

Database Systems

Lecture02 – Introduction to Relational Model

Beomseok Nam (남범석)

bnam@skku.edu

Relation

- Relation = Table

The diagram illustrates a relational table with four columns: *ID*, *name*, *dept_name*, and *salary*. Four arrows point from the text "attributes (or columns)" to the column headers. Three arrows point from the text "tuples (or rows or records)" to the first three data rows.

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000
45565	Katz	Comp. Sci.	75000
58583	Califieri	History	62000
76543	Singh	Finance	80000
76766	Crick	Biology	72000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000

Domain: Attribute Types

_domain

- **Domain** is the set of allowed values for each attribute

- Domain = data type + constraints
- Eg. score domain: 0~100

값이 목록

- Attribute values need to be **atomic**

Name_N_ID	Department
Alice 001	Comp. Sci.
Bob 002	Elec. Eng.
Charlie 003	NULL



Name	ID	Department
Alice	001	Comp. Sci.
Bob	002	Elec. Eng.
Charlie	003	NULL



- The special value **null** is a member of every domain

Relation Schema and Instance

- A_1, A_2, \dots, A_n are **attributes**
- $R = (A_1, A_2, \dots, A_n)$ is a **relation schema**

Example:

instructor = (*ID*, *name*, *dept_name*, *salary*)



- a **relation** is an instance, i.e., a set of **tuples**

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000



Relations are Unordered

- Order of tuples is irrelevant (tuples may be stored in an arbitrary order)
 - i.e., a relation is a set.
- Example: instructor relation with unordered tuples

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
22222	Einstein	Physics	95000
12121	Wu	Finance	90000
32343	El Said	History	60000
45565	Katz	Comp. Sci.	75000
98345	Kim	Elec. Eng.	80000
76766	Crick	Biology	72000
10101	Srinivasan	Comp. Sci.	65000
58583	Califieri	History	62000
83821	Brandt	Comp. Sci.	92000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
76543	Singh	Finance	80000

Database

- A database consists of multiple relations
- e.g.) a university database consists of

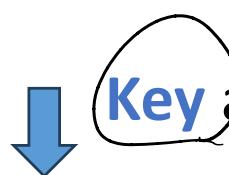
instructor relation

student relation

advisor relation

Key

- Q: What attribute plays the most significant role in the following relation?



Key attribute(s) uniquely identify a record in a table

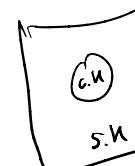
기본키

ID	name	dept_name	salary
22222	Einstein	Physics	95000
12121	Wu	Finance	90000
32343	El Said	History	60000
45565	Katz	Comp. Sci.	75000
98345	Kim	Elec. Eng.	80000
76766	Crick	Biology	72000
10101	Srinivasan	Comp. Sci.	65000
58583	Califieri	History	62000
83821	Brandt	Comp. Sci.	92000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
76543	Singh	Finance	80000

99999

Kim

(ID, name) \Rightarrow 주문번호 \Rightarrow key



\Rightarrow S.K D.C.K
+α

주문번호

superkey

Definition of Keys

- Let $K \subseteq R$
- K is a superkey of R if values for K are sufficient to identify a unique tuple in relation R .

- Example: $\{ID\}$ and $\{ID, name\}$ are both superkeys of *instructor*.

설명

= superset of key
= ↗설명의 키

- Minimal superkey is called candidate key

Example: $\{ID\}$ is a candidate key for *Instructor*

- One of the candidate keys is selected to be the primary key.

Relational Query Languages

- Two mathematical Query Languages

~~• Relational algebra~~ ← SQL

- More operational
- Useful for representing execution plans

~~• Relational calculus~~

- More declarative, i.e., describes what they want rather than how to compute it.
- Rather theoretical, so not covered in this class
- Eg. Find students whose SSN is 10.

$$\{t \mid t \in STUDENT \wedge t[\underline{ssn}] = 123\}$$

Relational Query Languages

■ Relational operators

- 5 fundamental relational operators
 - Union (\cup)
 - Selection (σ)
 - Projection (Π)
 - Cartesian Product (\times)
 - Set Difference (-)

5개의 표준 연산자

Union (\cup)

■ Relations r, s :

A	B
α	1
α	2
β	1

r

A	B
α	2
β	3

s

$r \cup s$:

A	B
α	1
α	2
β	1
β	3



Set difference (-)

- Relations r, s :

A	B
α	1
α	2
β	1

r

A	B
α	2
β	3

s

$r - s$:

A	B
α	1
β	1

Common tuple

$$= (\alpha, 2)$$

매우 유익

Selection (σ)

- Selection picks rows

- Relation r

- Select tuples with
 $A=B$ and $D > 5$

Condition

- $\sigma_{A=B \wedge D > 5}(r)$

$\sigma(r)$ unary

A	B	C	D
α	α	1	7
α	β	5	7
β	β	12	3
β	β	23	10

ignore

select

select

A	B	C	D
α	α	1	7
β	β	23	10

Selection (σ) - Example

- Find ‘Comp. Sci.’ major student named “Kim”

Student

ID	name	dept_name
10	Kim	Comp. Sci.
20	Lee	Biology
30	Kim	Electrical Eng.

σ \wedge /and 등 중 하나가 true
name = ‘Kim’ \wedge dept_name = ‘Comp. Sci.’ (Student)

Projection (Π)

- Projection picks columns and removes duplicates
- Relation r :

A	B	C
α	10	1
α	20	1
β	30	1
β	40	2

- Select A and C
 - Projection
 - $\Pi_{A, C}(r)$

$\xrightarrow{\text{Projection}}$

A	C
α	1
α	1
β	1
β	2

=

A	C
α	1
β	1
β	2

$\hookrightarrow \beta$ 는 삭제

Projection (Π) - Example

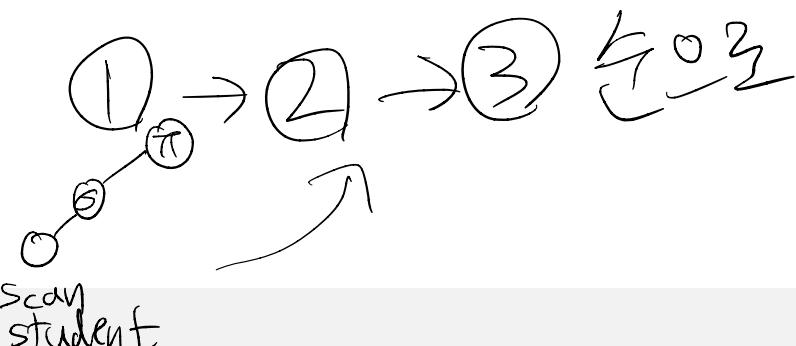
- Find ID of 'Comp. Sci.' major students.
 - Hint: relational operators can be cascaded

ID	name	dept_name
10	Kim	Comp. Sci.
20	Lee	Biology
30	Kim	Electrical Eng.

- $\Pi_{ID}(\sigma_{dept_name = 'Comp. Sci.'} (Student))$

③ ②

①

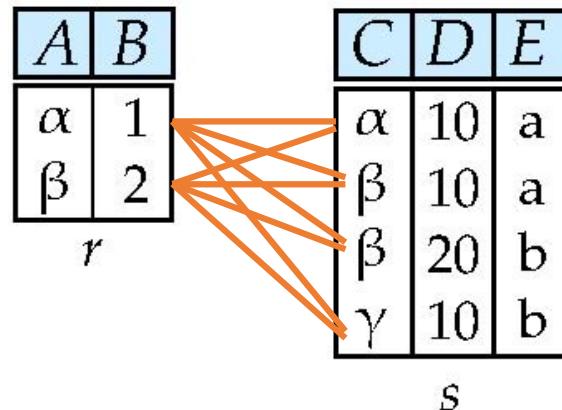


Cartesian Product (x)

$r \times s$

- Cartesian Product generates all possible pairs

- Relations r, s :



$r \times s$:

A	B	C	D	E
---	---	---	---	---

works as in nested loops:
for each row i in r
 for each row j in s
 output i, j

Cartesian Product (x)

- Cartesian Product generates all possible pairs

- Relations r, s :

Tuples of R

\cap

	A	B
α	1	
β	2	

	C	D	E
α	10	a	
β	10	a	
β	20	b	
γ	10	b	

$r \times s$:

	A	B	C	D	E
α	1		α	10	a
α	1		β	10	a
α	1		β	20	b
α	1		γ	10	b
β	2		α	10	a
β	2		β	10	a
β	2		β	20	b
β	2		γ	10	b

works as in nested loops:

for each row i in r

$n \times m$

for each row j in s
output i, j

Cartesian Product (x) - Example

여러개의 테이블을 연산할 때 오른쪽 사용

- Find name of students who take course 'swe3003'

Student

ID	name	dept_name
10	Kim	Comp. Sci.
20	Lee	Biology
...		

Take

ID	c_id	grade
10	swe2021	A
20	swe3003	B
30	sw4016	A

중복특성이 같아야 함 (결합성)

가장 먼저

ID	name	dept_name	ID	c_id	grade
10	Kim	Comp. Sci.	10	swe2021	A
10	Kim	Comp. Sci.	20	swe3003	B
10	Kim	Comp. Sci.	30	sw4016	A
20	Lee	Biology	10	swe2021	A
20	Lee	Biology	20	swe3003	B
20	Lee	Biology	30	sw4016	A

Student x Take

Cartesian Product (x) - Example

- Find name of students who take course 'swe3003'

Student

ID	name	dept_name
10	Kim	Comp. Sci.
20	Lee	Biology
...		

ID	c_id	grade
10	swe2021	A
20	swe3003	B
30	sw4016	A

Take

ID	name	dept_name	ID	c_id	grade
10	Kim	Comp. Sci.	10	swe2021	A
10	Kim	Comp. Sci.	20	swe3003	B
10	Kim	Comp. Sci.	30	sw4016	A
20	Lee	Biology	10	swe2021	A
20	Lee	Biology	20	swe3003	B
20	Lee	Biology	30	sw4016	A

$\sigma_{\text{Student.ID}=\text{Take.ID}} (\text{Student} \times \text{Take})$

Cartesian Product (x) - Example

- Find name of students who take course 'swe3003'

Student

ID	name	dept_name
10	Kim	Comp. Sci.
20	Lee	Biology
...		

Take

ID	c_id	grade
10	swe2021	A
20	swe3003	B
30	sw4016	A

student \bowtie Take

ID	name	dept_name	ID	c_id	grade
10	Kim	Comp. Sci.	10	swe2021	A
10	Kim	Comp. Sci.	20	swe3003	B
10	Kim	Comp. Sci.	30	sw4016	A
20	Lee	Biology	10	swe2021	A
20	Lee	Biology	20	swe3003	B
20	Lee	Biology	30	sw4016	A

$\sigma_{\text{Student.ID} = \text{Take.ID} \wedge \text{c_id} = \text{'swe3003'}} (\text{Student} \times \text{Take})$

Cartesian Product (x) - Example

- Find name of students who take course 'swe3003'

Student

ID	name	dept_name
10	Kim	Comp. Sci.
20	Lee	Biology
...		

ID	c_id	grade
10	swe2021	A
20	swe3003	B
30	sw4016	A

Take

ID	name	dept_name	ID	c_id	grade
10	Kim	Comp. Sci.	10	swe2021	A
10	Kim	Comp. Sci.	20	swe3003	B
10	Kim	Comp. Sci.	30	sw4016	A
20	Lee	Biology	10	swe2021	A
20	Lee	Biology	20	swe3003	B
20	Lee	Biology	30	sw4016	A

$\Pi_{\text{name}} (\sigma_{\text{Student.ID} = \text{Take.ID} \wedge \text{c_id} = \text{'swe3003'}} (\text{Student} \times \text{Take}))$

$\text{10} \quad \text{20}$
 $\text{10} \quad \text{20}$
 $\text{10} \quad \text{20}$

5 relational operators are enough! But,

- 5 fundamental relational operators are enough to write any query
 - Union (\cup)
 - Selection (σ)
 - Projection (Π)
 - Cartesian Product (\times)
 - Set Difference (-)
- But, derived operators, for convenience
 - Set intersection (\cap)
 - Join (\bowtie)
 - Rename (ρ)
 - Division (\div)
 - Group by (G)

Set Intersection of two relations (\cap)

- Relation r, s :

A	B
α	1
α	2
β	1

r

A	B
α	2
β	3

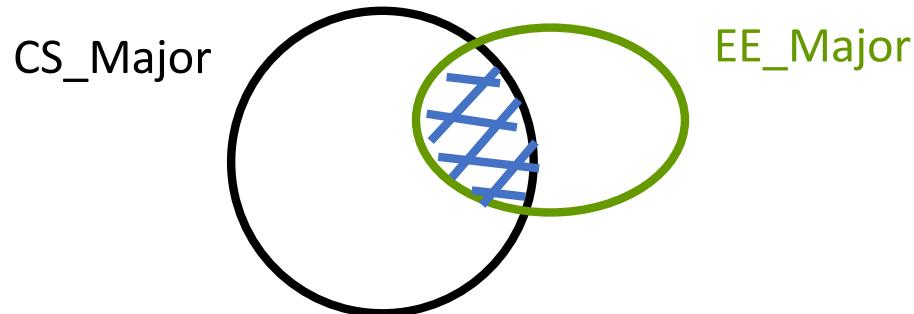
s

$r \cap s$

A	B
α	2

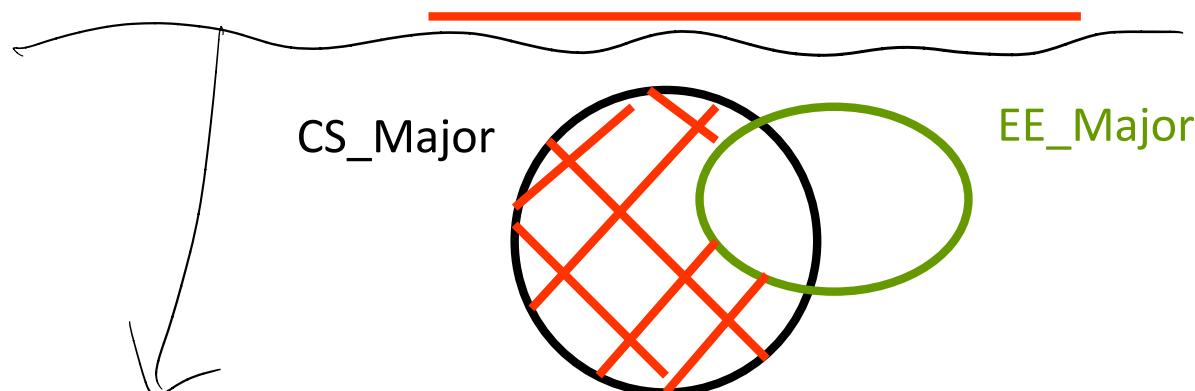
Observation

∧ 를 5가지 연산으로 바꾸는 방법



$$\blacksquare \text{ CS}_\text{Major} \cap \text{EE}_\text{Major} =$$

$$\text{CS}_\text{Major} - (\text{CS}_\text{Major} - \text{EE}_\text{Major})$$



So, \cap is not fundamental

Natural Join (\bowtie)

▪ Natural Join

- $R \bowtie S = \sigma_{R.a=S.a \wedge R.b=S.b \wedge \dots} (R \times S)$

$\begin{array}{c} \exists H \\ \sqsubseteq_0 \quad \underline{abcd} \quad \underline{ab} \rho \eta \end{array}$

- For each pair of tuples t_R from R and t_S from S ,
- If t_R and t_S have the same values on common attributes, add a tuple t to the result, where
 - t has the same value as t_R
 - t has the same value as t_S
 - Drop duplicate attributes

$\begin{array}{c} \exists \rho \\ \sqsubseteq_0 \quad \underline{abcd} \end{array}$

Natural Join (\bowtie) - Example

- Find name of students who take course 'swe3003'

Student

ID	name	dept_name
10	Kim	Comp. Sci.
20	Lee	Biology
30	Kim	Electrical Eng.

Take

ID	c_id	grade
10	swe2021	A
20	swe3003	B
30	sw4016	A

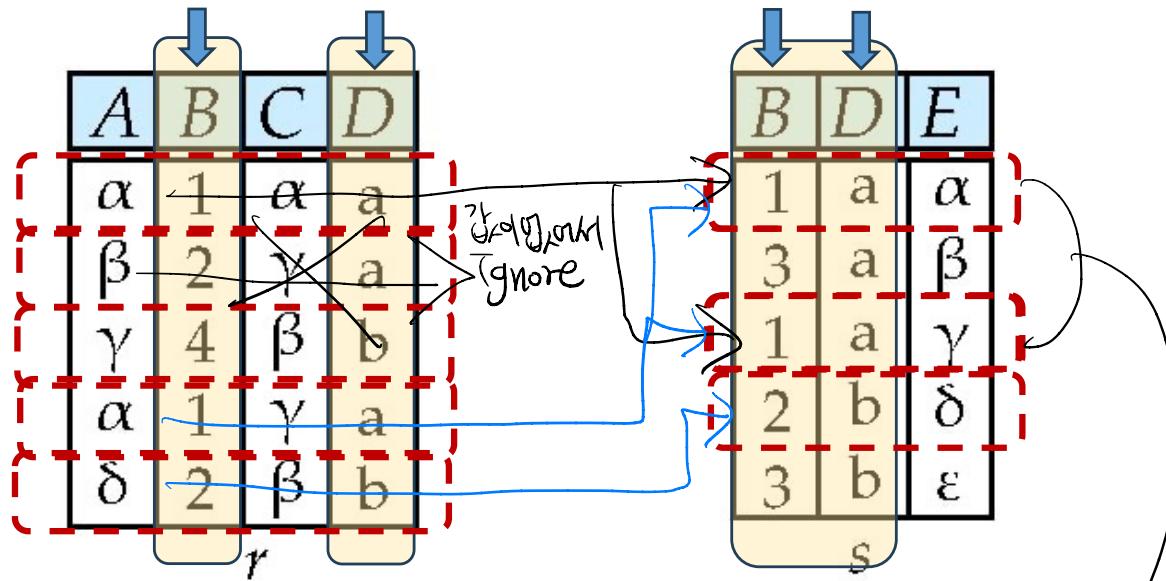
ID	name	dept_name	c_id	grade
10	Kim	Comp. Sci.	swe2021	A
20	Lee	Biology	swe3003	B
30	Kim	Electrical Eng.	sw4016	A

이 풀가
Student \bowtie Take

$$\Pi_{\text{name}} (\sigma_{\text{c_id}=\text{'swe3003'}} (\text{Student} \bowtie \text{Take}))$$

Natural Join with multiple common attributes

- Relations r, s:



- Natural Join

$r \bowtie s$

A	B	C	D	E
α	1	α	a	α
α	1	α	a	γ
α	1	γ	a	α
α	1	γ	a	γ
δ	2	β	b	δ

같은 특징 & 결합하는
행 찾기

rename (ρ)

$$\rho_{\text{AFTER}}(\underbrace{\text{BEFORE}}_{\text{before}})$$

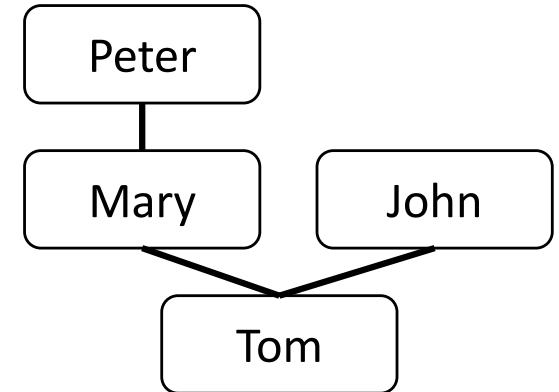
new before

- Q: why?
- A: The same table is used multiple times



rename (ρ)

- find the grand-parents of ‘Tom’,
given $\underbrace{\text{PC}(\text{parent-id}, \text{child-id})}$



parent	child
Mary	Tom
Peter	Mary
John	Tom

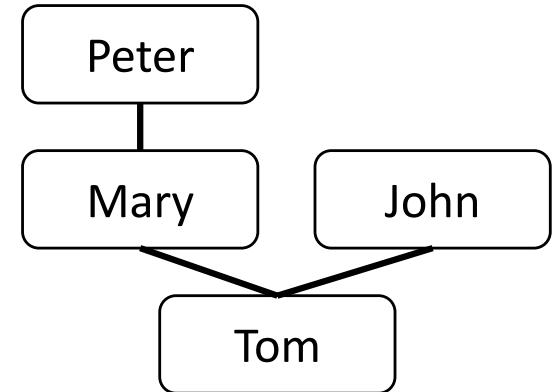
rename (ρ)

- find the grand-parents of ‘Tom’, given PC(parent-id, child-id)

parent	child
Mary	Tom
Peter	Mary
John	Tom

parent	child
Mary	Tom
Peter	Mary
John	Tom

copy!



- We need two different names for the same table - hence, the ‘rename’ op.

$\pi_{PC.parent}(\sigma_{PC1.child=PC.parent}(\rho_{PC1}(PC) \times PC))$

$\nwarrow_{PC.child='Tom'}$

John 조건을 만들기 위해
부모를 찾을 때(이중의)
이름을 연결함

Division (\div)

- Useful for expressing queries like:

Find students who have taken all CS courses.

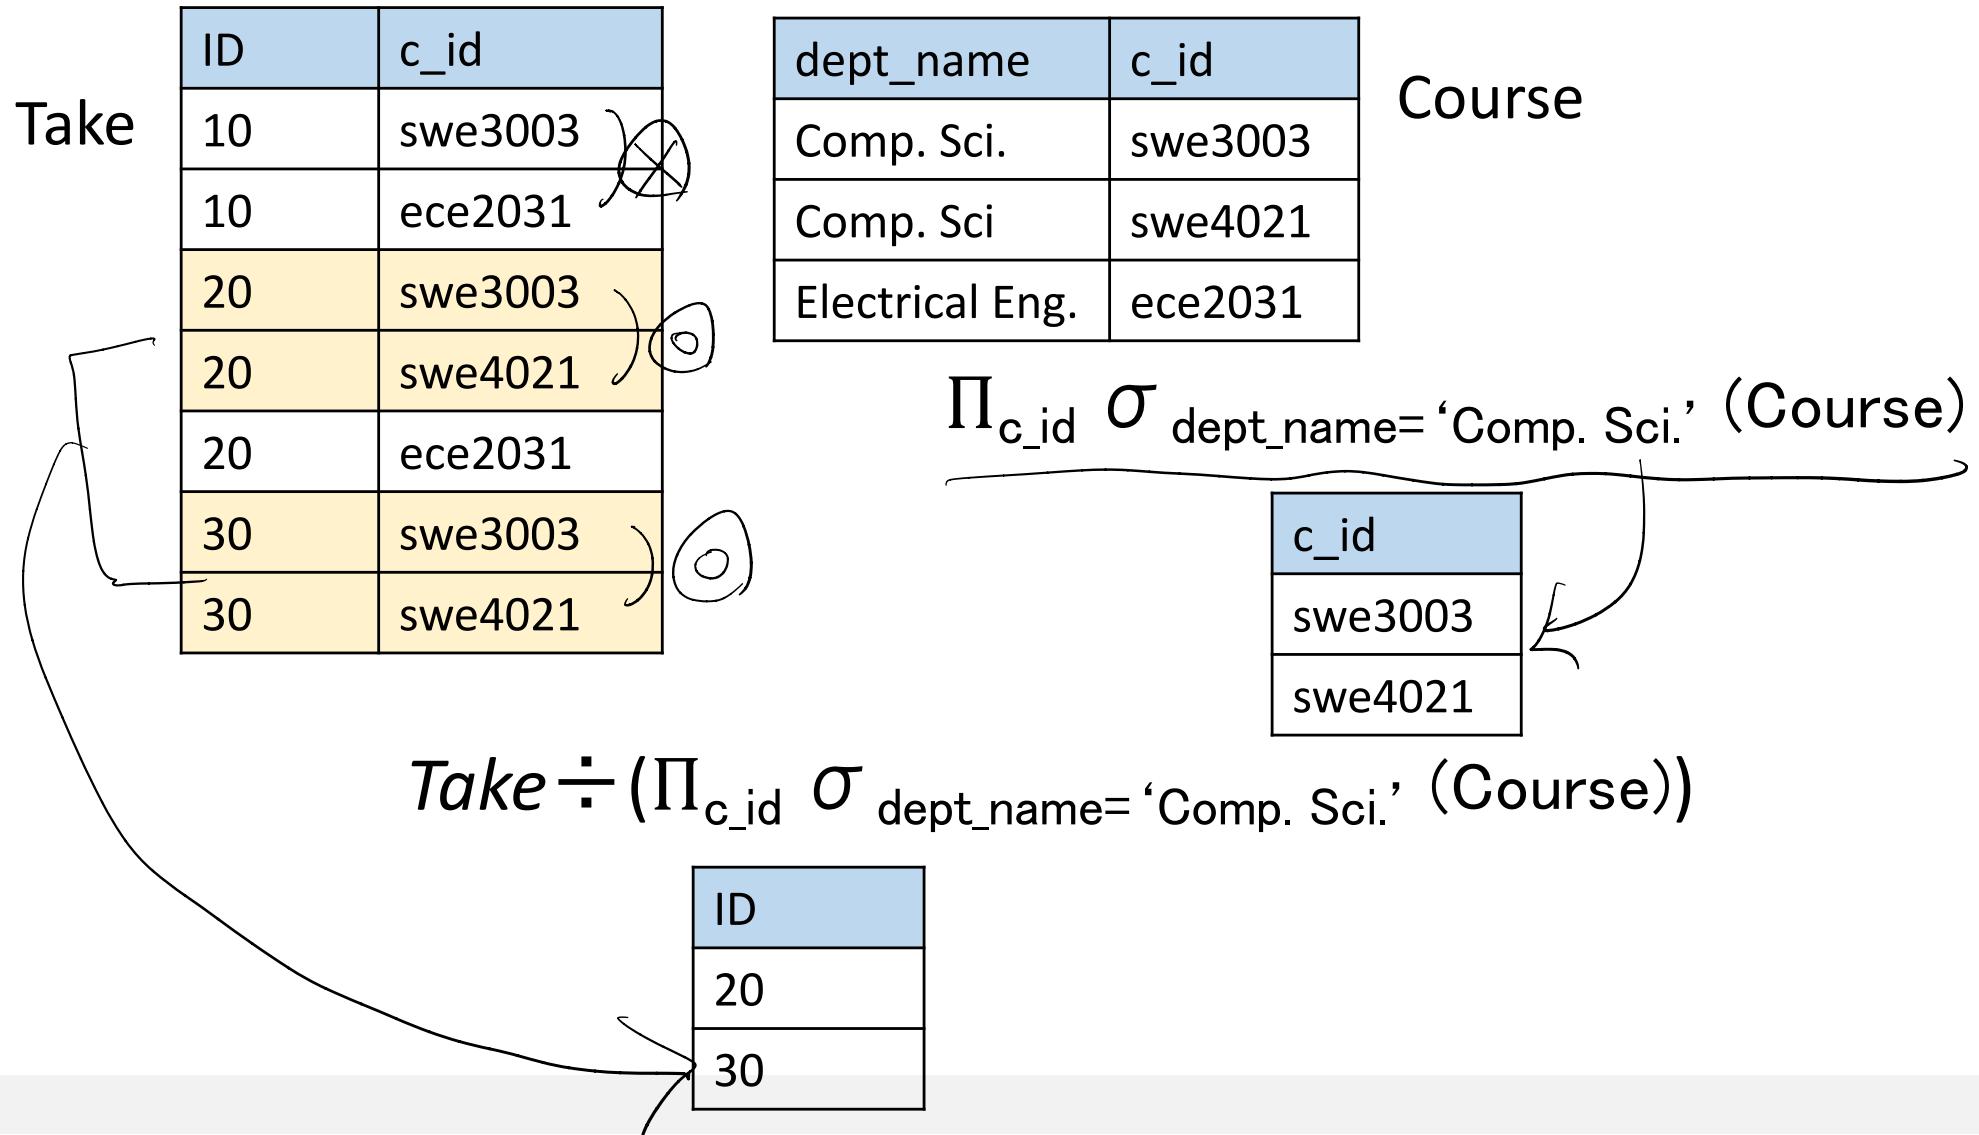
- Suppose R has 2 attributes (x, y) and S has (y) .
- Then $R \div S$ returns the set of x 's that match all y values in S .
- Example: $R = \text{Friend}(x, y)$. $S = \text{set of 10 students}$.
 - Then $R \div S$ returns the set of all x 's that are friends with all 10 students.



나는는 S 에(이)에 있는
모든값이 있는 R 에(이)에
남김

Division (\div) - Example

- Find ID of students who take all ‘Comp. Sci’ courses



Division (\div) - Example

- Find ID of students who take all ‘Comp. Sci’ courses

R	ID	proj_no
	10	p1
	10	p2
	10	p3
	10	p4
	20	p1
	20	p2
	30	p2
	40	p2
	40	p4

proj_no
p2

~~10~~ S_1

proj_no
p2
p4

~~20~~ S_2

proj_no
p1
p2
p4

~~30~~ S_3

ID
10
20
30
40

$R \div \textcircled{10}$ S_1

ID
10
40

$R \div \textcircled{20}$ S_2

ID
10

$R \div \textcircled{30}$ S_3

Observations

- Division (\div) is reverse of cartesian product
- It can be derived from the 5 fundamental operators

$$r \div s = \pi_{(R-S)}(r) - \pi_{(R-S)}[(\pi_{(R-S)}(r) \times s) - r]$$

Let's practice using the following schema

- branch (branch-name, branch-city, assets)
- customer (customer-name, customer-street, customer-only)
- account (account-number, branch-name, balance)
- loan (loan-number, branch-name, amount)
- depositor (customer-name, account-number)
- borrower (customer-name, loan-number)

Example Queries in Relational Algebra

loan (loan-number, branch-name, amount)

- Find all attributes of loans of over \$1200

$$\sigma_{\text{amount} > 1200} (\text{loan})$$

- ↗ one column
- Find the loan number for each loan of an amount greater than \$1200

$$\Pi_{\text{loan-number}} (\sigma_{\text{amount} > 1200} (\text{loan}))$$

Example Queries in Relational Algebra

depositor (customer-name, account-number)

borrower (customer-name, loan-number)

- Find the names of all customers who have a loan, an account, **or** both, from the bank

빌려거나 빌려온 사람 모두

$$\prod_{\text{customer-name}} (\text{borrower}) \cup \prod_{\text{customer-name}} (\text{depositor})$$

- Find the names of all customers who have a loan **and** an account at bank.

빌리고 빌려온 사람

$$\prod_{\text{customer-name}} (\text{borrower}) \cap \prod_{\text{customer-name}} (\text{depositor})$$

Example Queries in Relational Algebra

loan (loan-number, branch-name, amount)

borrower (customer-name, loan-number)

depositor (customer-name, account-number)

- Find the names of customers who have a loan at Suwon branch.

$\prod \text{customer-name} (\sigma_{\text{branch-name} = \text{"Suwon"}} (\text{borrower} \bowtie \text{loan}))$

$\prod \text{customer-name} (\sigma_{\text{b-name} = \text{"Suwon"}} (\text{loan}) \bowtie \text{borrower})$

- Find the names of all customers who have a loan at the Suwon branch but do not have an account at any branch of the bank.

$\prod \text{customer-name} (\sigma_{\text{branch-name} = \text{"Suwon"}} (\text{borrower} \bowtie \text{loan}))$

- $\prod \text{customer-name} (\text{depositor})$

$$a * (b + c) = a * b + a * c$$

commutative associative rule