

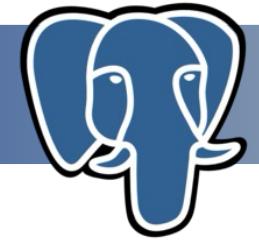
# Database Systems

## Lecture03 & 04 – Ch 3. Introduction to SQL

Beomseok Nam (남범석)

[bnam@skku.edu](mailto:bnam@skku.edu)

# PostgreSQL



PostgreSQL

- A Relational Database Management System
- PostgreSQL (pronounced “POST-gress” and QL is silent)
- ANSI-SQL:2008 compatible (Most Oracle-compatible)

# DBA creates Users and Databases

- When PostgreSQL is installed, DBA account ‘postgres’ is created
- ‘postgres’ can login to PostgreSQL DBMS by running ‘psql’

```
postgres@swui:~$ psql  
psql (14.7 ...)  
postgres=#
```

- To create a new user, DBA (postgres) runs the following statements in psql

```
postgres=> CREATE USER db20333999 with  
encrypted password 'changethis';
```

- To create a database, DBA runs the following statements in psql

```
postgres=> CREATE DATABASE db20333999;  
postgres=> GRANT ALL PRIVILEGES on database  
db20333999 TO db20333999;
```

# User Login to PostgreSQL

- Users can login to psql as follows

```
2020333999@swye:~$ psql -h localhost -U db20333999
Password for user db20333999: changethis <ENTER>
psql (14.7 (...))
SSL connection (protocol: TLSv1.3, ...)
Type "help" for help.
```

```
db20333999=>
```

# PostgreSQL: psql Command Line options

- -A, --no-align: set output to non-aligned, no padding
- -c sql, --command sql: execute the sql command
- **-d name, --dbname name**: name of database
- -f name, --file name: use name as the source of commands
- -o name, --output name: put the output in name
- -q, --quiet: suppress welcome messages
- -t, --tuples-only: suppress print column names, result row counters, etc
- -?, --help: get command line help

# psql Meta-commands

- `\l`: list existing databases
- `\c db_name`: connect to a database name
- `\d`: list all tables (display)
- `\d table`: show information about a *table*
- `\copy table`: copy a *table* to/from a file
- `\e`: edit the query buffer
- `\p`: display the query buffer
- `\g`: execute the query buffer (go)
- `\r`: resets (clears) the query buffer
- `\w name`: write the query buffer to *name*

# List & Choosing Databases in PostgreSQL

- Use the `\l` command to find out which databases currently exist on the server:

```
db20333999 => \l
```

List of databases

Name	Owner	Encoding	Collate	Ctype	privileges
db16311252	postgres	UTF8	en_US.UTF-8	en_US.UTF-8	
db16311648	postgres	UTF8	en_US.UTF-8	en_US.UTF-8	
db20333999	postgres	UTF8	en_US.UTF-8	en_US.UTF-8	
...					

```
db20333999=> \c db20333999
```

```
psql (14.7 (...))
```

```
SSL connection (protocol: TLSv1.3, ...)
```

```
You are now connected to database "db20333999" ...
```

```
db20314193=>
```

# psql Meta-commands

- `\i file_name`: read *file* into query buffer (input)
- `\q`: quit the program
- `\! Command`: execute the Linux *command*
- `\h command`: display help on command
- `\?:` help on meta-commands

# Creating a Table

- Once you choose a database, you can display tables in the database:

```
db20333999 => \dt
```

```
Did not find any relations.
```

- An empty set indicates that you have not created any table yet.

# Loading Sample Data

- Create a text file `pet.txt` containing one record per line.
- Values must be separated by commas or tabs, and given in the order in which the columns were listed in the CREATE TABLE statement.
- Then load the data via the COPY Command.

# Sample Data File

Fluffy	Harold	cat	f	1993-02-04	\N
Claws	Gwen	cat	m	1994-03-17	\N
Buffy	Harold	dog	f	1989-05-13	\N
Fang	Benny	dog	m	1990-08-27	\N
Bowser	Diane	dog	m	1979-08-31	1995-07-29
Chirpy	Gwen	bird	f	1998-09-11	\N
Whistler	Gwen	bird	\N	1997-12-09	\N
Slim	Benny	snake	m	1996-04-29	\N

To Load pet.txt:

db20333999 => COPY myTable FROM 'pet.txt' (FORMAT CSV,  
DELIMITER ('\t'));

# Run SQL source code

- You can create a SQL source file `test.sql` containing SQL statements using a text editor and run it in psql shell.

To Run test.sql:

```
db20333999 => \i test.sql;
```

# University Database Example

- in ~swe3003/sample\_db

To Create and populate University Database:

```
db20333999 => \i /home/swe3003/sample_db/DDL.sql;  
db20333999 => \i /home/swe3003/sample_db/data.sql;
```

DDL, SQL  
3 Entity 테이블  
Data.sql  
설계 Entity 테이블

# Chapter 3: Introduction to SQL

# DDL: Data Definition Language

- The SQL **data-definition language (DDL)** allows the specification about relations, including:
  - The schema for each relation.
  - The domain of values associated with each attribute.
  - Integrity constraints
  - And as we will see later, also other information such as
    - The set of indexes to be maintained for each relations.
    - Security and authorization information for each relation.
    - The physical storage structure of each relation on disk.

# Domain Types in SQL

- **char(n).** Character string, with fixed length  $n$ .
- **varchar(n).** Variable length string, with maximum length  $n$ .
- **int.** Integer
- **smallint.** Small integer
- **numeric(p,d).** Fixed point number
  - with precision of  $p$  digits, with  $d$  digits to the right of decimal point.
- **real.** Floating point number
- **float(n).** Floating point number
  - with user-specified precision of at least  $n$  digits.

# Create Table

```
create table TableName (Attribute1 DataType1,  
                          Attribute2 DataType2,  
                          ...,  
                          Attributen DataTypen,  
                          [integrity-constraint],  
                          ...))
```

- Example:

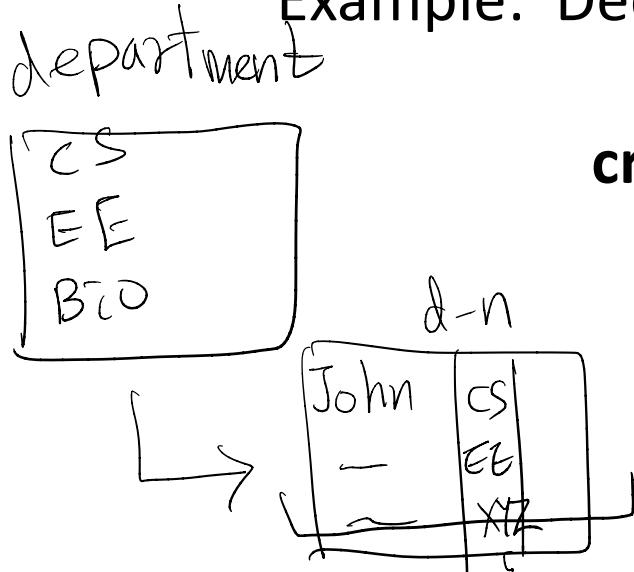
```
create table instructor (  
    ID              char(5),  
    name            varchar(20) not null,  
    dept_name      varchar(20),  
    salary         numeric(8,2))
```

↑ name or null  
0123456789EX

# Integrity Constraints in Create Table

- **not null** → 제작자를 말하는 데이터입력시
- **primary key ( $A_1, \dots, A_n$ )** → 조건을 만족시키지 못하면 알림을 가능
- **foreign key ( $A_m, \dots, A_n$ ) references r**

Example: Declare `dept_name` as the primary key for `department`



create table `instructor` (  
    `ID`           `char(5)`,  
    `name`       `varchar(20) not null`,  
    `dept_name`   `varchar(20)`,  
    `salary`       `numeric(8,2)`,  
    **primary key (`ID`)**,  
    **foreign key (`dept_name`) references department**)

primary key automatically ensures not null

Foreign Key 외래키  
외래키에 해당하는 특장에 사용되는  
다른 테이블에  
값이 존재하지 않으면 해당테이블에 입력가능

ALTER TABLE

저장되는 값을  
설정위해

department table의  
primary key이자

# And a Few More Relation Definitions

- **create table student (**  
*ID*           **varchar(5),**  
*name*       **varchar(20) not null,**  
*dept\_name*   **varchar(20),**  
*tot\_cred*     **numeric(3,0),**  
**primary key (ID),**  
**foreign key (dept\_name) references department );**

- **create table takes (**  
*ID*           **varchar(5),**  
*course\_id*   **varchar(8),**  
*sec\_id*       **varchar(8),**  
*semester*     **varchar(6),**  
*year*          **numeric(4,0),**  
*grade*        **varchar(2),**  
**primary key (ID, course\_id, sec\_id, semester, year),**  
**foreign key (ID) references student,**  
**foreign key (course\_id, sec\_id, semester, year) references section );**

상당히 많은 대가로 한 번도 조건

Primary Key

## And more still

- `create table course (  
 course_id varchar(8) primary key,  
 title varchar(50),  
 dept_name varchar(20),  
 credits numeric(2,0),  
 foreign key (dept_name) references department) ;`

# Drop and Alter Table Constructs

## ■ **drop table** *student*

- Deletes the table schema and its contents

## ■ **delete from** *student* -- DML

- Deletes all contents of table, but retains table schema

## ■ **alter table**

### • **alter table** *r* add *A D*

- *A* is the name of the attribute to be added to relation *r*
- *D* is the domain of *A*.
- All existing tuples in the relation get *null* for the new attribute.

### • **alter table** *r* drop *A*

*↳ SQL standard*

- Dropping of attributes is not supported by many databases

# DML (Data-Manipulation Language)

## ■ DML

- select / insert / update / delete
- set operations
- ordering
- aggregate functions
- nested subqueries

SUM, MAX, MIN

R. Al Redundant

# DML (Data-Manipulation Language)

- A typical SQL query :

```
select A1, A2, ..., An
from r1, r2, ..., rm
where P
```

$\rightarrow$  Set  
Relation

- Note: The result of an SQL query is also a relation.

# The **select** Clause

- The **select** clause lists the attributes
  - corresponds to the projection operation (  $\Pi$  )
- Example: find the names of all instructors:

```
select name  
from instructor
```

- NOTE: SQL are case insensitive
  - i.e., upper- or lower-case letters do not matter.
  - E.g.  $Name \equiv NAME \equiv name$

## The **select** Clause (Cont.)

- SQL allows duplicates in relations
- To eliminate duplicates, use **distinct** after select.

↑  
distinct

keyword

- Find all department names of instructor, and remove duplicates

```
select distinct dept_name  
from instructor
```

↑  
distinct

- The keyword **all** allows duplicates to be retained.

```
select all dept_name  
from instructor
```

# The **Select** Clause (Cont.)

- An asterisk ( \* ) denotes “all attributes”

```
select *  
from instructor
```

- The **select** clause can use arithmetic expressions

- +, -, \*, and /

- E.g.:

```
select ID, name, salary/12  
from instructor
```

Salary / 12 한 달

# The **where** Clause

- The **where** specifies conditions that the result must satisfy
  - boolean operators (**and or not ...**)
  - comparison operators (<, >, =, ...)
  - and more...
- Find all instructors in Comp. Sci. dept with salary > 80000

```
select name  
from instructor  
where dept_name = 'Comp. Sci.' and  
      salary > 80000
```

# The **from** Clause

- The **from** clause lists the relations

- Corresponds to the Cartesian product ( X )

**select \***

**from instructor, teaches**

Generates Cartesian product *instructor X teaches*, with all attributes from both relations

{ 127 }

select : 선택

where : 조건

from : 어떤 틀에 대해서, (,) 는 X 곱집합(의)

## Join ( $\bowtie$ )

대기 대상

R<sub>0</sub>att, S<sub>0</sub>att - R<sub>0</sub>att

- $R \bowtie S = \sigma_{R.common\_attrs=S.common\_attrs} (R \times S)$

↳ 조건부 결합

- Find the names of instructors and the course ID of the courses they taught.

```
select name, course_id  
from instructor, teaches  
where instructor.ID = teaches.ID
```

# Join ( $\bowtie$ ) - Example

- Find the course ID, semester, year and title of each course offered by the Comp. Sci. department

*select section.course\_id, semester, year, title  
from section, course  
where section.course\_id = course.course\_id  
and dept\_name = 'Comp. Sci.'*

Filter condition

natural Join

고정  
X 흔에  
가장 가능 학생수  
EC 오거나  
가장이 or.

section
course_id
sec_id
semester
year
building
room_no
time_slot_id

course
course_id
title
dept_name
credits

# Natural Join

- Natural join retains only one copy of common column
  - **select \***  
**from *instructor* natural join *teaches*;**

→ 광동 특유의 자동차로 차(21)

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>	<i>course_id</i>	<i>sec_id</i>	<i>semester</i>	<i>year</i>
10101	Srinivasan	Comp. Sci.	65000	CS-101	1	Fall	2009
10101	Srinivasan	Comp. Sci.	65000	CS-315	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	CS-347	1	Fall	2009
12121	Wu	Finance	90000	FIN-201	1	Spring	2010
15151	Mozart	Music	40000	MU-199	1	Spring	2010
22222	Einstein	Physics	95000	PHY-101	1	Fall	2009
32343	El Said	History	60000	HIS-351	1	Spring	2010
45565	Katz	Comp. Sci.	75000	CS-101	1	Spring	2010
45565	Katz	Comp. Sci.	75000	CS-319	1	Spring	2010
76766	Crick	Biology	72000	BIO-101	1	Summer	2009
76766	Crick	Biology	72000	BIO-301	1	Summer	2010

# Natural Join Example

- List the names of instructors along with the course ID of the courses that they taught.

- **select name, course\_id  
from instructor, teaches  
where instructor.ID = teaches.ID;**

$\Rightarrow \oplus =$

○ 거냥 naturalJoin

○ select,

- **select name, course\_id  
from instructor natural join teaches;**

# Danger in Natural Join

- Danger in natural join: beware of unrelated attributes with same name which get equated incorrectly

*natural join*

Instructor			Teaches			
ID	name	dept_name	ID	c_id	semester	year
10	Nam	Comp. Sci.	10	swe3003	Spring	2025
20	Alice	Electrical Eng.	20	swe3003	Fall	2024
			10	swe3021	Spring	2024

List the names of  
instructors along with the  
titles of courses that  
they teach

*natural join*

*OK (?)*

Course		
c_id	title	dept_name
swe3003	DB	'Comp. Sci'
swe3021	Multicore	'Comp. Sci'

*Alice → swe3003*

*Comp. Sci*

*Electrical*

*- "IH"*

**select name, title**

**from instructor natural join teaches natural join course;**

# Danger in Natural Join

Instructor

ID	name	dept_name
10	Alice	Comp. Sci.
20	Bob	Electrical Eng.

Teaches

ID	c_id	semester	year
10	swe3003	Spring	2025
20	swe3003	Fall	2024
10	swe3004	Spring	2024

List the names of  
instructors along with the  
the titles of courses that  
they teach

Course

c_id	title	dept_name
swe3003	DB	'Comp. Sci'
swe3004	OS	'Comp. Sci'

**select name, title**

**from instructor natural join teaches natural join course;**

name	title
Alice	DB
Alice	OS

Uh oh, ( Bob, DB ) is missing because  
course.dept\_name is different from  
instructor.dept\_name

Bob  
DB

# Danger in Natural Join

Instructor

ID	name	dept_name
10	Alice	Comp. Sci.
20	Bob	Electrical Eng.

Teaches

ID	c_id	semester	year
10	swe3003	Spring	2025
20	swe3003	Fall	2024
10	swe3004	Spring	2024

List the names of  
instructors along with the  
the titles of courses that  
they teach

Course

c_id	title	dept_name
swe3003	DB	'Comp. Sci'
swe3004	OS	'Comp. Sci'

Correct version:

**select name, title**

**from instructor natural join teaches, course**

**where teaches.course\_id = course.course\_id;**

cartesian 사용 후  
직접 C\_id 확인  
하는게 좋다

이거

# Rename ( $\rho$ ) Operation

- as clause renames relations and attributes :

*old-name as new-name*

- E.g.

- **select** *ID, name, salary/12 as monthly\_salary*  
**from** *instructor*

할수는 있지만  
아직은 가능하지 않음

- Keyword as is optional and may be omitted

*instructor as T*  $\equiv$  *instructor T*

- Keyword as must be omitted in Oracle DBMS

# String Operations

- string-matching operators: = and **like**
  - percent % - matches any substring.
  - underscore \_ - matches any single character.

- Find the names of all instructors whose name includes the substring “dar”.

```
select name  
from instructor  
where name like '%dar%'
```

like  
✓  
dar로 시작하는 걸 찾기 → %dar

.. 끝나는 걸 찾기 → dar%

dar를 포함하는

%를 포함한 문자열

- Match the string “100 %”

like '100 \%' escape '\'

# String Operations (Cont.)

- Patterns are case sensitive.
- Pattern matching examples:
  - ‘Intro%’ matches any string beginning with “Intro”.
  - ‘%Comp%’ matches any string containing “Comp”
  - ‘\_\_\_’ matches any string of exactly three characters.  
*3개의 문자로 구성된*
  - ‘\_\_\_%’ matches any string of at least three characters.  
*최소 3개 이상*
- SQL supports a variety of string operations such as
  - concatenation (using “||”)      *'abc' || 'def' → 'abcdef'*
  - converting from upper to lower case (and vice versa)
    - UPPER(), LOWER()
  - finding string length, extracting substrings, etc.
    - LENGTH(), SUBSTRING(str, position, length)  
*선택한 문자열 추출 가능*

# Sorting: **order by** Clause

- List in alphabetic order the names of all instructors

```
select name  
from instructor  
order by name
```

- **asc** for ascending order (default)

ID	name	dept_name	salary
20	Alice	Comp. Sci.	8,000
30	Charlie	Biology	null
40	Bob	Comp. Sci.	6,000



name
Alice
Bob
Charlie

- **desc** for descending order

- ```
select name  
from instructor  
order by name desc
```



| name    |
|---------|
| Charlie |
| Bob     |
| Alice   |

# Sorting: **order by** Clause

- Can sort on multiple attributes (lexicographic sorting)

- select dept\_name, name  
from instructor  
order by dept\_name, name**

select 키워드로 order by 가능

{st}

| ID | name    | dept_name  | salary |
|----|---------|------------|--------|
| 20 | Alice   | Comp. Sci. | 8,000  |
| 30 | Charlie | Biology    | null   |
| 40 | Bob     | Comp. Sci. | 6,000  |

| dept_name  | name    |
|------------|---------|
| Biology    | Charlie |
| Comp. Sci. | Alice   |
| Comp. Sci. | Bob     |

null은 가장 뒤로 정렬됨

# Range query and tuple comparison

- SQL includes a **between** comparison operator 비교
- Example: Find the names of all instructors with salary between \$90,000 and \$100,000 (that is,  $\geq \$90,000$  and  $\leq \$100,000$ )
  - **select name  
from instructor  
where salary between 90000 and 100000**
- Tuple comparison
  - **select name, course\_id  
from instructor, teaches  
where (instructor.ID, dept\_name) =  
(teaches.ID, 'Biology');**

# Set Operations

- Set operations **union**, **intersect**, and **except**
  - Each of the above operations automatically eliminates duplicates
- To retain all duplicates use the corresponding multiset operations **union all**, **intersect all** and **except all**.

# Set Operations

- Find courses that were offered in Fall 2024 or in Spring 2025

*(select c\_id from course where semester = 'Fall' and year = 2024)*  
*union*  
*(select c\_id from course where semester = 'Spring' and year = 2025)*

- Find courses that were offered in Fall 2024 and in Spring 2025

*(select c\_id from section where semester = 'Fall' and year = 2024)*  
*intersect*  
*(select c\_id from section where semester = 'Spring' and year = 2025)*

- Find courses that were offered in Fall 2024 but not in Spring 2025

*(select c\_id from section where semester = 'Fall' and year = 2024)*  
*except*  
*(select c\_id from section where semester = 'Spring' and year = 2025)*

# Null Values

- Tuples may have a null value, denoted by *null*
- *null* signifies an **unknown** value or that a value does not exist.
- The result of any arithmetic expression involving *null* is *null*
  - Example:  $5 + \text{null}$  returns null
- The predicate **is null** can be used to check for null values.
  - Example: Find all instructors whose salary is null.

```
select name  
from instructor  
where salary is null
```

|  | Salary |
|--|--------|
|  | null   |
|  | ○      |

null값은 확정하기 어렵네  
사칙연산을 하면 null임  
unknown value ↗

# Null Values and Three Valued Logic

- Any comparison with *null* returns *unknown*
  - Example:  $5 < \text{null}$  or  $\text{null} < > \text{null}$  or  $\text{null} = \text{null}$
- Three-valued logic using the truth value *unknown*:
  - OR:  $(\text{unknown or true}) = \text{true}$ ,  
 $(\text{unknown or false}) = \text{unknown}$   
 $(\text{unknown or unknown}) = \text{unknown}$
  - AND:  $(\text{true and unknown}) = \text{unknown}$ ,  
 $(\text{false and unknown}) = \text{false}$ ,  
 $(\text{unknown and unknown}) = \text{unknown}$
  - NOT:  $(\text{not unknown}) = \text{unknown}$
  - “*P is unknown*” evaluates to true if predicate *P* evaluates to *unknown* *null*
- Result of **where** clause predicate is treated as *false* if it evaluates to *unknown*

# Aggregate Functions

- These functions operate on the multiset of values of a column of a relation, and return a value

**avg:** average value

**min:** minimum value

**max:** maximum value

**sum:** sum of values

**count:** number of values

Select 다음 구절에 사용 가능

# Aggregate Functions (Cont.)

- Find the average salary of instructors in the Computer Science department

- ```
• select avg (salary)
  from instructor
  where dept_name= 'Comp. Sci.';
```

avg는  
null(값이 없는 행)  
 $\Rightarrow \text{avg} = 200$

	salary
	100
	300
	null

- Find the total number of instructors who teach a course in the Spring 2014 semester

- ```
• select count (distinct ID)
  from teaches
  where semester = 'Spring' and year = 2014
```

중복제거

- Find the number of tuples in the *course* relation

- ```
• select count (*)
  from course;
```

전체

Count(salary)는 null(?)까지  
Count

# Aggregate Functions – group by

- Find the average salary of instructors in each department

- `select dept_name, avg (salary)`

- `from instructor`

- `group by dept_name;`

aggregate : 합계

합계(합계)는 잘됨만 가능하다.

- Note: departments with no instructor will not be shown

ID	name	dept_name	salary
76766	Crick	Biology	72000
45565	Katz	Comp. Sci.	75000
10101	Srinivasan	Comp. Sci.	65000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000
12121	Wu	Finance	90000
76543	Singh	Finance	80000
32343	El Said	History	60000
58583	Califieri	History	62000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
22222	Einstein	Physics	95000

dept_name	avg_salary
Biology	72000
Comp. Sci.	77333
Elec. Eng.	80000
Finance	85000
History	61000
Music	40000
Physics	91000

구하기 사용할 때  
avg를 쓰면 됩니다.

# Aggregation (Cont.)

- Attributes in **select** clause outside of aggregate functions must appear in **group by** list

- /\* erroneous query \*/

```
select dept_name, ID, avg(salary)  
from instructor  
group by dept_name;
```

dept-nameG(dept-name, avg(salary))  
(Instructor)

$$G_{1, G_2, \dots, G_n} \mathcal{G}_{F_1(A_1), F_2(A_2), \dots, F_n(A_n)}(E)$$

$E$  is any relational-algebra expression

- $G_1, G_2, \dots, G_n$  is a list of attributes on which to group (can be empty)
- Each  $F_i$  is an aggregate function
- Each  $A_i$  is an attribute name

# Aggregate Functions – **having** Clause

- Find the names and average salaries of all departments whose average salary is greater than 42000

```
select dept_name, avg (salary)  
from instructor  
group by dept_name  
having avg (salary) > 42000;
```

group을 한 후에 필터 적용됨 (where는 group 전에 적용됨)

Note: predicates in the **having** clause are applied after the formation of groups whereas predicates in the **where** clause are applied before forming groups

Salary > 1000 → 남아 있는 부서들만  
↑ avg 계산  
where Salary > 1000

# Null Values and Aggregates

- The following query

```
select sum (salary), avg (salary)  
from instructor
```

- ignores tuples with salary=null
- Result is *null* if there is no non-null amount

ID	name	dept_name	salary
20	Alice	Comp. Sci.	8,000
30	Charlie	Biology	null
40	Bob	Comp. Sci.	6,000



sum(salary)	avg(salary)
14,000	7,000

- All aggregate operations except **count(\*)** ignore tuples with null values on the aggregated attributes

# Nested subquery

- SQL allows nesting of subqueries.
- A **subquery** is a **select-from-where** expression within another query.
- Common use cases of subqueries
  - set membership test (IN , NOT IN)
  - set comparisons (SOME, ALL)
  - set cardinality

# Example Query

- Find IDs of students who take DB class

```
select student_id      → 20  
from Take             → 40  
where course_id IN (  
    select course_id  
    from Course  
    where title = 'DB'  
)
```

sub query

먼저 실행

```
select course_id  
from Course  
where title = 'DB'
```

Student

ID	name
20	Alice
30	Charlie
40	Bob

Take

ID	c_id
20	swe3003
20	swe3004
30	swe3004
40	swe3003
40	swe3021

Course

c_id	title
swe3003	DB
swe3004	OS
swe3021	Multicore

sub query

to include

join

with

join

join

join

join

null

join

join

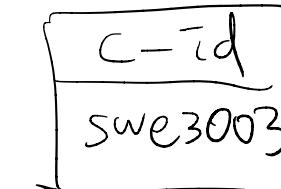
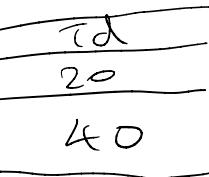
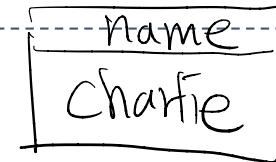
# Example Query

- Find IDs of students who **do not** take DB class

```
select student_name  
from Student  
where student_id NOT IN (  
    select student_id  
    from Take  
    where course_id IN (  
        select course_id  
        from Course  
        where title = 'DB'  
    )  
)
```

ID	name
X 20	Alice
0 30	Charlie
X 40	Bob

ID	c_id
20	swe3003
20	swe3004
30	swe3004
40	swe3003
40	swe3021



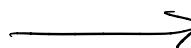
c_id	title
swe3003	DB
swe3004	OS
swe3021	Multicore



# Example Query

- Find the total number of (distinct) students take course taught by the instructor with ID 10101

```
select count (distinct ID)
from takes
where c_id IN
    (select course_id
     from teaches
     where teaches.ID= 10101);
```



ID
20
40



c_id
swe3003

- Note: This query can be written in a much simpler manner.

Take	ID	c_id
	20	swe3003
	20	swe3004
	30	swe3004
	40	swe3003
	40	swe3021

Teaches

ID	c_id
10101	swe3003
20202	swe3004
30303	swe3021

# Set Comparison

- Find names of instructors with salary greater than that of some (at least one) instructor in the Biology department.

```
select distinct T.name  
from instructor as T, instructor as S  
where T.salary > S.salary and S.dept_name = 'Biology';
```

Cartesian

product ( $T \times S$ )

T.name	dept	T.salary	S.name	S.dept	S.salary
Alice	Comp. Sci.	8,000	Alice	Comp. Sci.	8,000
Alice	Comp. Sci.	8,000	Charlie	Biology	2,000
Alice	Comp. Sci.	8,000	Bob	Comp. Sci.	6,000
Alice	Comp. Sci.	8,000	David	Biology	7,000
Charlie	Biology	2,000	Alice	Comp. Sci.	8,000
Charlie	Biology	2,000	Charlie	Biology	2,000
Charlie	Biology	2,000	Bob	Comp. Sci.	6,000
Charlie	Biology	2,000	David	Biology	7,000
Bob	Comp. Sci.	6,000	Alice	Comp. Sci.	8,000
...					

name	dept	salary
Alice	Comp. Sci.	8,000
Charlie	Biology	2,000
Bob	Comp. Sci.	6,000
David	Biology	7,000

$8,000 > 2,000$

$8,000 > 7,000$

# Set Comparison

- Find names of instructors with salary greater than that of some (at least one) instructor in the Biology department.
- Same query using  $> \text{some}$  clause

```
select name          인수선택  
      from instructor    문제를 풀 때마다 같은 문장  
      where salary > some (select salary  
                            from instructor  
                            where dept_name = 'Biology');
```

name	dept	salary
Alice	Comp. Sci.	8,000
Charlie	Biology	2,000
Bob	Comp. Sci.	6,000
David	Biology	7,000

$> \text{some}$

name	dept	salary
Charlie	Biology	2,000
David	Biology	7,000

# Definition of **some** Clause

- $F <\text{comp}> \text{some } r \Leftrightarrow \exists t \in r \text{ such that } (F <\text{comp}> t)$

Where  $\text{comp}$  can be:  $<$ ,  $\leq$ ,  $>$ ,  $=$ ,  $\neq$

(5 < some	<table border="1"><tr><td>0</td></tr><tr><td>5</td></tr><tr><td>6</td></tr></table>	0	5	6	) = TRUE	(read: 5 < some tuple in the relation)
0						
5						
6						
(5 < some	<table border="1"><tr><td>0</td></tr><tr><td>5</td></tr></table>	0	5	) = FALSE	$(= \text{some } R) \equiv \text{in } R$ [ <i>존재(적)성자</i> ]	
0						
5						
(5 = some	<table border="1"><tr><td>0</td></tr><tr><td>5</td></tr></table>	0	5	) = TRUE		
0						
5						
(5 \neq some	<table border="1"><tr><td>0</td></tr><tr><td>5</td></tr></table>	0	5	) = TRUE (since $0 \neq 5$ )	$(\neq \text{some } R) \equiv \text{not in } R$ <i>비존재성자</i> <i>5가 not in</i>	
0						
5						

# Example Query

- Find the names of all instructors whose salary is greater than the salary of **all** instructors in the Biology department.

```
select name
from instructor
where salary > all (select salary
              from instructor
              where dept_name = 'Biology');
```

name	dept	salary
Alice	Comp. Sci.	8,000
Charlie	Biology	2,000
Bob	Comp. Sci.	6,000
David	Biology	7,000

> all

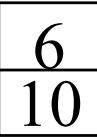
name	dept	salary
Charlie	Biology	2,000
David	Biology	7,000

→ 2000, 7000 를 더해 9000

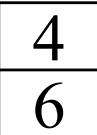
# Definition of **all** Clause

- $F <\text{comp}> \text{all } r \Leftrightarrow \forall t \in r (F <\text{comp}> t)$

(5 < all  ) = false

(5 < all  ) = true

(5 = all  ) = false

(5 ≠ all  ) = true (since  $5 \neq 4$  and  $5 \neq 6$ )

전부이다 = 속해 있지 않다  
 $(\neq \text{all}) \equiv \text{not in}$

However,  $(= \text{all}) \neq \text{in}$

각지각지

- $= \text{some } R \equiv \text{in } R$
- $\neq \text{some } R \equiv \text{not in } R$
- $\neq \text{all } R \equiv \text{not in } R$
- $= \text{all } R \equiv \text{in } R$

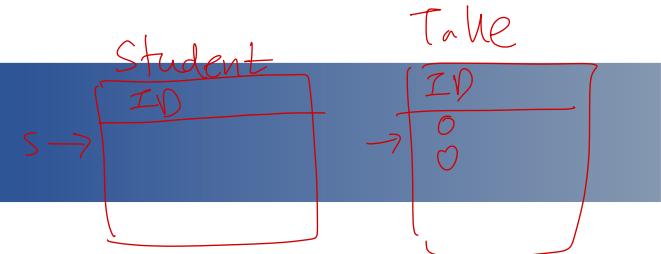
# Test for Empty Relations: **exists**

- The **exists** construct returns the value **true** if the argument subquery is nonempty.
- **exists**  $r \Leftrightarrow r \neq \emptyset$
- **not exists**  $r \Leftrightarrow r = \emptyset$

not exists   
select \_\_\_\_  
from \_\_\_\_  
where exists ( select \*  
from T  
\_\_\_\_ )



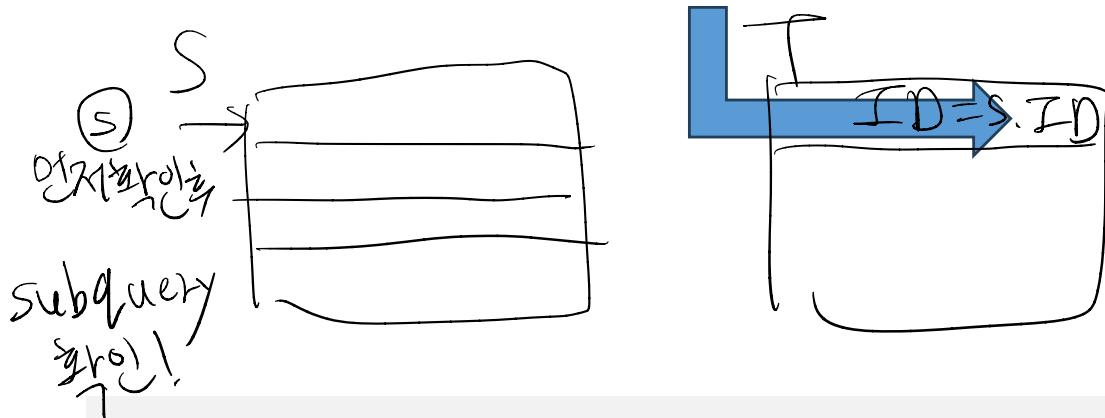
# Correlation Variables



- **correlated subquery** is a subquery that depends on the outer query for its values.

- Unlike a **regular subquery** (which runs independently and returns a fixed result), a **correlated subquery** is executed **once for each row** in the outer query.
- Eg. Find names of student who have taken at least one course

```
select distinct name
from Student natural join Take;
```



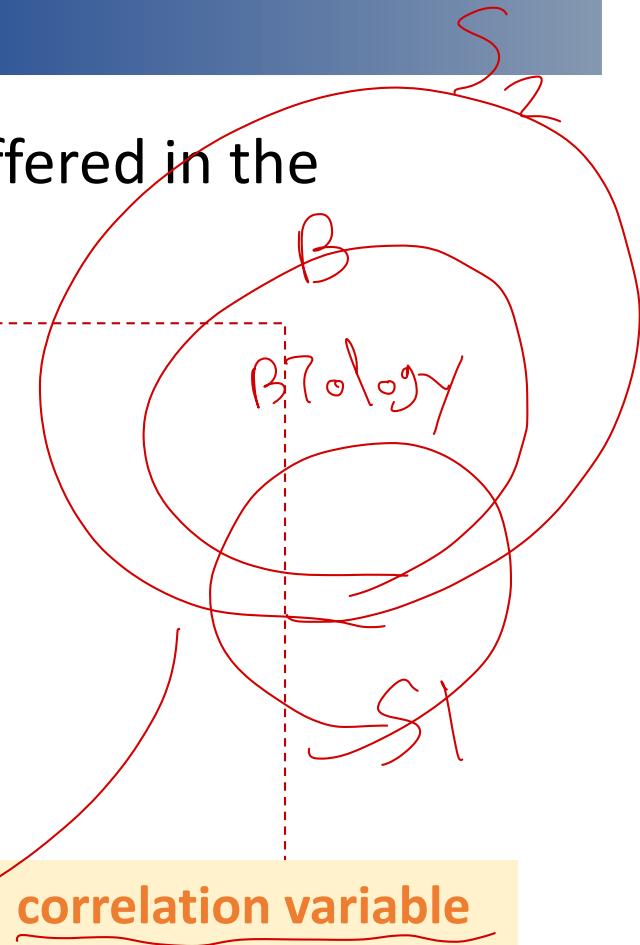
```
select distinct name
from Student s
where EXISTS (
    select 1
    from Take t
    where t.ID = s.ID
);
```

S가 원래(?) 있는  
외부 쿼리  
여기까지는 외부 쿼리  
L  
내부 쿼리  
인덱스가 다 '1'로 바뀜  
S : correlation variable

# Test for all: **not exists**

- Find all students who have taken **all** courses offered in the Biology department.

```
select distinct S.ID, S.name  
from student as S student table 청약2021-22  
where not exists ( (select course_id  
                    from course  
                    where dept_name = 'Biology')  
                   not exist( ) → True  
                   except 자기집합  
                   (select T.course_id  
                     from takes as T  
                     where S.ID = T.ID));  
not exist( ~ ) → False
```



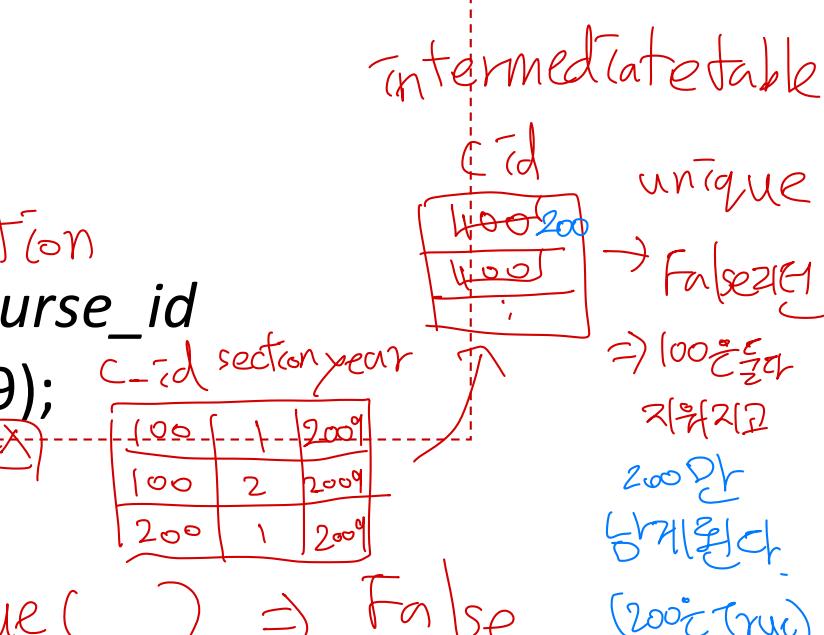
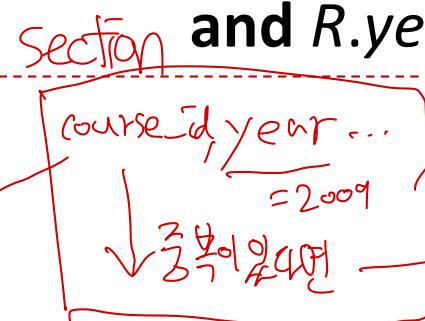
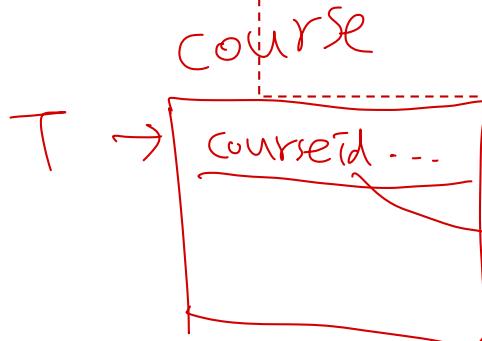
- If 'ALL' is included in the query description, consider 'not exists' and 'except'.
- Note that  $X - Y = \emptyset \Leftrightarrow X \subseteq Y$
- Note:* Cannot write this query using = **all** and its variants

# Test for Absence of Duplicate Tuples

- The **unique** construct tests whether a subquery has any duplicate tuples in its result.  
• (Evaluates to “true” on an empty set)
- Find all courses that were offered at most once in 2009

```
select T.course_id  
from course as T  
where unique (select R.course_id
```

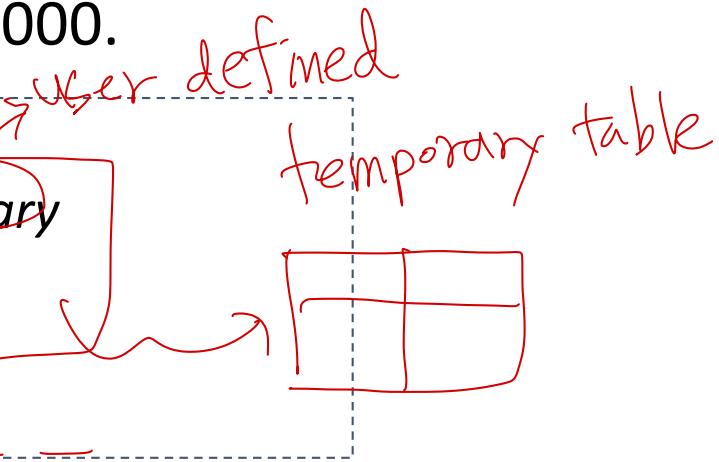
```
        from section as R  
        where T.course_id = R.course_id  
              and R.year = 2009);
```



# Subqueries in the From Clause

- SQL allows a subquery expression to be used in the **from** clause
- Find the average instructors' salaries of those departments where the average salary is greater than \$42,000.

```
select dept_name, avg_salary  
from (select dept_name, avg(salary) as avg_salary  
      from instructor  
     group by dept_name)  
where avg_salary > 42000;
```



or

```
with dept_avg(dept_name,avg_sal)  
select dept_name, avg_salary  
from (select dept_name, avg(salary)  
      from instructor  
     group by dept_name)  
       as dept_avg(dept_name, avg_salary)  
where avg_salary > 42000;
```

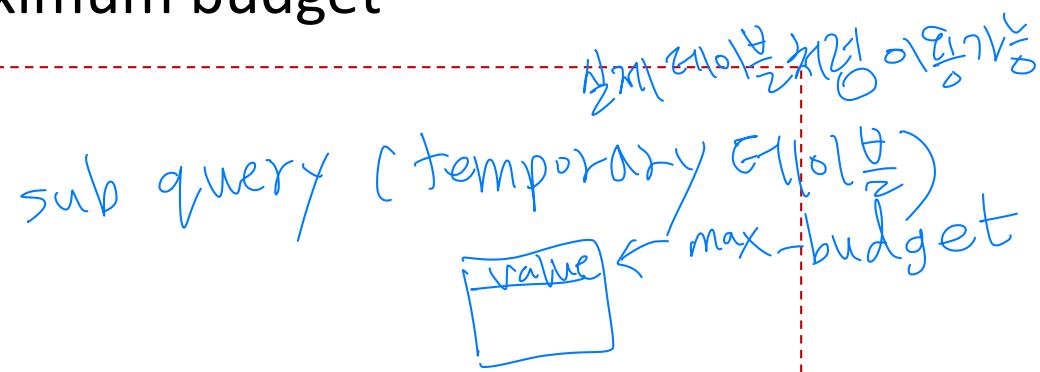
or  
select  
from instructor  
group by d\_n  
having avg\_sal > 42000

select **문장이후의 AS3**  
선행하는 테이블 name이 위치.

# with Clause

- The **with** clause provides a way of defining a temporary **view** whose definition is available only to the query in which the **with** clause occurs.
- Find all departments with the maximum budget

```
→ G(10) ↴ name  
with max_budget (value) as  
    (select max(budget)  
     from department)  
select budget  
from department, max_budget  
where department.budget = max_budget.value;
```



# Complex Queries using **with** Clause

- With clause is very useful for writing complex queries
- Supported by most database systems, with minor syntax variations
- Find all departments where the total salary is greater than the average of the total salary at all departments

```
with dept_total(dept_name, value) as  
    (select dept_name, sum(salary)  
     from instructor  
     group by dept_name),  
     dept_total_avg(value) as  
        (select avg(value)  
         from dept_total)  
select dept_name  
from dept_total, dept_total_avg  
where dept_total.value >= dept_total_avg.value;
```

Handwritten notes:

- A blue bracket groups the first two lines of the query (the WITH clause).
- A blue bracket groups the first three lines of the query (the WITH clause and the first part of the SELECT clause).
- A blue circle highlights the word "dept\_total" in the WHERE clause.

# Scalar Subquery

- **Scalar subquery** is used where a single value is expected

- E.g. `select dept_name,`

```
(select count(*)
     from instructor
    where department.dept_name = instructor.dept_name)
   as num_instructors
  from department;
```

单一值  
Single Value

- E.g. `select name`

```
from instructor
```

```
where salary * 10 >
```

```
(select budget from department
```

```
where department.dept_name = instructor.dept_name)
```

单一值  
Single Value

- Runtime error if subquery returns more than one result tuple

# Modification of the Database

- Deletion of tuples from a given relation
- Insertion of new tuples into a given relation
- Updating values in some tuples in a given relation

# Modification of the Database – delete from

- Delete all instructors

**delete from** *instructor*

Arsh

- Delete all instructors from the Finance department

**delete from** *instructor*  
**where** *dept\_name*= 'Finance';

- Delete all tuples in the *instructor* relation for those instructors associated with a department located in the Watson building.

```
delete from instructor  
where dept_name in (select dept_name  
                 from department  
                 where building = 'Watson');
```

## delete from (Cont.)

- Delete all instructors whose salary is less than the average salary of instructors

```
delete from instructor  
where salary < (select avg (salary) from instructor);
```

= scalar subquery → 틀리게

subquery인지 → 예상 결과값  
→ avg 재귀적 계산법

오류 찾는법  
원인 찾기  
원인 찾기

- Problem: as we delete tuples from deposit, the average salary changes
- Solution used in SQL:
  - First, compute **avg** salary and find all tuples to delete
  - Next, delete all tuples found above (without recomputing **avg** or retesting the tuples)

# Modification of the Database – **insert into**

- Add a new tuple to *course*

**insert into** *course*

**values** ('CS-437', 'Database Systems', 'Comp. Sci.', 4);

- or equivalently

**insert into** *course* (*course\_id*, *title*, *dept\_name*, *credits*)

**values** ('CS-437', 'Database Systems', 'Comp. Sci.', 4);

- Add a new tuple to *student* with *tot\_creds* set to null

**insert into** *student*

**values** ('3003', 'Green', 'Finance', *null*);

## insert into (Cont.)

- Add all instructors to the *student* relation with tot\_creds set to 0

~~insert into student~~

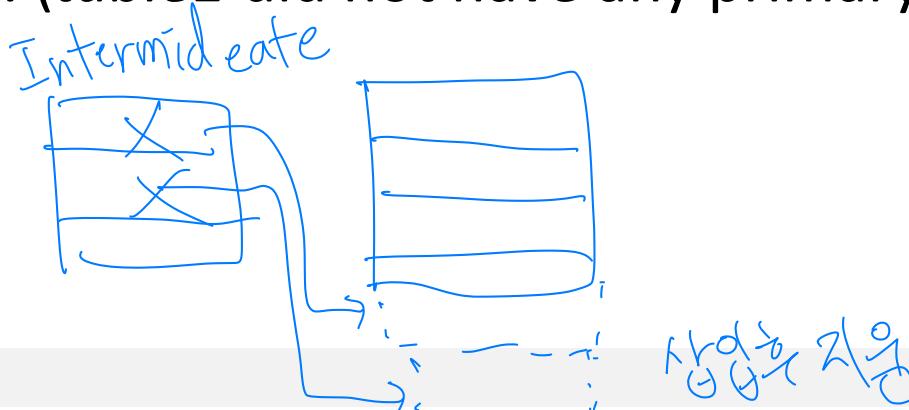
~~select ID, name, dept\_name, 0  
from instructor~~

- The **select from where** statement is evaluated fully before any of its results are inserted into the relation

- otherwise queries like

~~insert into table1 select \* from table1~~

would cause problems. (*table1* did not have any primary key defined)



# Modification of the Database – **update set**

- Increase salaries of instructors whose salary is over \$100,000 by 3%, and all others receive a 5% raise

- Write two **update** statements:

**update** *instructor*  
**set** *salary* = *salary* \* 1.03  
**where** *salary* > 100000;

②

MH  
OC

**update** *instructor*  
**set** *salary* = *salary* \* 1.05  
**where** *salary* <= 100000;

①

salary = 99999 인 사람은

① → ② 하면 됨 이유는?

- What if we change the order of the two statements?
  - The order is important
- Can be done better using the **case** statement (next slide)

# case Statement for Conditional Updates

- Same query as before but with case statement

```
update instructor  
set salary = case  
    when salary <= 100000 then salary * 1.05  
    else salary * 1.03  
end
```

SQL Key word

if else 조건문

# Updates with Scalar Subqueries

- Recompute and update `tot_creds` value for all students

```
update student S
  set tot_cred = ( select sum(credits)
                     from takes natural join course
                   where S.ID= takes.ID and
                         takes.grade <> 'F' and
                         takes.grade is not null);
```

*correlation value*

- Sets `tot_creds` to null for students who have not taken any course
- Instead of `sum(credits)`, use:

```
case
  when sum(credits) is not null then sum(credits)
  else 0
end
```