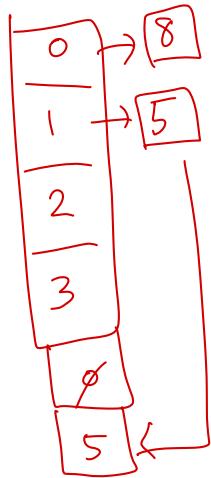


Next



5 삽입 → collision

8 삽입 → collision

Database Systems

Lecture13 – Chapter 14: Indexing

Beomseok Nam (남범석)

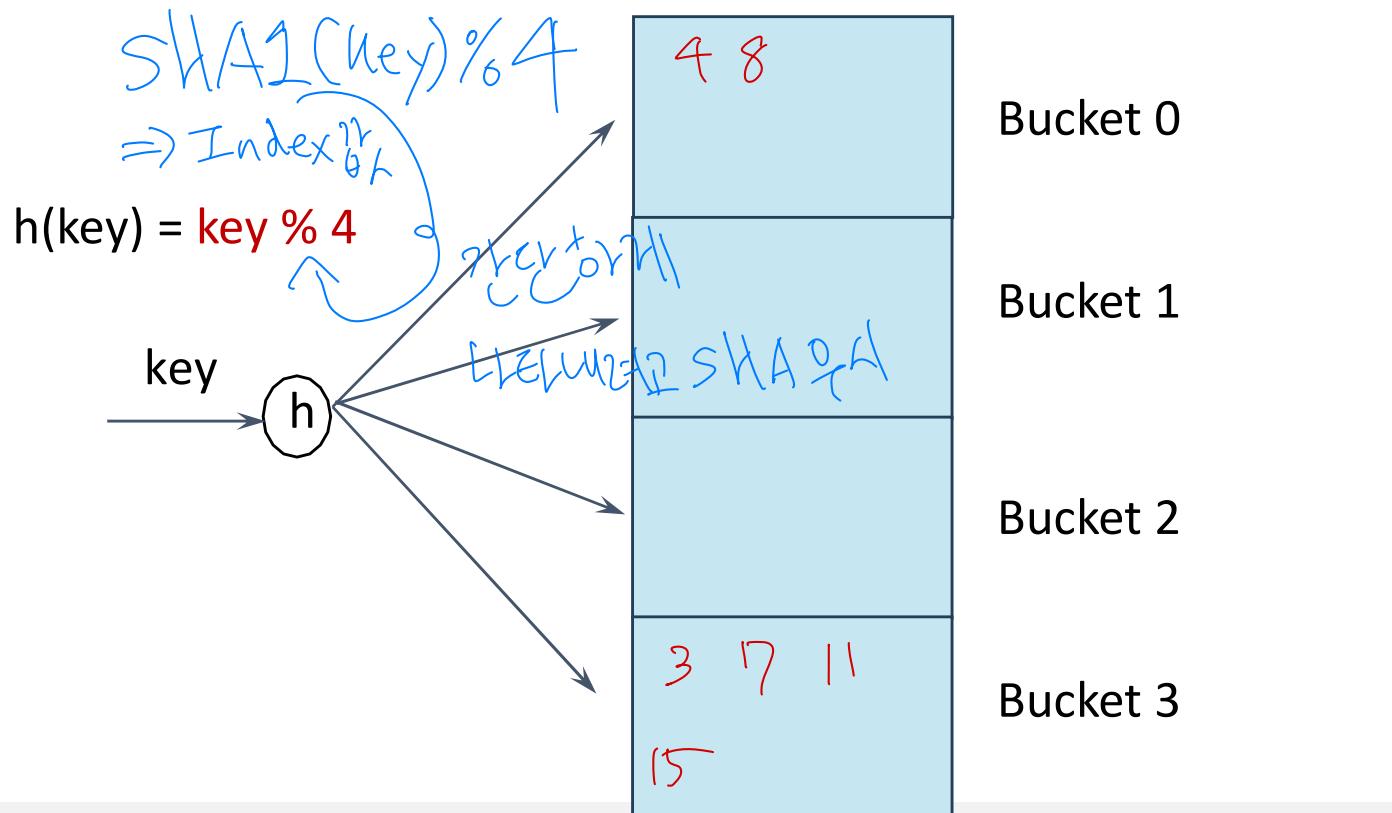
bnam@skku.edu

Static Hashing

- Bucket: Storage unit holding index entries
 - (typically 4KB disk block)
 - Hash Function (h): Maps search keys \rightarrow bucket addresses
 - Different keys may map to the same bucket (**collisions**)
 - Entire bucket needs to be searched **sequentially**
 - Hash Index: Buckets store **pointers to records**
 - Hash File : Buckets store **records**
- Hash 함수
 $key \% 10 \rightarrow$ Index
112
↑
0 10 20 30
= 같은 버킷 저장
버킷에 더 많은 데이터
충돌하는 현상 \rightarrow collision

Static Hashing

- Fixed-size array of buckets (e.g., $\text{HT} = \text{new Bucket}[\text{size}]$;)
- $h(\text{key}) \rightarrow$ get bucket index in array
 - Often uses a **cryptographic hash function** (e.g., SHA-256) followed by modulo: $H(\text{key}) \{ \text{return } \text{SHA256}(\text{key}) \% \text{ size}; \}$
 - Let's assume $H(\text{key}) = \text{key \% size}$ for simplicity



Handling of Bucket Overflows

- Bucket overflow occurs because of Hash Collision

[1] Insufficient buckets

- This type of bucket overflow cannot be eliminated

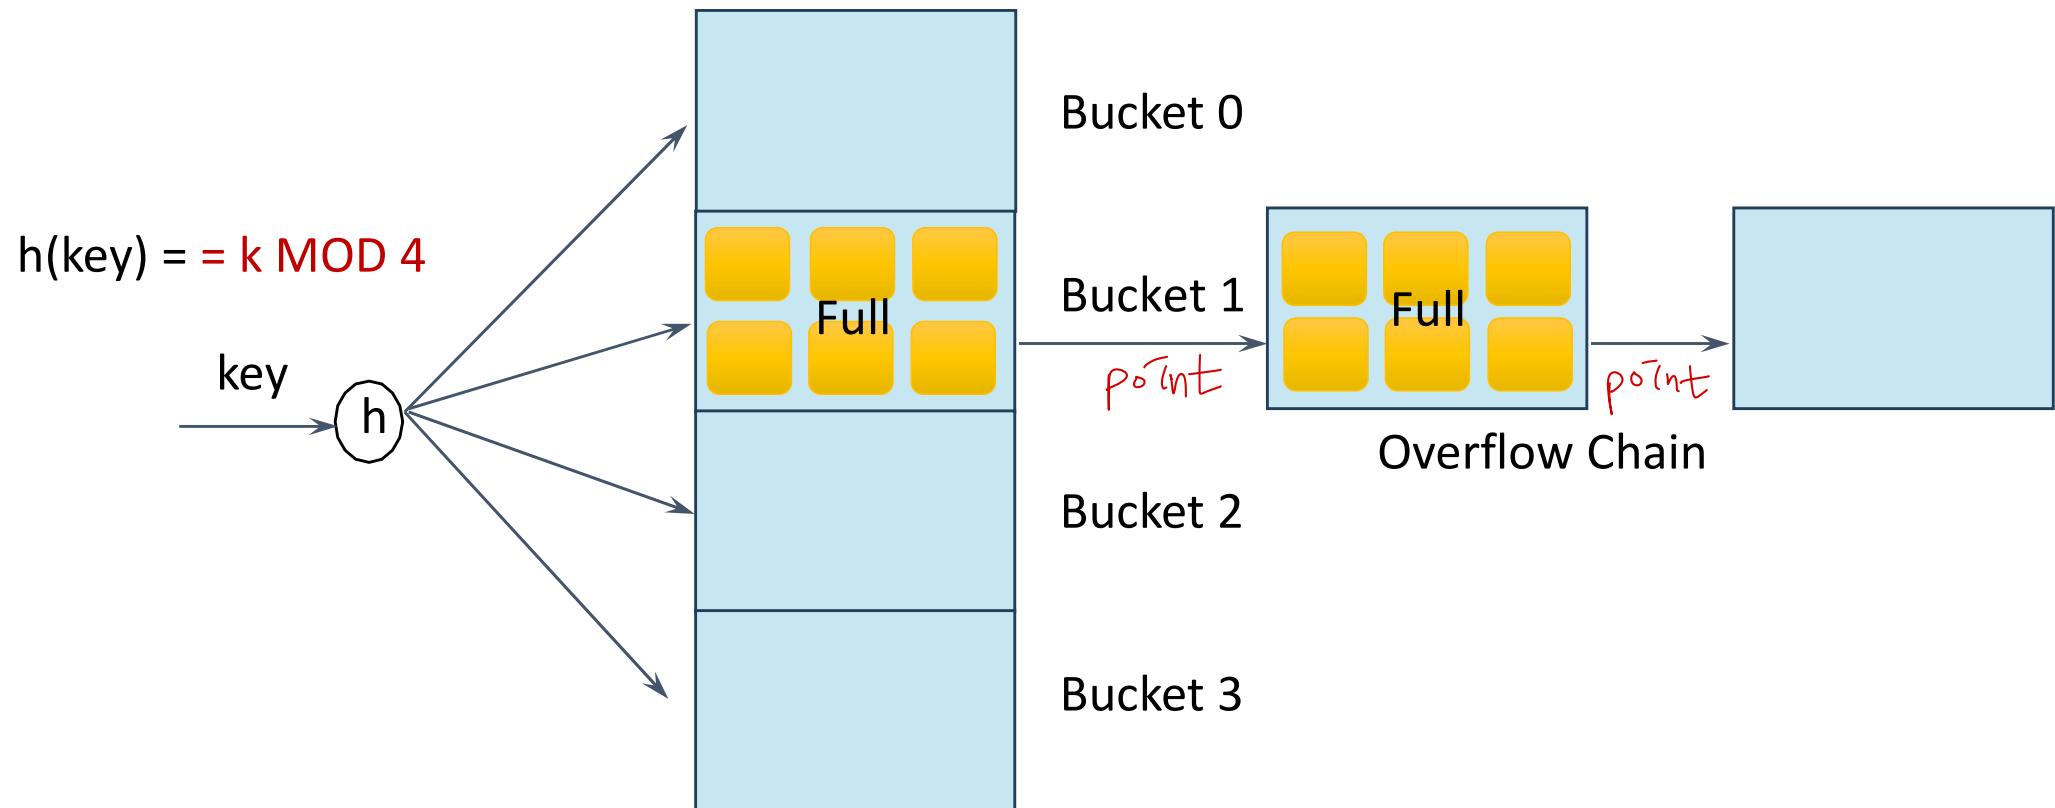
[2] Skew in distribution of records.

- Skew occurs due to two reasons:
 - Hash function produces skewed distribution of keys
- Overflow is handled by either
 - Overflow Chain**
 - or
 - Closed Hashing (Open Addressing)**

Handling of Bucket Overflows

■ Overflow Chain (*open Hashing*)

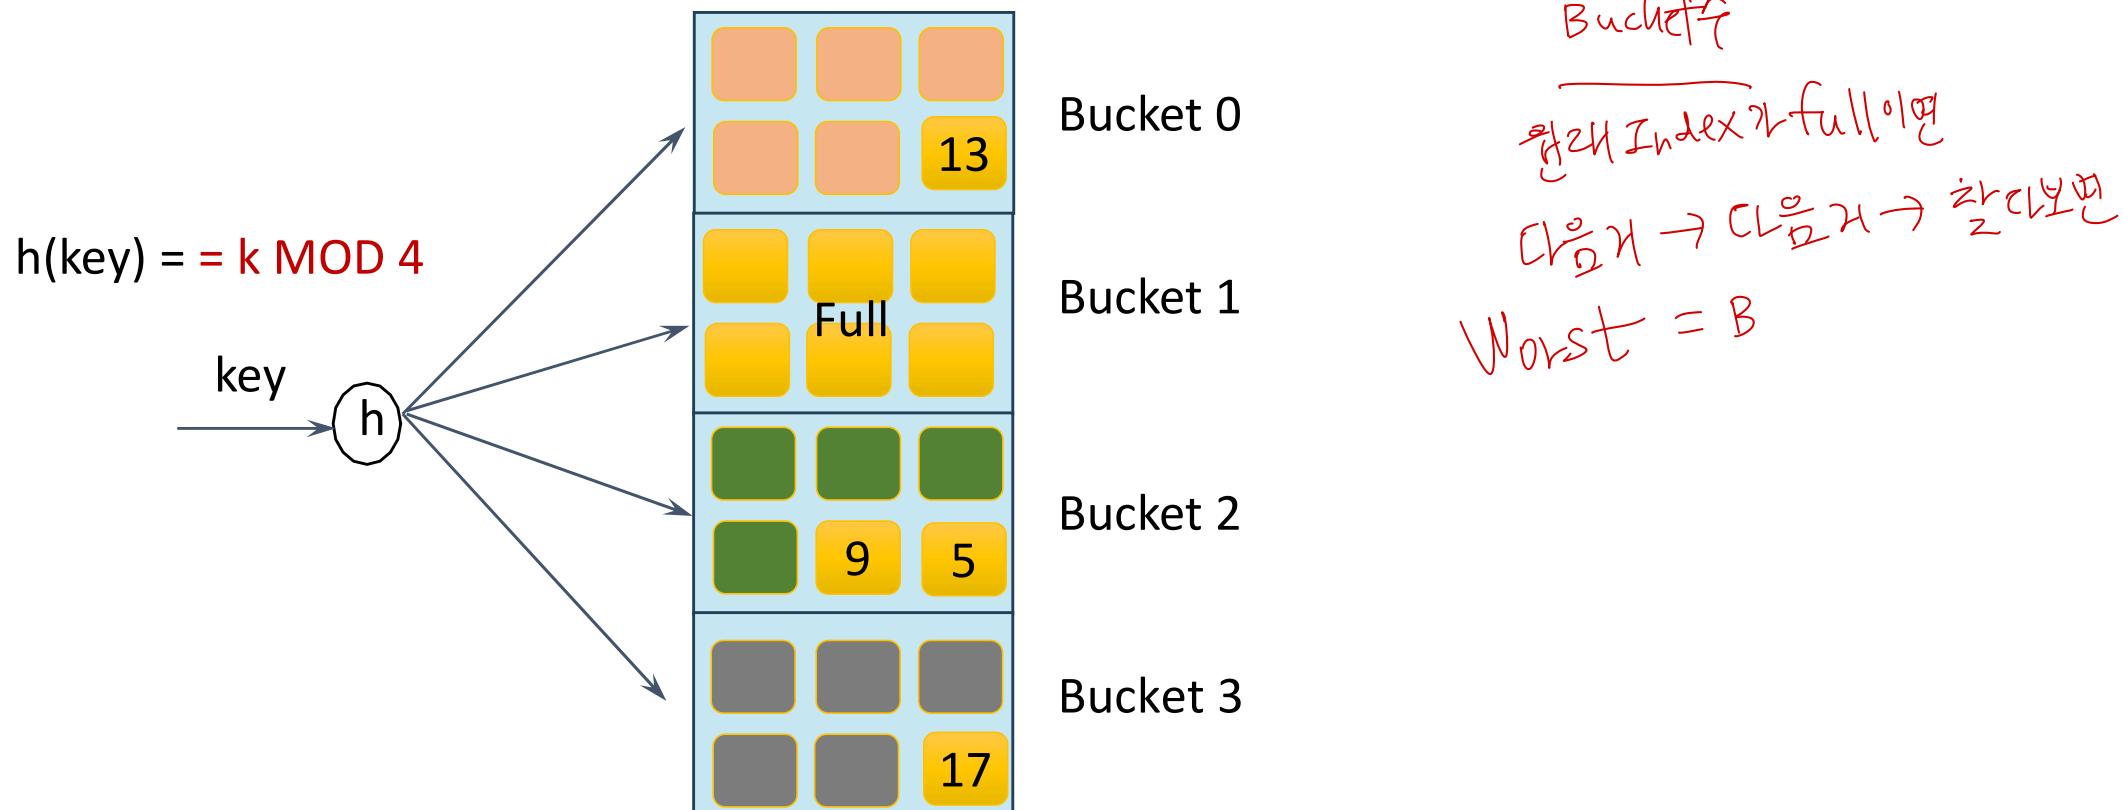
- Overflow buckets are chained together in a linked list.



Handling of Bucket Overflows (Closed Hashing)

- **Closed Hashing (a.k.a Open Addressing)**

- **Linear probing**: Use the next bucket (in cyclic order) that has space.
↳ Next Free Space 가 저장 $\Rightarrow O(B)$



Problem: Worst Search Complexity = $O(N)$. \rightarrow Limit Probing Distance.

Example of Hash File

- Hash file of *instructor* file, using *dept_name* as key

- E.g. $h(\text{Music}) = 1$
 - $h(\text{History}) = 2$
 - $h(\text{Physics}) = 3$
 - $h(\text{Elec. Eng.}) = 3$

The given hash function fails to provide a uniform distribution.

→ Static Hashing is a Bad Choice

bucket 0			

bucket 1

bucket 2			
32343	El Said	History	80000
58583	Califieri	History	60000

bucket 3			
22222	Einstein	Physics	95000
33456	Gold	Physics	87000
98345	Kim	Elec. Eng.	80000

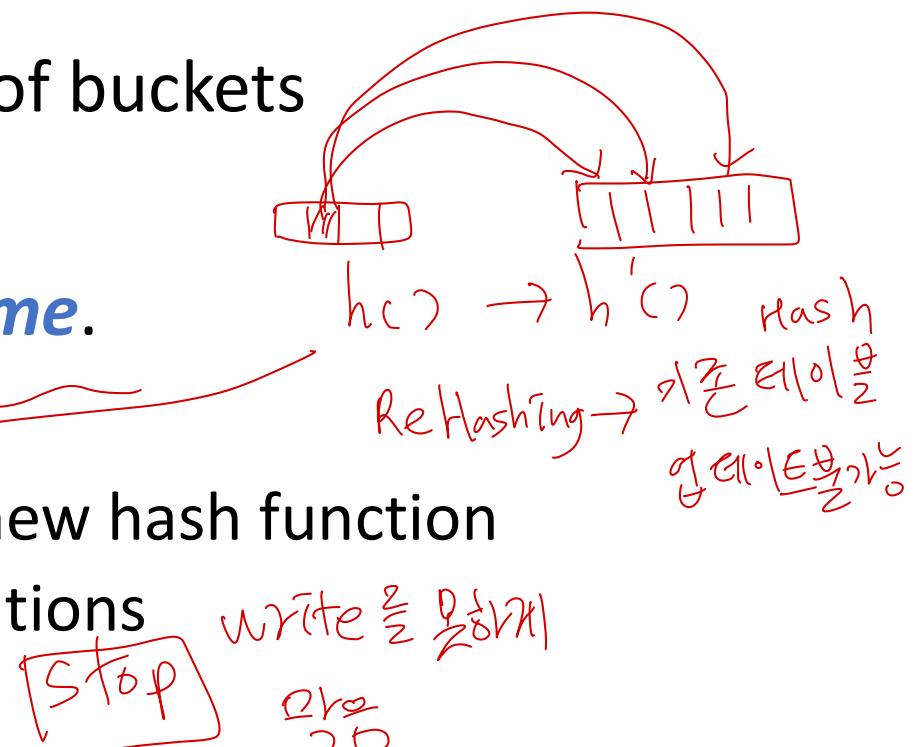
bucket 4			
12121	Wu	Finance	90000
76543	Singh	Finance	80000

bucket 5			
76766	Crick	Biology	72000

bucket 6			
10101	Srinivasan	Comp. Sci.	65000
45565	Katz	Comp. Sci.	75000
83821	Brandt	Comp. Sci.	92000

Deficiencies of Static Hashing

- Static Hashing → Fixed-size array of buckets
- *Databases grow or shrink with time.*
- Solution: **Rehashing**
 - periodic re-organization using a new hash function
 - Expensive, disrupts normal operations
- Better solution: *allow the number of buckets to be modified dynamically.*



정지
정지

Dynamic Hashing

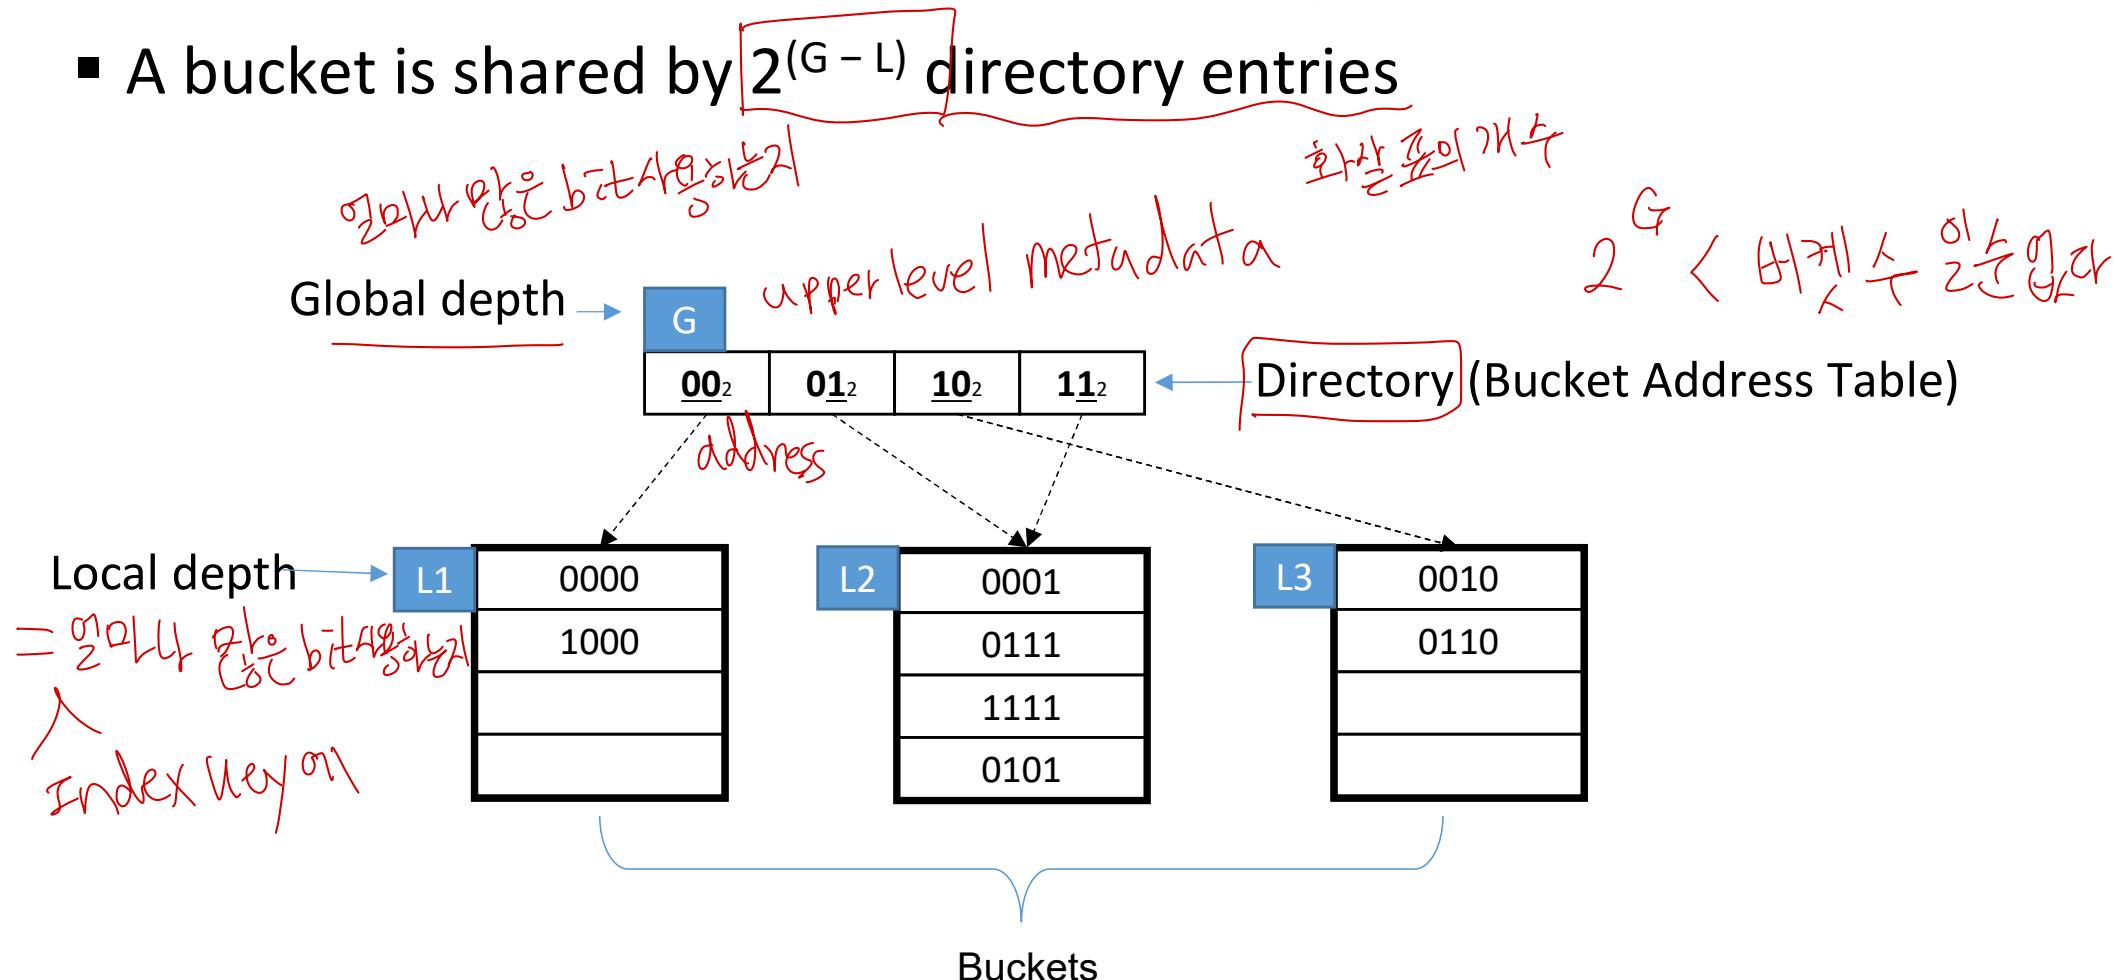
- Trick lies in how hash function is adjusted!
- **Extendable Hashing**
 - Designed for **disk-based** hash tables
 - Buckets may be **shared by multiple hash values**
 - Uses a **directory (array)** of pointers to buckets
 - Idea: Add a level of indirection!
- **Linear Hashing**
 - Do rehashing in an incremental manner

Extendable Hashing

- Use **prefix (or postfix)** bits of the hash key
- Let ***Global Depth G = # of bits used***
 - Number of buckets $\leq 2^G$
 - **Initially, G = 0**, grows/shrinks with DB size
- Buckets do not need to be stored contiguously on disk.
 - Need Directory (i.e., Bucket Address Table)
 - directory size = 2^G
- Dynamic Behavior
 - **Multiple directory entries** may point to the **same bucket**
 - ***Local Depth:*** # of bits used by each bucket
 - Buckets split/merge as data changes

Extendable Hash Structure

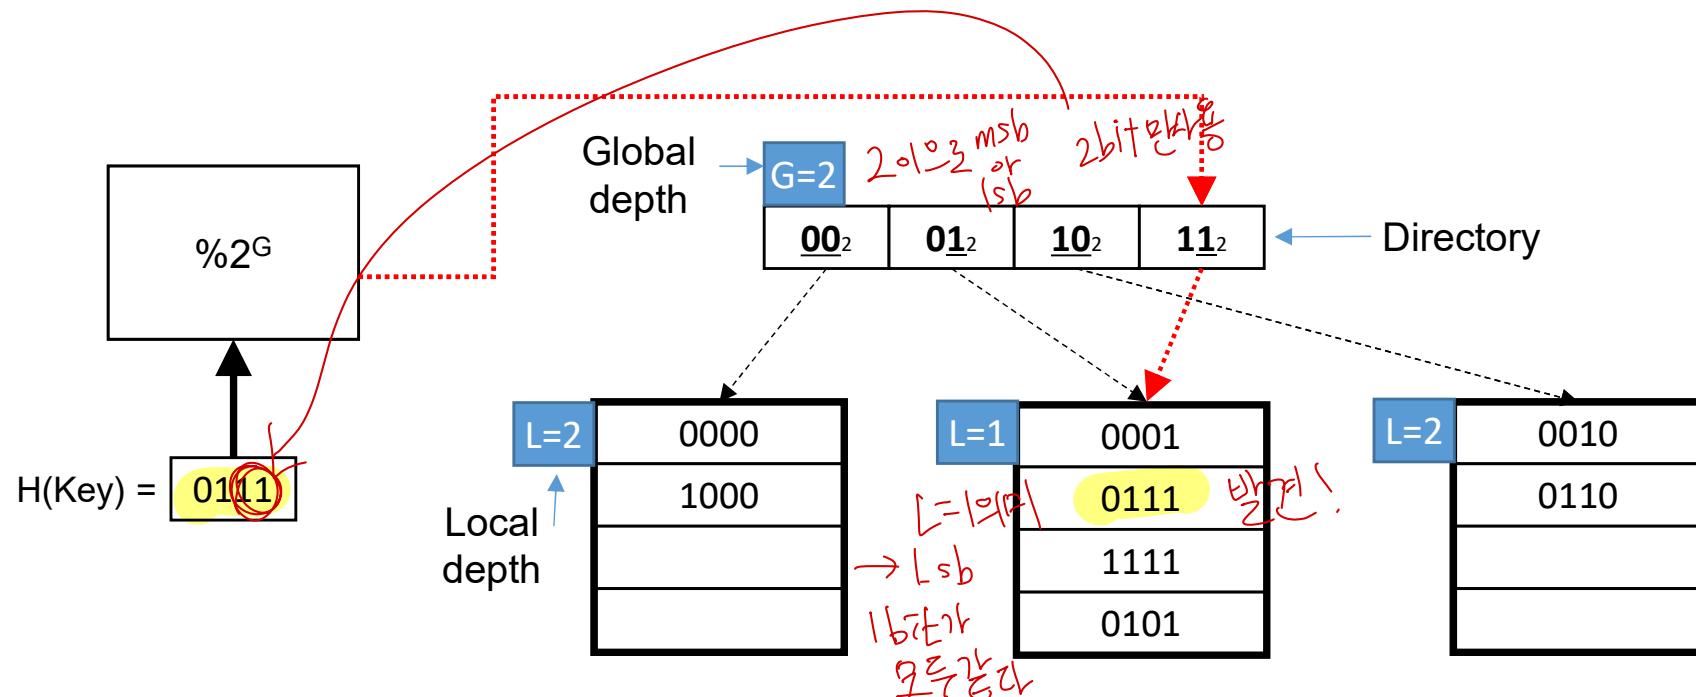
- **Global Depth G** = # of bits used to look up the directory
- **Local Depth L** = # of bits used to index keys in each bucket
- A bucket is shared by $2^{(G - L)}$ directory entries



In this example, $L1 = L3 = G = 2$, whereas $L2 = G - 1 = 1$

Extendable Hashing – How it Works

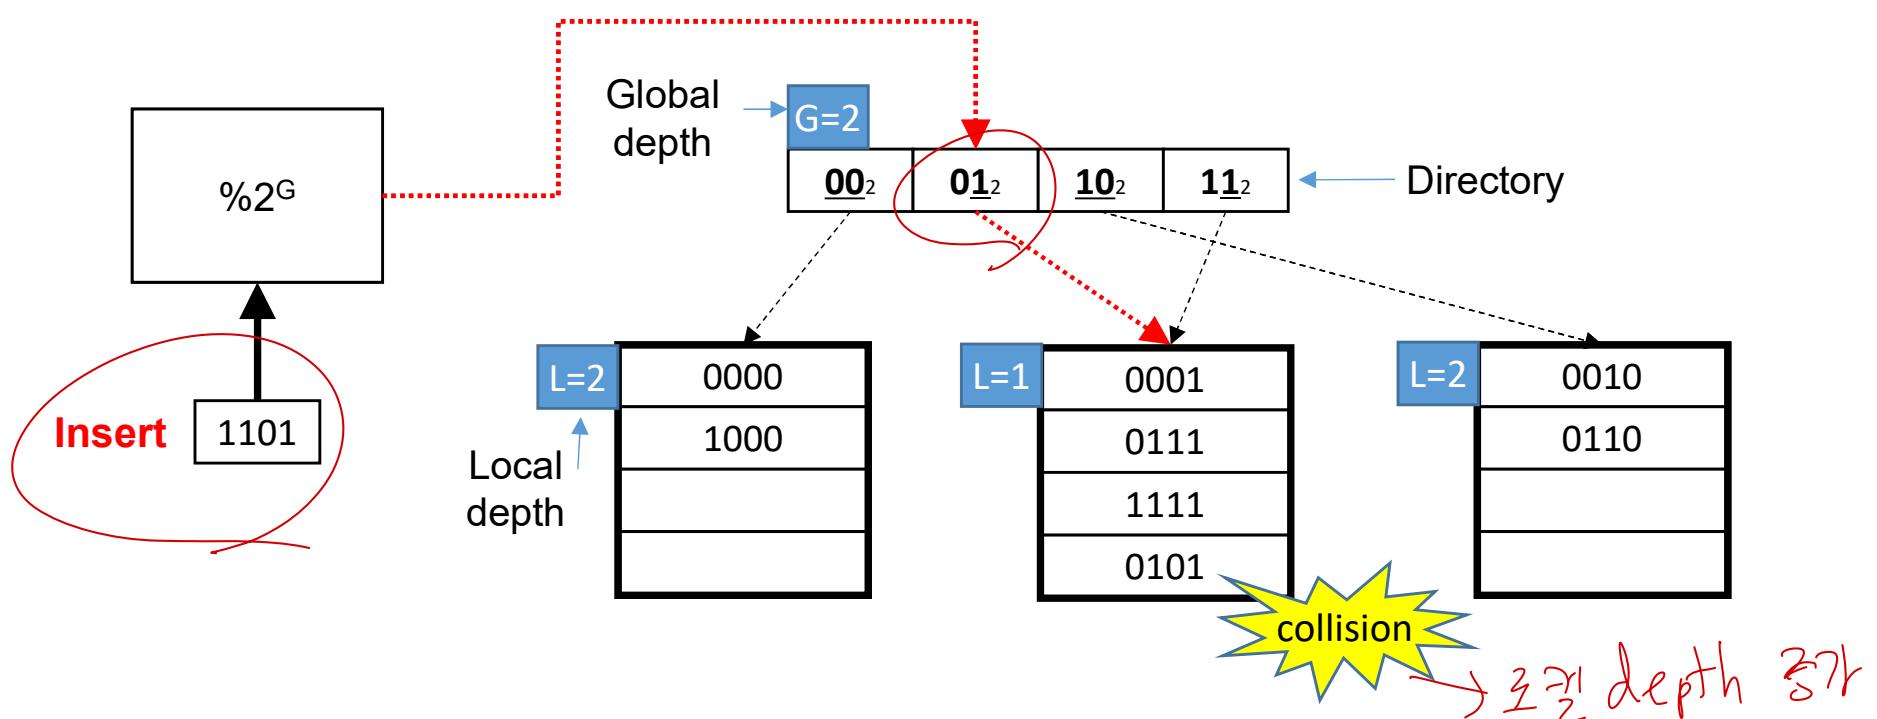
- Directory is an array of size 4 (2^2), so 2 bits are needed.
- We will use postfix in the examples
 - global depth = # of least significant bits of $h(K)$;
- If $h(K) = 5$ = binary 0101 , it is in bucket pointed to by 01 .
- If $h(K) = 7$ = binary 0111 , it is in bucket pointed to by 11



Extendable Hashing – How it Works

■ Bucket Split

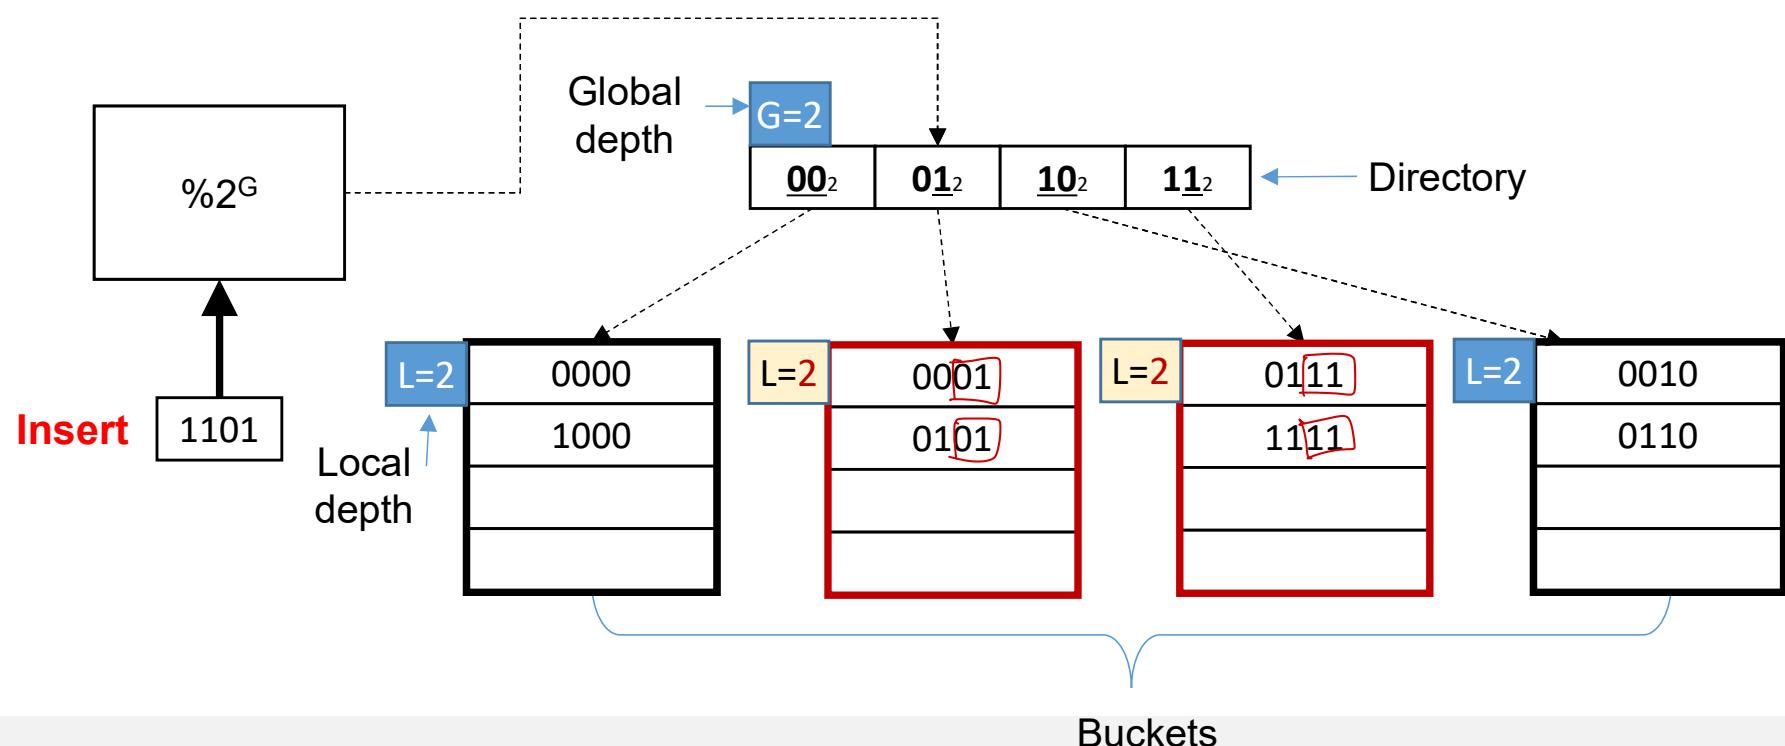
- if $G > L$ (more than one pointer to bucket)
 - allocate a new bucket, and set $L = L+1$
 - Update the directory to point to the new bucket
 - move records in the overflow bucket to the new bucket
 - Further splitting is required if the bucket is still full



Extendable Hashing – How it Works

▪ Bucket Split

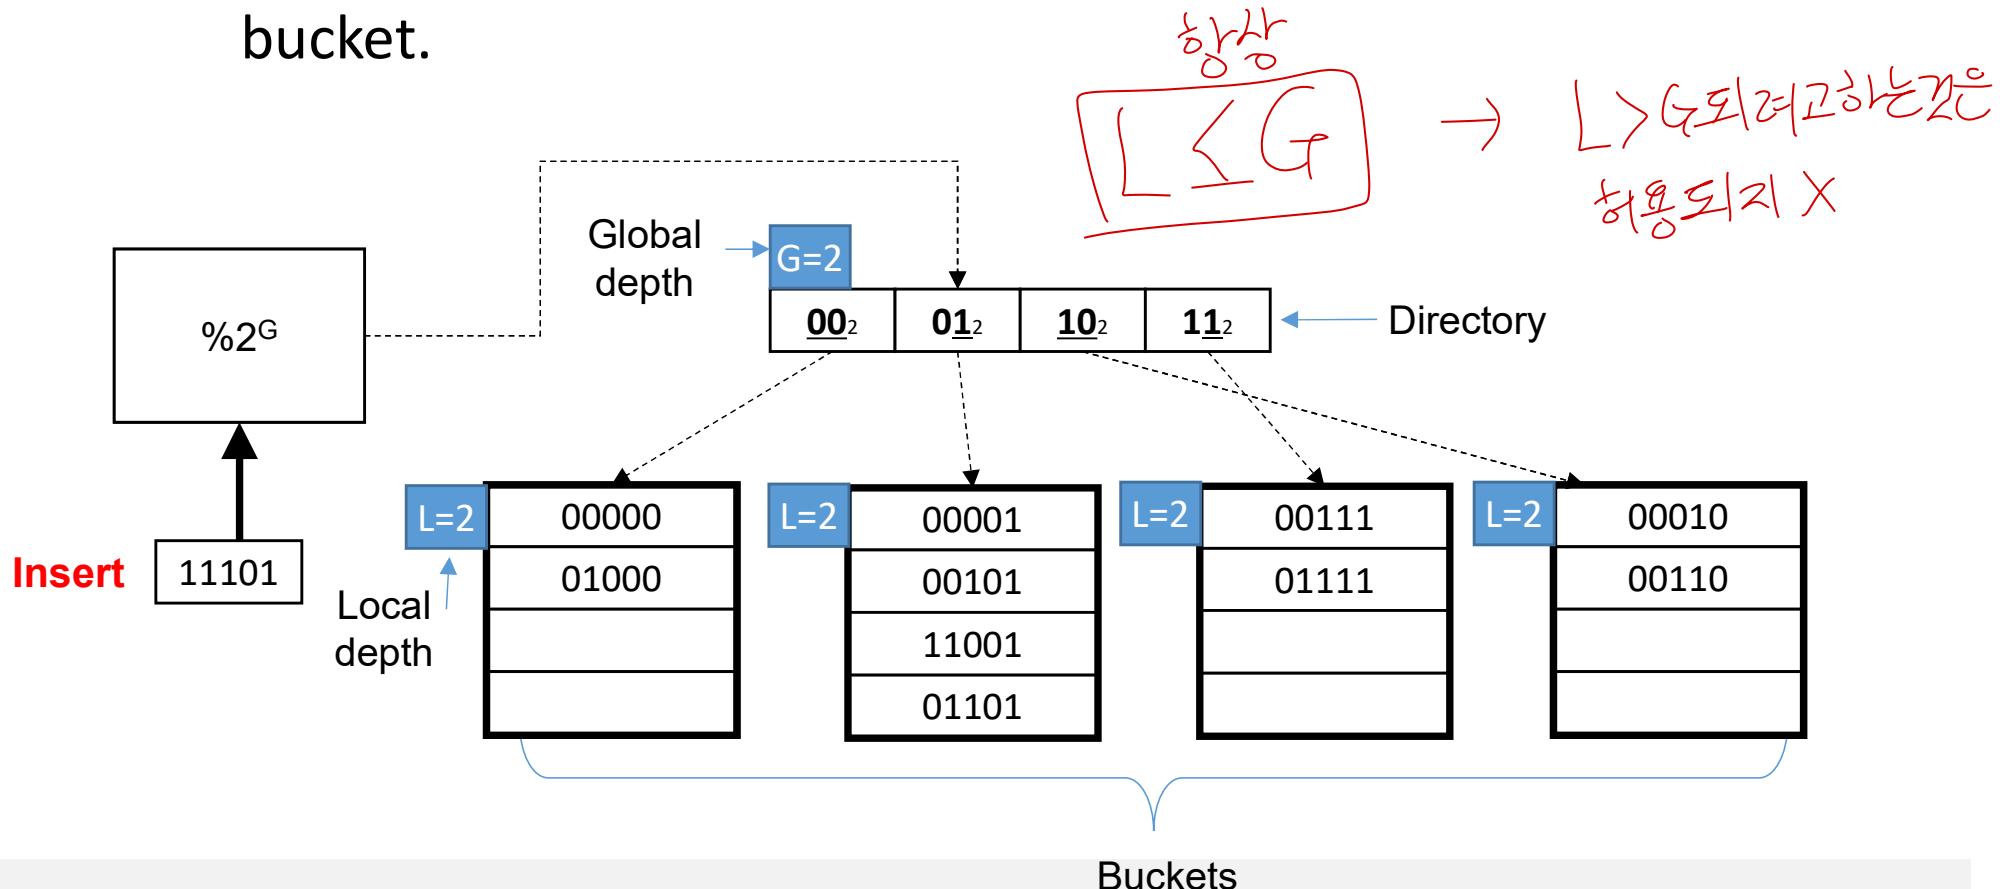
- if $G > L$ (more than one pointer to bucket)
 - allocate a new bucket, and set $L = L+1$
 - Update the directory to point to the new bucket
 - move records in the overflow bucket to the new bucket
 - Further splitting is required if the bucket is still full



Extendable Hashing – How it Works

- Bucket Split → *Directory Doubling*

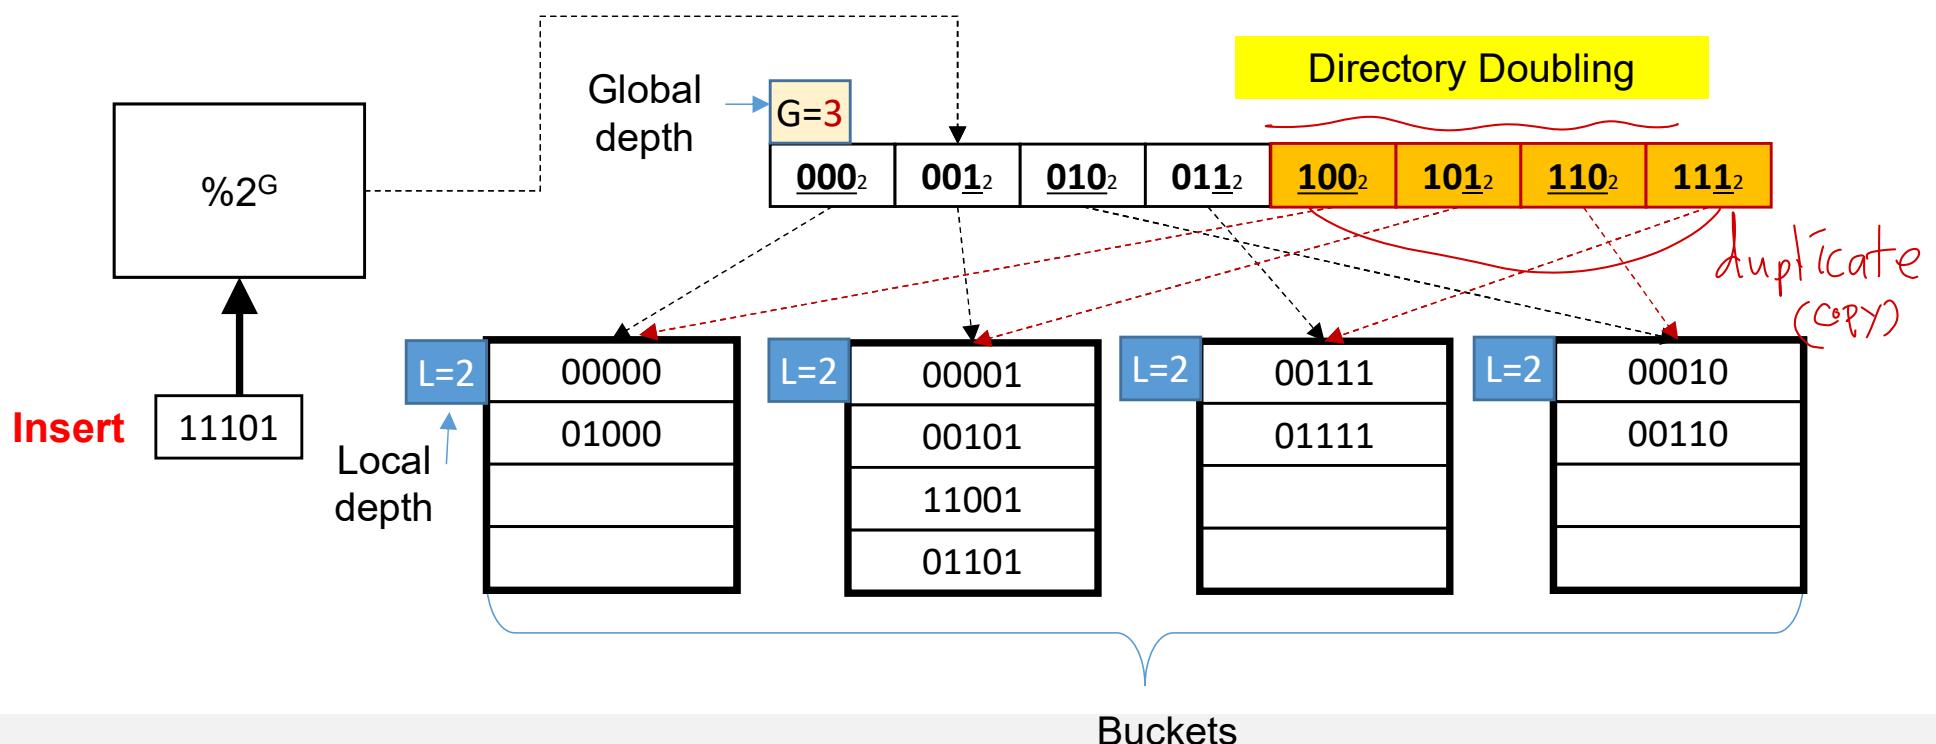
- If $G = L$ (only one pointer to bucket)
 - increment G and double the size of directory.
 - Now $G > L$.
 - replace each entry by two entries that point to the same bucket.



Extendable Hashing – How it Works

- Bucket Split → *Directory Doubling*

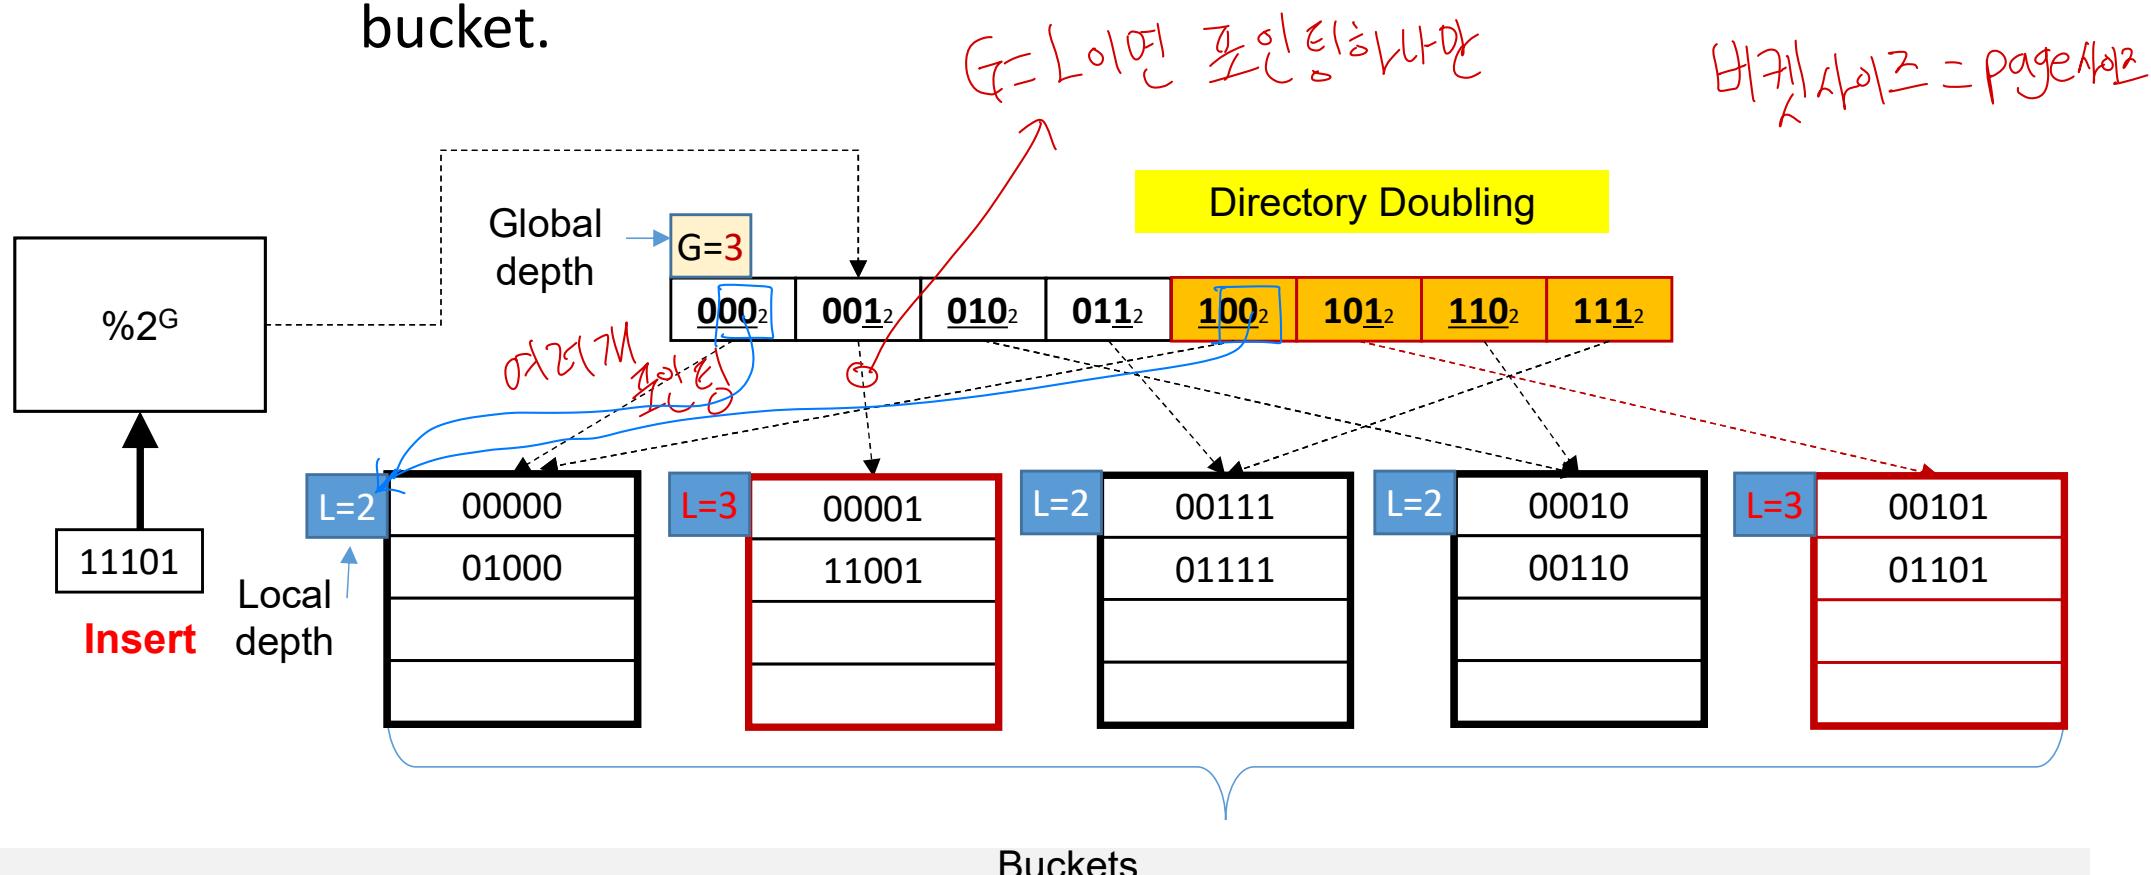
- If $G = L$ (only one pointer to bucket)
 - increment G and double the size of directory.
 - Now $G > L$.
 - replace each entry by two entries that point to the same bucket.



Extendable Hashing – How it Works

- Bucket Split → *Directory Doubling*

- If $G = L$ (only one pointer to bucket)
 - increment G and double the size of directory.
 - Now $G > L$.
 - replace each entry by two entries that point to the same bucket.



Extendable Hashing - Pros and Cons

- Benefits of extendable hashing:

- Hash performance does not degrade with growth of file
- Minimal space overhead

- Disadvantages of extendable hashing

- Extra level of indirection to find records
- Directory doubling is expensive → Directory \neq B^+ tree!
- Directory can become large
 - Solution: B^+ -tree structure for bucket address table

- Linear hashing is an alternative mechanism

- Allows incremental growth of its directory
- At the cost of more bucket overflows

Linear Hashing – a lazy approach

- No directory → uses temporary overflow pages
bucket address چی지 ?
- Round-robin splitting:
 - split the bucket pointed by Next, then advance Next
- Use a family of hash functions h_0, h_1, h_2, \dots
- $h_{\text{Level}}(\text{key}) = h(\text{key}) \bmod (2^{\text{Level}}N)$
 - N = initial # buckets (must be a power of 2)
 - h is some hash function
- h_{i+1} doubles the range of h_i (similar to directory doubling)



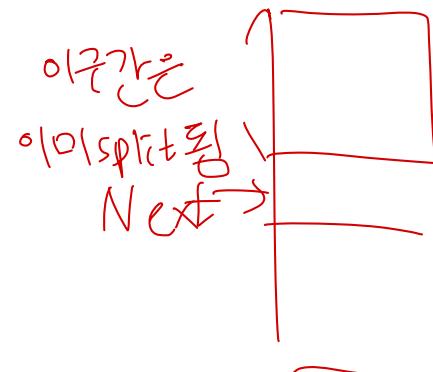
다차면 늘리고!

Linear Hashing - algorithm

- Algorithm works in 'rounds'.
- **Level** = Current round number

- At start of a round

- Total buckets = $N * 2^{Level}$ buckets (denoted as N_{Level})
- Buckets 0 to $Next - 1 \rightarrow$ already split →
- $Next$ to $N_{Level} - 1 \rightarrow$ not yet split →



- End of round when $Next = N_{Level}$

- Move to the next round:

`Level++;`

`Next = 0;`

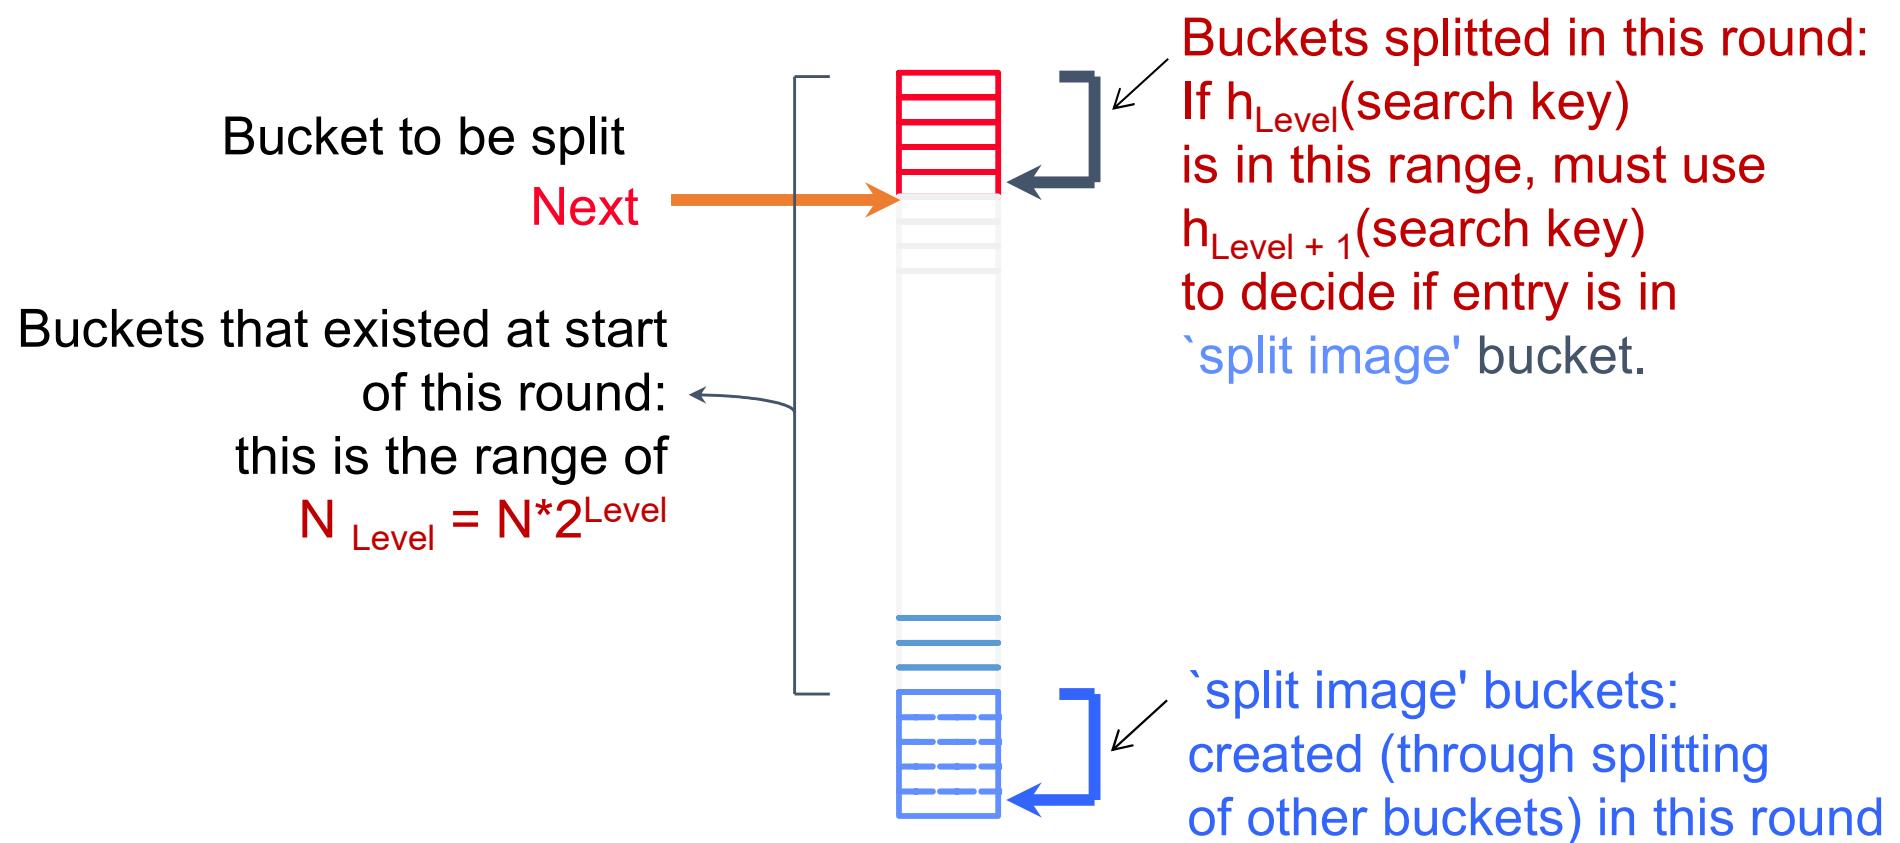
Linear Hashing - Insert

- Find target bucket using current hash (h_{Level})
- If there's space → insert and done.
- If no space:
 - Add a temporary overflow page and insert there.
 - Then, split the bucket pointed by **Next** bucket.
 - *This is likely NOT the bucket that just overflowed!!!*
 - Increment **Next**.
 - After split, use $h_{\text{Level+1}}$ to **re-distribute** entries.
- **⚠️** Buckets are split **in round-robin**, not by need → **Long overflow chains may develop!**

$O(N)$
↳ overflow chain bucket ↳

Linear Hashing – How it Works

■ Overview

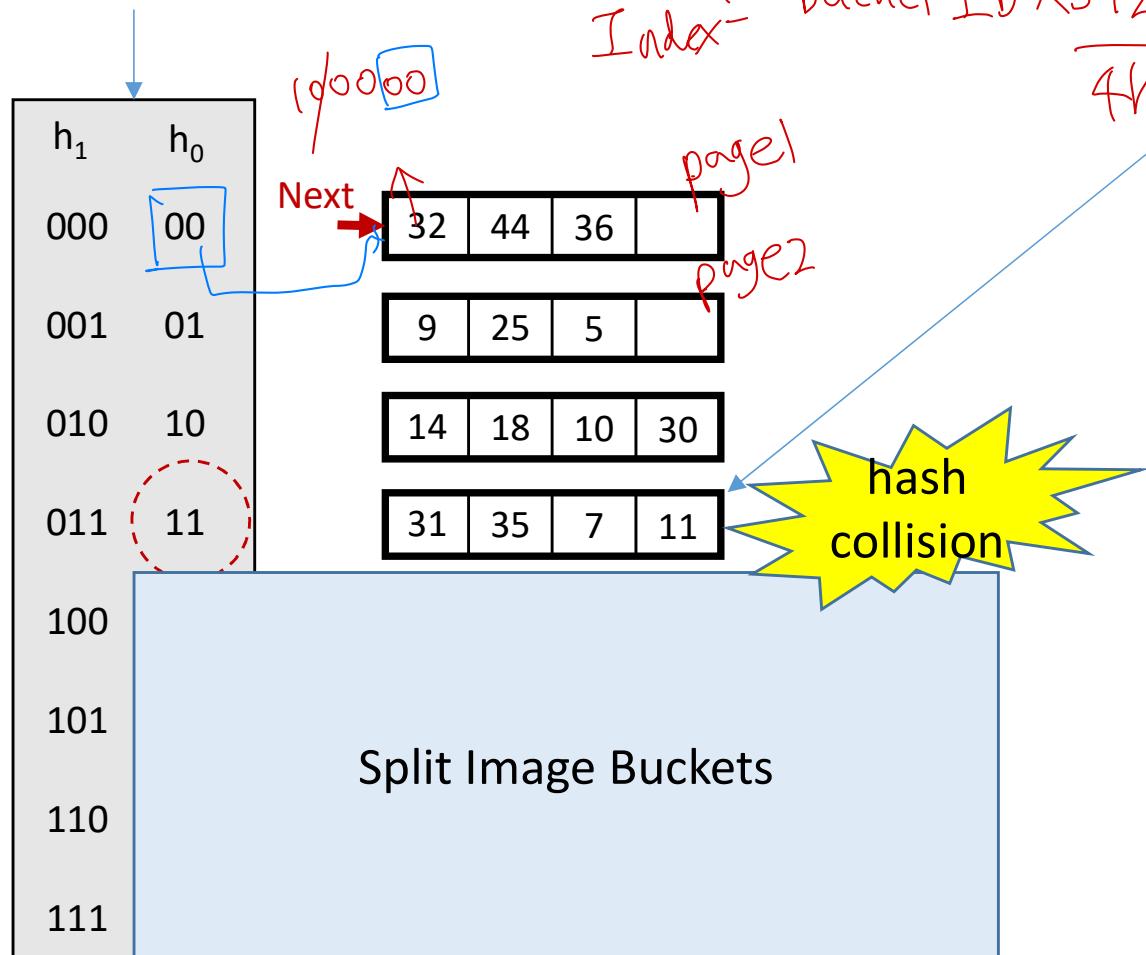


Linear Hashing – How it Works

- Insert 43 ($101011_{(2)}$) (Level 0, N=4)

$$43\%N = 11_{(2)}$$

This is for illustration only!

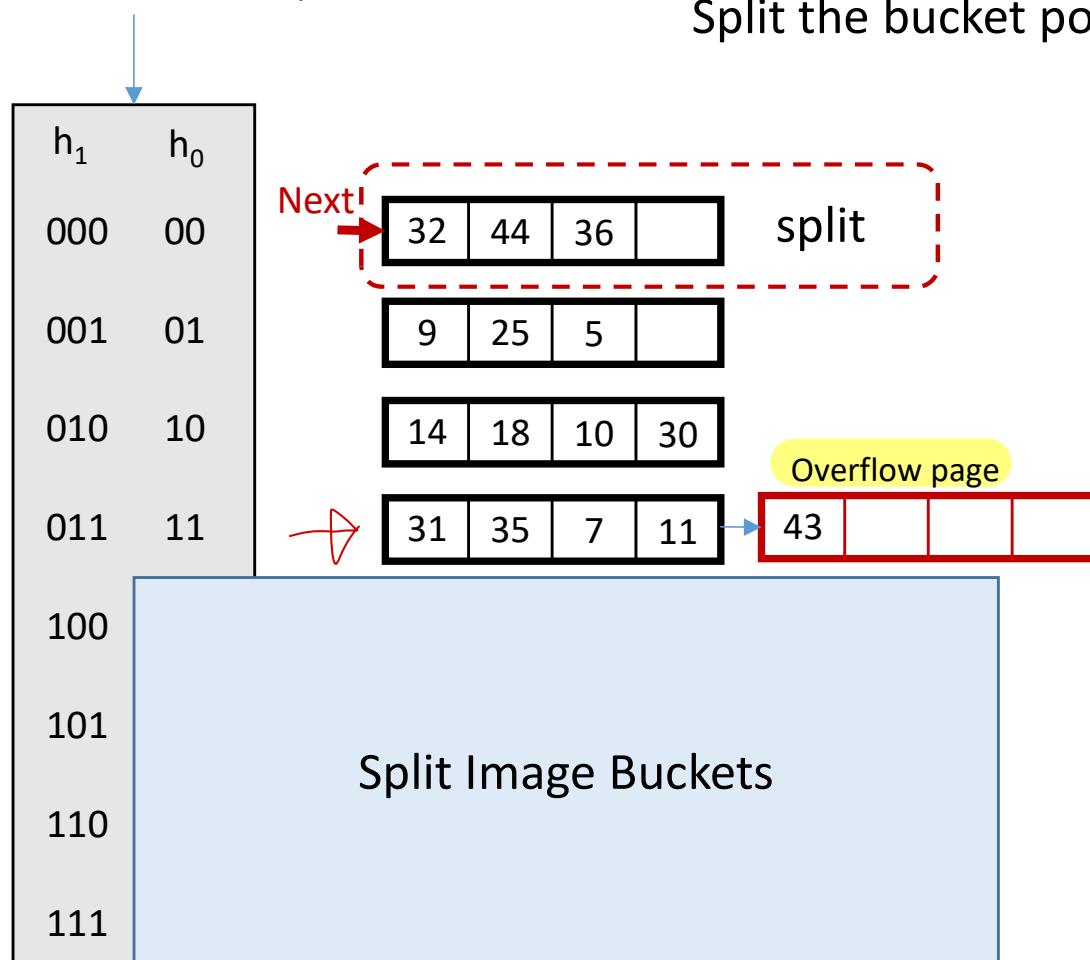


Linear Hashing – How it Works

- Insert 43 ($101011_{(2)}$) (Level 0, N=4)

$$43 \% N = 11_{(2)}$$

This is for illustration only!



Split the bucket pointed by Next

$$\begin{aligned}32 &= 100000_{(2)} \\44 &= 101100_{(2)} \\36 &= 100100_{(2)}\end{aligned}$$

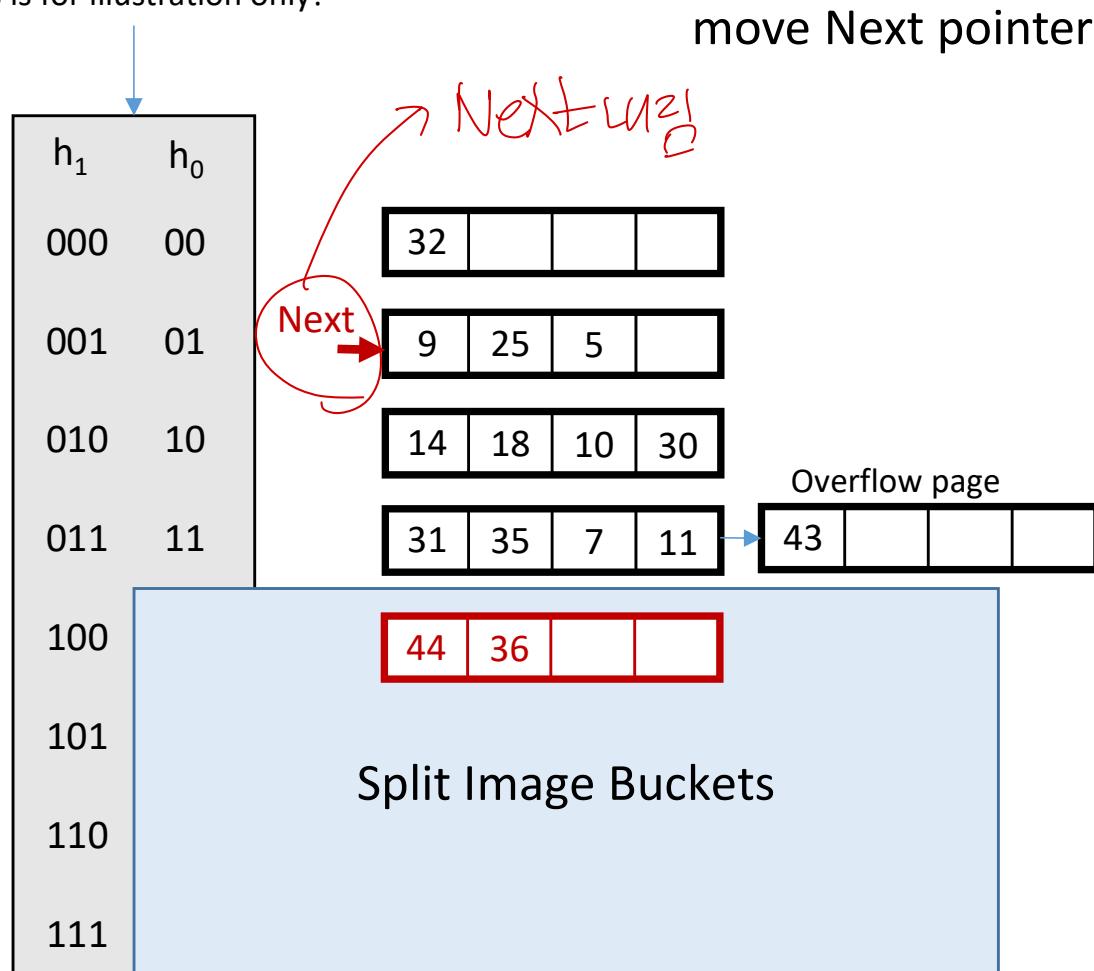
Collision 발생 \rightarrow Next pointer 가리키는 버킷 split

Linear Hashing – How it Works

- Insert 43 ($101011_{(2)}$) (Level 0, N=4)

$$43 \% N = 11_{(2)}$$

This is for illustration only!



$$32 = 100\textcolor{red}{000}_{(2)}$$

$$44 = 101\textcolor{red}{100}_{(2)}$$

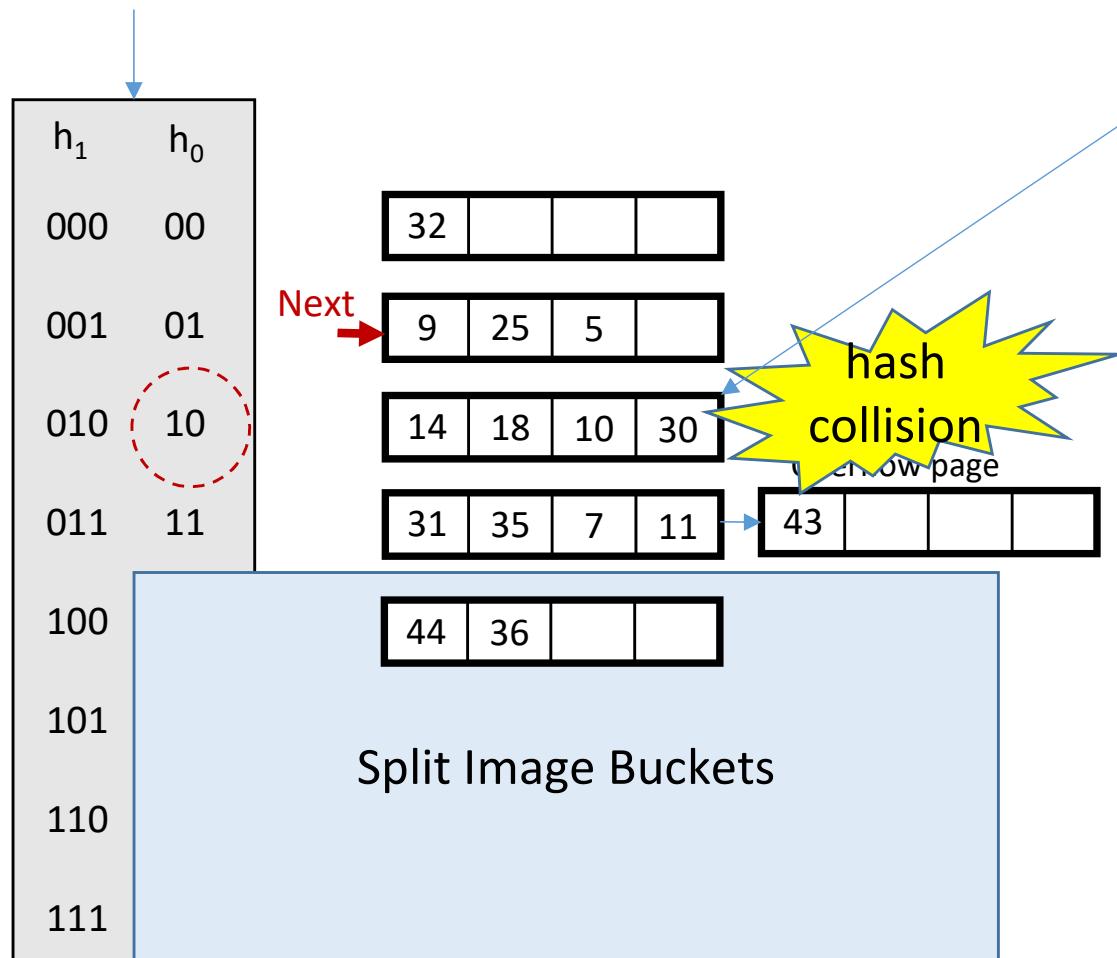
$$36 = 100\textcolor{red}{100}_{(2)}$$

Linear Hashing – How it Works

- Insert 6 ($000110_{(2)}$) (Level 0, N=4)

$$6\%N = 10_{(2)}$$

This is for illustration only!

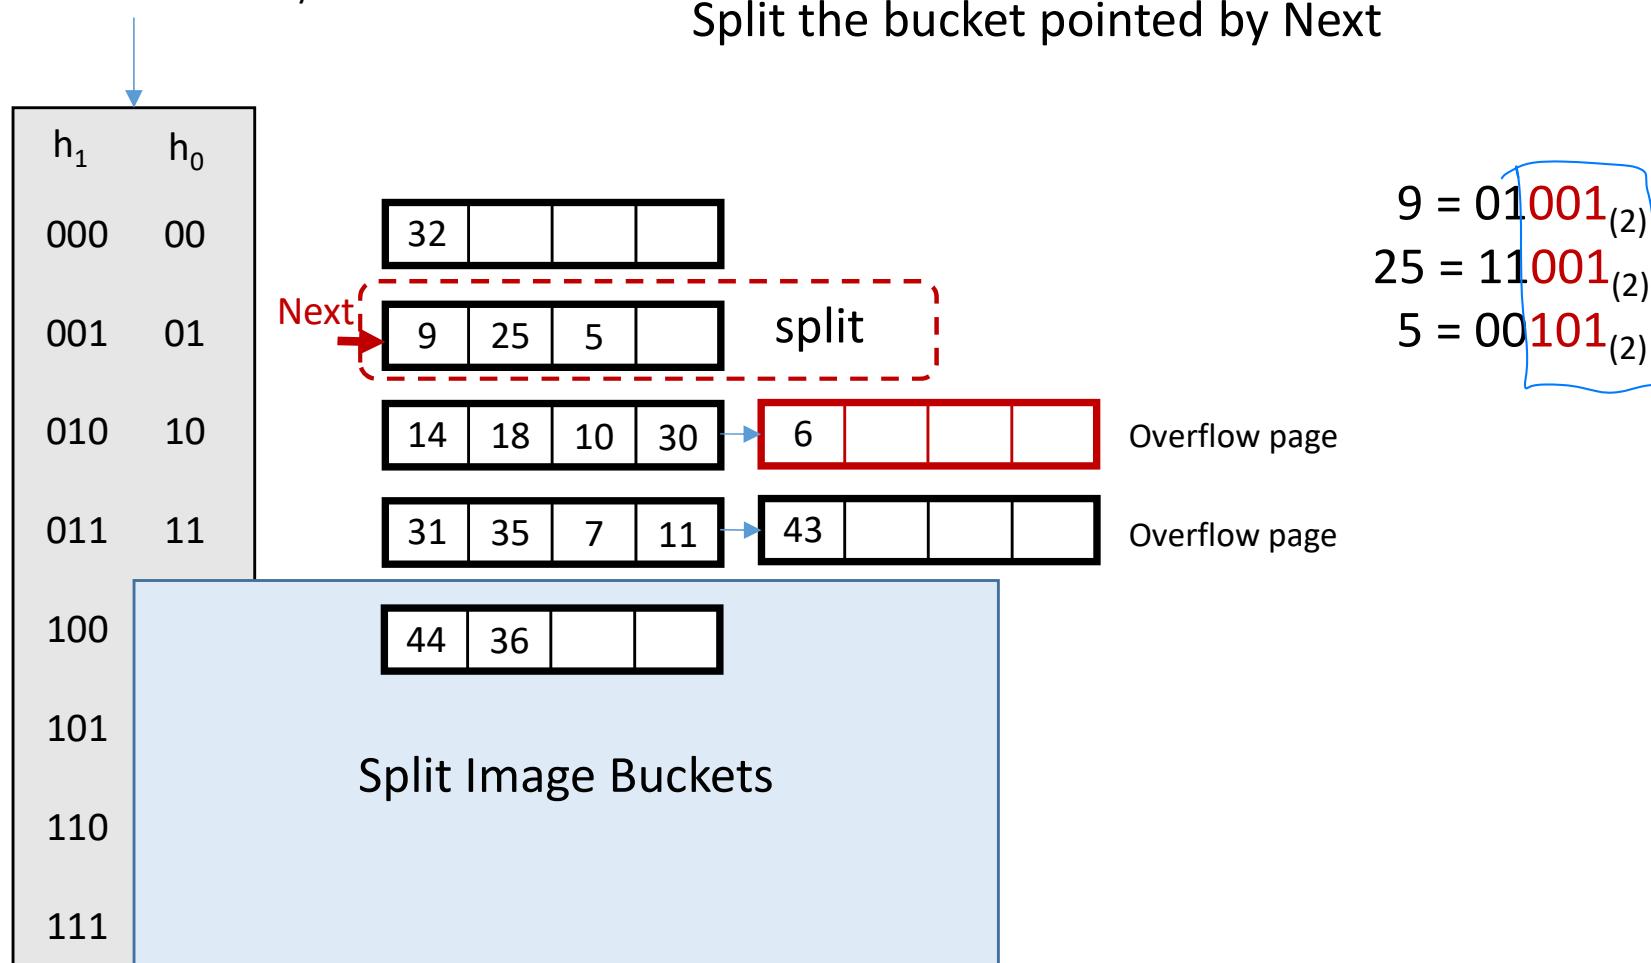


Linear Hashing – How it Works

- Insert 6 ($000110_{(2)}$) (Level 0, N=4)

$$6\%N = 10_{(2)}$$

This is for illustration only!

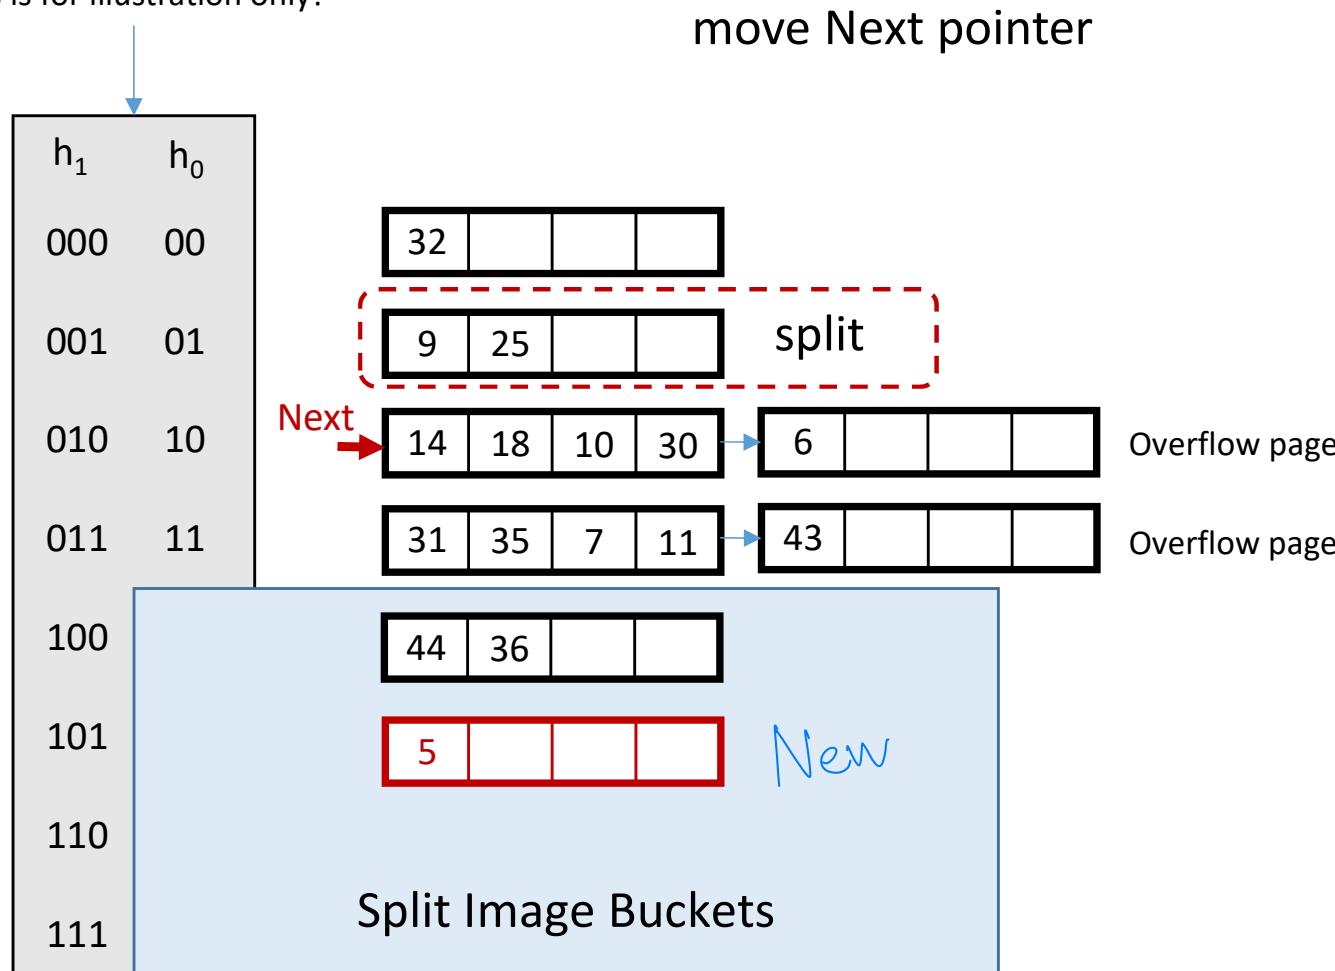


Linear Hashing – How it Works

- Insert 6 ($000110_{(2)}$) (Level 0, N=4)

$$6\%N = 10_{(2)}$$

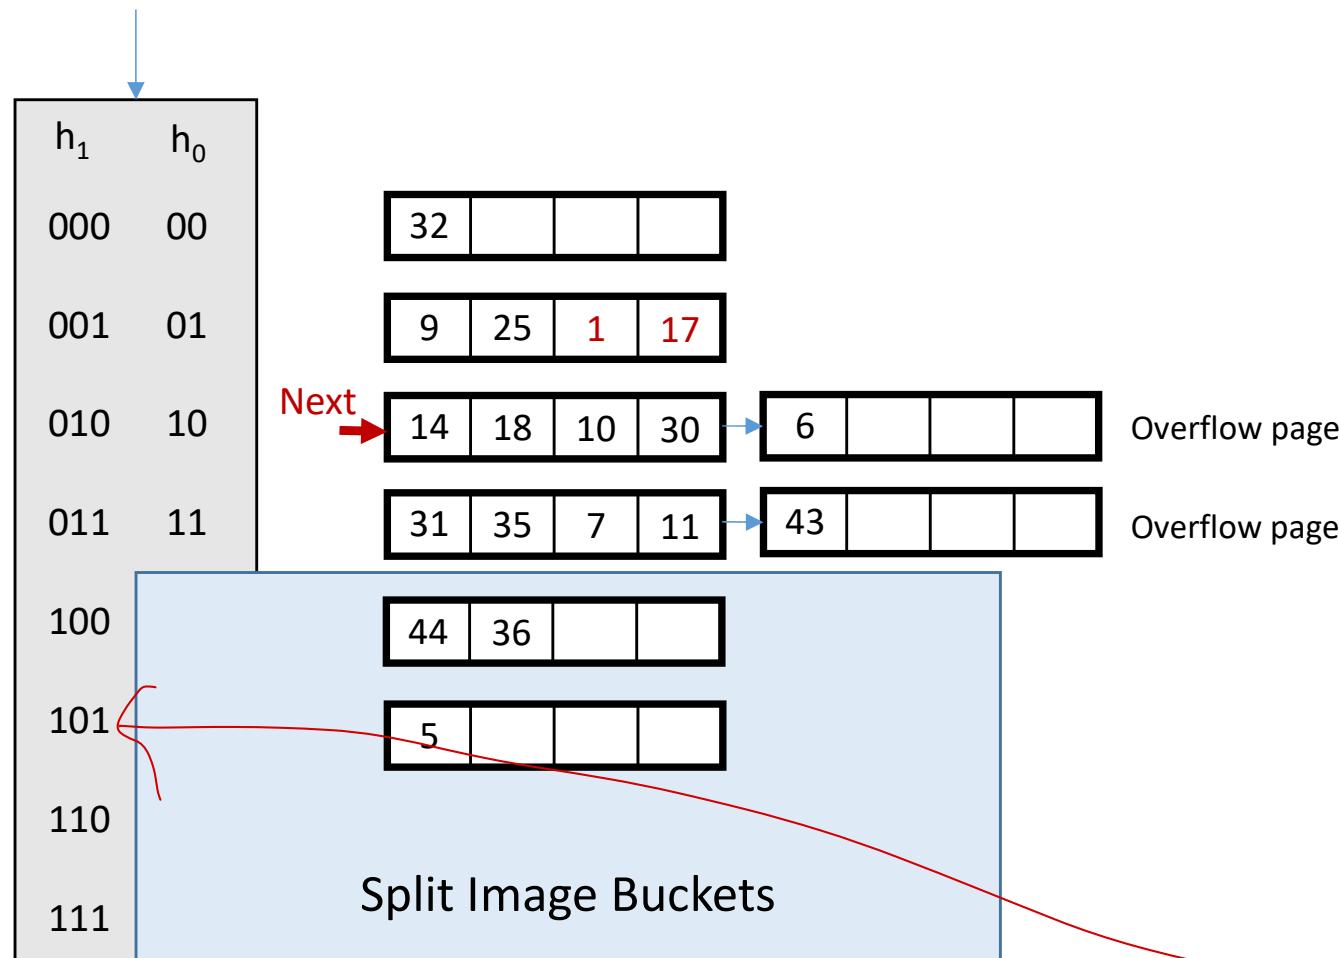
This is for illustration only!



Linear Hashing – How it Works

- Insert 1 & 17 ($000001_{(2)}$, $010001_{(2)}$) (Level , N=4) $1\%N = 01_{(2)}$

This is for illustration only!



$$25 = 1[00] \quad) \text{비교}$$

Next=10

2bit 확인 01 < 10

→ previous buckets

3bit 확인은 아니야
→ New(split 된) 버켓을

호환하지

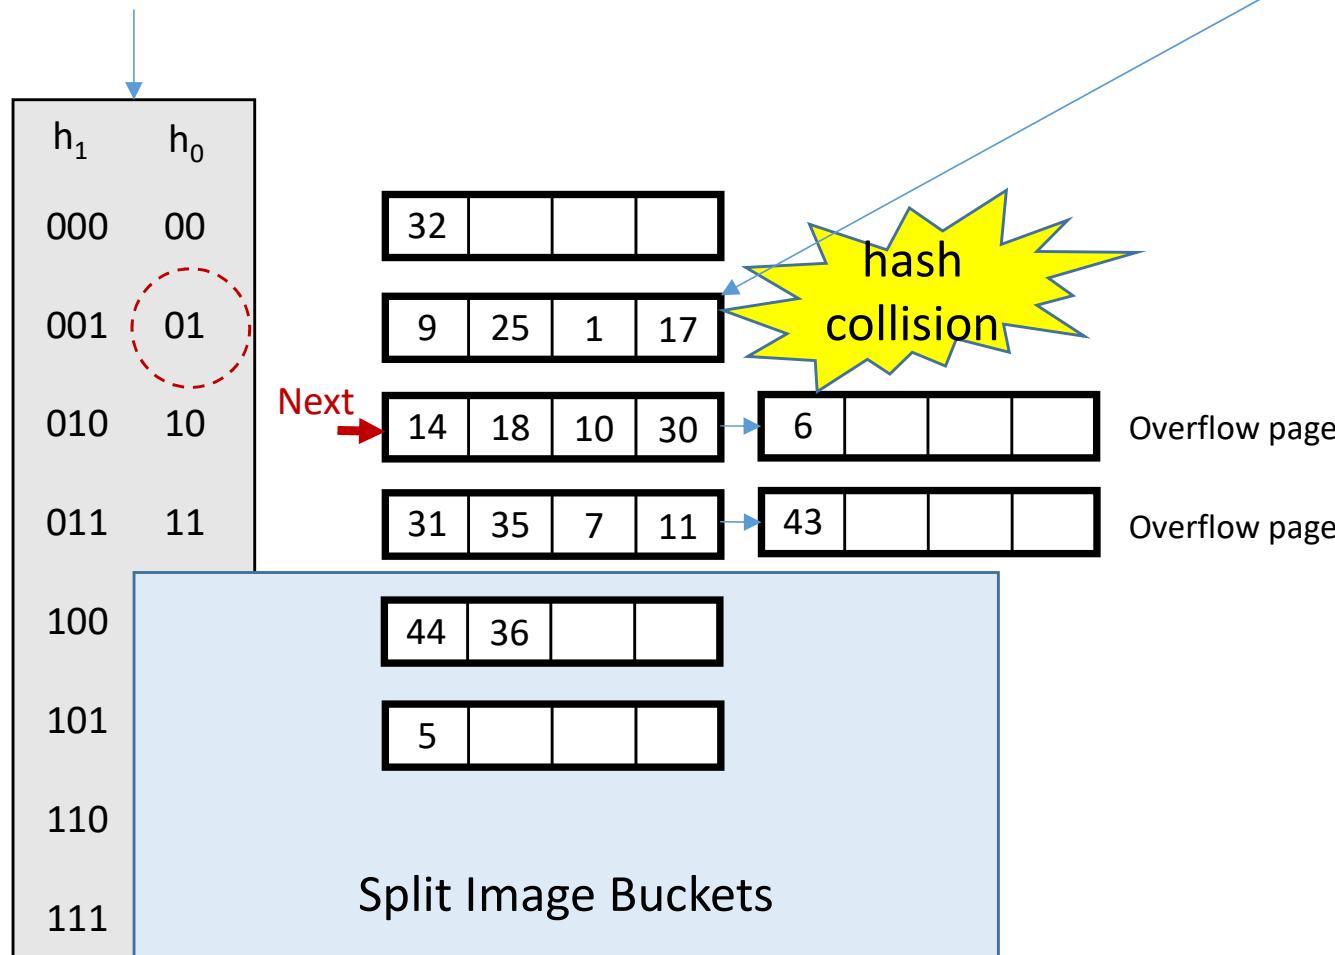
$$29 = 1[0] \quad)$$

Linear Hashing – How it Works

- Insert 33 ($100001_{(2)}$) (Level 0, N=4)

$$33\%N = 01_{(2)}$$

This is for illustration only!

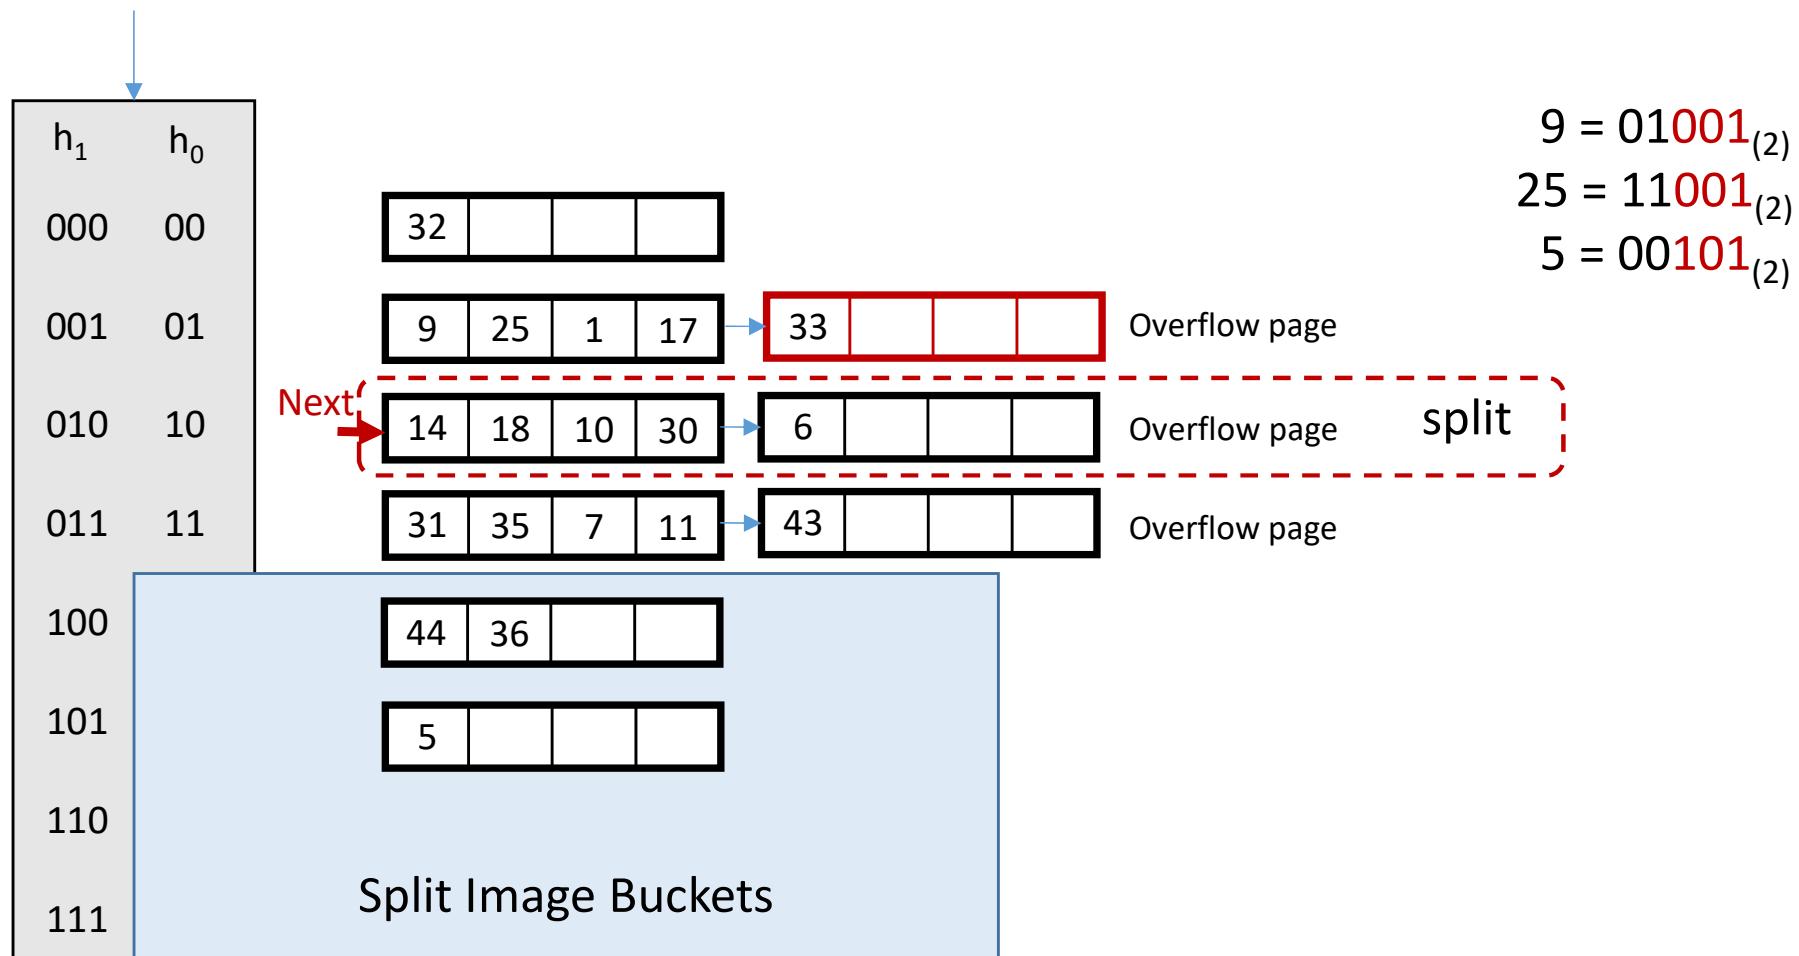


Linear Hashing – How it Works

- Insert 33 ($100001_{(2)}$) (Level 0, N=4)

$$33\%N = 01_{(2)}$$

This is for illustration only!

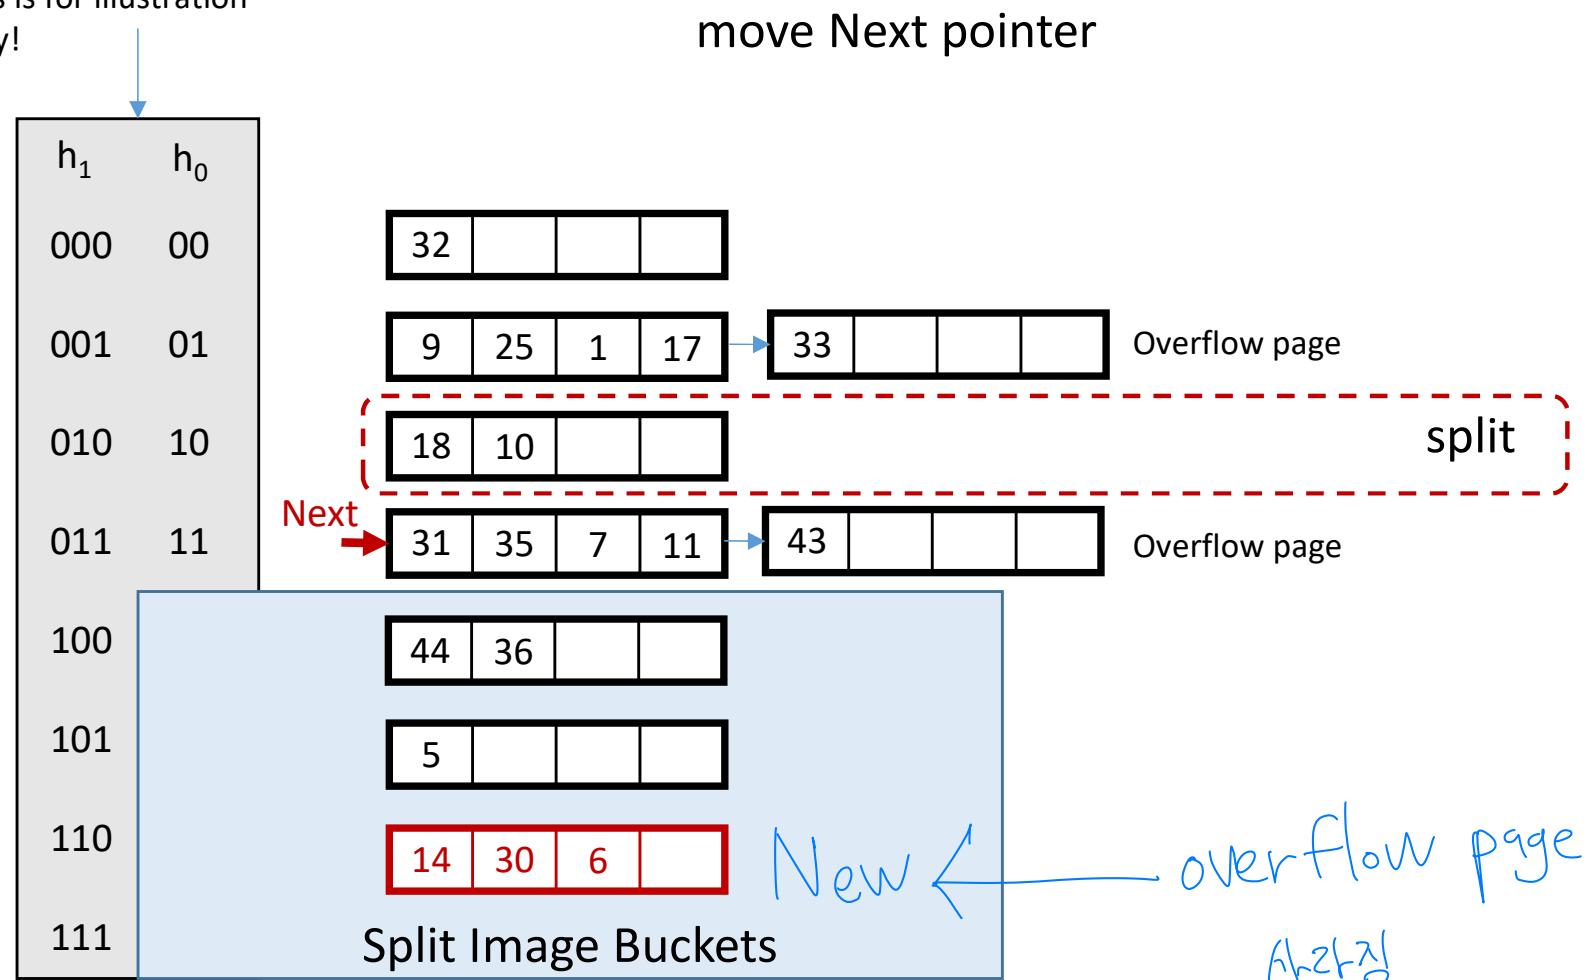


Linear Hashing – How it Works

- Insert 33 ($100001_{(2)}$) (Level 0, N=4)

$$33 \% N = 01_{(2)}$$

This is for illustration only!



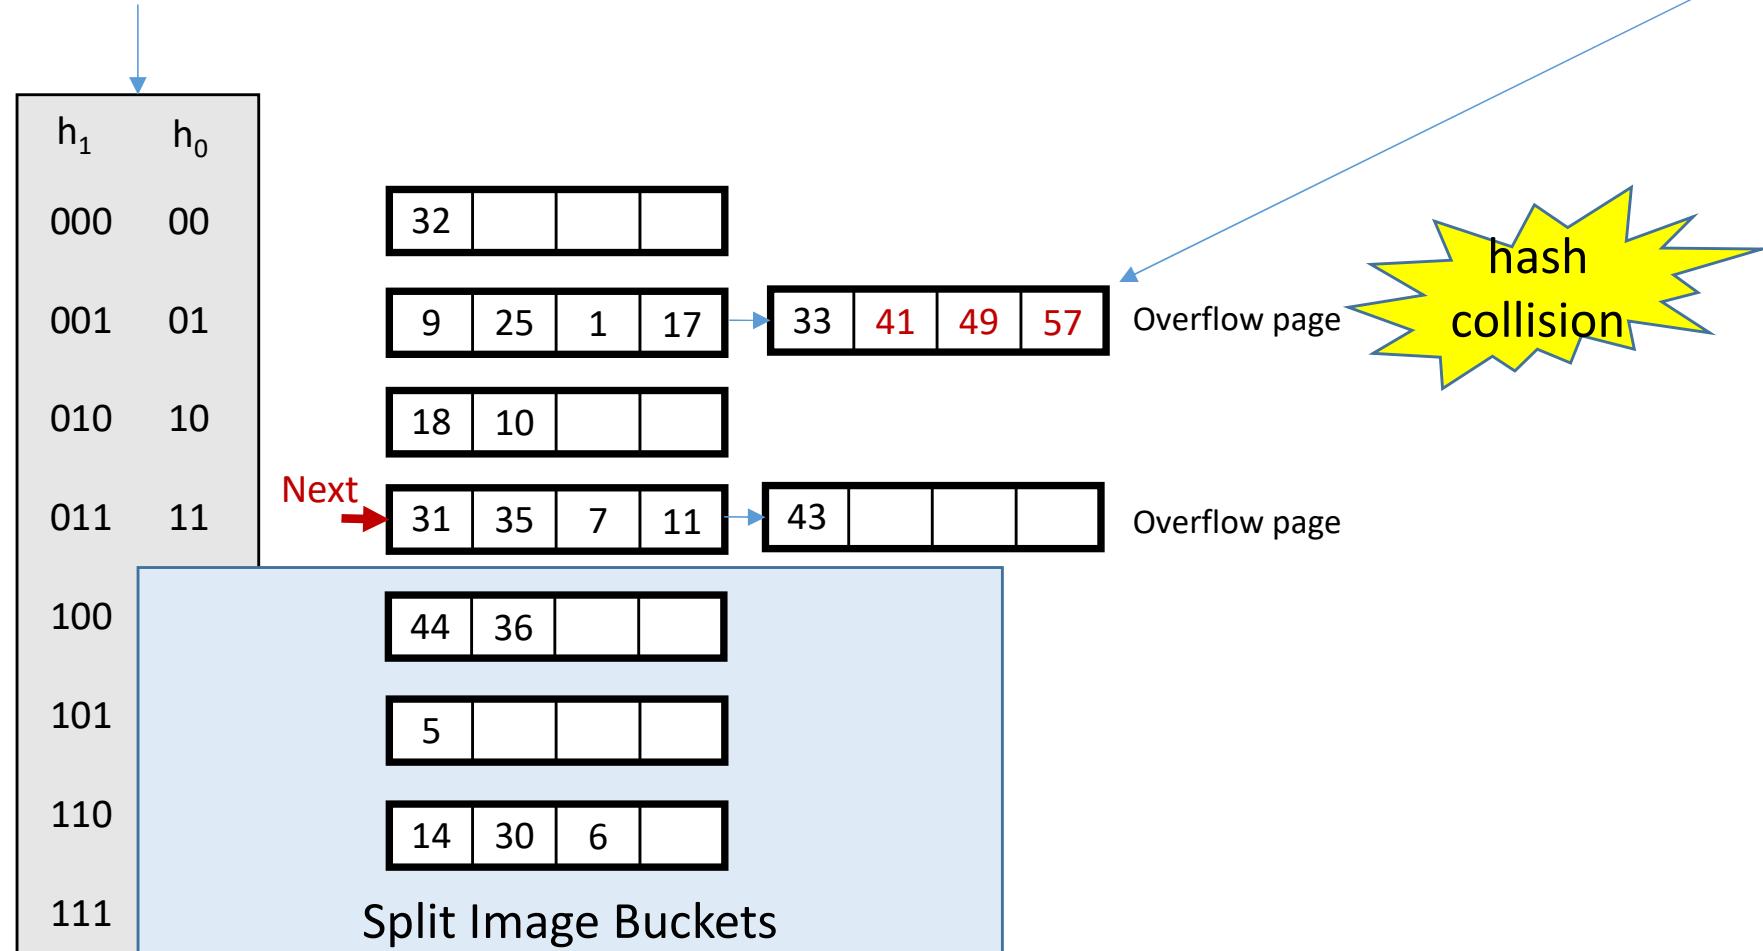
14 = 01	110	(2)
18 = 10	010	(2)
10 = 01	010	(2)
30 = 11	110	(2)
6 = 00	110	(2)

사각점

Linear Hashing – How it Works

- Insert 41, 49, 57, 65 ($101001_{(2)}$, $110001_{(2)}$, $111001_{(2)}$, $1000001_{(2)}$)

This is for illustration only!

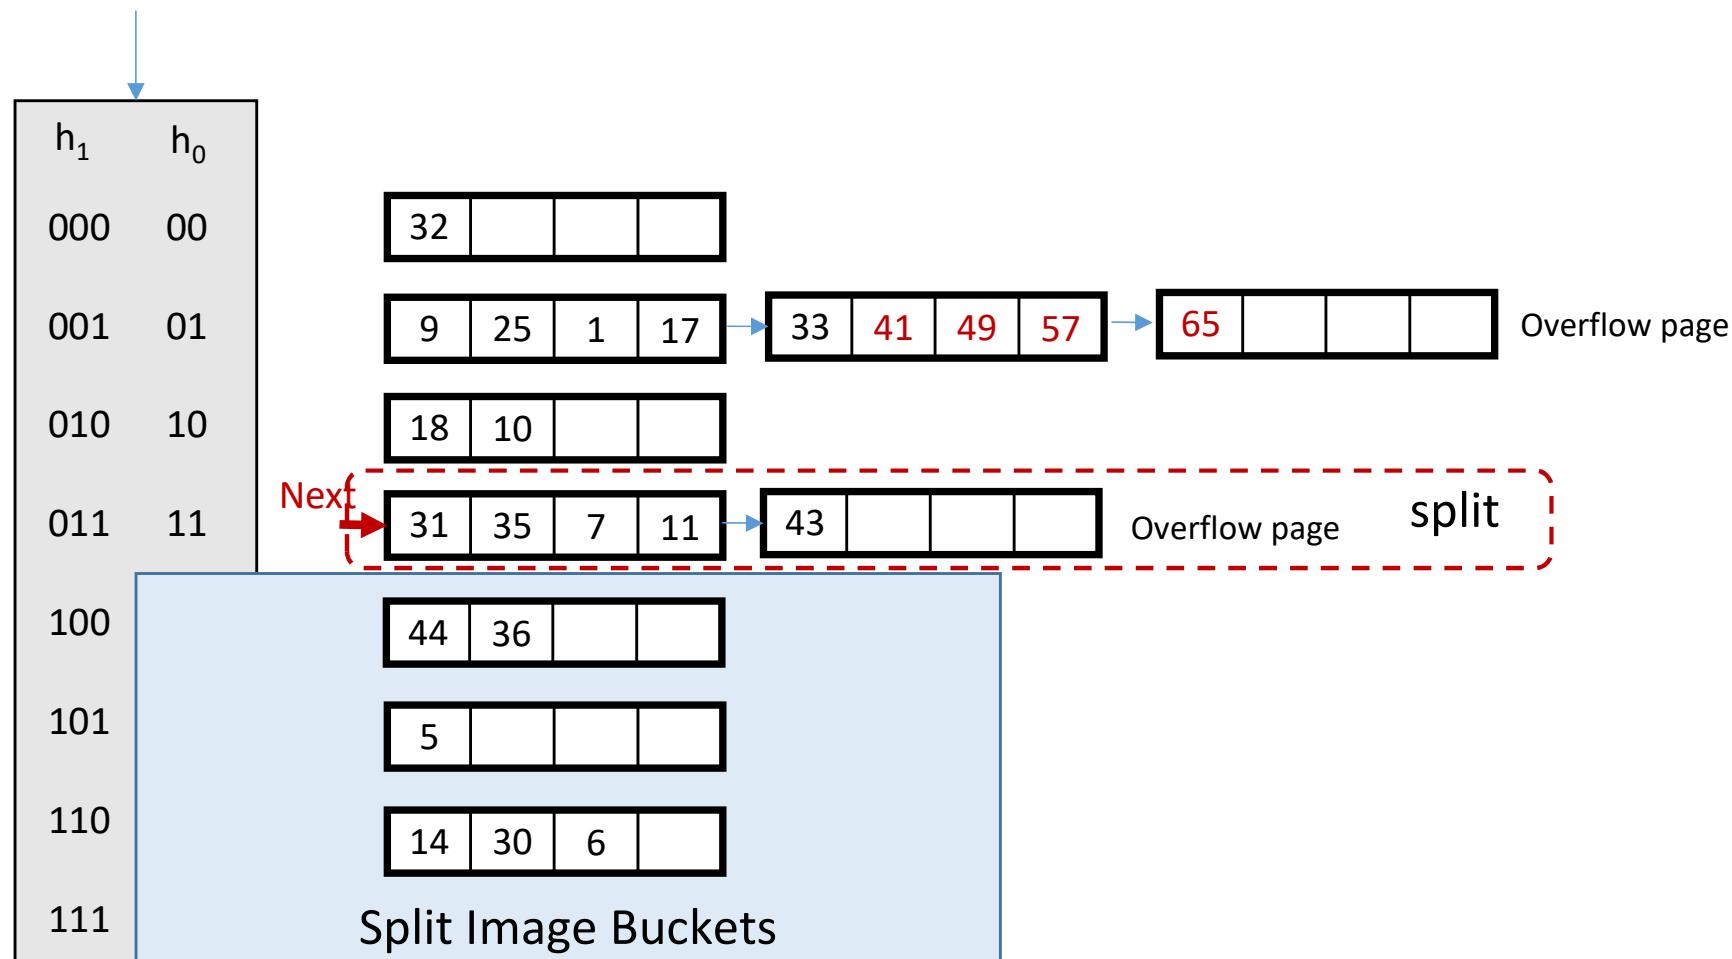


Linear Hashing – How it Works

- Insert 41, 49, 57, 65 ($101001_{(2)}$, $110001_{(2)}$, $111001_{(2)}$, $1000001_{(2)}$)

This is for illustration only!

$$65\%N = 01_{(2)}$$

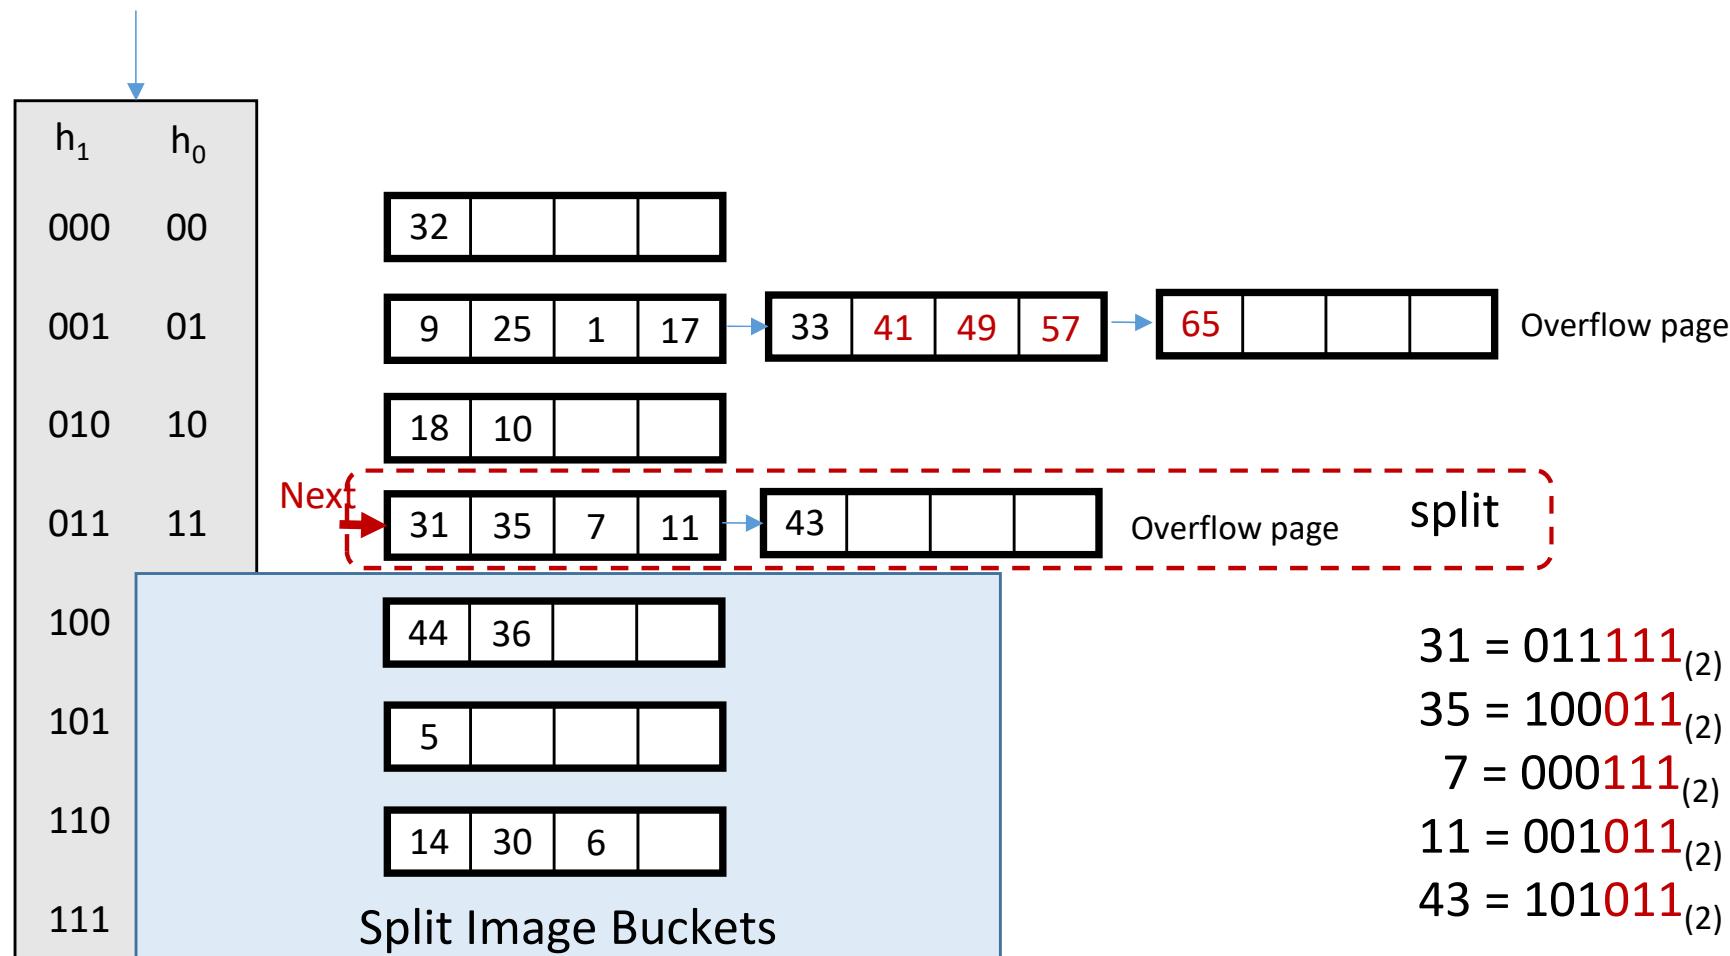


Linear Hashing – How it Works

- Insert 41, 49, 57, 65 ($101001_{(2)}$, $110001_{(2)}$, $111001_{(2)}$, $1000001_{(2)}$)

This is for illustration only!

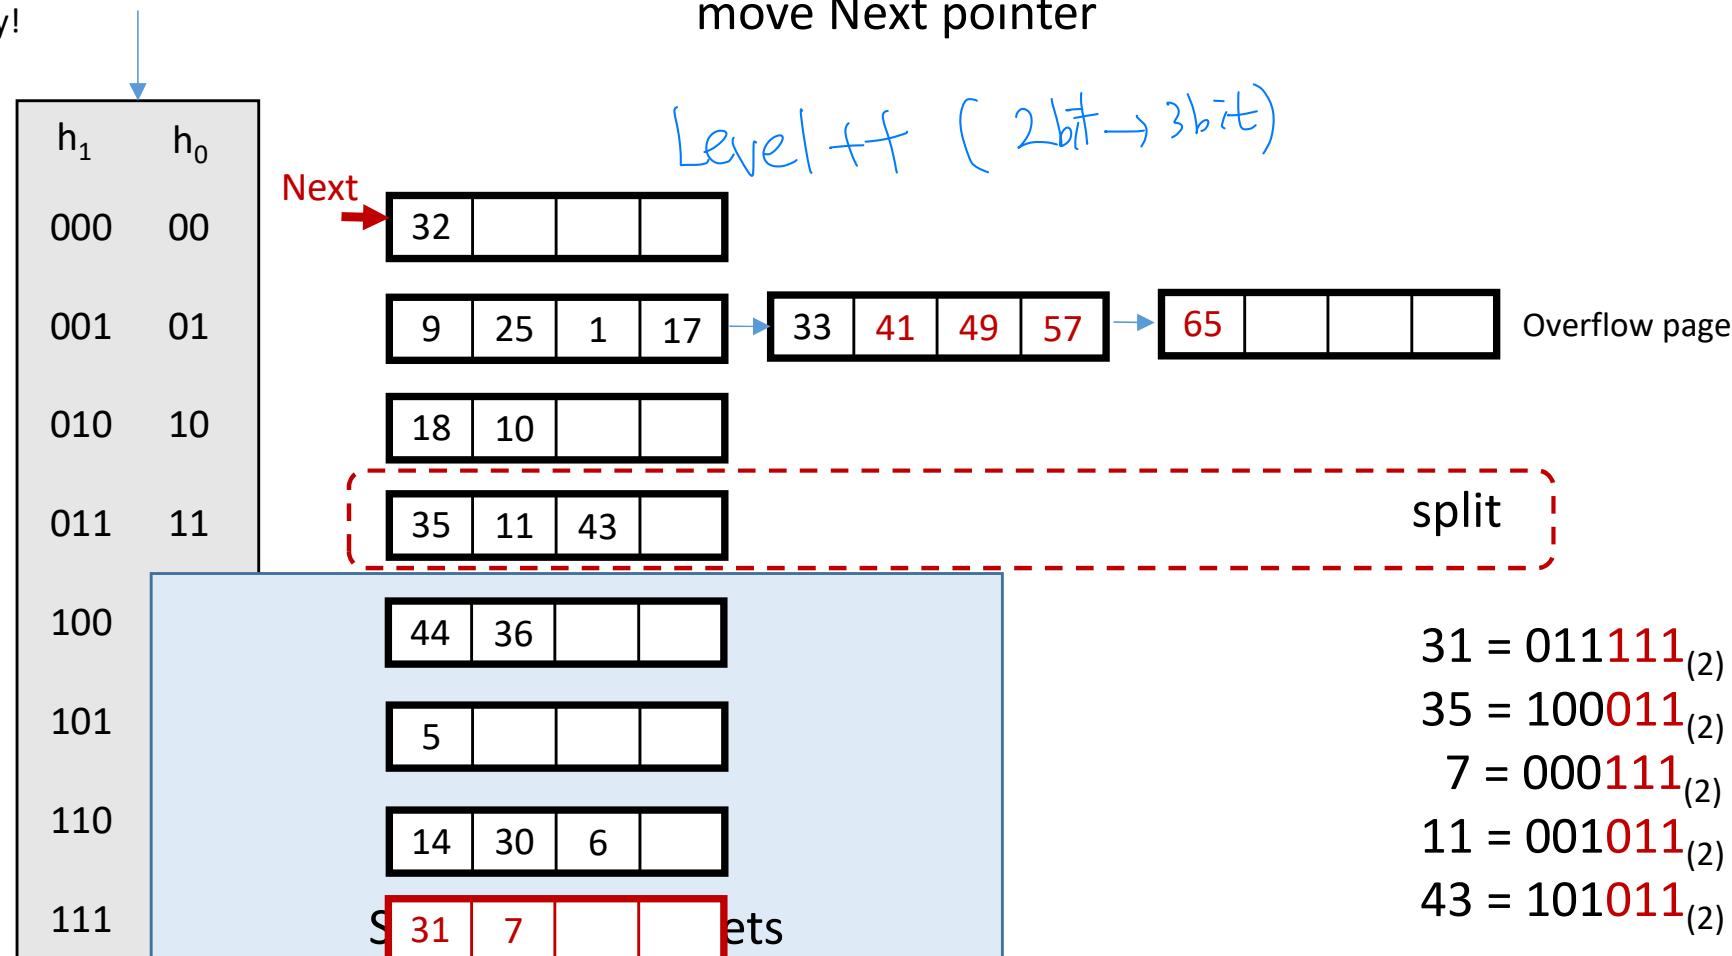
$$65\%N = 01_{(2)}$$



Linear Hashing – How it Works

- Insert 41, 49, 57, 65 ($101001_{(2)}$, $110001_{(2)}$, $111001_{(2)}$, $1000001_{(2)}$)

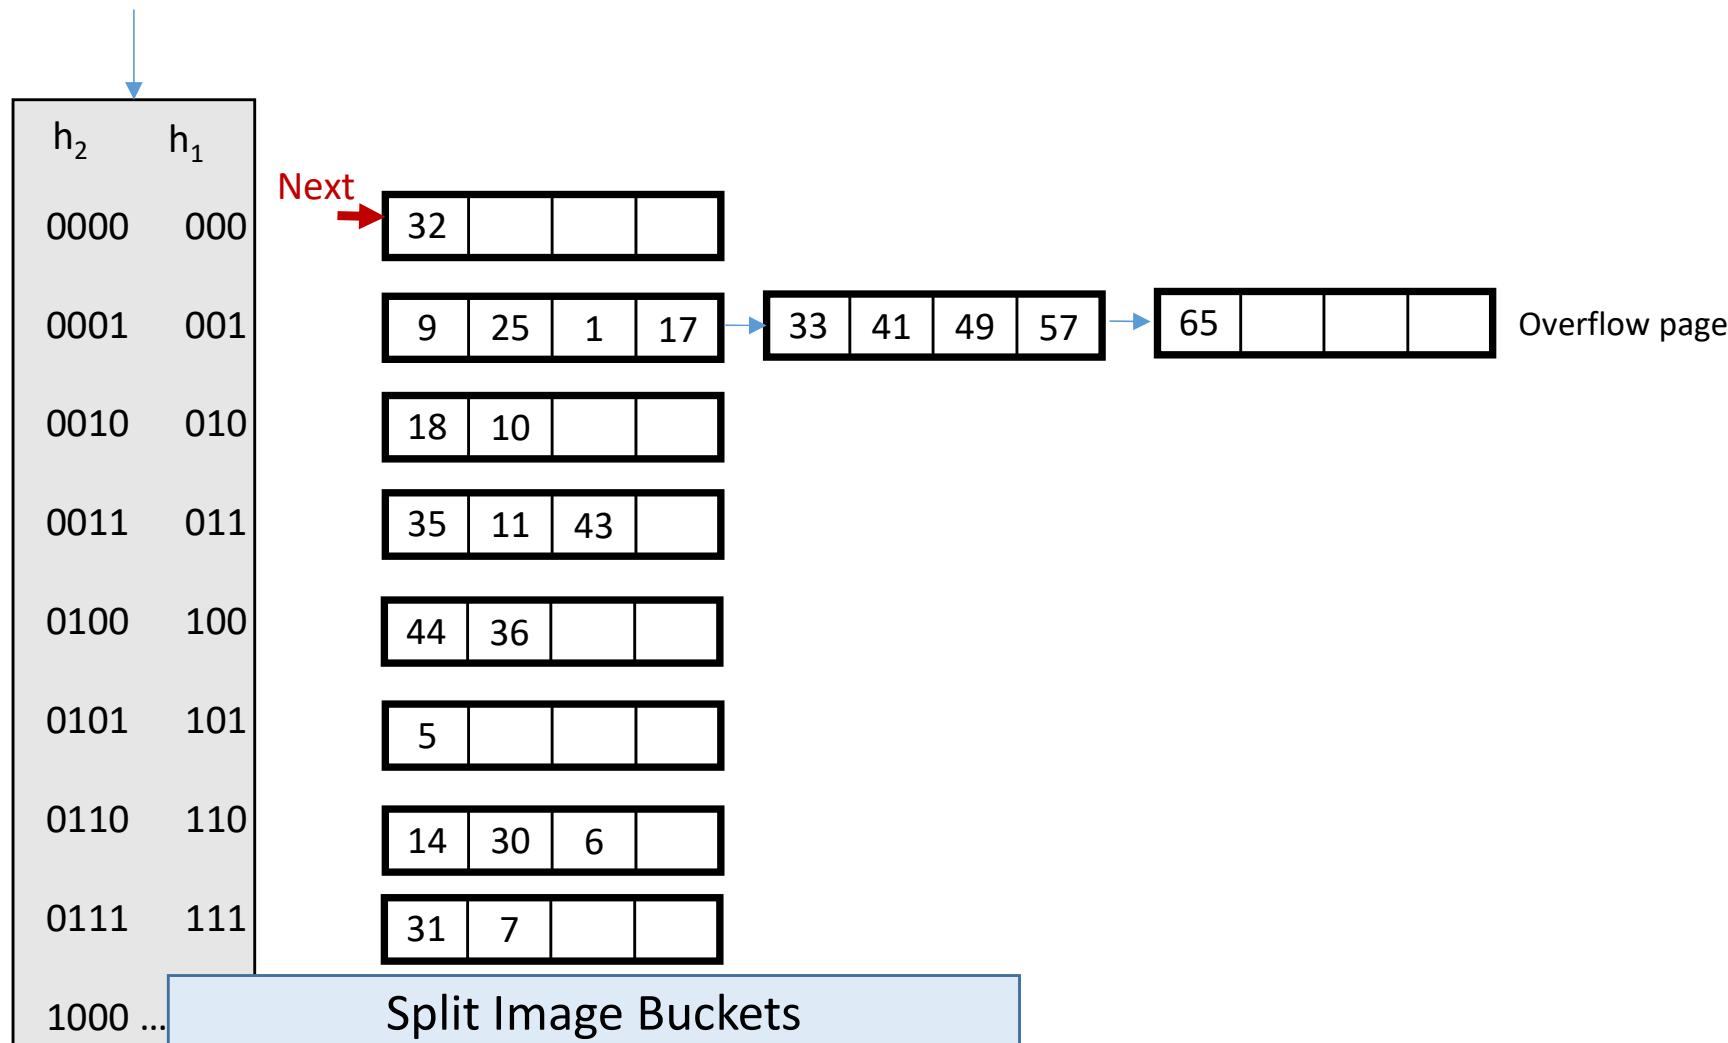
This is for illustration
only!



Linear Hashing – How it Works

- Level 1, N=8

This is for illustration only!



Linear Hashing Search Algorithm

- Compute $h_{Level}(key)$ to find a target bucket for key
- If $h_{Level}(r) \geq Next \rightarrow after$
 - Bucket that has not been split this round
 - key belongs to that bucket for sure
- Else $\rightarrow previous$ (이전 스플릿됨) \rightarrow additional bit 확인!
 - Buckets has already been split $2bit \rightarrow 3bit$
 - Must recompute $h_{Level+1}(r)$ or $h_{Level}(r) + N_{Level}$
 - Key will go to either
 - $h_{Level}(r)$ or
 - $h_{Level}(r) + N_{Level}$

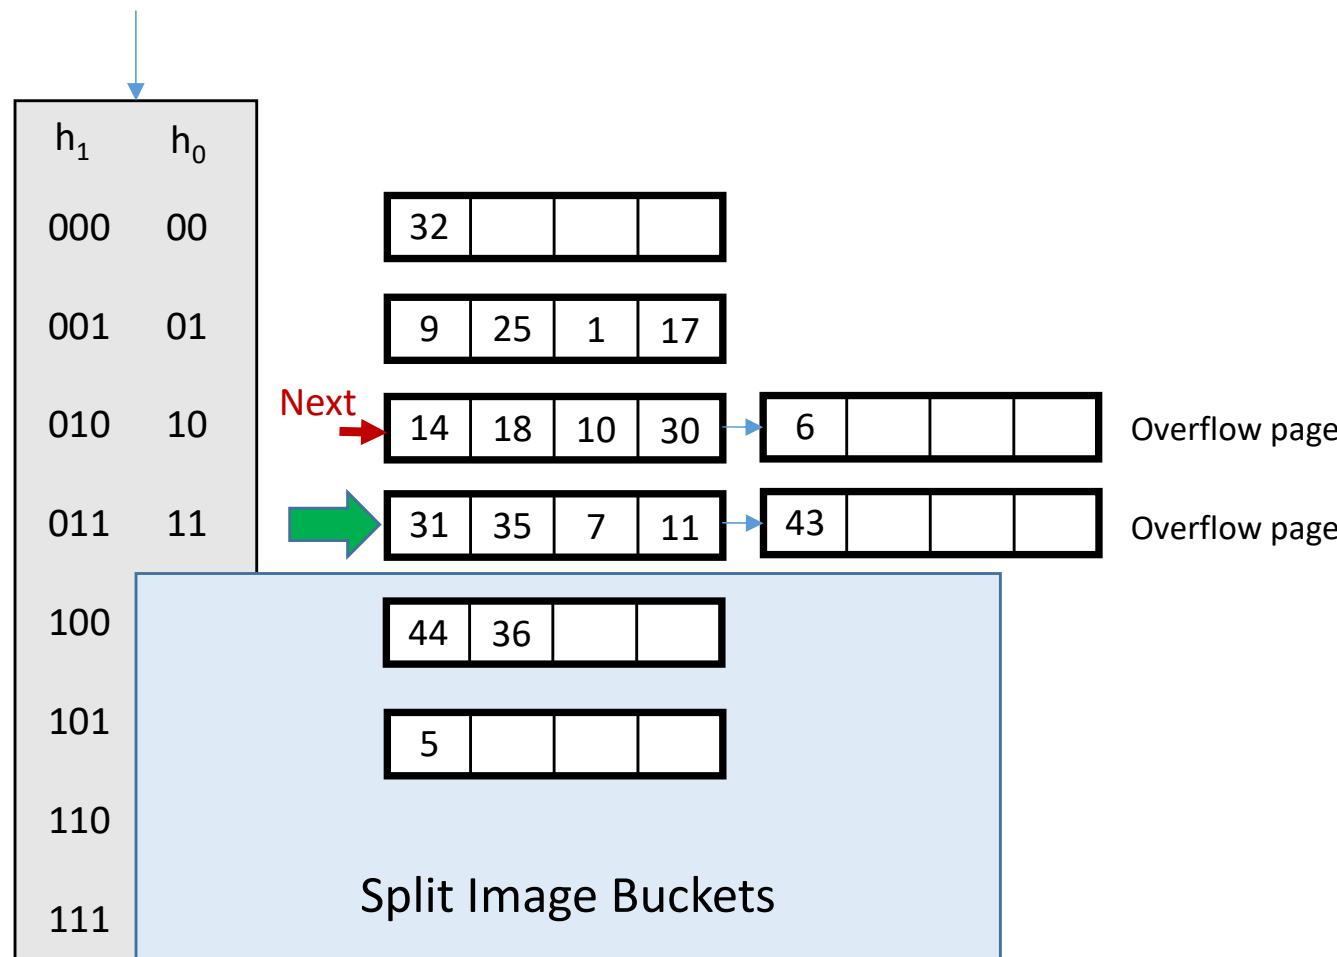
Linear Hashing – How it Works

- Search 7 ($00111_{(2)}$) (Level , N=4)

$$7\%N = 11_{(2)} \geq \text{Next}$$

This is for illustration only!

bucket $11_{(2)}$ has not split yet



Linear Hashing – How it Works

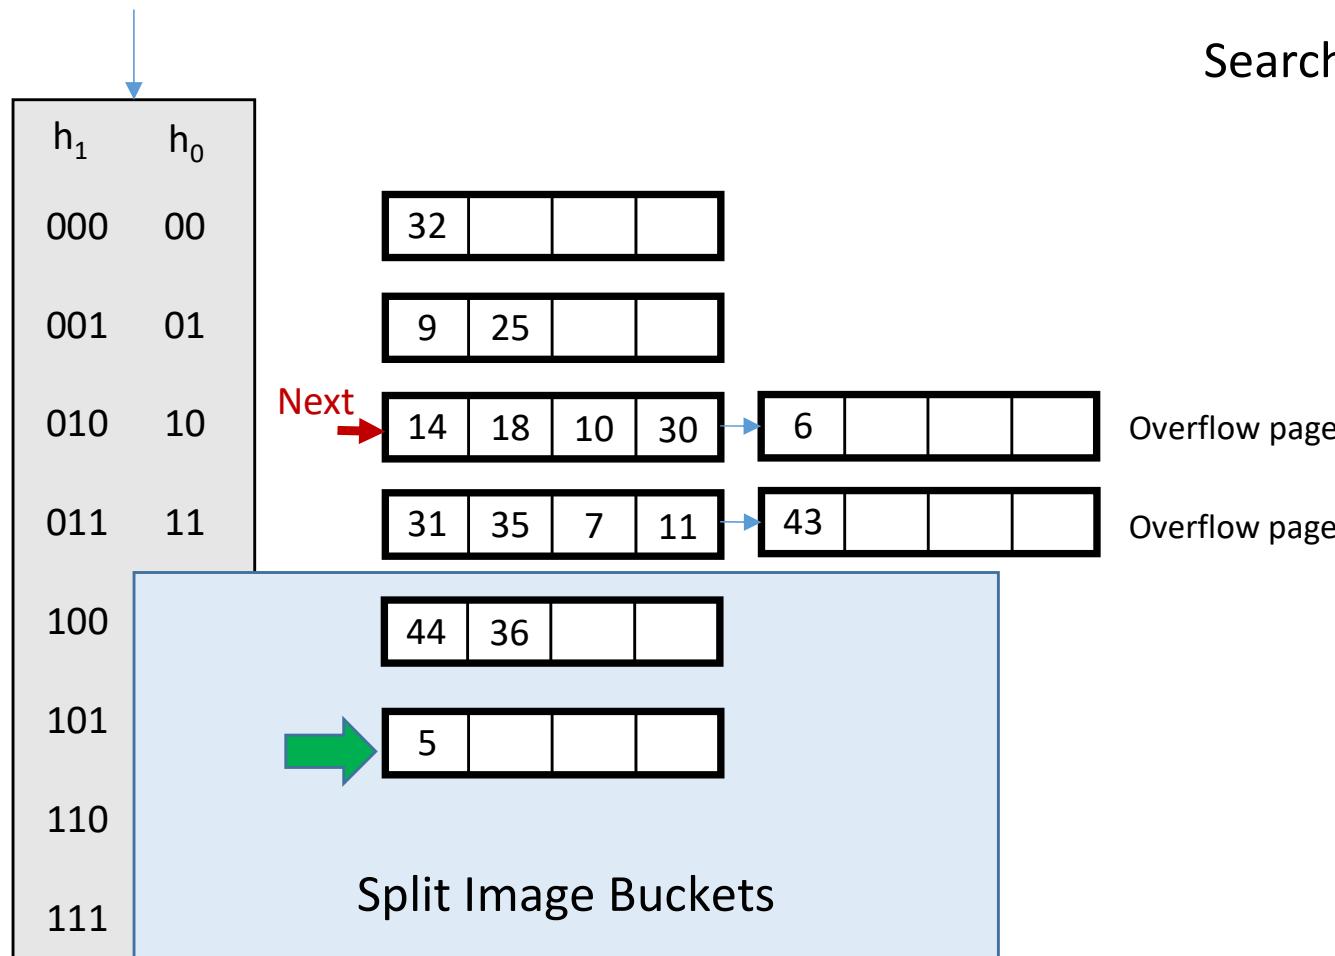
- Search 5 ($000101_{(2)}$) (Level , N=4)

$$5\%N = 01_{(2)} < \text{Next}$$

This is for illustration only!

bucket $01_{(2)}$ has split

Search bucket $5\%(2^{\text{level}}N) = 101_{(2)}$



Comparison of Ordered Indexing and Hashing

- Cost of periodic re-organization
*비교 (< <)
다른 알고리즘*
- Relative frequency of insertions and deletions
- Is it desirable to optimize average access time at the expense of worst-case access time?
- Expected type of queries:
 - Hashing is generally better for point query
 - If range queries are common, ordered indices are to be preferred
- In practice:
 - PostgreSQL supports hash indices, but discourages use due to poor performance
 - Oracle supports static hash organization, but not hash indices
 - SQLServer supports only B+-trees*Hash based Index 사용 권장 X overflow 가능*

Multiple-Key Access

- Use multiple indices for certain types of queries.
- Example:

```
select ID
```

```
from instructor
```

```
where dept_name = "Finance" and salary = 80000
```

Index → 1 번째 조건 먼저 체크 → 2 번째 체크

- Possible strategies for processing query using indices on single attributes:
 1. Use index on *dept_name* to find instructors with department name Finance; test *salary* = 80000
 2. Use index on *salary* to find instructors with a salary of \$80000; test *dept_name* = "Finance".
 3. Take intersection of both sets

교집합

Multiple-Key Access and Bitmap Index

다중 키의 DBMS 가속화

(OH $\in O(n^2)$)

- Bitmap indices are designed for efficient querying on multiple keys
- Records are assumed to be numbered sequentially from, say, 0
 - Given a number n it must be easy to retrieve record n
 - Particularly easy if records are of fixed size
- Applicable on attributes that take on a small number of distinct values
 - E.g., gender, country, state, ...
 - E.g., income-level (income broken up into a small number of levels such as 0-9999, 10000-19999, 20000-50000, 50000- infinity)
- A bitmap is simply an array of bits

Bitmap Index

- a bitmap index has a bitmap for each value of the attribute
 - Bitmap has as many bits as records
 - In a bitmap for value v, the bit for a record is 1 if the record has the value v for the attribute, and is 0 otherwise

Example

Ex) 8bit size

record number	ID	gender	income_level
0	76766	m 1800	L1 100
1	22222	f	L2 1000
2	12121	f	L1 2000
3	15151	m 4000	L4 500
4	58583	f	L3 1500

gender Index



Bitmaps for gender

m	10010
f	01101

5bit \rightarrow 5bit

unidentified 2. 01110
(조금 다른 경우)
bitwise and 연산자

Bitmaps for income_level

L1	10100
L2	01000
L3	00001 \rightarrow L1 male
L4	00010
L5	00000

Index

0 ~ 400 / 401 ~ 800 / ... 범위 생성

Bitmap Indices (Cont.)

- Bitmap indices are useful for queries on multiple attributes
 - not particularly useful for single attribute queries
- Queries are answered using bitmap operations
 - Intersection (and)
 - Union (or)
- Each operation takes two bitmaps of the same size and applies the operation on corresponding bits to get the result bitmap
 - E.g., $100110 \text{ AND } 110011 = 100010$
 $100110 \text{ OR } 110011 = 110111$
 $\text{NOT } 100110 = 011001$
 - Males with income level L1: $10010 \text{ AND } 10100 = 10000$
 - Can then retrieve required tuples.
 - Counting number of matching tuples is even faster

bit wise 연산

Bitmap Indices (Cont.)

- Bitmap indices generally very small compared with relation size
 - E.g., if record is 100 bytes, space for a single bitmap is 1/800 of space used by relation.
 - If number of distinct attribute values is 8, bitmap is only 1% of relation size

Creation of Indices

- Example

create index takes_pk Index name
on takes (ID,course_ID, year, semester, section)

primary key \rightarrow Index D₂₀
non-primary key \rightarrow createIndex

drop index takes_pk
Index ~~exists~~

- Most database systems allow specification of type of index, and clustering.
- Indices on primary key created automatically by all databases
- Some database also create indices on foreign key attributes
 - Index on foreign key would be useful for this query:
 - $takes \bowtie \sigma_{name='Shankar'}(student)$
- Indices can greatly speed up lookups, but impose cost on updates