

# Database Systems

## Lecture16 – Chapter 17: Transactions

Beomseok Nam (남범석)

[bnam@skku.edu](mailto:bnam@skku.edu)

# Outline

- Transaction Concept
- Transaction State
- Concurrent Executions
- Serializability
- Recoverability
- Implementation of Isolation
- Transaction Definition in SQL
- Testing for Serializability.

# Transaction Concept

단위 일관성

## ■ Transaction

- a sequence of operations performed as a **single logical unit** of work that satisfies **ACID** principles.

## ■ E.g., transaction to transfer \$50 from account A to account B:

1. **read(A)**      *(기존의 값을)*
2. **A := A - 50**
3. **write(A)**
4. **read(B)**
5. **B := B + 50**
6. **write(B)**

6개  
작업  
가지  
한번에  
실행되어야

If the transaction fails after step 3 but before step 6, \$50 will be “lost” leading to an inconsistent database state

## ■ Two main issues to deal with:

- **Failures**: hardware failures and system crashes
- **Concurrent execution** of multiple transactions

# ACID Principles

## ■ **Atomicity.** All or Nothing.

- Either all operations of the transaction are properly reflected in the database or none are.

## ■ **Consistency.** Serializability

- Concurrent execution of multiple transactions preserves the consistency of the database.

## ■ **Isolation.**

- Each transaction must be unaware of other concurrently executing transactions.
- Intermediate transaction results must be hidden from other concurrently executed transactions.

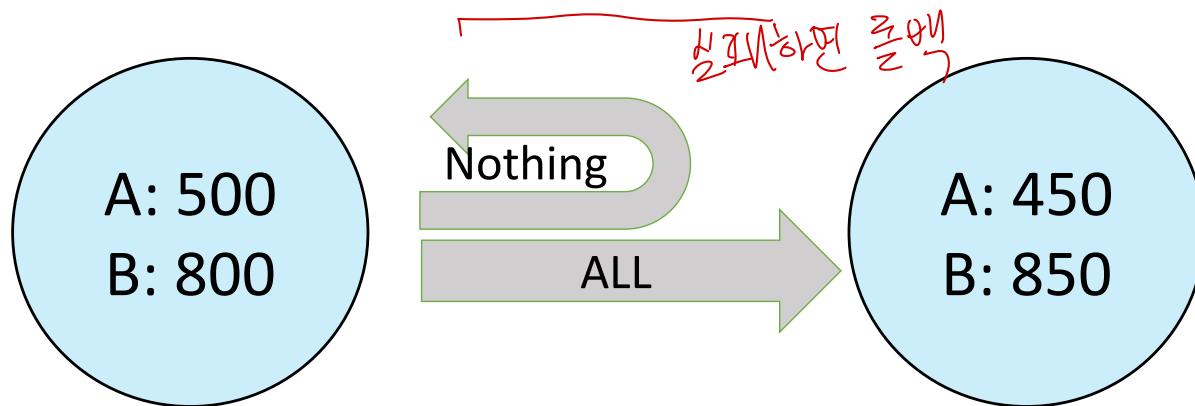
## ■ **Durability.**

- After a transaction completes successfully, the changes it has made to the database persist, even if there are system failures.

# Atomicity

## ■ Atomicity = ALL or NOTHING

- All operations within the transaction are completed successfully, or none are. If any of the operations fails, the entire transaction is rolled back.



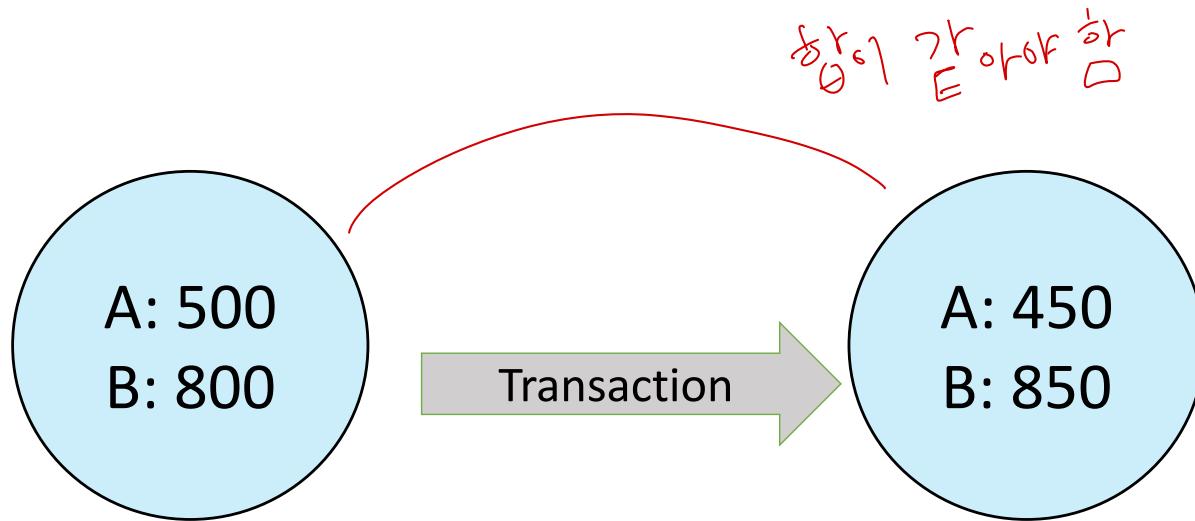
1. **read(A)**
2.  $A := A - 50$
3. **write(A)**
4. **read(B)**
5.  $B := B + 50$
6. **write(B)**

# Consistency

일관성

## ▪ Consistency requirement :

- The database moves from one valid state to another.



## ▪ consistency requirements include

- ① Explicit integrity constraints such as primary keys and foreign keys
- ② Implicit integrity constraints such as
- the sum of A and B is unchanged by the execution of the transaction

27/2  
2023

# Isolation

- ▶ **Isolation requirement:** Prevent any concurrent transaction  
from accessing the partially updated DB

다른 트랜잭션이 들어와도 올바른  
트랜잭션의 결과를 보호해야 함

T1	T2
read(A)	
A := A-50	
write (A)	
	read(A), read(B), print(A+B)
read(B)	
B := B+50	
write(B)	

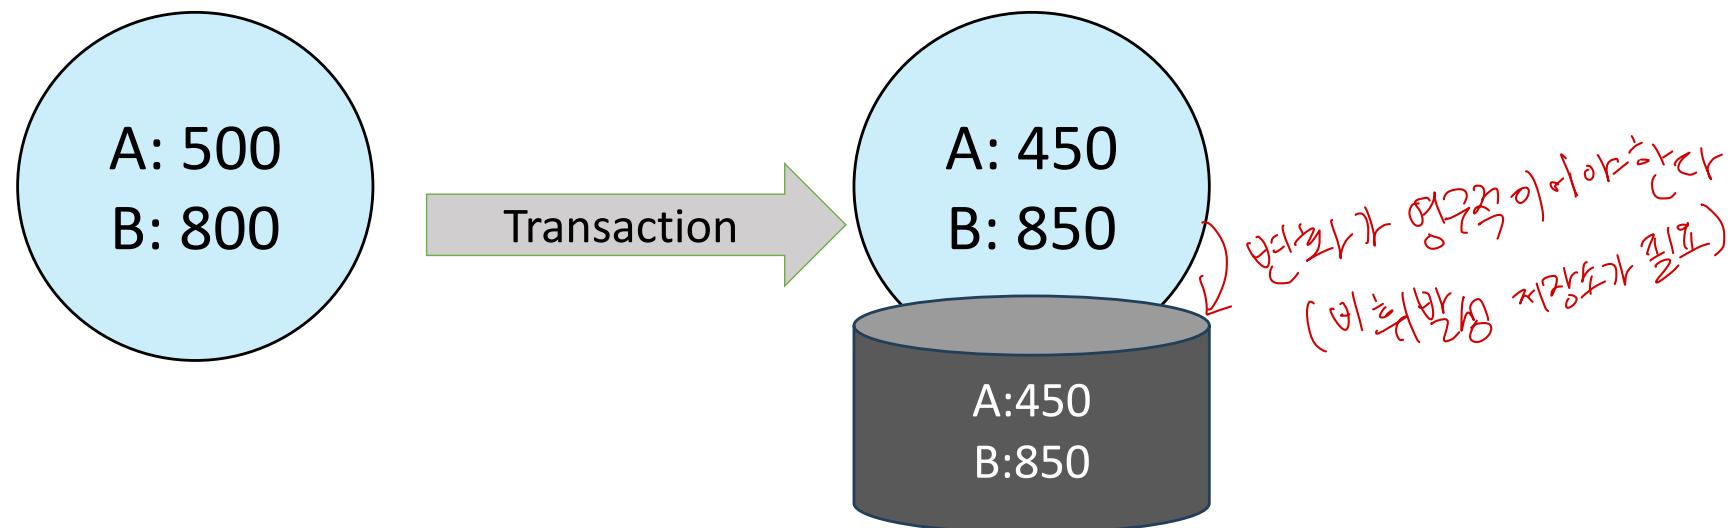
T2 will see an inconsistent DB  
(the sum  $A + B$  will be less than it should be).

- ▶ Isolation can be ensured trivially by running transactions **serially**, i.e., one after the other.
- ▶ However, executing multiple transactions sequentially has performance problems.

# Durability

## ▪ Durability requirement

- Once a transaction is committed, its changes are permanent, even in the case of a system failure.



# Isolation: Concurrent Executions

- Multiple transactions are allowed to run concurrently.

- Advantages

- Increased resource utilization → higher **throughput**
    - E.g., one transaction uses the CPU while another performs I/Os.
  - Reduced waiting time → low **response time** (latency)

## ▪ Concurrency control schemes

- control the interaction among the concurrent transactions

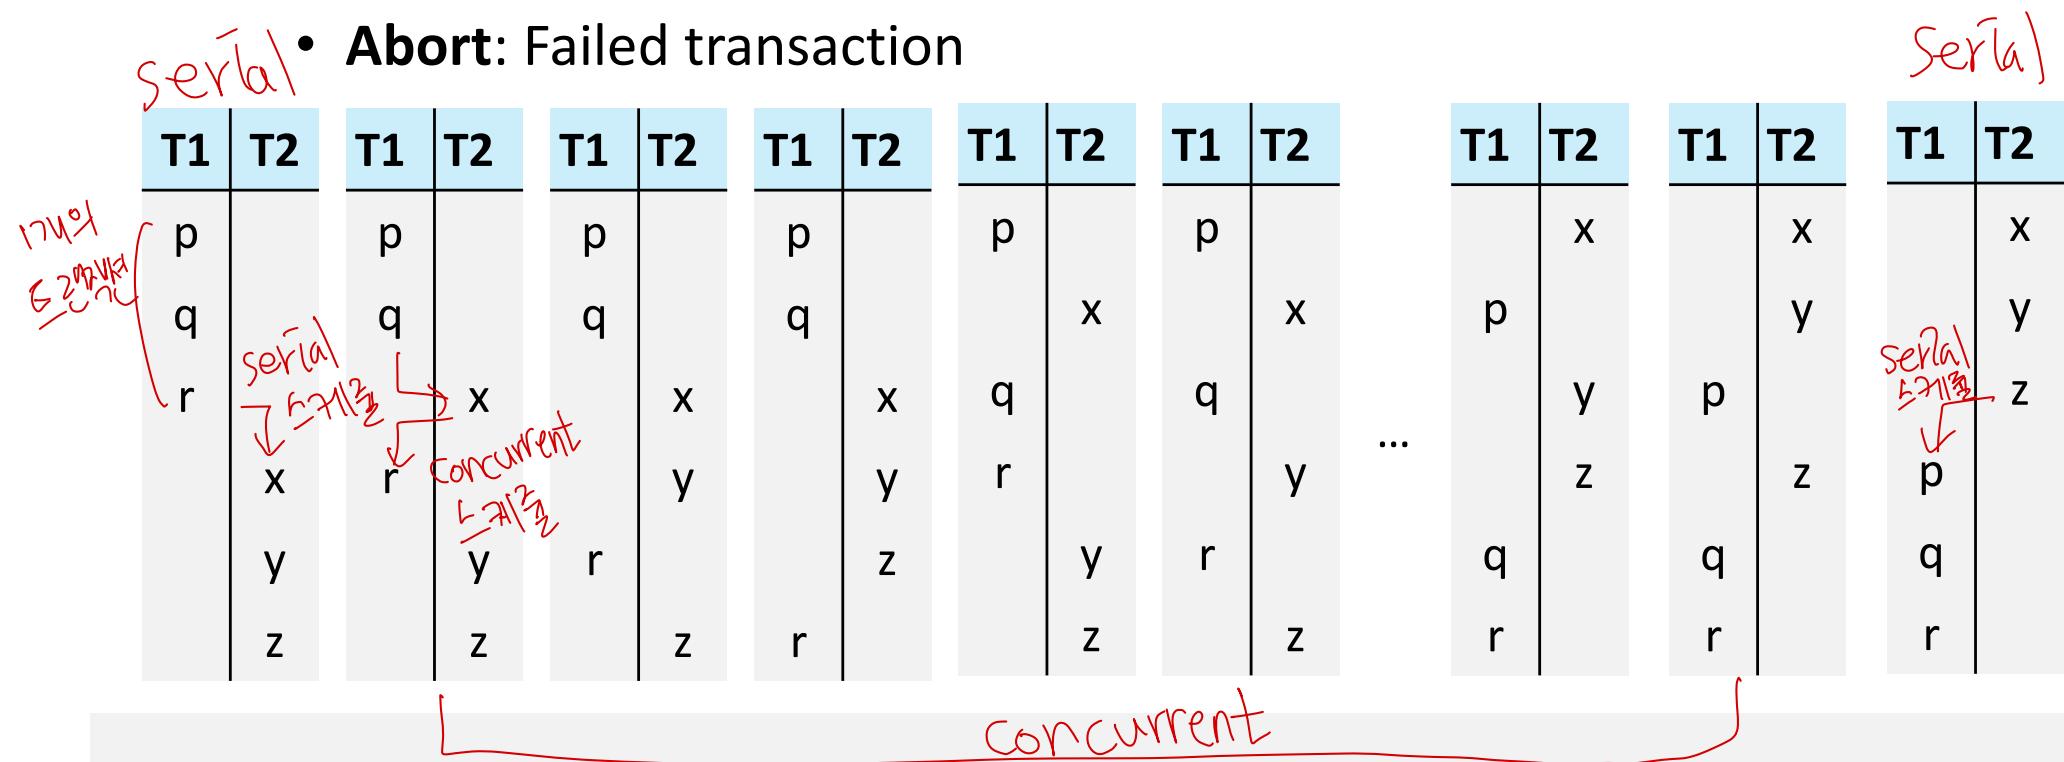
# Schedules

## ■ Schedule 스케줄링

- Interleave instructions of concurrent transactions
  - Must preserve the order of instructions that appear in each individual transaction.

- Last instruction of transaction

- **Commit:** Successful transaction
  - **Abort:** Failed transaction



# Serial Schedule 1

- Let  $T_1$  transfer \$50 from  $A$  to  $B$ , and  $T_2$  transfer 10% of the balance from  $A$  to  $B$ .
- A serial schedule in which  $T_1$  is followed by  $T_2$ :

$T_1$	$T_2$
$A: 400$ $B: 800$ ↓ $A: 350$ $B: 850$	read ( $A$ ) $A := A - 50$ write ( $A$ ) read ( $B$ ) $B := B + 50$ write ( $B$ ) commit

DBMS가 트랜잭션의 순서를  
결정 고려하지 않음  
각 순서가 충돌 가능성  
applicability 확인

↓  
 $A: 315$   
 $B: 885$   
 $\Rightarrow 1200$

## Serial Schedule 2

- A serial schedule where  $T_2$  is followed by  $T_1$

$T_1$	$T_2$
	read ( $A$ ) $A', 400$
	$temp := A * 0.1$ $B; 800$
	$A := A - temp$
	write ( $A$ ) $A', 360$
$A', 310$	read ( $B$ )
$B, 890$	$B := B + temp$ $B', 8400$
$\Rightarrow 1200$	write ( $B$ )
	commit
read ( $A$ )	
$A := A - 50$	
write ( $A$ )	
read ( $B$ )	
$B := B + 50$	
write ( $B$ )	
commit	

## Serializable Schedule 3

↑ serializable 결과가 serial처럼 보이게

- This is not a serial schedule, but it is *equivalent* to Schedule 1

$T_1$	$T_2$
$A: 400$ $B: 800$ read ( $A$ ) $A := A - 50$ write ( $A$ ) $350$  read ( $B$ ) $800$ $B := B + 50$ write ( $B$ ) $850$ commit	read ( $A$ ) $350$ $temp := A * 0.1$ $A := A - temp$ write ( $A$ ) $A: 315$  read ( $B$ ) $850$ $B := B + temp$ write ( $B$ ) $B: 885$ commit

the sum  $A + B$  is preserved.

↑  $1200$ 이卫  
↓  $2160$  Schedule 1

## Schedule 4 (Not Serializable)

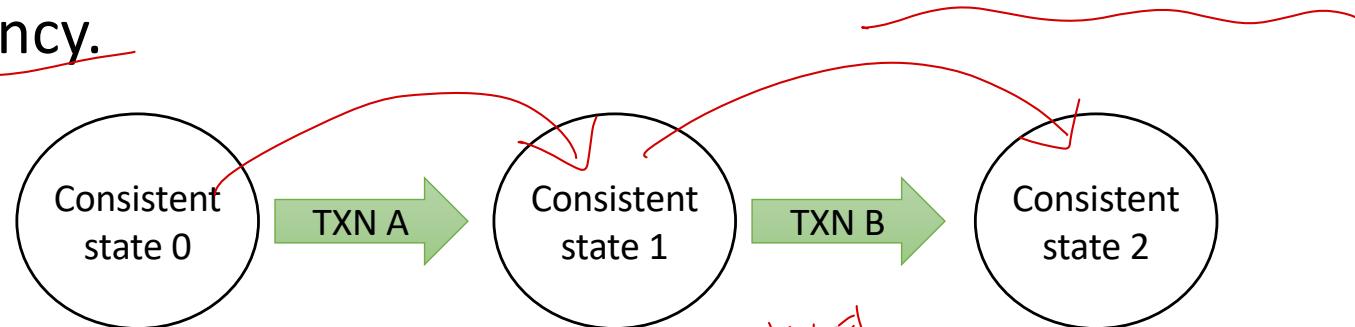
- The following schedule does not preserve the value of  $(A + B)$ .

$T_1$	$T_2$
$40^o$ read ( $A$ ) $350$ $A := A - 50$ <i>019210101, Disk 0121012101X</i>	$read(A) \quad 40^o$ $temp := A * 0.1 \quad 360$ $A := A - temp$ $write(A)$ $read(B) \quad 80^o$
$A: 350$ write ( $A$ ) $\rightarrow$ overwritten $800$ read ( $B$ ) $B := B + 50$ $850$ write ( $B$ ) commit	$B := B + temp \quad B: 840$ $write(B)$ $commit \quad B: 840$

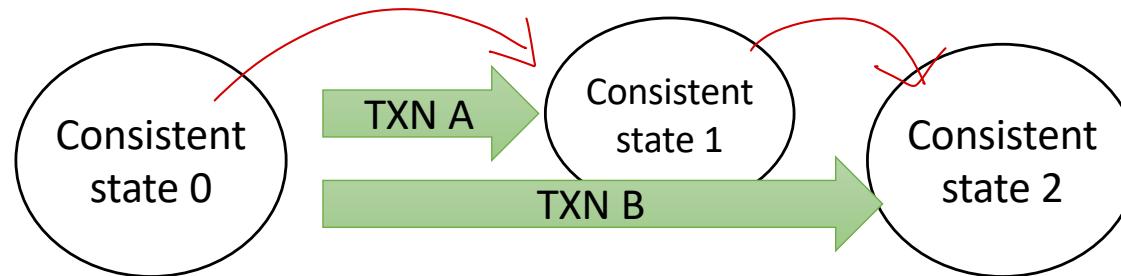
$1190 \neq 1200 \sim \text{non-Serializable} \rightarrow \text{try}$

# Serializability

- Serial execution of a set of transactions preserves database consistency.



- A schedule is **serializable** if it is equivalent to a serial schedule.



- Different forms of serializability:
  1. **Conflict serializability**
  2. **View serializability**

## *Simplified view of transactions*

- We ignore operations other than **read** and **write** instructions
- We assume that transactions **perform computations using local buffers** in between reads and writes.
- Modified data is **exposed** to other transactions **only when write is called.**

# Conflicting Instructions

- Instructions  $I_a$  and  $I_b$  of transactions  $T_a$  and  $T_b$  **conflict** if and only if at least one of the instructions write a shared data item.

읽는것은 문제 없음

- $I_a = \text{read}(Q)$ ,  $I_b = \text{read}(Q)$ .  $I_a$  and  $I_b$  don't conflict.
- $I_a = \text{read}(Q)$ ,  $I_b = \text{write}(Q)$ . They conflict. 같은게 old / new 달라짐
- $I_a = \text{write}(Q)$ ,  $I_b = \text{read}(Q)$ . They conflict
- $I_a = \text{write}(Q)$ ,  $I_b = \text{write}(Q)$ . They conflict

- If two contiguous *instructions* do not conflict, we can swap them
- If a schedule  $S$  can be transformed into a schedule  $S'$  by swapping non-conflicting instructions, we say that  $S$  and  $S'$  are **conflict equivalent**.

冲突이 발생하지 않도록 스왑하는방법

# Conflict Serializability

- Schedule  $S$  is **conflict serializable** if it is conflict equivalent to a serial schedule
  - Schedule 1 can be transformed into a serial schedule 2 by series of swaps of non-conflicting instructions.

$T_1$	$T_2$
read ( $A$ ) write ( $A$ )	read ( $A$ ) write ( $A$ )
↑ read ( $B$ ) write ( $B$ )	↓ read ( $B$ ) write ( $B$ )

**Conflict Serializable** Schedule 1

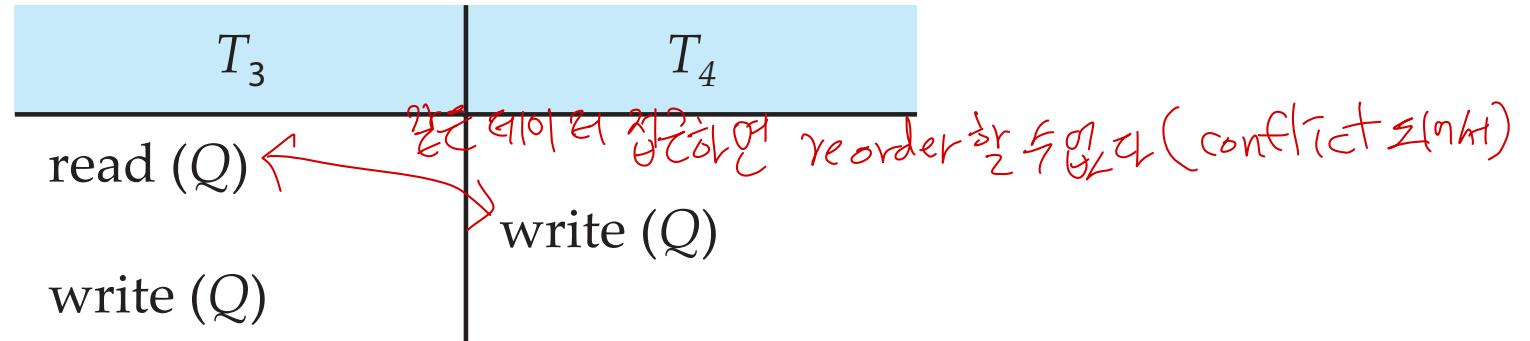
$T_1$	$T_2$
read ( $A$ ) write ( $A$ ) read ( $B$ ) write ( $B$ )	
	read ( $A$ ) write ( $A$ ) read ( $B$ ) write ( $B$ )

Serial Schedule 2

*↳ OWN  
Serializable  
Serializable  
→*

## Conflict Serializability (Cont.)

- Example of a schedule that is **not** conflict serializable:



- We are unable to swap instructions in the above schedule to obtain either the serial schedule  $\langle T_3, T_4 \rangle$ , or the serial schedule  $\langle T_4, T_3 \rangle$ .

# View Serializability

- Two schedules  $S$  and  $S'$  with the same set of transactions are **view equivalent** if the following three conditions hold

같은 3가지 조건을 만족하면

## 1. Same initial Reads

- If  $T_a$  reads the initial value of a data item  $Q$  in  $S$ , it must do the same in  $S'$ .

초기 Read( $Q$ ) 가 같은 드레인 같음  
2개의 스케줄 같음

## 2. Same Read-From

- If  $T_a$  reads  $Q$  written by transaction  $T_b$  in  $S$ ,  $T_i$  must read the same value of  $Q$  written by  $T_b$  in  $S'$ . ( $W(Q) \rightarrow R(Q)$ )

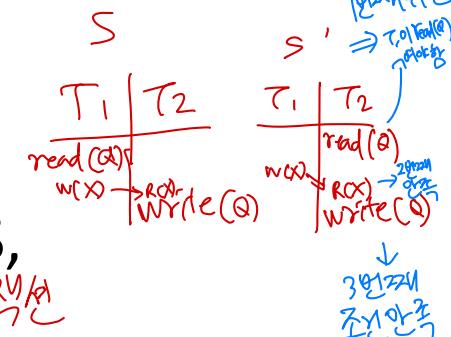
2개의 허브 동작  
 $(W \rightarrow R)$ 의 동작  
같아야 함

## 3. Same Final Write

- Final write( $Q$ ) must be done by the same transaction in both  $S$  and  $S'$ .

초기 Write( $Q$ ) 가 같은 드레인  
2개의 스케줄 같음

- $S$  is **view serializable** if it is view equivalent to a serial schedule.
- Every conflict serializable schedule is also view serializable.



# Q1: Is this a view serializable schedule?

T1	T2	T3
$R(Q)$ $W(Q)$	$W(Q)$ $R(Q)$	$R(Q)$ $W(Q)$

Serial Schedule

T1	T2	T3
$R(Q)$		
	$W(Q)$	
	$R(Q)$	
		$R(Q)$
		$W(Q)$

View-Serializable?

- In both schedules, T1 reads the initial value of Q
- In both schedules, T3 performs the final write

## Q2: Is this a view serializable schedule?

T1	T2	T3
R(Q) W(Q)		

Serial Schedule

The diagram shows a serial schedule for three transactions (T1, T2, T3) over three time steps. In the first step, T1 performs a read (R(Q)) and a write (W(Q)). In the second step, T2 performs a write (W(Q)) and a read (R(Q)). In the third step, T3 performs a read (R(Q)) and a write (W(Q)). Red arrows indicate the sequence of operations: from T1's write to T2's read, and from T2's write to T3's read.

T1	T2	T3
R(Q) W(Q)		

View-Serializable?

The diagram shows a schedule for view-serializability. It includes the same transaction operations as the serial schedule: T1 reads R(Q) and writes W(Q), T2 reads R(Q) and writes W(Q), and T3 reads R(Q) and writes W(Q). Handwritten red annotations include 'No! 2 번째 원칙' (No! Second rule) pointing to T2's write, and '1,3 번째 원칙' (1,3rd rule) pointing to T3's write.

- In both schedules, T1 reads the initial value of Q
- In both schedules, T3 performs the final write

### Q3: Is this a view serializable schedule?

T1	T2	T3
R(Q)		
W(Q)		
	W(Q)	
		W(Q)

Serial Schedule

- ① True
- ② W→R : Read  $\neq$  Write → True
- ③ True

View conflict

Read conflict

→ DB의 일관성이 있는지

T1	T2	T3
R(Q)		
	W(Q)	
	W(Q)	
		W(Q)

View-Serializable? Yes

But, conflict Serializable? No

why? conflict가 발생

(R-W, W-R, W-W)

reorder 발생

( R-R 같은 경우 )

오류 발생

- Above is a schedule which is view-serializable but *not* conflict serializable.
- Every view serializable schedule (that is not conflict serializable) has blind writes.

# Blind Write

- **Blind Write** occurs when a transaction **writes a value to a data item (without reading it)**.

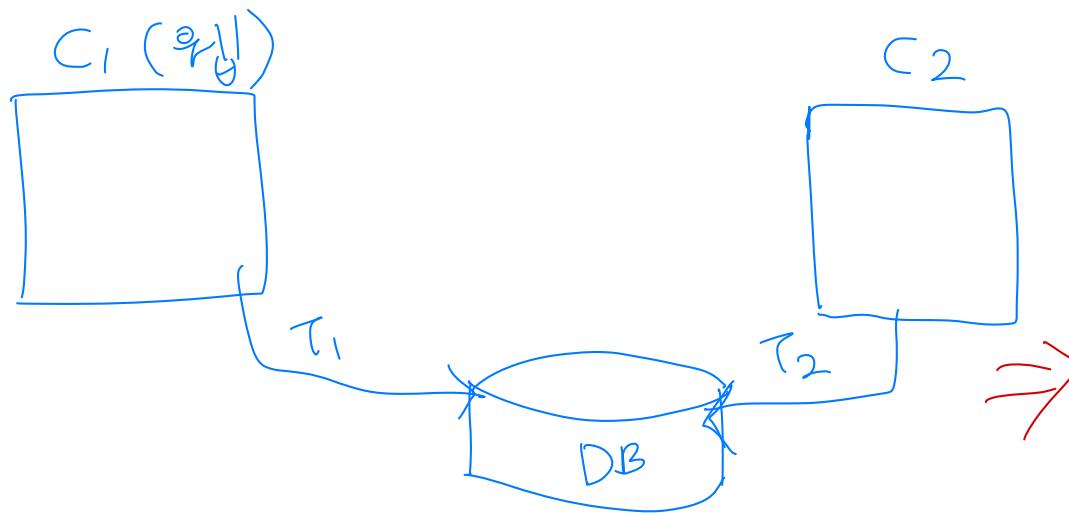
- E.g.

UPDATE customers

이 아이템을  
읽지 않고 저장

SET reward\_points = 0

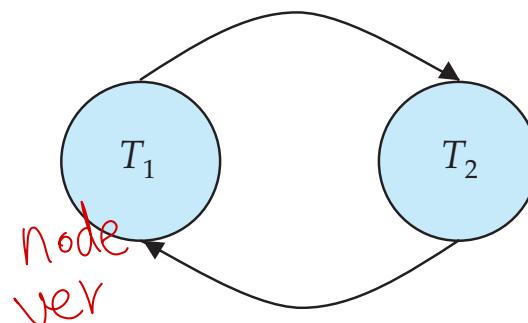
WHERE customer\_id = 123;



시간상으로  $T_1$ 을 먼저 실행 (도보다)  
하지만 네트워크상의 문제로  $T_2$ 보다 더 이  
는게 도착함  $\Rightarrow$  그대로 DBMS는  
신경쓰지 않음 (순서가 중요하지만,  
업데이트 차수가 충돌하지 않는다)

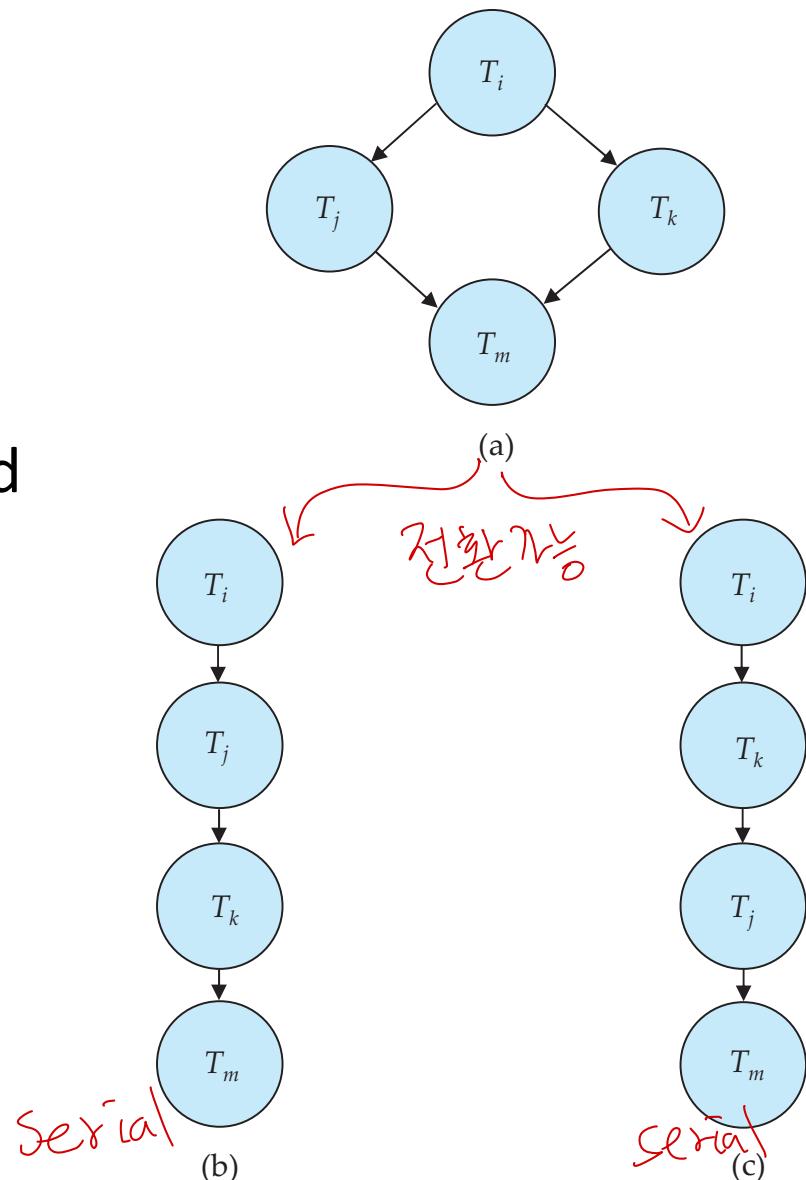
# Testing for Serializability

- Consider some schedule of a set of transactions  $T_1, T_2, \dots, T_n$
- **Precedence graph** — a direct graph where the vertices are the transactions (names).
- We draw an arc from  $T_a$  to  $T_b$  if the two transactions conflict and  $T_a$  accessed the conflicting data item earlier than  $T_b$ .
- We may label the arc by the item that was accessed.
- Example of a precedence graph



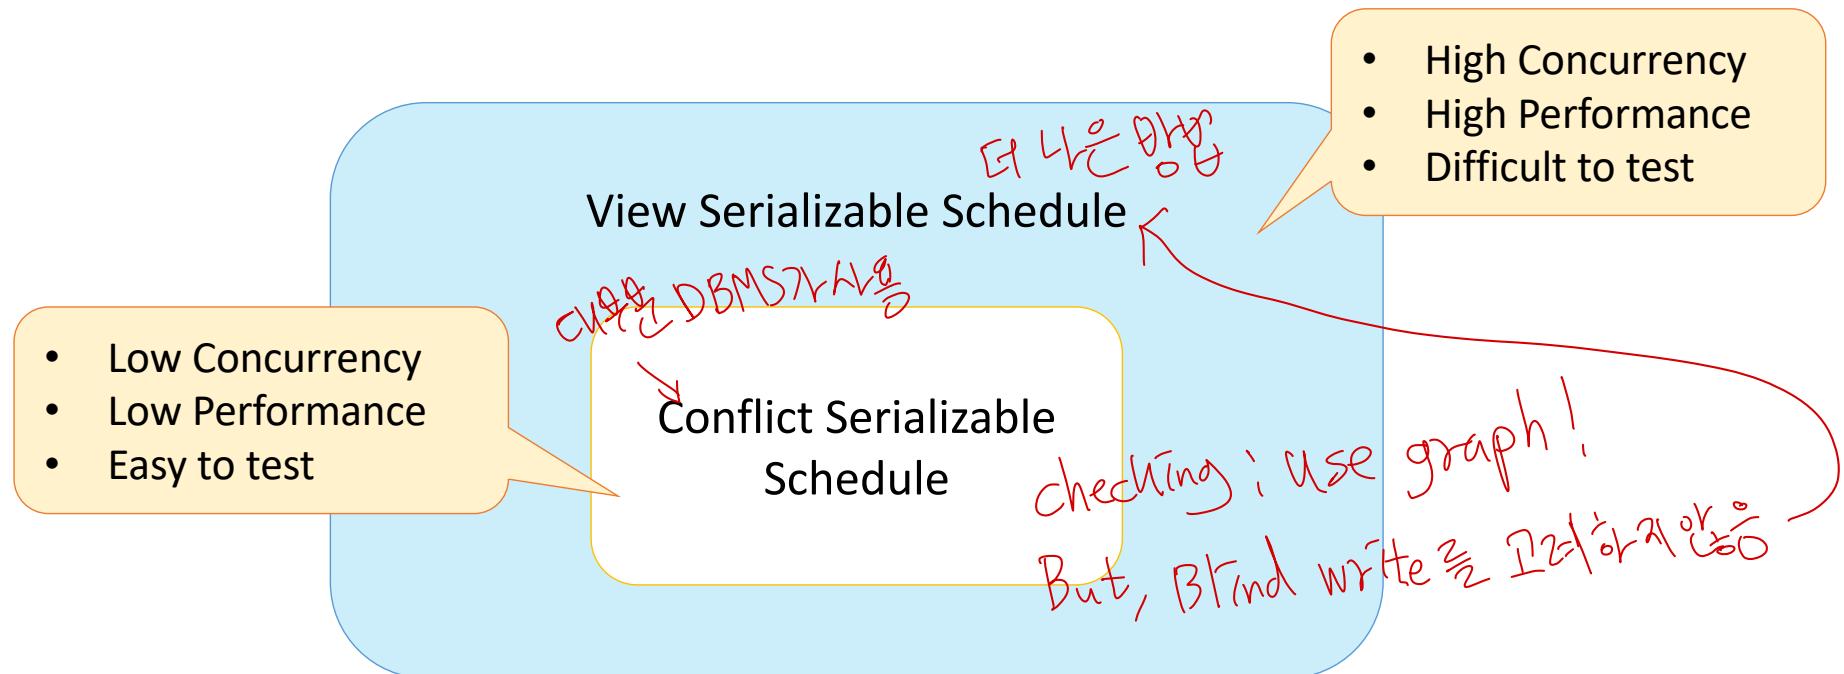
# Test for Conflict Serializability

- A schedule is **conflict serializable** if and only if its precedence graph is **acyclic**.  
*=not cyclic*
- Cycle-detection algorithms exist
- If precedence graph is acyclic, the **serializability order** can be obtained by a **topological sorting** of the graph.
  - This is a linear order consistent with the partial order of the graph.



# Conflict Serializability vs. View Serializability

- Conflict serializability is more popular in DBMS in practice
- But, conflict serializability actually permit only a subset of serializable schedules.
- The view serializability is expensive (*NP-complete*) to test.
  - Practical algorithms that just check some **sufficient conditions**



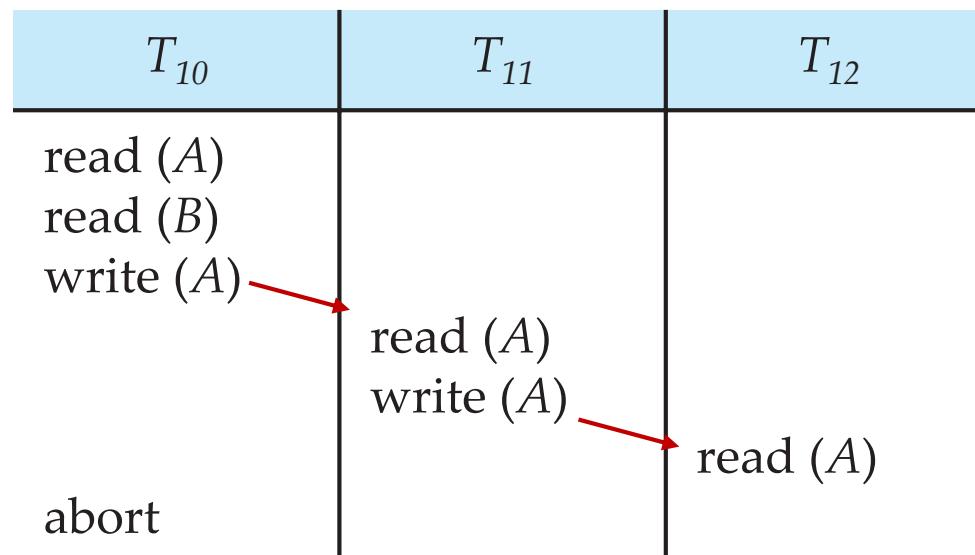
# Recoverable Schedules

- Schedules must be recoverable
- Recoverable schedule** — if a transaction  $T_b$  reads a data written by a transaction  $T_a$ , then  $T_a$  must commit before  $T_b$ .
- The following schedule is not recoverable

$T_8$	$T_9$
read ( $A$ ) write ( $A$ )	read ( $A$ ) commit
read ( $B$ ) <b>abort</b>	<p>↑ 연결된 <math>A</math>를 client가 읽음 But, <math>A</math>를 연결한 <math>T_8</math>이 abort(폐지)됨 ⇒ Not recoverable</p> <p>This is called <b>"Dirty Read"</b></p> <p>commit하지 않은 트랜잭션의 쓴 데이터 읽은 경우</p>

# Cascading Rollbacks

- **Cascading rollback** – a transaction failure may lead to a series of transaction rollbacks.
- Consider the following schedule where none of the transactions has yet committed



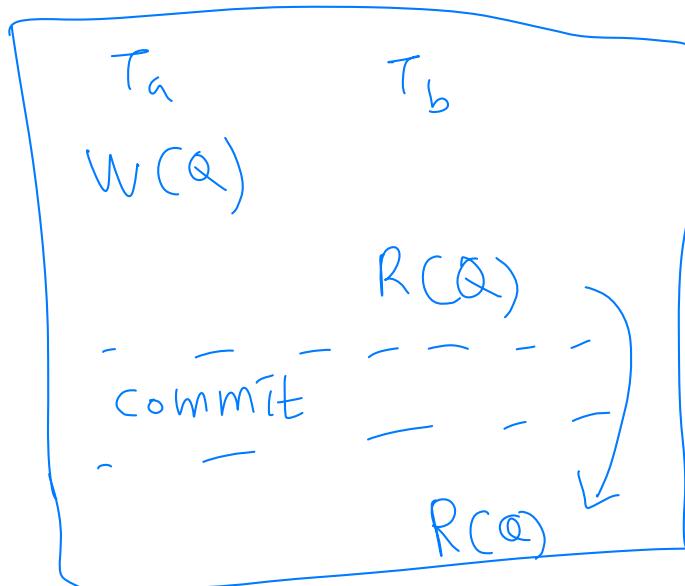
This schedule is recoverable.

If  $T_{10}$  fails,  $T_{11}$  and  $T_{12}$  must also be rolled back.

- Can lead to the undoing of a significant amount of work (overhead ↑↑)

# Cascadeless Schedules

- Cascadeless schedules — prevents cascading rollbacks;
  - If  $T_b$  reads a data written by  $T_a$ , the commit operation of  $T_a$  appears before the read operation of  $T_b$ .
  - Read committed data only!
- Every Cascadeless schedule is also recoverable
- It is desirable to restrict the schedules to those that are cascadeless



concurrent 트랜잭션을 제한할 때  
부여해야 하는  
(자유롭게 제한해야 하기)  
concurrency 를 해친다  
→ 성능 ↓

여기로 이동 시키면  
 $T_a$ 가  
( $T_a$ 가 commit한 후에  $R(Q)$ 를  
한다는 규칙을 정하면 충분히  
막을 수 있다)  
= Cascadeless 스케줄링

# Concurrency Control

- Goal: A DBMS must make scheduling decisions that are
  - either conflict or view serializable, and *consistency 관점*
  - recoverable and preferably cascadeless *→ performance 관점*
    - ↳ consistency 관점
    - 성능 관점으로 병렬 실행하기 힘들
- DBMS scheduler should provide a high degree of concurrency
  - We study concurrency control protocols in Chapter 18.

# ~~Weak Levels of Consistency~~

~~Tradeoff of consistency~~

- Some applications live with weak levels of consistency
  - Tradeoff accuracy for performance

- **On Line Transaction Processing –OLTP**

- Transactions need to be serializable
    - Short simple transactions
    - Transactions access only a small fraction of the database

- **On Line Analytic Processing –OLAP**

- Transactions do not need to be serializable
  - Uses database to guide strategic decisions.
    - E.g., A CEO wants to get an approximate total balance of all customers' accounts
    - Transactions access a large fraction of the database

# Phenomena caused by Concurrent Transactions

Inconsistency 일치

## ■ dirty read

- A transaction reads data written by an uncommitted transaction.

Read(Q) → 100  
read(A) → 200  
이후  
결과 같지 다른현상

## ■ nonrepeatable read

- A transaction runs the same point query twice and finds that a data has been modified by another transaction.

## ■ phantom read

- A transaction runs the same range query twice and finds that the result set has new records due to another recently-committed transaction.

## ■ serialization anomaly

serialize 할 수 없는 경우

- The result of successfully committing a group of transactions is inconsistent with all possible orderings of running those transactions one at a time.

# Phantom Phenomenon

- E.g., Transaction 1:

```
select ID, name from instructor where salary > 90000
```

- E.g., Transaction 2:

```
insert into instructor values ('11111', 'James', 'Marketing', 100000)
```

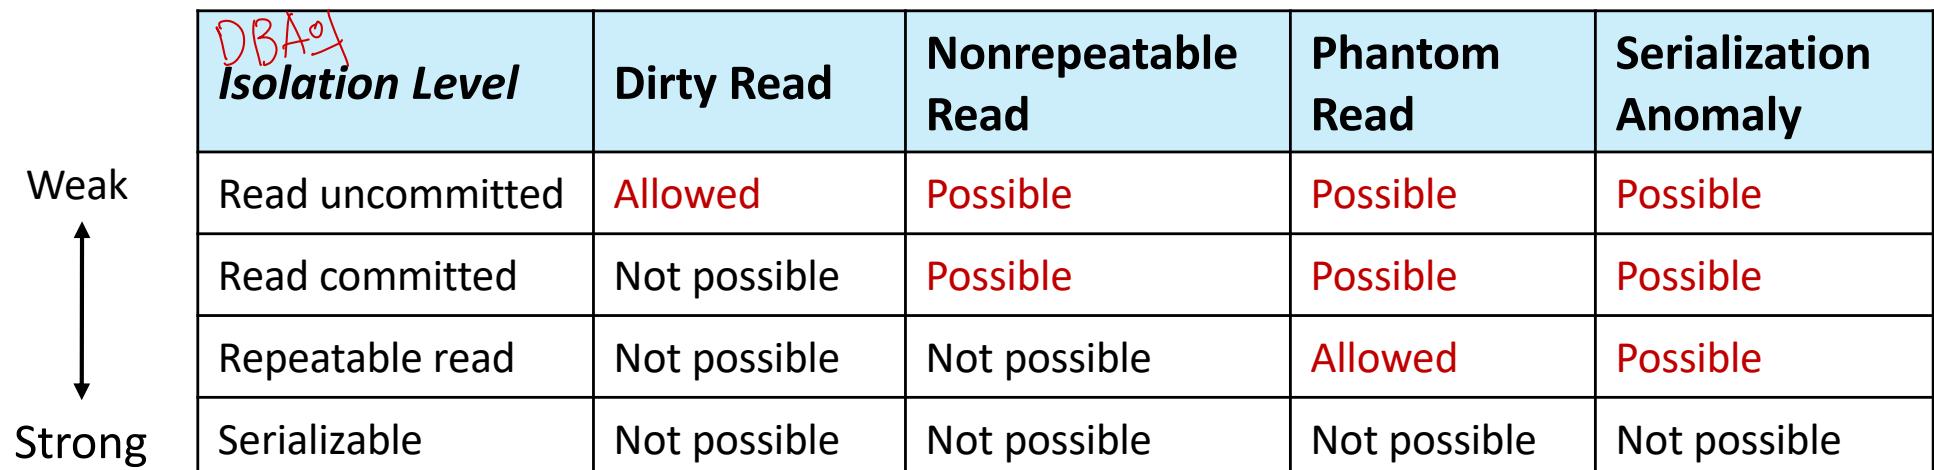
- Suppose

- T1 starts, finds tuples salary > 90000 using index and locks them
- And then T2 executes.
- Do T1 and T2 conflict? Does tuple level locking detect the conflict?
- Instance of the **phantom phenomenon**

조건에 맞는 데이터  
↓ 삽입 했지만  
T1은 이미 데이터를  
지정하지 않았다  
T2는 새로운  
데이터로 바뀐  
⇒ 패턴 현상

# Transaction Isolation Level

- **Serializable** — default in SQL-92, **SQLite**  
*DBA는 Serializable는 강한 일관성을 제공하는, 상당히 강한 weak level의 consistency를 제공합니다.*
- **Repeatable read** — default in MySQL
  - Repeated reads of same record must return same value.
  - However, a transaction may not be serializable – it may find some records inserted by a transaction but not find others.
- **Read committed** — default in PostgreSQL & Oracle
  - Successive reads may return different (but committed) data.
- **Read uncommitted** — even uncommitted records may be read.



<i>DBA의 Isolation Level</i>	Dirty Read	Nonrepeatable Read	Phantom Read	Serialization Anomaly
Read uncommitted	Allowed	Possible	Possible	Possible
Read committed	Not possible	Possible	Possible	Possible
Repeatable read	Not possible	Not possible	Allowed	Possible
Serializable	Not possible	Not possible	Not possible	Not possible

# Repeatable Read vs. Read Committed

T1	T2	T3
TS_start	TS_start	
	$W(Q \leftarrow 1)$	
	TS_commit	
$1 \leftarrow R(Q):$		
$1 \leftarrow R(Q):$	$\begin{array}{l} \text{H1} \leftarrow \text{2} \\ \text{H1} \leftarrow \text{3} \end{array}$	
$1 \leftarrow R(Q):$	$\begin{array}{l} \text{H1} \leftarrow \text{2} \\ \text{H1} \leftarrow \text{3} \\ \text{H2} \leftarrow \text{1} \end{array}$	
$1 \leftarrow R(Q):$	$\begin{array}{l} \text{H1} \leftarrow \text{2} \\ \text{H1} \leftarrow \text{3} \\ \text{H2} \leftarrow \text{1} \\ (\text{Repeatable} \\ \text{Read}) \end{array}$	
TS_commit		TS_start
		$W(Q \leftarrow 2)$
		TS_commit

Repeatable Read

T1	T2	T3
TS_start	TS_start	
	$W(Q \leftarrow 1)$	
	TS_commit	
$1 \leftarrow R(Q):$		
$1 \leftarrow R(Q):$	$\begin{array}{l} \text{H1} \leftarrow \text{2} \\ \text{H1} \leftarrow \text{3} \end{array}$	
$1 \leftarrow R(Q):$	$\begin{array}{l} \text{H1} \leftarrow \text{2} \\ \text{H1} \leftarrow \text{3} \\ \text{H2} \leftarrow \text{1} \end{array}$	
$2 \leftarrow R(Q):$	$\begin{array}{l} \text{H1} \leftarrow \text{2} \\ \text{H1} \leftarrow \text{3} \\ \text{H2} \leftarrow \text{1} \\ \text{Commit}(T2) \end{array}$	
TS_commit		TS_start
		$W(Q \leftarrow 2)$
		TS_commit

Read Committed

Phantom Read!!!  
 (a.k.a. Non-Repeatable Read)  
 This problem happens in PostgreSQL & Oracle

↑ Default

# Phantom Phenomenon in Repeatable Read

T1	T2
<p>TS_start</p> <p>select name from instructor where ID='1'; → empty table</p> <p>select name from instructor where ID='1'; → empty table // GOOD (Repeatable Read)</p> <p>update instructor set salary = 2000 where ID = '1';</p> <p>select name from instructor where ID='1'; → 'James' Phantom Read here!!!!</p> <p>TS_commit</p>	<p>TS_start</p> <p>insert into instructor values ('1', 'James', 'CS', 1000);</p> <p>TS_commit</p> <p><i>T<sub>2</sub>) write or T<sub>1</sub>) write to R<sub>1</sub>) Read<sub>2</sub>) 2nd to R<sub>2</sub>) Cr.</i></p>

This anomaly happens in MySQL

# Transaction Definition in SQL

- A transaction in SQL ends by:
    - **Commit work** commits current transaction
    - **Rollback work** causes current transaction to abort.
  - **Auto Commit**
    - Each SQL statement commits implicitly
    - Implicit commit can be turned off by a database directive
      - E.g., in JDBC -- `connection.setAutoCommit(false);`
  - Isolation level can be set at database level
  - Isolation level can be changed at start of transaction
    - E.g. In SQL
      - SET TRANSACTION ISOLATION LEVEL SERIALIZABLE
    - E.g. in JDBC - `connection.setTransactionIsolation(Connection.TRANSACTION_SERIALIZABLE)`
- auto commit*
- consistency level을 설정 가능*
- 事务  
提交  
回滚*

# Implementation of Isolation Levels

DBMS7Y  
serializability  
부정합법성의 증명

## ■ Locking

- Lock on whole database vs lock on items
- How long to hold lock?
- Shared vs exclusive locks

## ■ Timestamps

- Transaction timestamp assigned e.g. when a transaction begins
- Data items store two timestamps
  - Read timestamp
  - Write timestamp
- Timestamps are used to detect out of order accesses

## ■ Multiple versions of each data item

- Allow transactions to read from a “snapshot” of the database