

Database Systems

Lecture18 – Chapter 18: Concurrency Control (Part I: Lock-based Concurrency Control)

Beomseok Nam (남범석)

bnam@skku.edu

Last Class

2023-7-8

- Serializable schedule in any DBMS
 - “Conflict serializable” schedule
- Swapping method
 - Swap non-conflicting operations
- Dependency graph
 - Topological sorting
- *real system*
 - ✓ Only if **the entire schedule is given**, we can determine whether the schedule is serializable or not.

Scheduling at runtime

- Q: What if we do not know the entire schedule ahead of time?
- We need a way to guarantee that all execution schedules are serializable.
- Solution:
 - Use **locks**
 - Use **timestamps**
 - Use **versions**

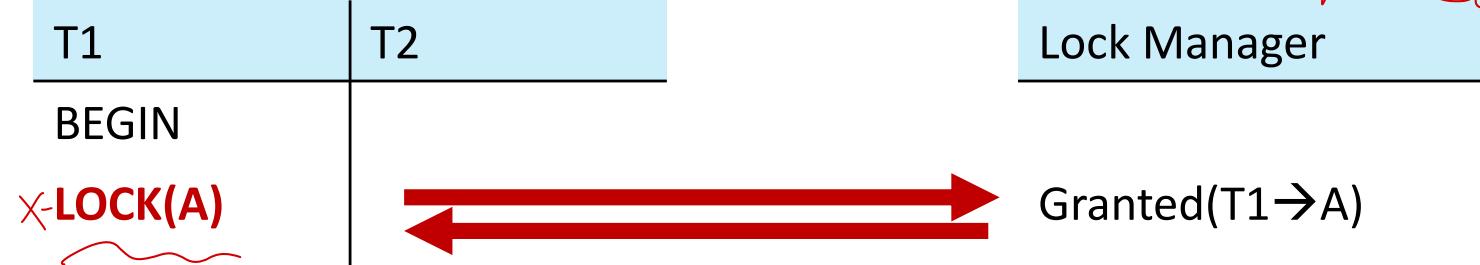
Lock Types

- S-LOCK: Shared locks for reads (=read lock)
- X-LOCK: Exclusive locks for both reads and writes
- Lock-compatibility matrix (锁兼容性矩阵)

	S	X
S	true	false
X	false	false

- Transactions request or release locks.
- Lock manager grants or blocks requests.
- Lock manager updates its internal lock-table.
- → It keeps track of what transactions hold what locks and what transactions are waiting to acquire any locks.

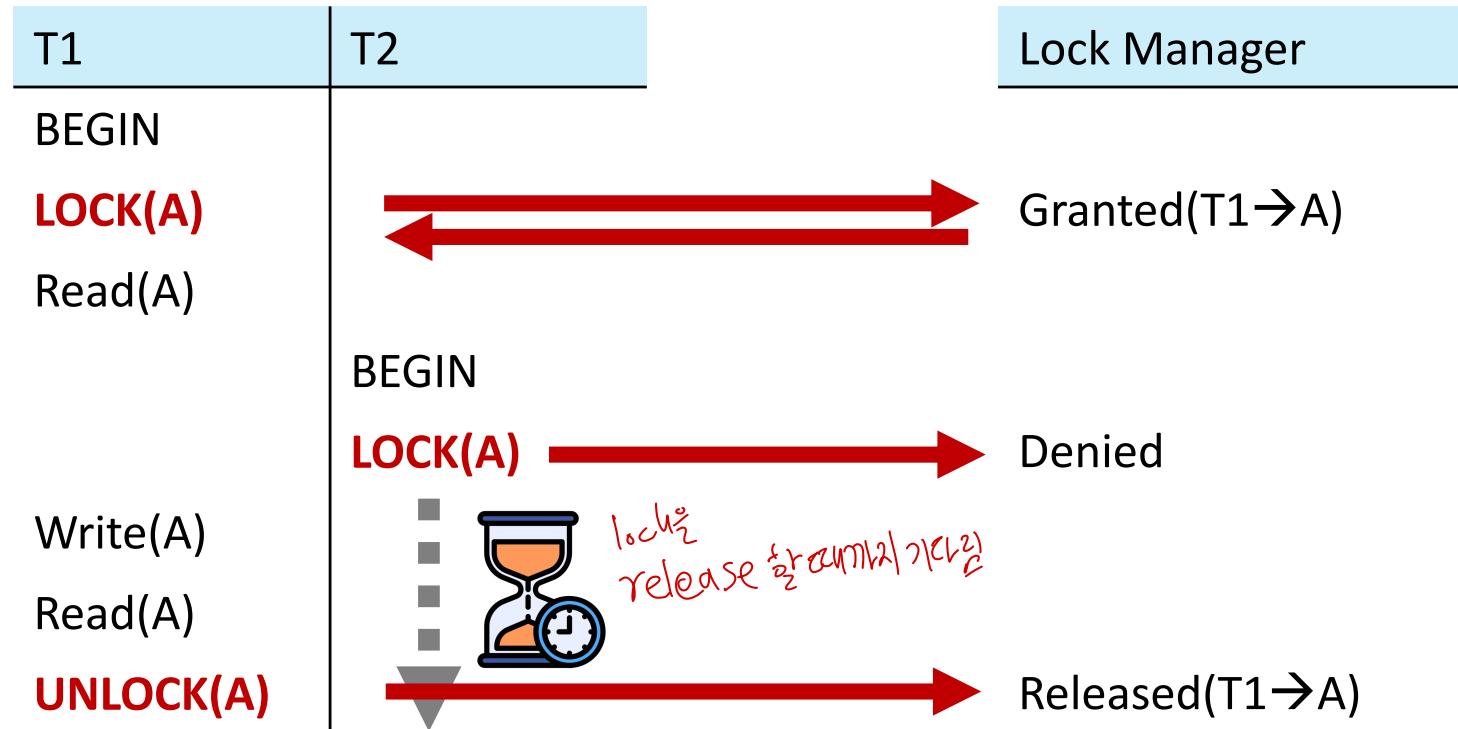
Executing with Locks



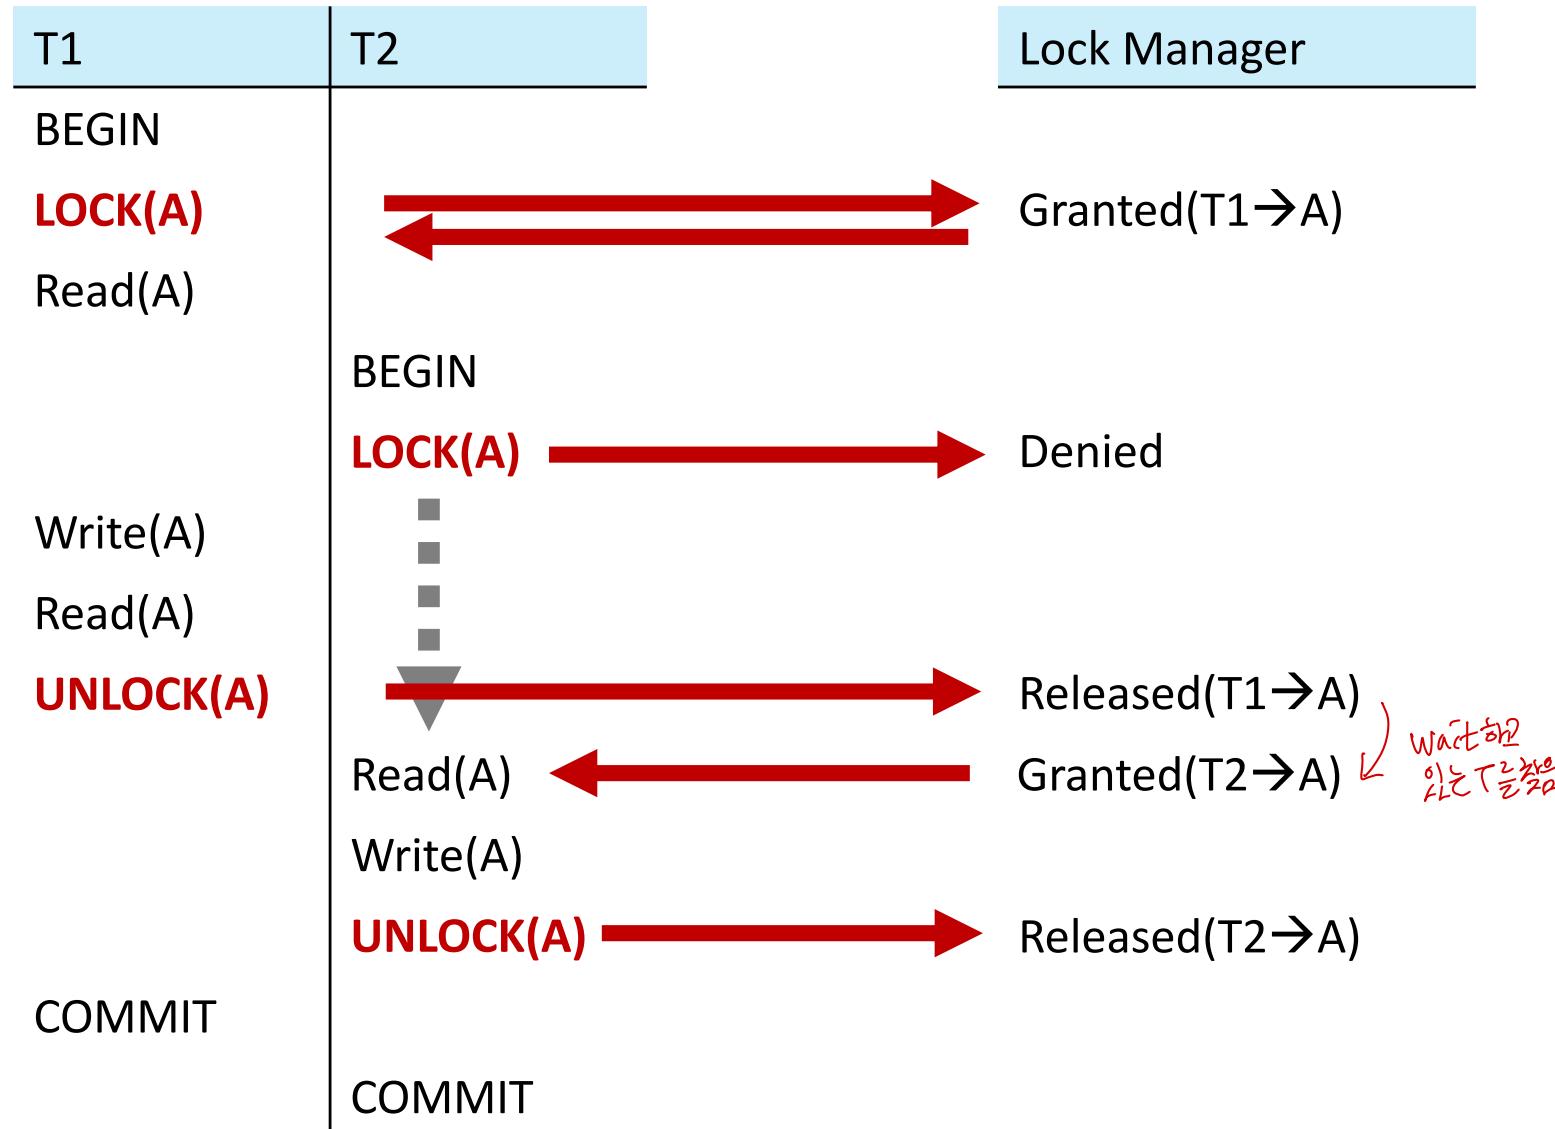
Executing with Locks



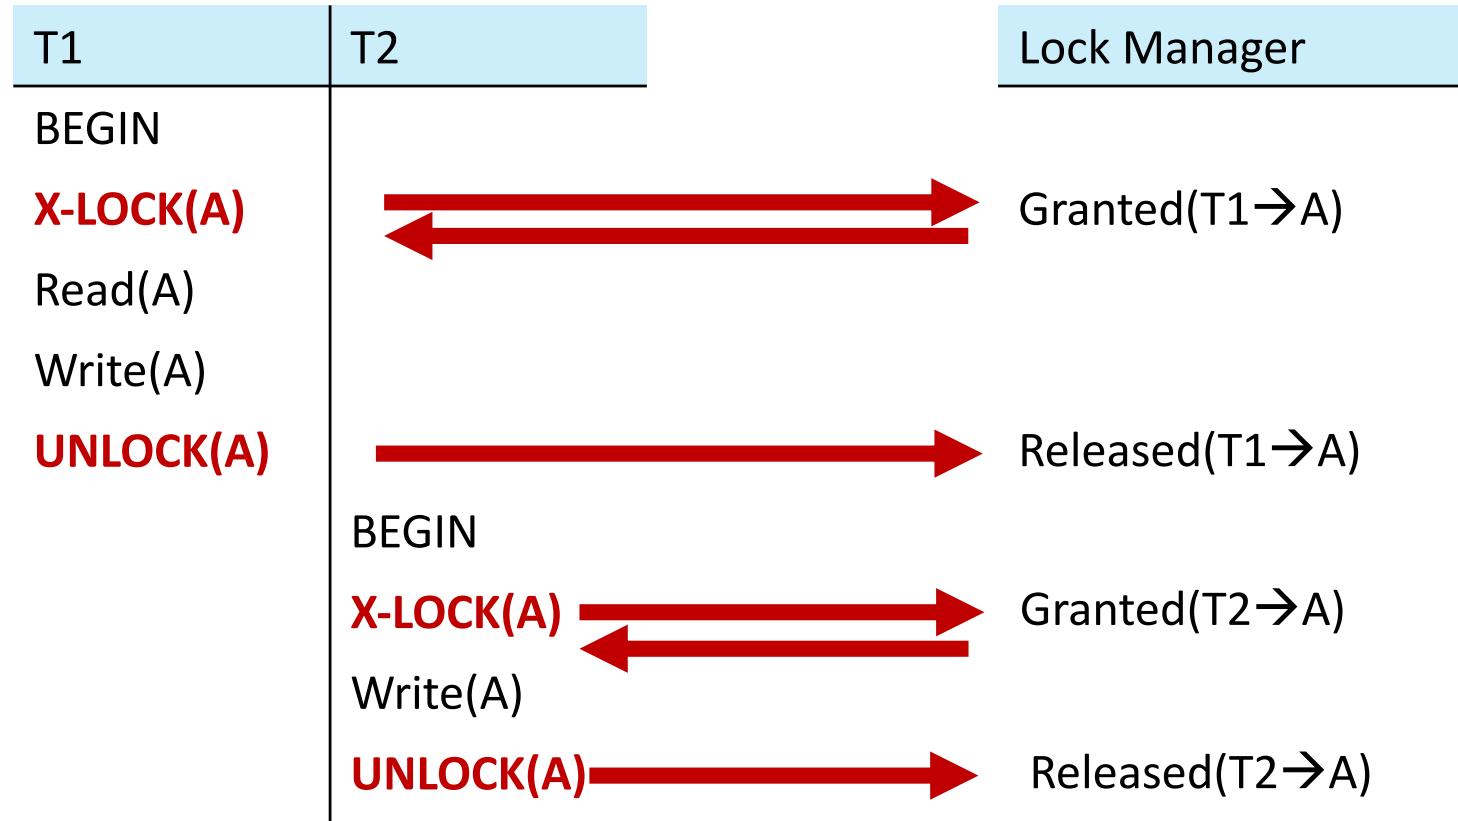
Executing with Locks



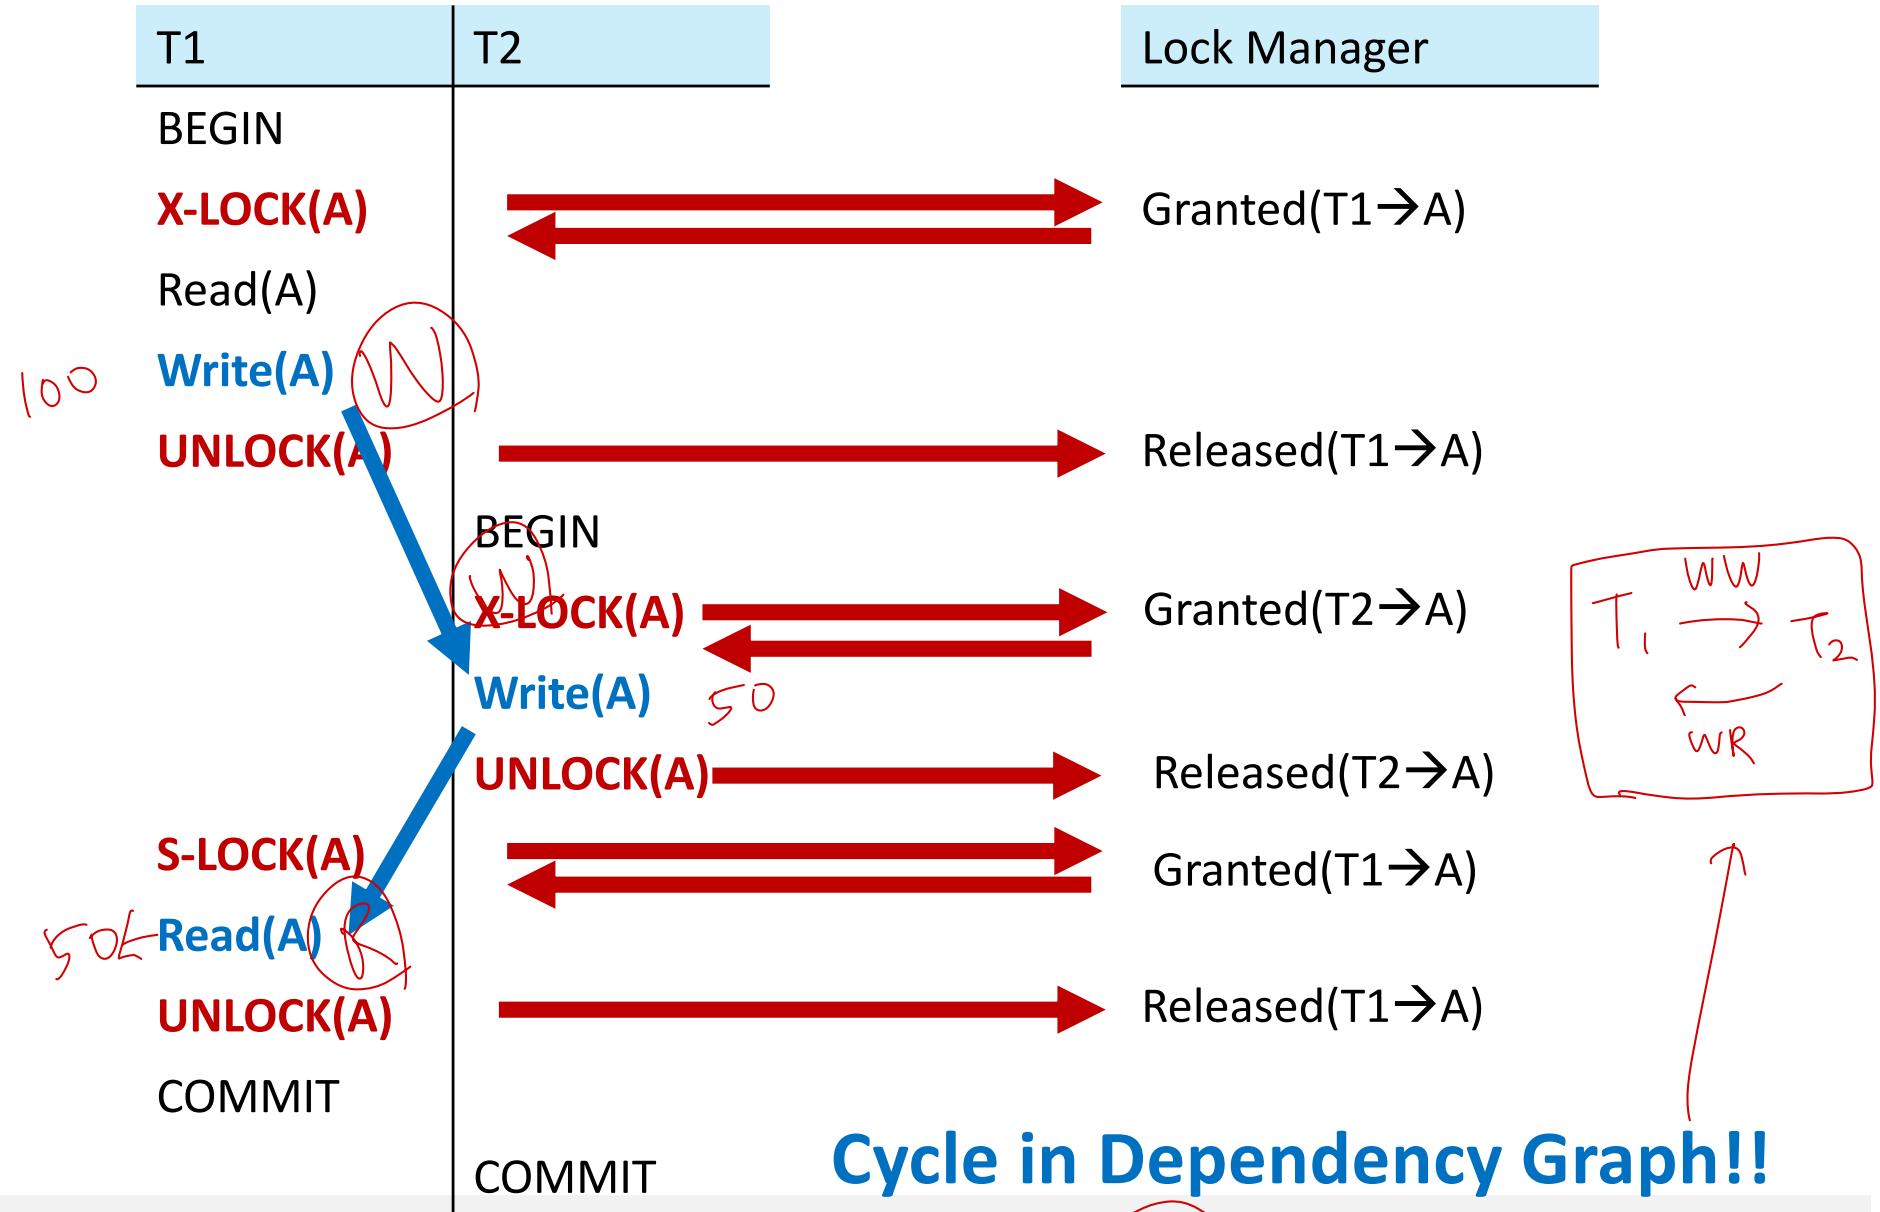
Executing with Locks



What if releasing locks early?



What if releasing locks early?



What if releasing locks early? – Another Example

T1	T2
X-LOCK(A)	
Read(A)	
900 A=A-100	S-LOCK(A) 
Write(A)	↓ Read(A) 900  UNLOCK(A)
UNLOCK(A)	S-LOCK(B)
X-LOCK(B)	
Wait ↓	
Read(B) 	
1100 B=B+100	
Write(B)	
UNLOCK(B)	
1100 COMMIT	Print A+B → 1900  
	Commit

Initial Database

A=1000, B=1000

T2 Output

A+B=1900

Locking alone is not sufficient!!

Lock만으로는 충분하지 않다

Two-Phase Locking Protocol

Two-Phase Locking (2PL) Protocol

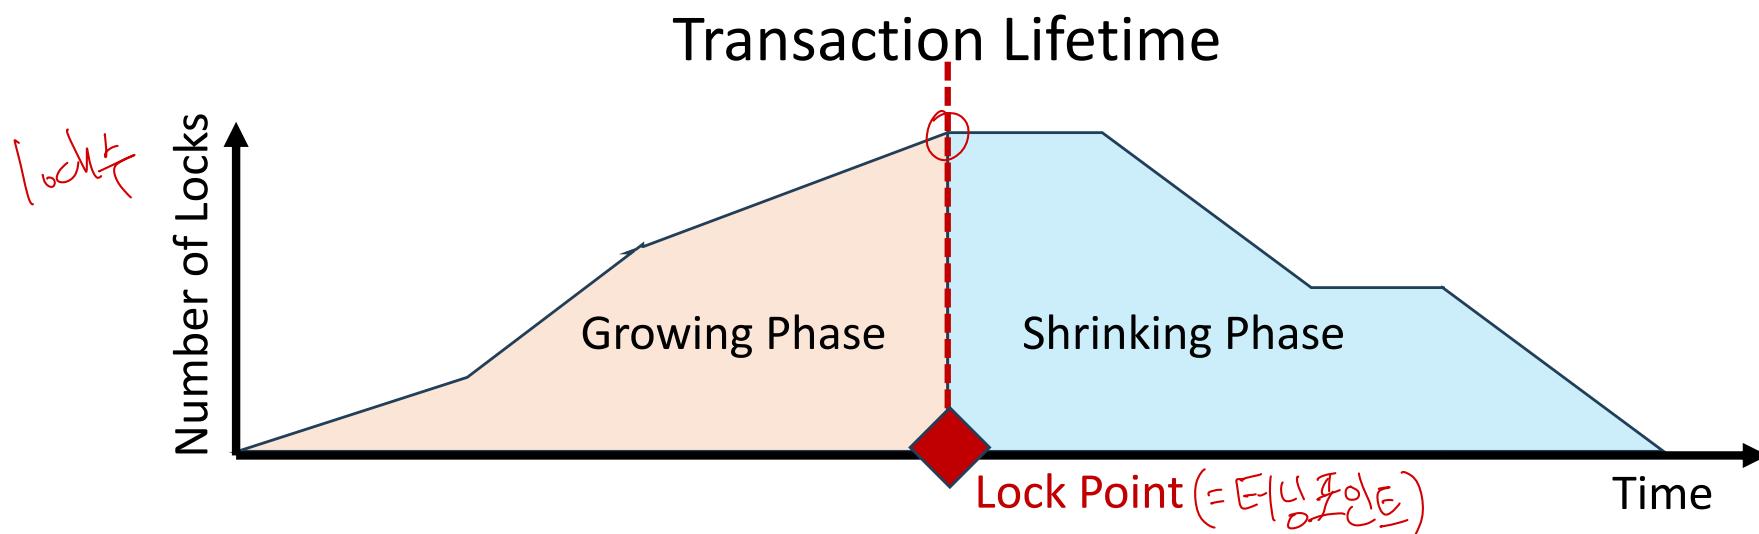
■ Phase 1: Growing Phase

- Transaction may obtain locks
- Transaction may not release locks

→ lock을 갖고 있는 흐름

■ Phase 2: Shrinking Phase

- Transaction may release locks
- Transaction may not obtain locks



- **Lock point:** a time when a transaction acquired its final lock

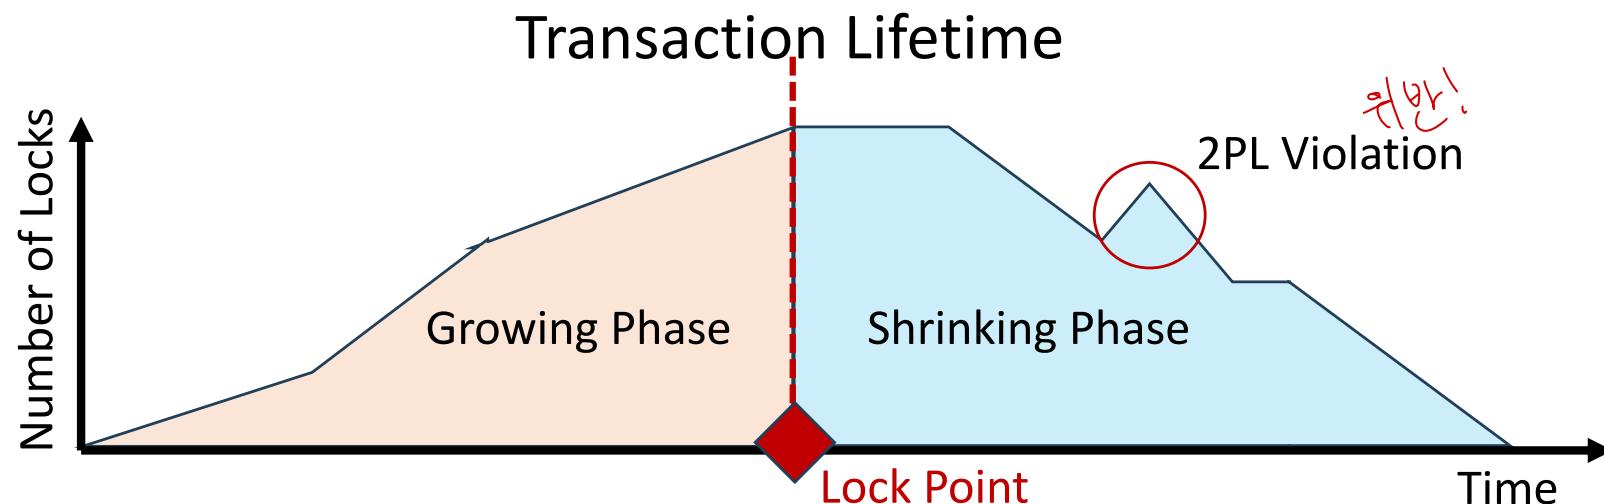
Two-Phase Locking (2PL) Protocol

- Phase 1: **Growing Phase**

- Transaction may obtain locks
- Transaction may not release locks

- Phase 2: **Shrinking Phase**

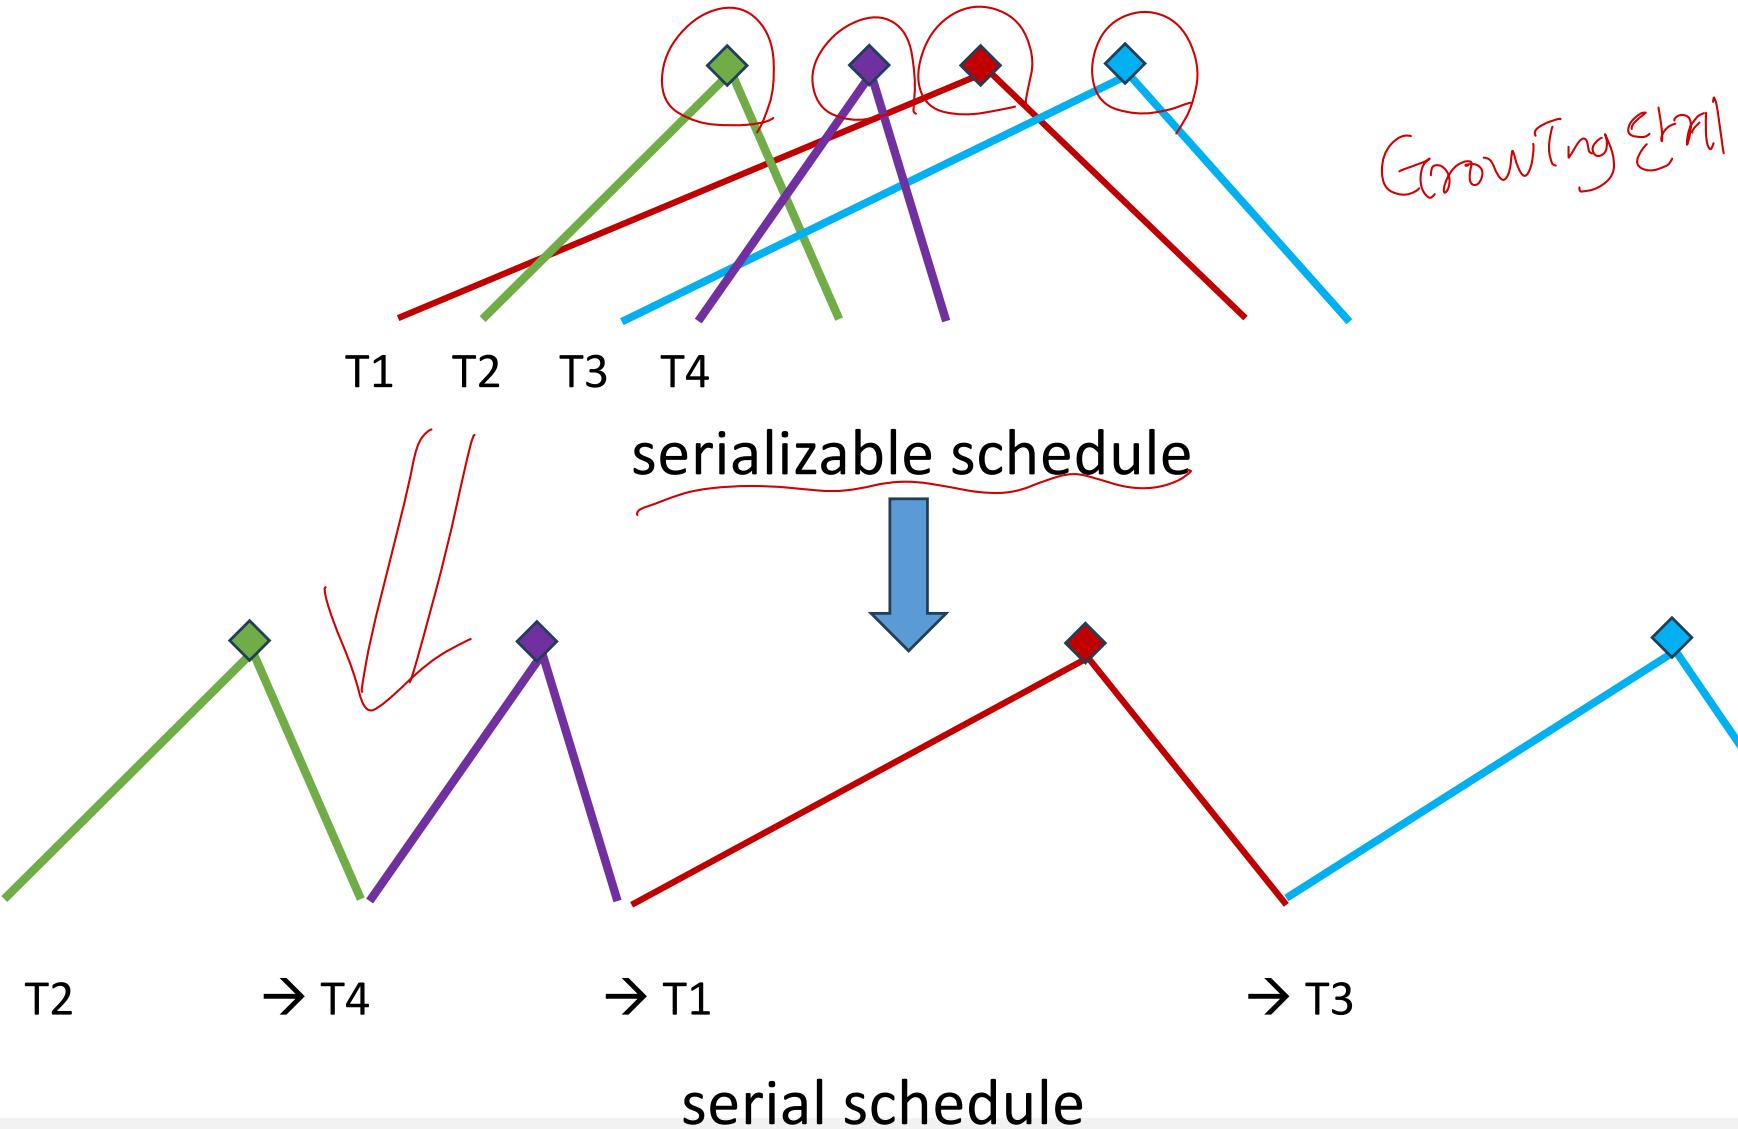
- Transaction may release locks
- Transaction may not obtain locks



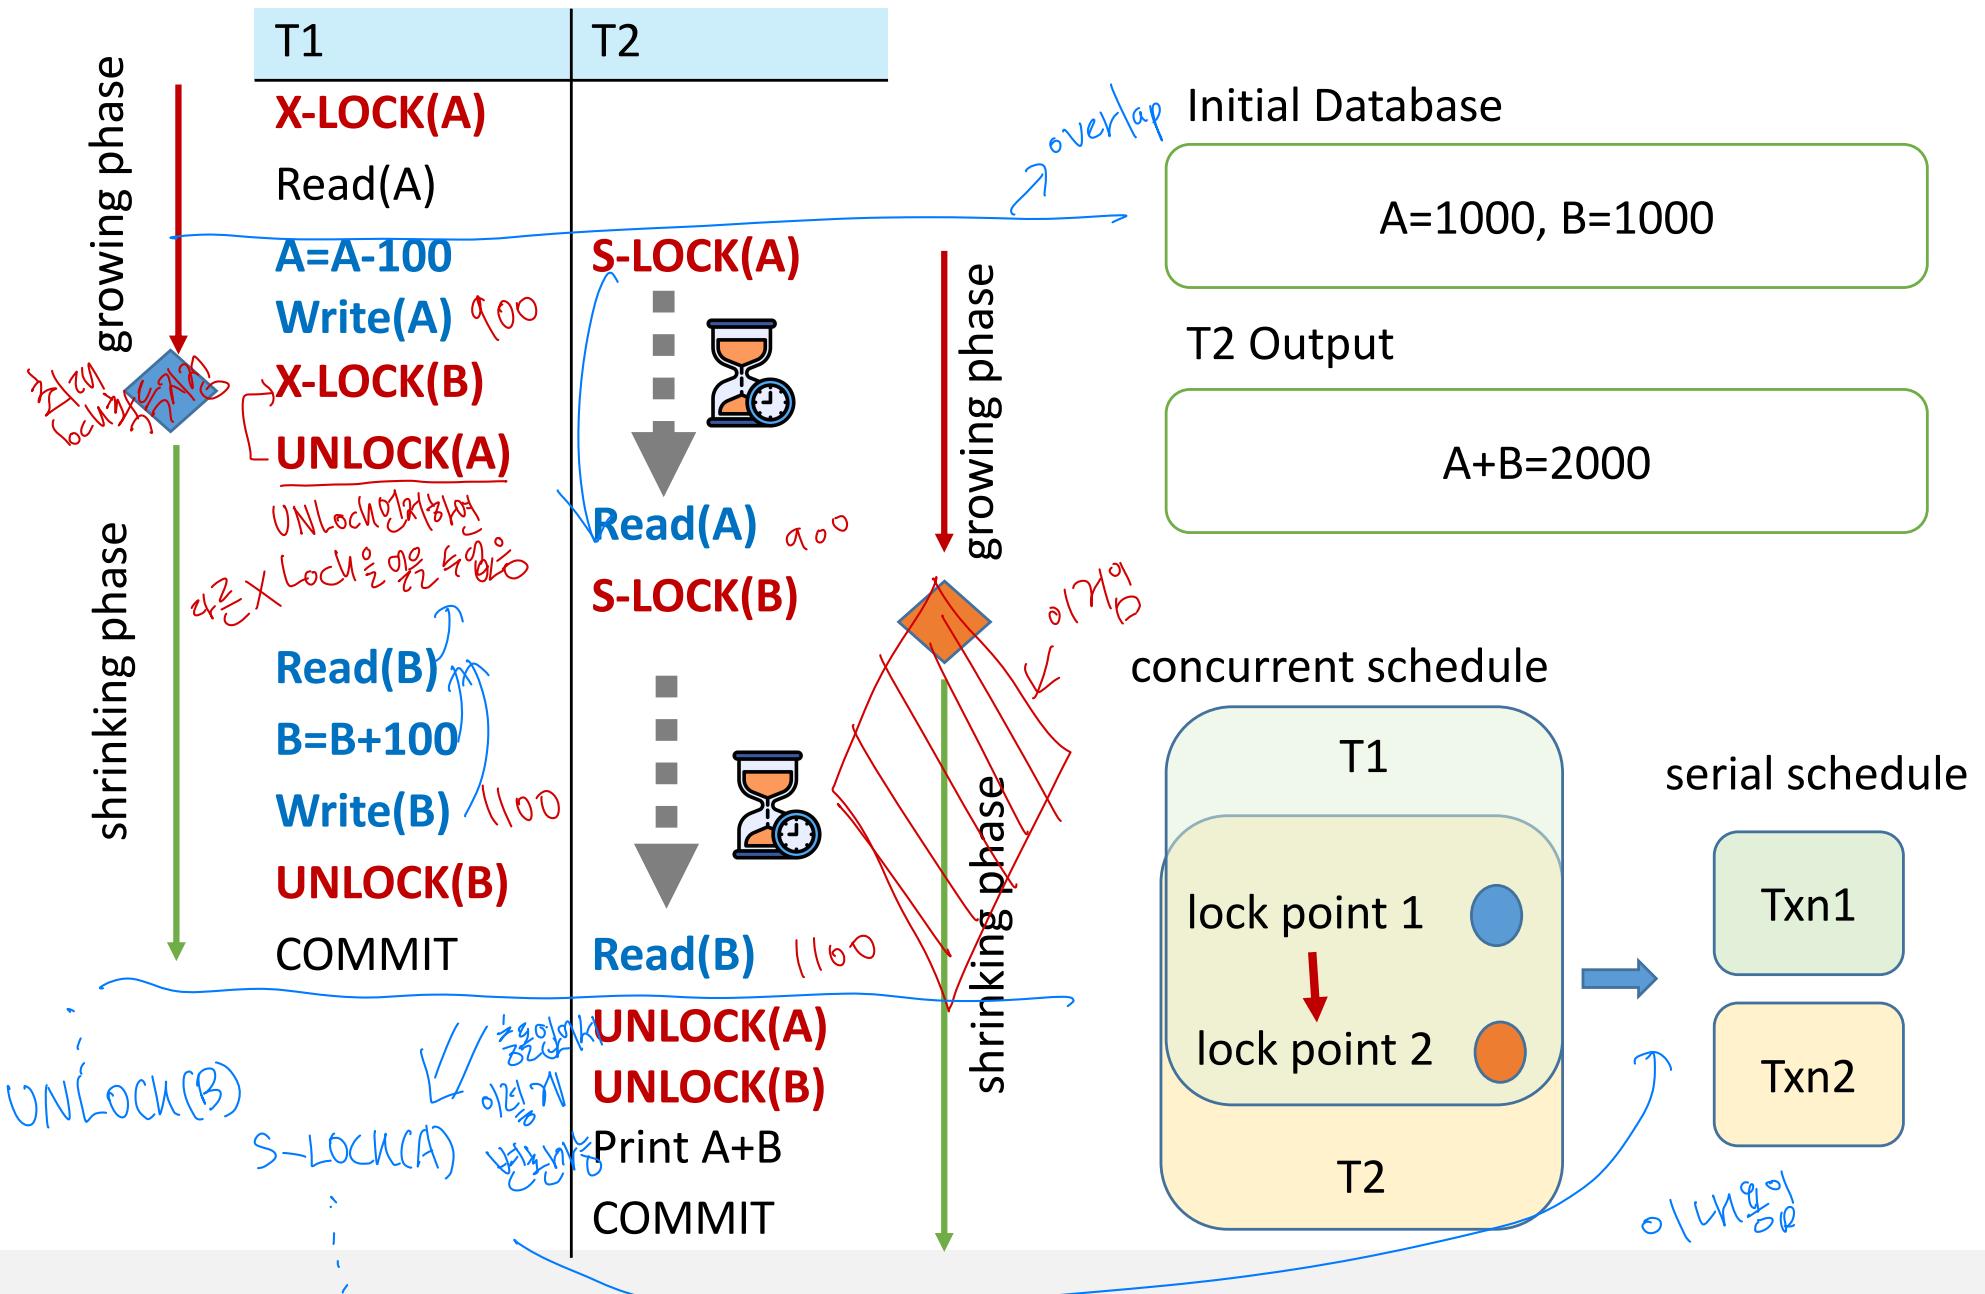
- **Lock point:** a time when a transaction acquired its final lock

Two-Phase Locking (2PL) Protocol

- The order of transactions = order of lock points



Two-Phase Locking (2PL) - Example



Cascading Rollback in 2PL

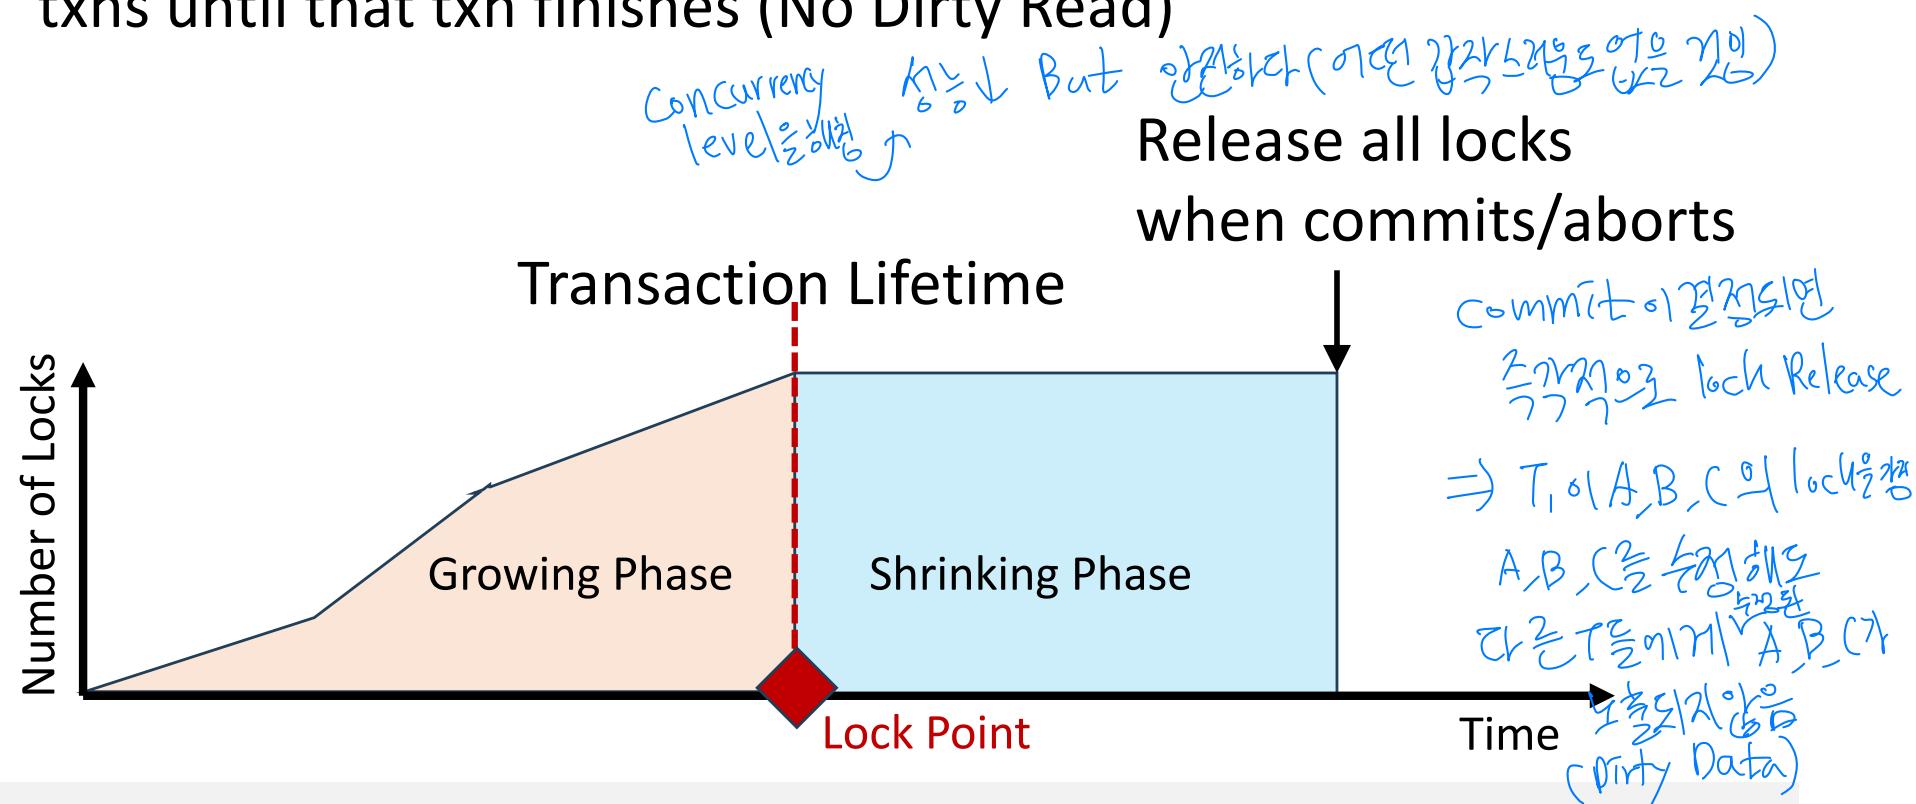
T1	T2	T3
BEGIN		
X-LOCK(A)		
Write(A)	BEGIN	
	X-LOCK(A) → blocked	
	↓ wait	
UNLOCK(A)	→ granted	
	Read(A)	
	Write(A)	BEGIN
	...	X-LOCK(A) → blocked
	...	↓ wait
	UNLOCK(A)	→ granted
		Read(A)
		Write(A)
Write(B)		
...		
ABORT		

T_2 and T_3 also abort
when T_1 aborts \Rightarrow 2차소상작업

Strict Two-Phase Locking Protocol (2PL)

- Strict 2PL (a.k.a Rigorous 2PL) prevents cascading roll-back

- Transaction is only allowed to release locks after it commits/aborts.
- A value written by a txn is not read or overwritten by other txns until that txn finishes (No Dirty Read)

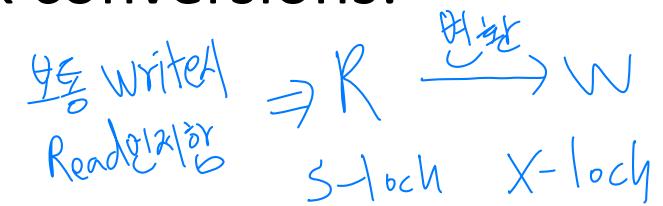


Lock Conversions

- Two-phase locking protocol with lock conversions:

- Growing Phase:

- can acquire a lock-S on item
 - can acquire a lock-X on item
 - can **convert** a lock-S to a lock-X (**upgrade**)



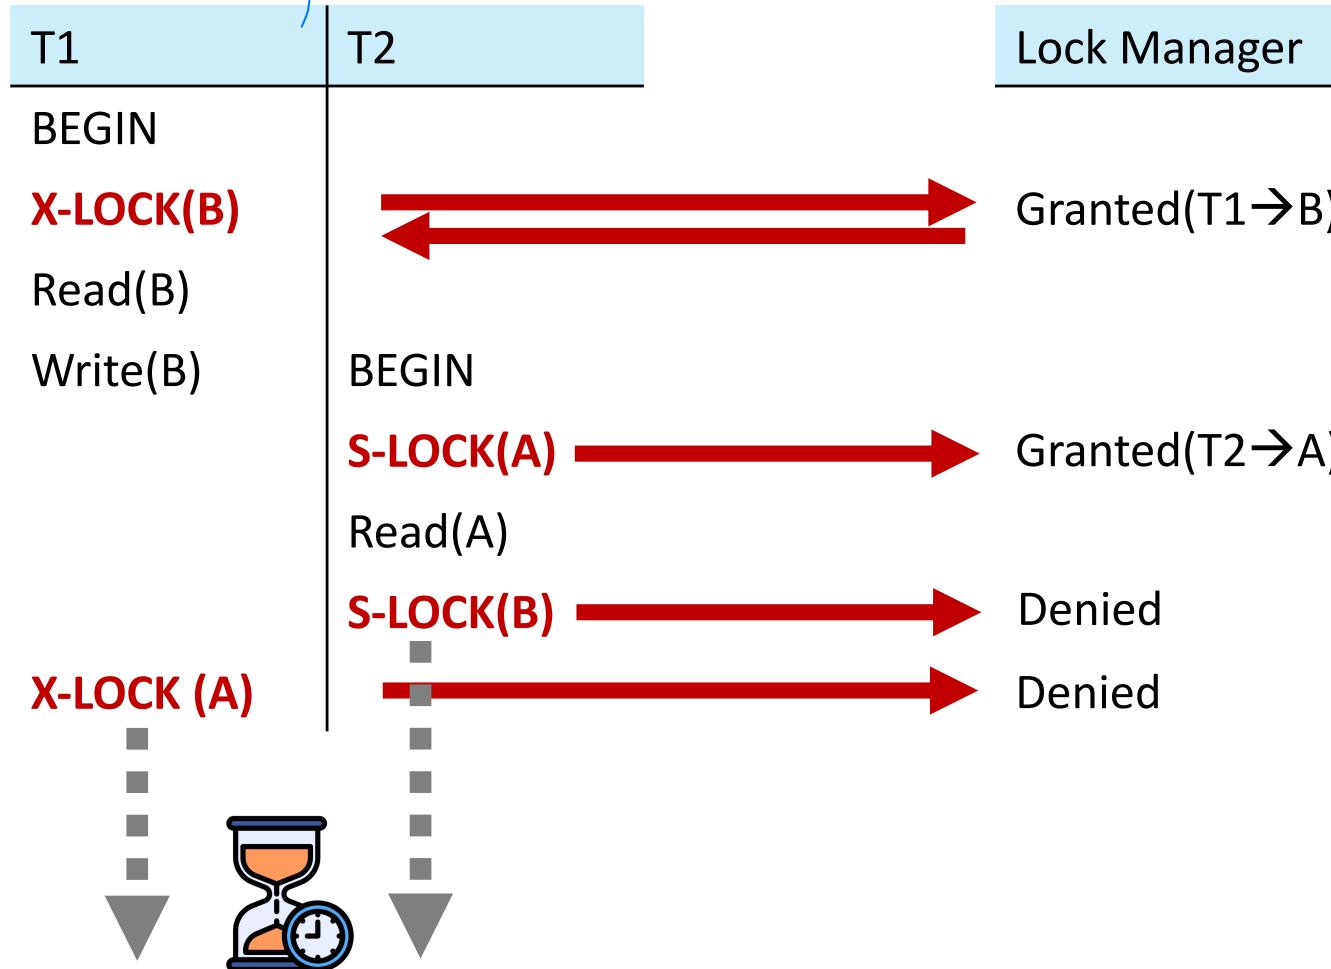
- Shrinking Phase:

- can release a lock-S
 - can release a lock-X
 - can convert a lock-X to a lock-S (**downgrade**)

- This protocol ensures serializability

Deadlock in 2PL

모든 기장자체로 고착상태(운제점)



	T ₁	T ₂
A	S Denied	S
B	X	S Denied
	↓	↓

AB는 각각 다른 흐름에 들어온다.

Neither T_1 nor T_2 can make progress.

Either T_1 or T_2 must be rollbacked and release its lock.

Two ways of dealing with Deadlocks

■ Approach #1:

- **Deadlock Detection**

- If a deadlock is detected, select and kill a victim txn

→ optimistic approach

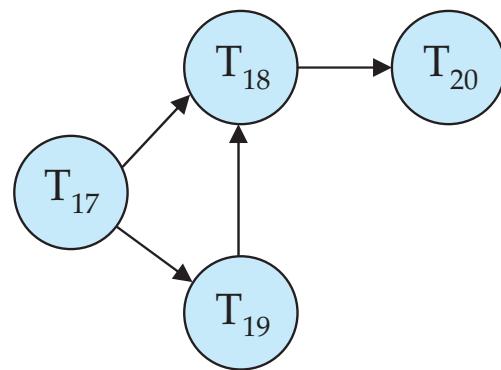
■ Approach #2:

- **Deadlock Prevention**

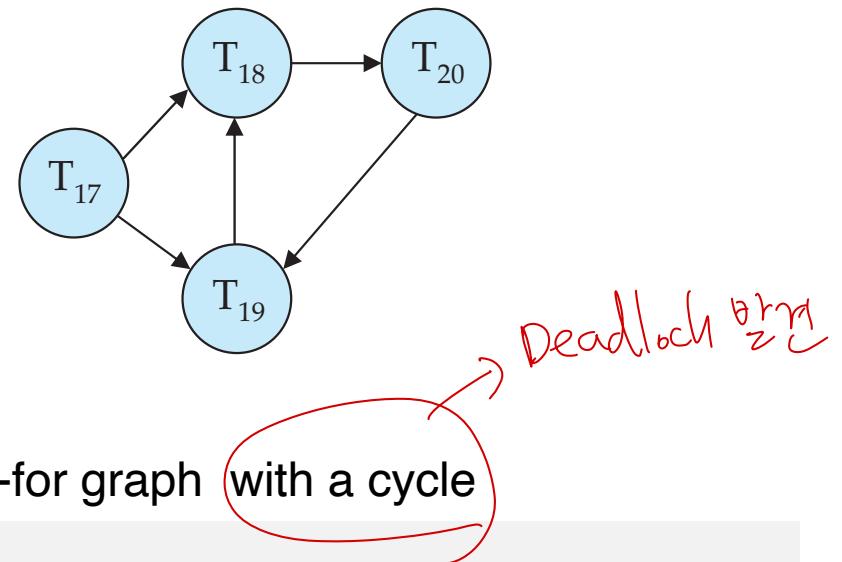
→ Deadlock 발생 가능성을 미리 방지
= pessimistic approach

Deadlock Detection

- The DBMS creates a **waits-for** graph to keep track of what locks each txn is waiting to acquire:
 - Nodes are transactions
 - **Edge** from T_i to T_j if T_i is waiting for T_j to release a lock.
 $T \rightarrow J$
- The system periodically checks for cycles in waits-for graph and then decides how to break it.



Wait-for graph without a cycle



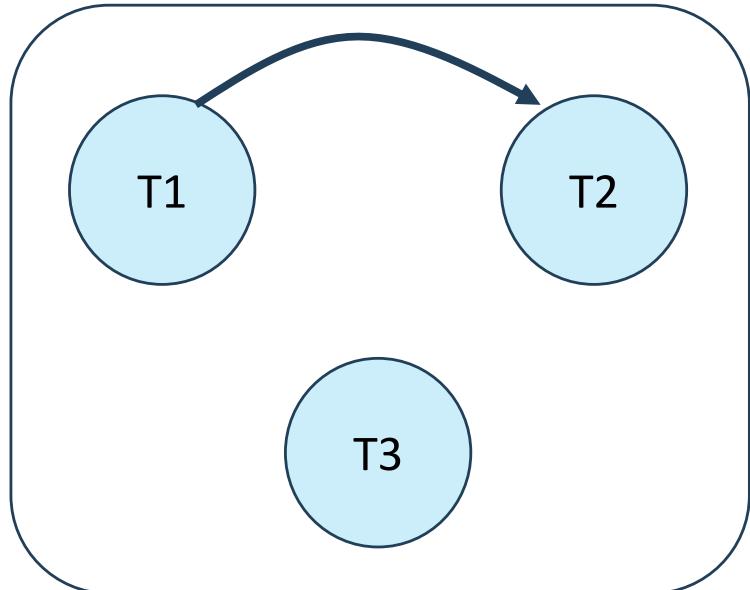
Wait-for graph **with a cycle**

Deadlock Detection

T1	T2	T3
BEGIN	BEGIN	BEGIN
S-LOCK(A)	X-LOCK(B)	S-LOCK(C)
S-LOCK(B)	X-LOCK(C)	X-LOCK(A)

*Denied
(wait)*

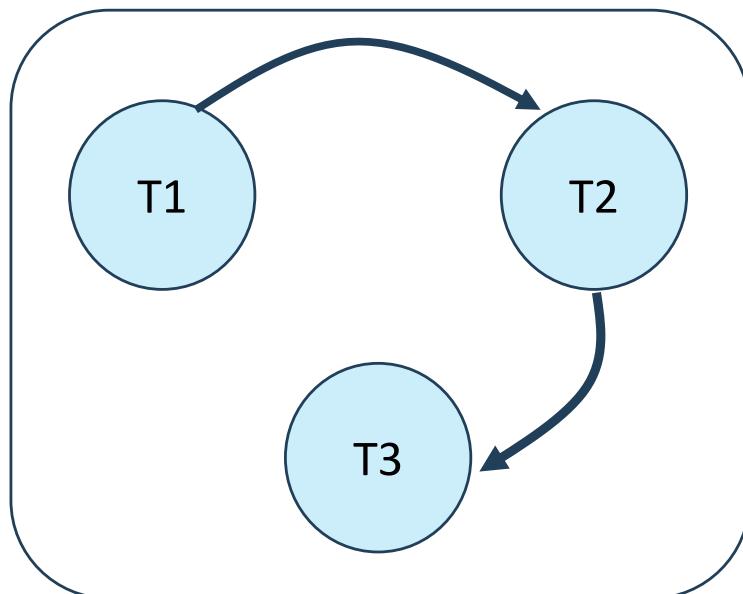
Waits-for Graph



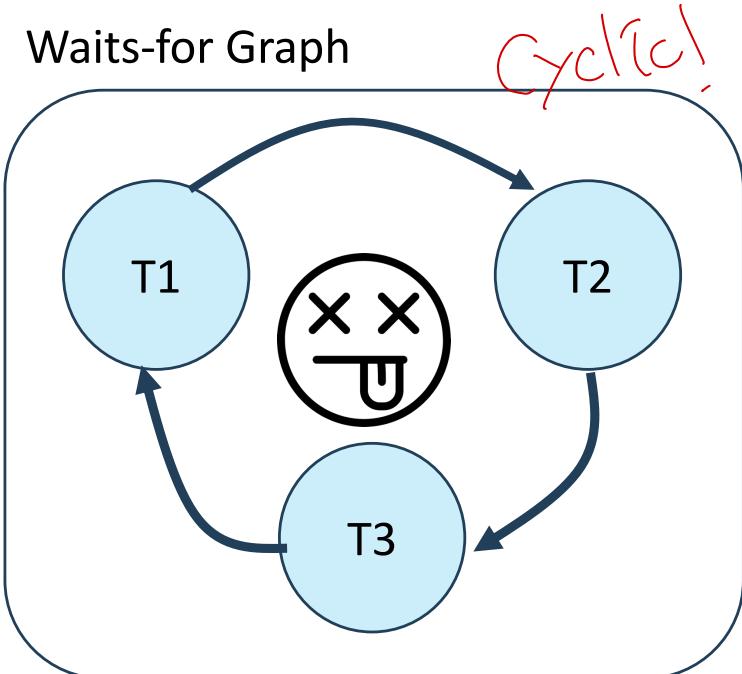
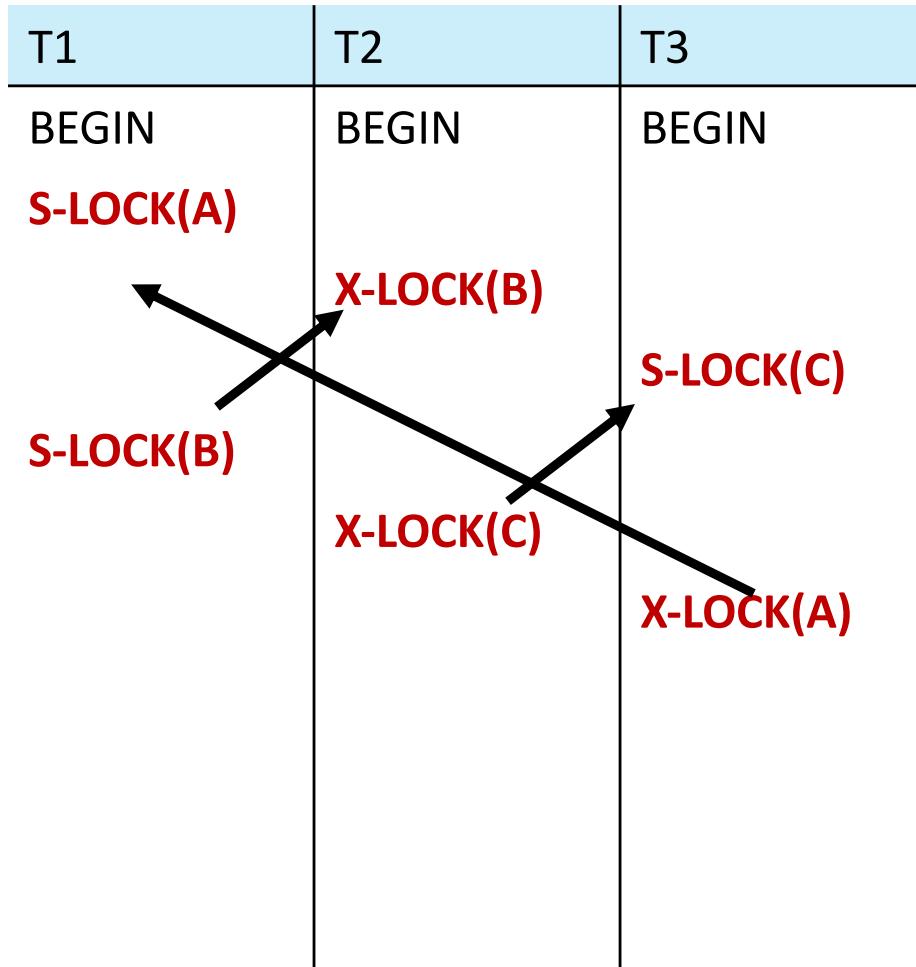
Deadlock Detection

T1	T2	T3
BEGIN	BEGIN	BEGIN
S-LOCK(A)	X-LOCK(B)	S-LOCK(C)
S-LOCK(B)	X-LOCK(C)	X-LOCK(A)

Waits-for Graph



Deadlock Detection



Deadlock Handling

- When the DBMS detects a deadlock, it will select a “victim” txn to rollback to break the cycle.
- Selecting the proper victim depends on a lot of different parameters....
 - By age (lowest timestamp)
 - By progress (least/most queries executed)
 - By the # of items already locked
 - By the # of txns that we have to rollback with it
- We also should consider the # of times a txn has been restarted in the past to prevent starvation.

Deadlock Prevention

- When a txn tries to acquire a lock that is held by another txn, the DBMS **kills** one of them **to prevent** a deadlock.
- This approach does not require a *waits-for* graph or detection algorithm.
→ *각각 언제 시작되었는지 알 필요*

Deadlock Prevention

- Assign priorities based on timestamps:

→ Older Timestamp = Higher Priority (e.g., $T_1 > T_2$)

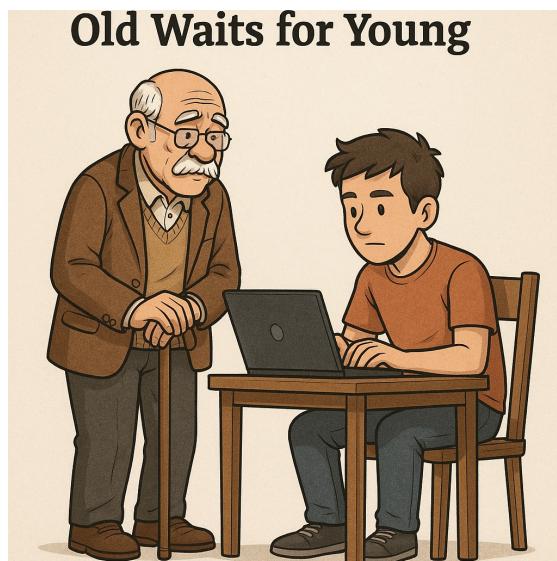
old txn



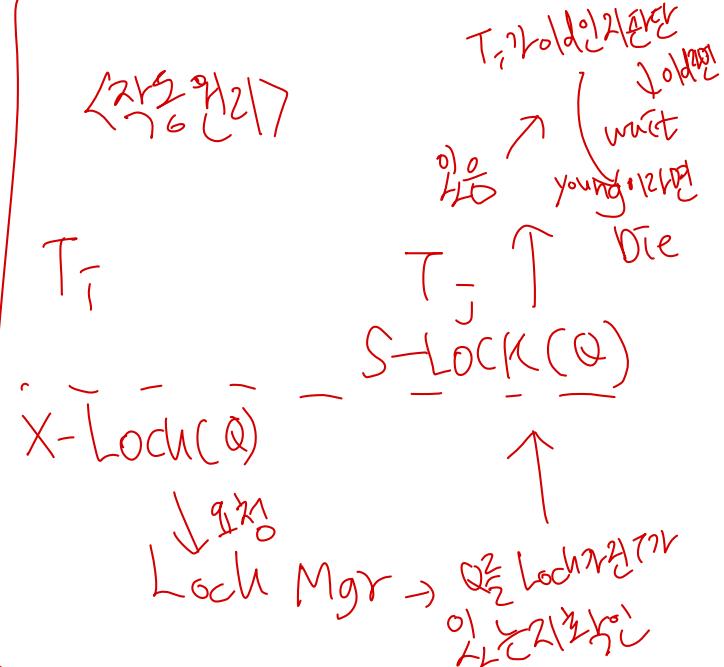
- Wait-Die

→ If requesting txn has higher priority than holding txn, then requesting txn waits for holding txn.

→ Otherwise requesting txn aborts.



Transaction A가 먼저 시작
→ Start Time 을 기준(더 작은게 old)



Deadlock Prevention

- Assign priorities based on timestamps:

→ Older Timestamp = Higher Priority (e.g., $T_1 > T_2$)

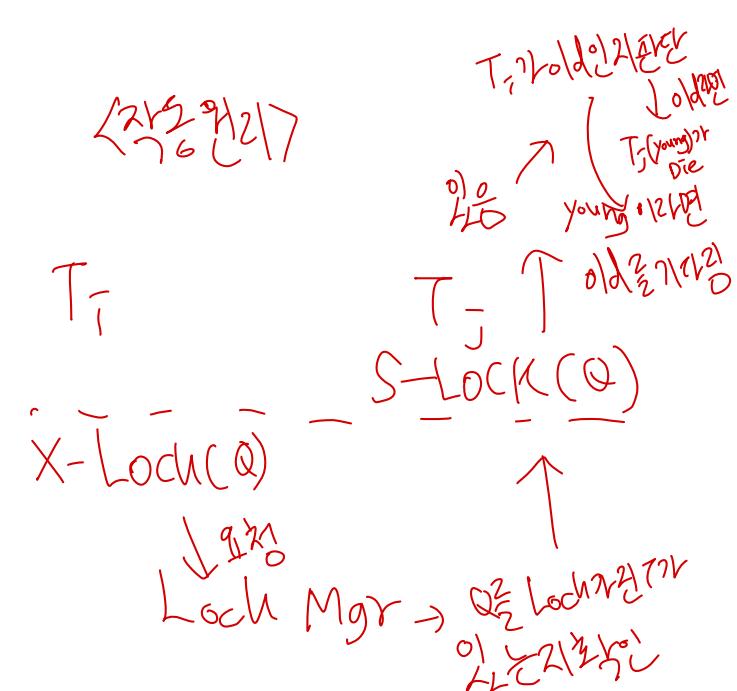
old txn kill young txn

Old txn

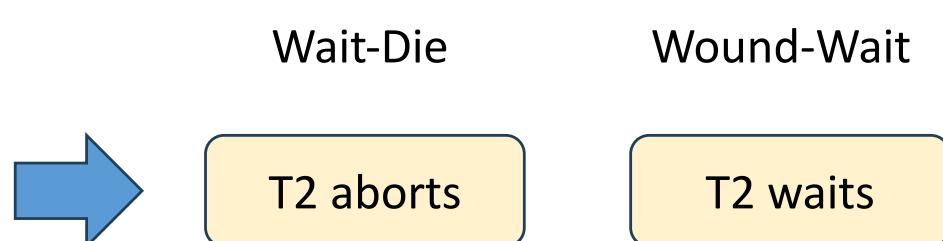
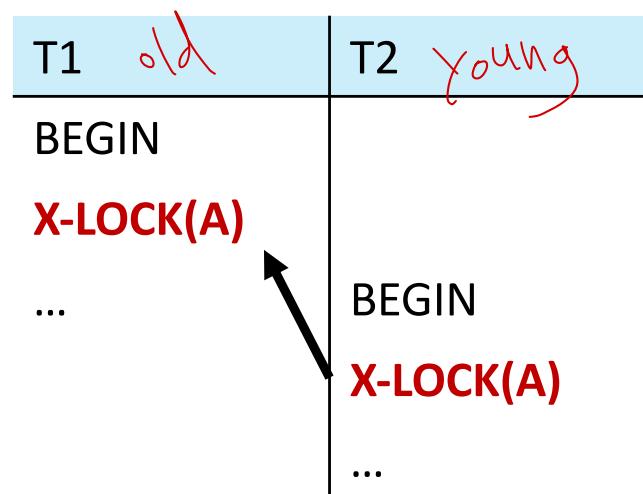
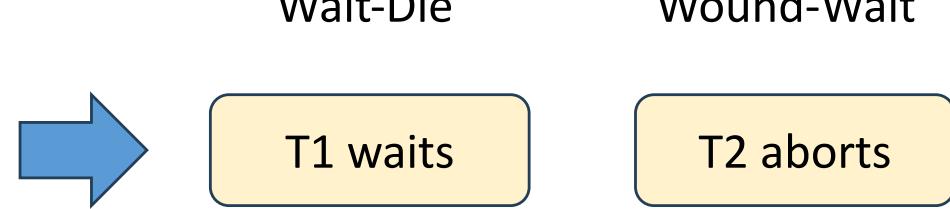
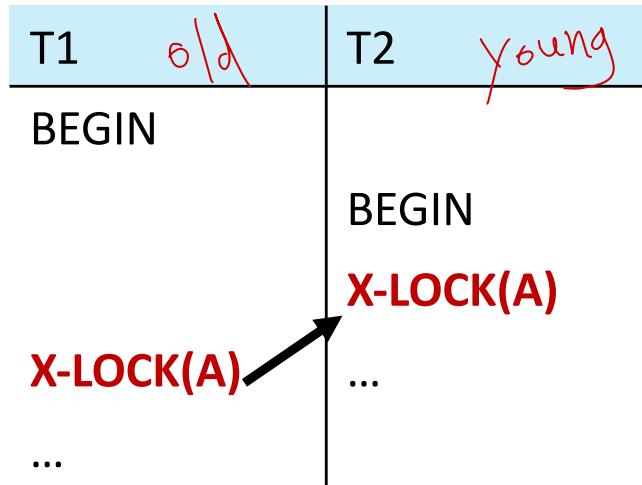
- **Wound-Wait** (“Young Waits for Old”)

→ If requesting txn has higher priority than holding txn, then holding txn aborts and releases lock.

→ Otherwise requesting txn waits.



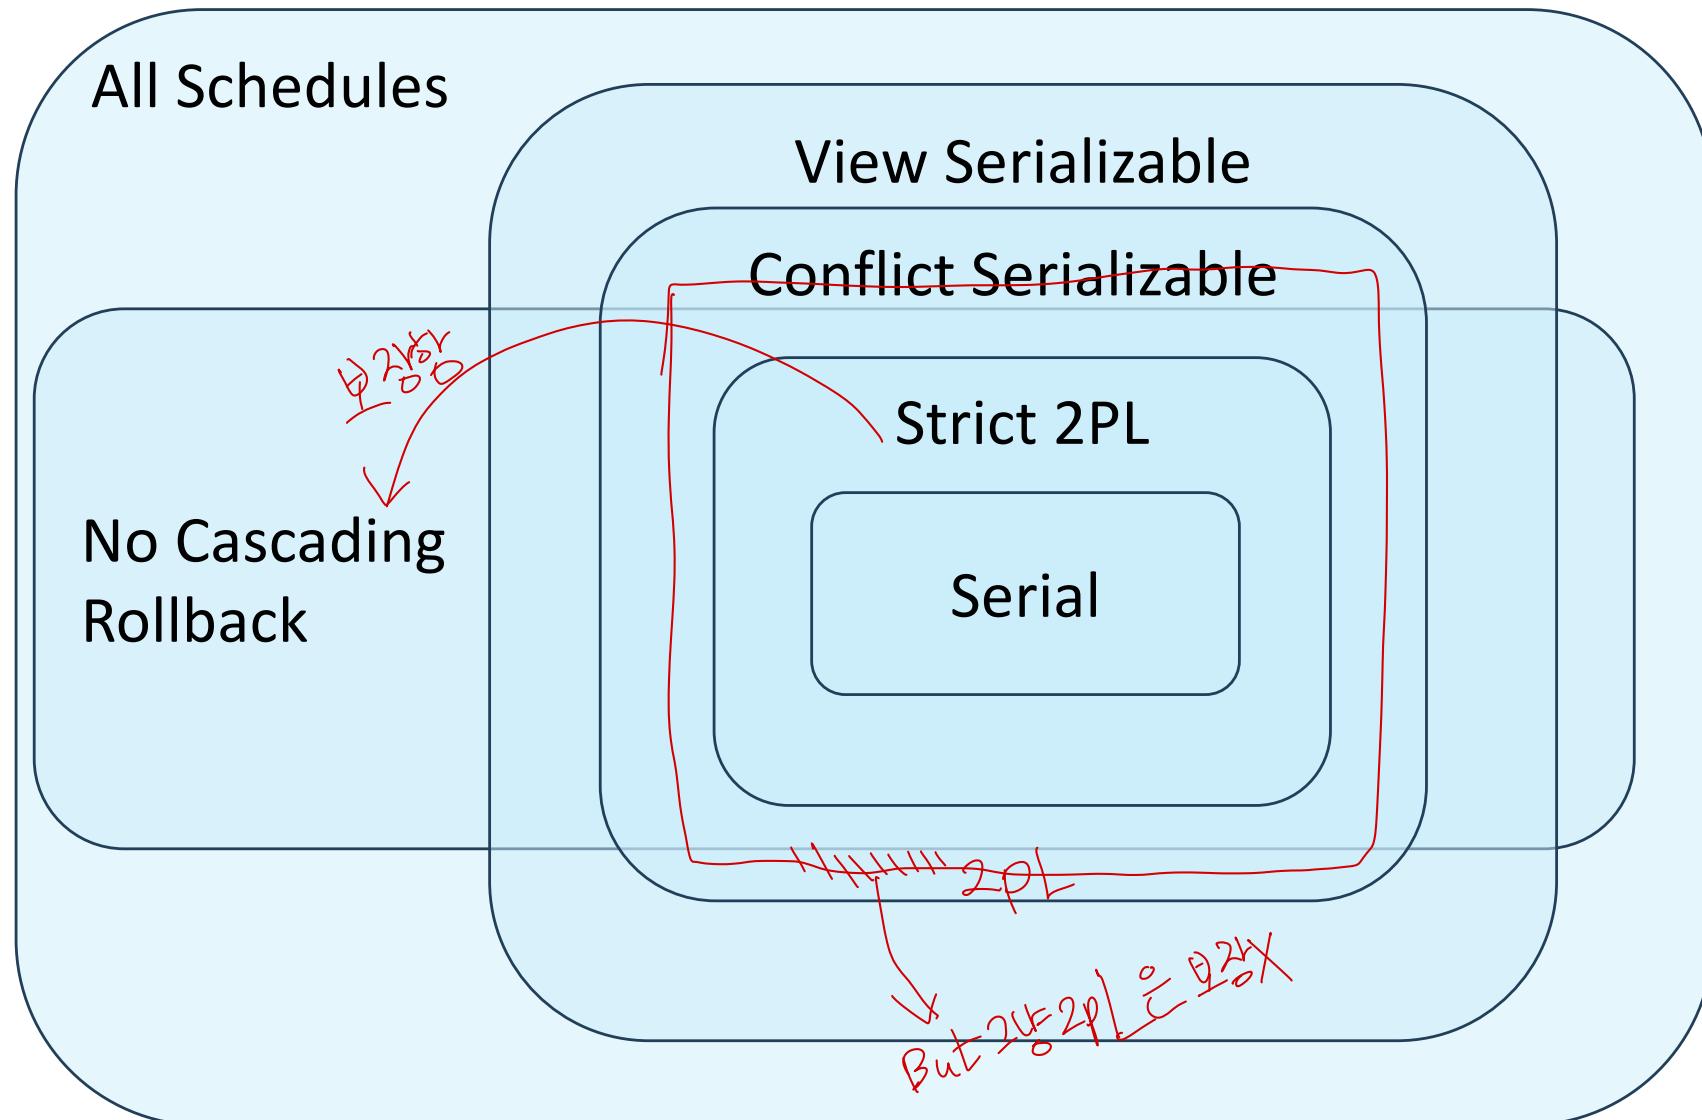
Deadlock Prevention



Deadlock Prevention

- Why do these schemes guarantee no deadlocks?
 - Only one “type” of direction allowed when waiting for a lock.
 ↑
 상수는 향상할 수 있다
 No cyclic locks
- When a txn restarts, what is its (new) priority?
 - Its original timestamp to prevent it from getting starved for resources.
 ↑
 Starvation (기다려야 하는 경우)
 or starvation of lock
 (≈ repeated abortion)

Universe of Schedules



Lock Granularities

세분화 (범위)

- When a txn wants to acquire a “lock”, the DBMS can decide the granularity (i.e., scope) of that lock.
→ Attribute? Tuple? Page? Table?
- The DBMS should ideally obtain fewest number of locks that a txn needs.
 - Coarse-grained → 테이블 기준 lock (범위가 큼)
< 범위가 큼 : 동시성이 잘 아질 수 있다 (overhead ↓)
< 범위가 작은 : 동시성이 높아지고, 성능이 좋아진다 (overhead ↑)
관리하는 cost
- Trade-off between Concurrency versus Overhead.
→ Fewer Locks, Larger Granularity vs. More Locks, Smaller Granularity.

Database Lock Hierarchy

The levels, starting from the top level are

- Database
- Table
- Page
- Tuple
- Attribute (rare)

update
from
where
→ 이 구조는 S IX 가 있음

<Level>

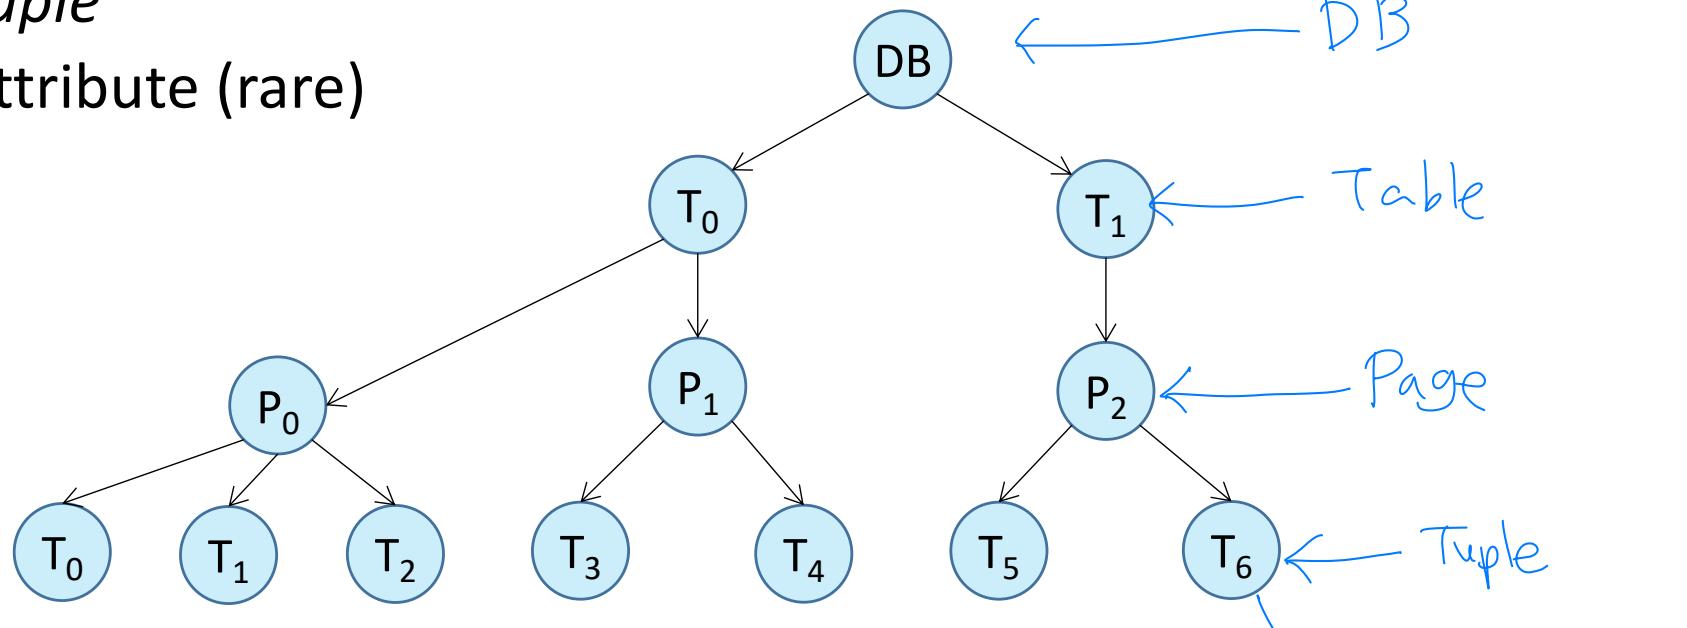
DB

Table

Page

Tuple

Attribute



ex) T_3 를 업데이트 $\rightarrow T_3 : X\text{lock}$

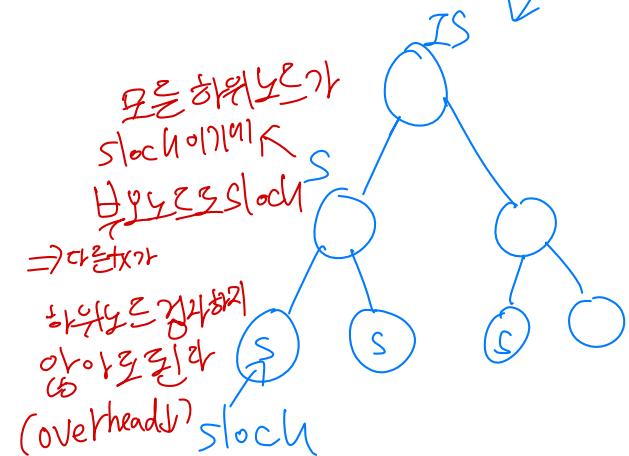
$\Rightarrow P_1, T_0, DB$ 모두 IX lock 걸림

\Rightarrow Tx2가 T_3, T_4 를 끌기 요청

\Rightarrow DB, $T_6 : IS, P, : S$ 가 도로 IX vs S 는 끌림
 \therefore Page 단계에서 false 되면 ↳ 차원 노드를 방문하지 않고 경쟁

Intention Lock Modes

- Intention locks allow a higher level node to be locked in S or X mode without having to check all descendant nodes.
- If a node is locked in an intention mode, then some txn is doing explicit locking at a lower level in the tree



Intention Lock Modes

- **intention-shared (IS)**: indicates explicit locking at a lower level of the tree but only with **shared** locks. \Rightarrow 하위 구조에锁을 가진 tx가 존재한다.
- **intention-exclusive (IX)**: indicates explicit locking at a lower level with **exclusive** locks \Rightarrow xlock 가진 tx 존재
- **shared and intention-exclusive (SIX)**: the subtree rooted by that node is locked explicitly in **shared** mode and explicit locking is being done at a lower level with **exclusive-mode** locks. \Rightarrow 하위 구조가 모두 lock \rightarrow 그 중에 몇몇이 xlock인 경우

Compatibility Matrix with Intention Lock Modes

- The compatibility matrix for all lock modes is:

	IS	IX	S	SIX	X
IS	true = comparable	true	true	true	false
IX	true	true	false	false	false
S	true	false	true	false	false
SIX	true	false	false	false	false
X	false	false	false	false	false

1) 흥미로운 관계
 1) S - IX (SIX)
 ⇒ False
 왜? S는 모든 하위노드 True
 하지만 IX는 대개 무조건
 X여야해서 충돌!
 2) IX - IX (무조건 관계)
 원도 나눠 같으면 됨
 → 그게 아니면 바로 하위노드에서
 차지하기가 성공 X → True

IS : lock 끼리는 관찰다 → True
 IX : 같은쪽에 X Lock을 얻으면 된다. 혹은 전족이어도 무조건 관계에서는 관찰다.
 S : 관찰다 → True
 SIX : IS와 비슷한 이유 → True
 X : X는 하위 모든 노드가 X Lock을 가져야하는데 이미 원쪽 노드에
 Lock을 가지고 있으니 충돌 → True

Multiple Granularity Locking Scheme

- Each txn obtains appropriate lock at highest level of the database hierarchy.

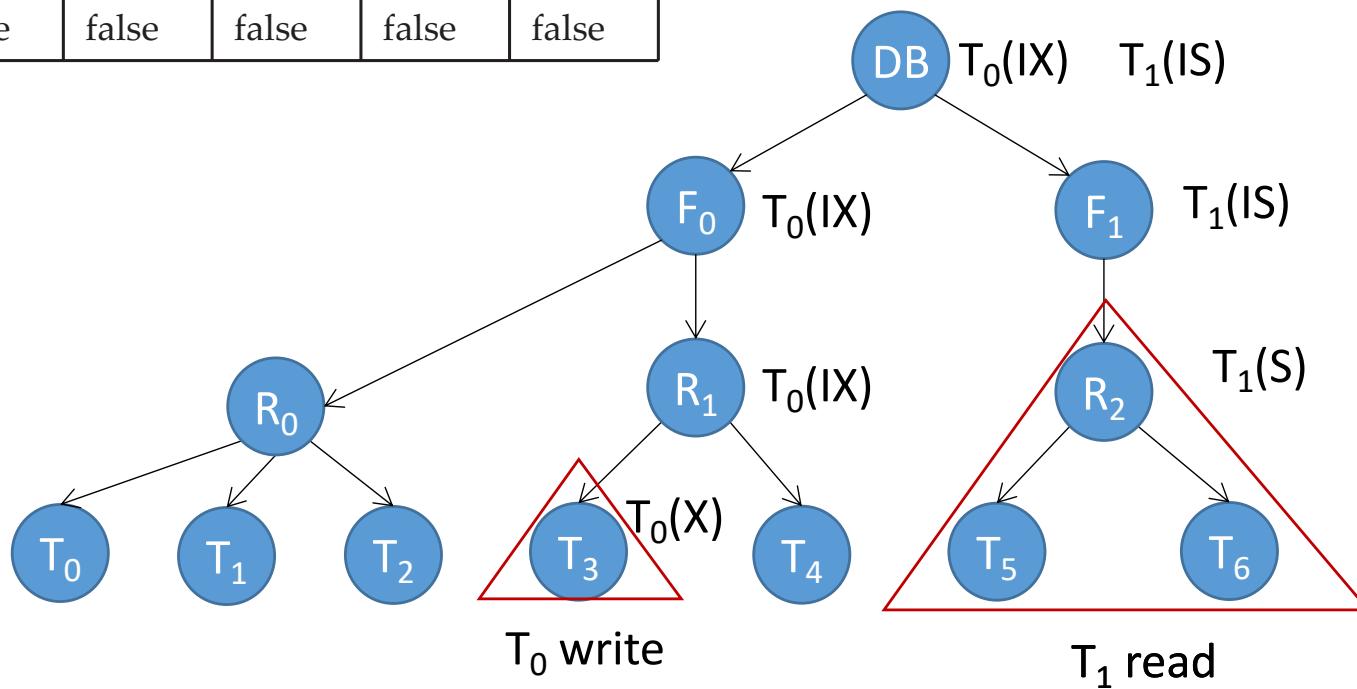
자식

부모

- To get S or IS lock on a node, the txn must hold at least IS on parent node.
- To get X, IX, or SIX on a node, must hold at least IX on parent node.

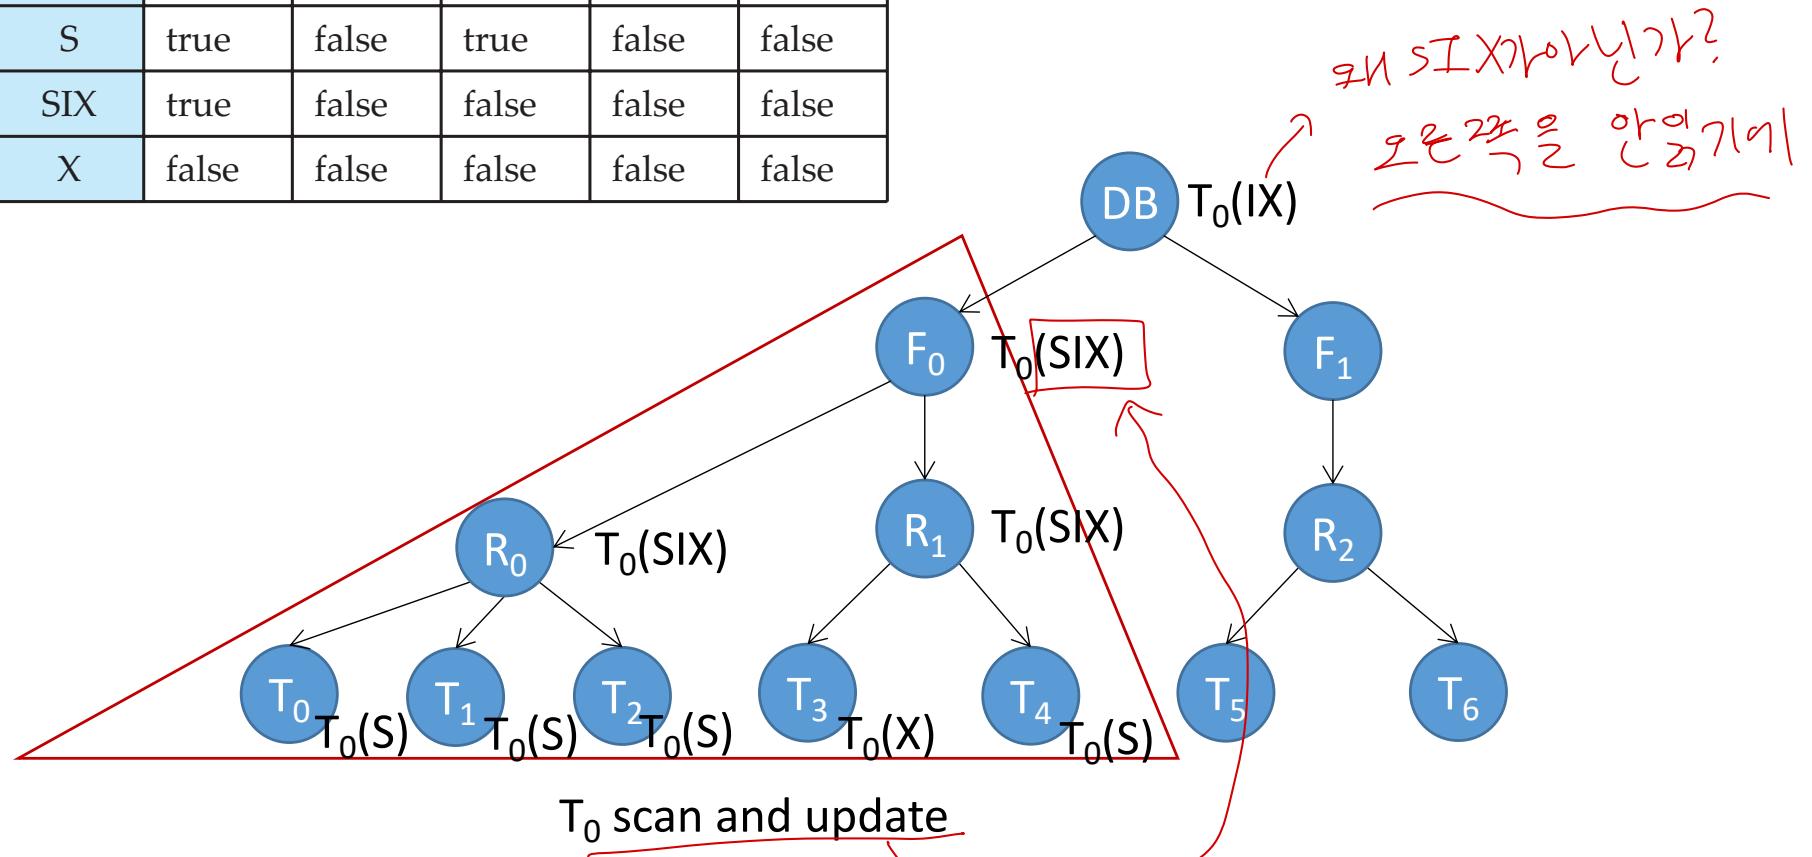
MGL Example

	IS	IX	S	SIX	X
IS	true	true	true	true	false
IX	true	true	false	false	false
S	true	false	true	false	false
SIX	true	false	false	false	false
X	false	false	false	false	false



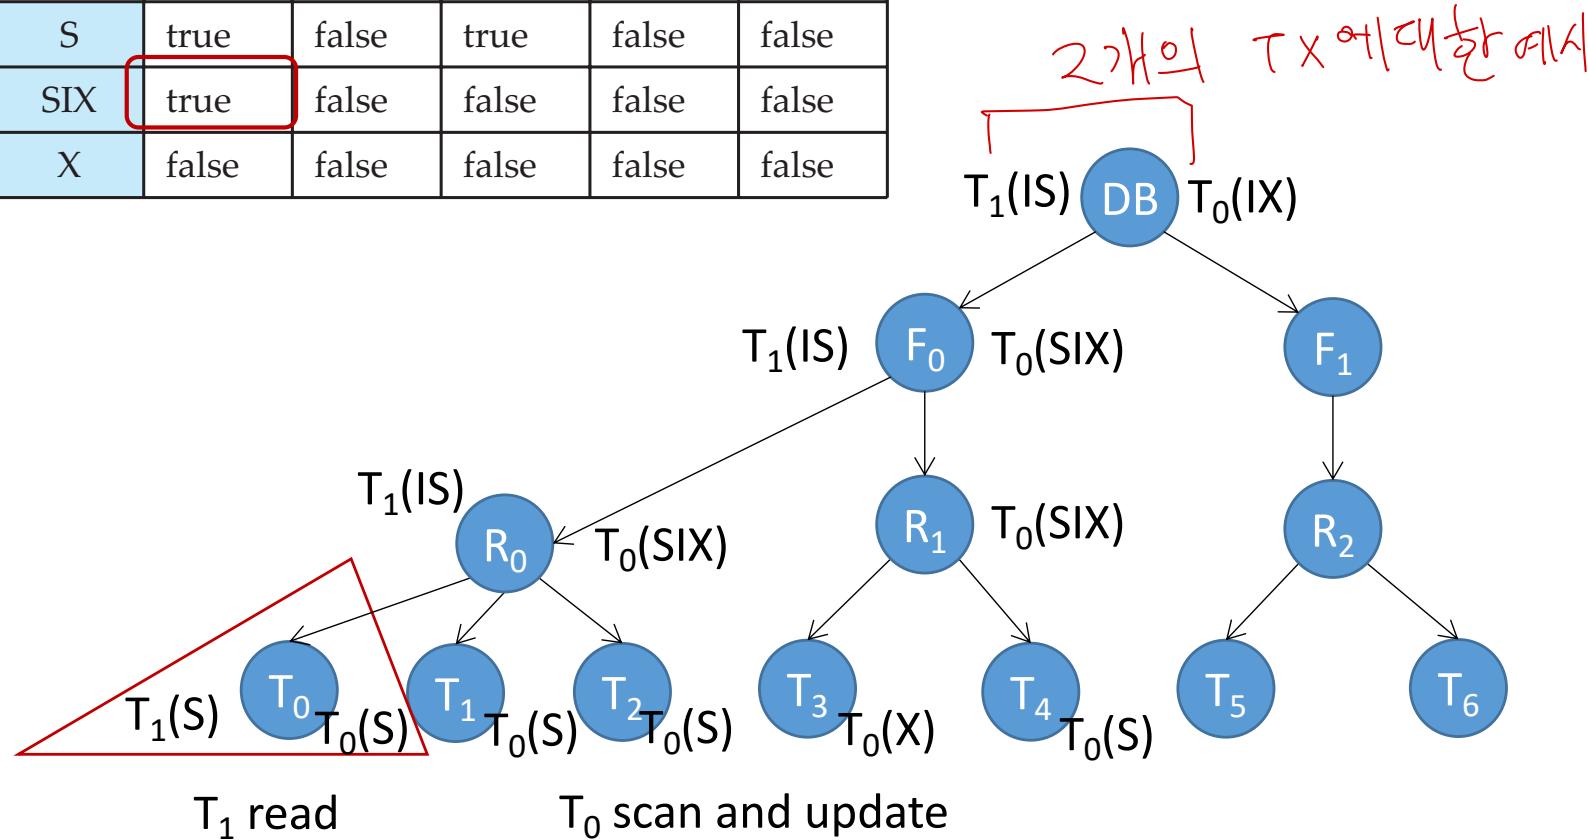
MGL Example

	IS	IX	S	SIX	X
IS	true	true	true	true	false
IX	true	true	false	false	false
S	true	false	true	false	false
SIX	true	false	false	false	false
X	false	false	false	false	false



MGL Example

	IS	IX	S	SIX	X
IS	true	true	true	true	false
IX	true	true	false	false	false
S	true	false	true	false	false
SIX	true	false	false	false	false
X	false	false	false	false	false

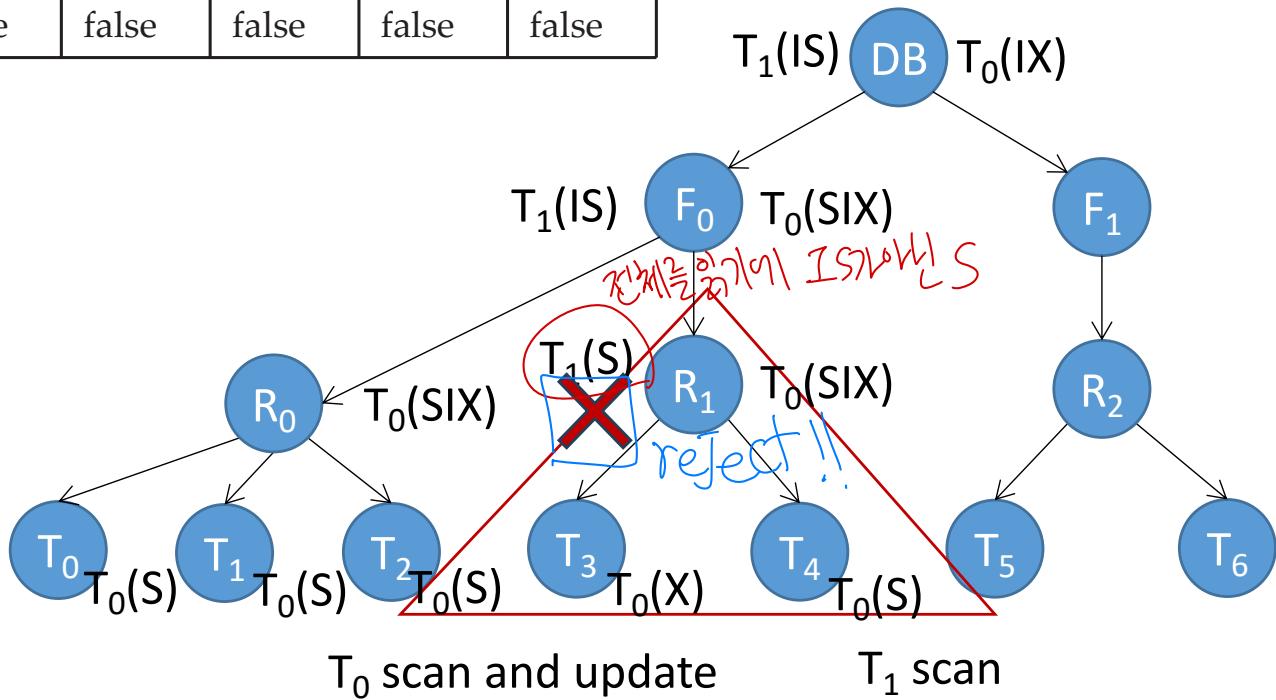


MGL Example

	IS	IX	S	SIX	X
IS	true	true	true	true	false
IX	true	true	false	false	false
S	true	false	true	false	false
SIX	true	false	false	false	false
X	false	false	false	false	false

S : entire access
 IS : partial access

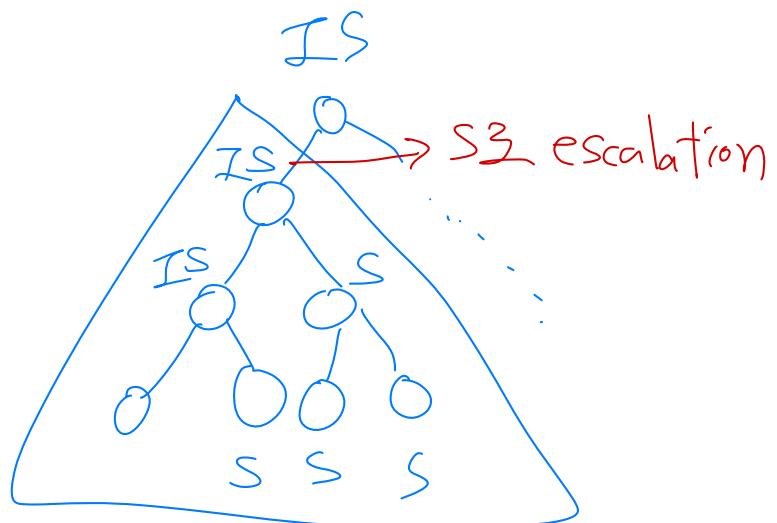
($X \in DB \setminus T_2$)



Lock Escalation

- The DBMS can automatically switch to coarser-grained locks when a txn acquires too many low-level locks.
- This reduces the number of requests that the lock manager must process.

하나의 테이블에 대한 많은 가볍은 풀을 잡을 때



↓
S로 바뀌면 Table 전체를 잡아옴
⇒ 최적화 기법

Locking in Practice

- Applications typically do not acquire a txn's locks manually (i.e., explicit SQL commands). → DBMS가 알아서 할 테니 app 개발자는 신경 X
- Sometimes you need to provide the DBMS with hints to help it to improve concurrency.
 - Update a tuple after reading it.
 - Skip any tuple that is locked.
- Explicit locks are also useful when doing major changes to the database.

SELECT...FOR UPDATE

- Perform a SELECT and then sets an exclusive lock on the matching tuples.
- Can also set shared locks:
 - Postgres: FOR SHARE
 - MySQL: LOCK IN SHARE MODE

판정현상과 함께 "for update" 키워드 사용
→ DBMS에게 판정현상
방지에 대한 힌트를 주는 것

