

Database Systems

Lecture19 – Chapter 18: Concurrency Control (Part II: Timestamp-based Protocol)

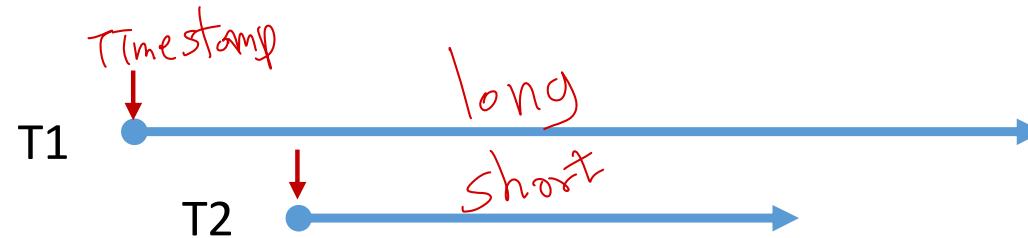
Beomseok Nam (남범석)
bnam@skku.edu

Observation

- If *conflicts* between transactions are *rare* and most transactions are *short-lived*, then acquiring locks adds unnecessary overhead.
 ⇒ expensive
 (∴ OS가 잠금을 요청하는걸)
 lock
 - Concurrency control protocol could be optimized for the no-conflict case.

Timestamp-Based Protocols

- Each transaction T_i has a unique timestamp $TS(T_i)$.
 - *Timestamp* is assigned when *a transaction starts*



- Newer transactions have timestamps greater than earlier ones
 - $TS(\text{old } T1) < TS(\text{new } T2)$ (TS: timestamp)
- Timestamp could be based on a logical counter, not real time.
= Tx 시작 시점

논리적인 증가값 (Tx id와 같은)

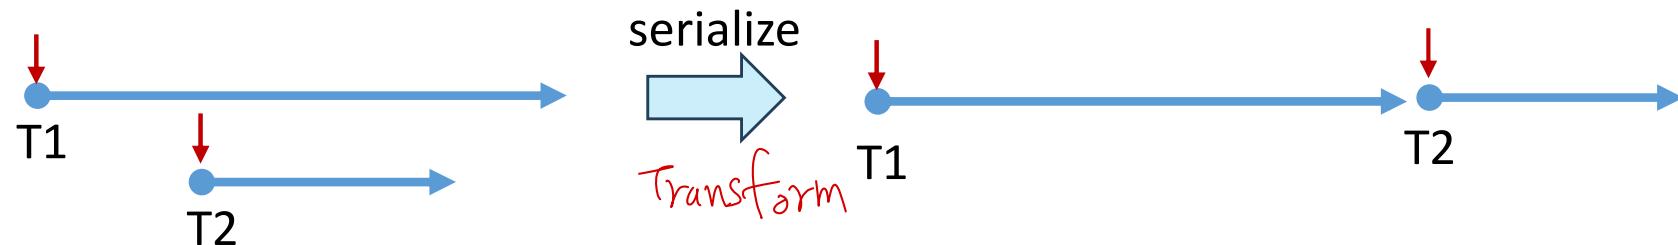
- Timestamp-Based Protocols
 - Timestamp-Ordering Protocol (TSO)
 - Validation-based Protocol (OCC)

최적화 가능

Timestamp-Based Protocols

- **Timestamp-Ordering Protocol (TSO)** serializes transactions such that

- timestamp order = serializability order
- If $TS(T1) < TS(T2)$, the schedule is equivalent to the serial schedule where $T1 \rightarrow T2$



- **Validation-based Protocol (OCC)** serializes transactions such that

- validation time order = serializability order

Timestamp-Ordering Protocol

The **timestamp ordering (TSO) protocol**

→ can be table, page, tuple..

- Every object Q is tagged with timestamp of the last txn that successfully did read/write:
 - **W-timestamp(Q)** : Write timestamp on Q
 - **R-timestamp(Q)** : Read timestamp on Q
- Rules for Read/Write
 - Conflicting operations are executed in timestamp order
 - Out of order operations trigger transaction rollback

Timestamp-Ordering Protocols (Cont.)

- Suppose a transaction T_i issues a **read(Q)**

W-TS (Q)

R-TS (Q)

T_i 가 young 일 때

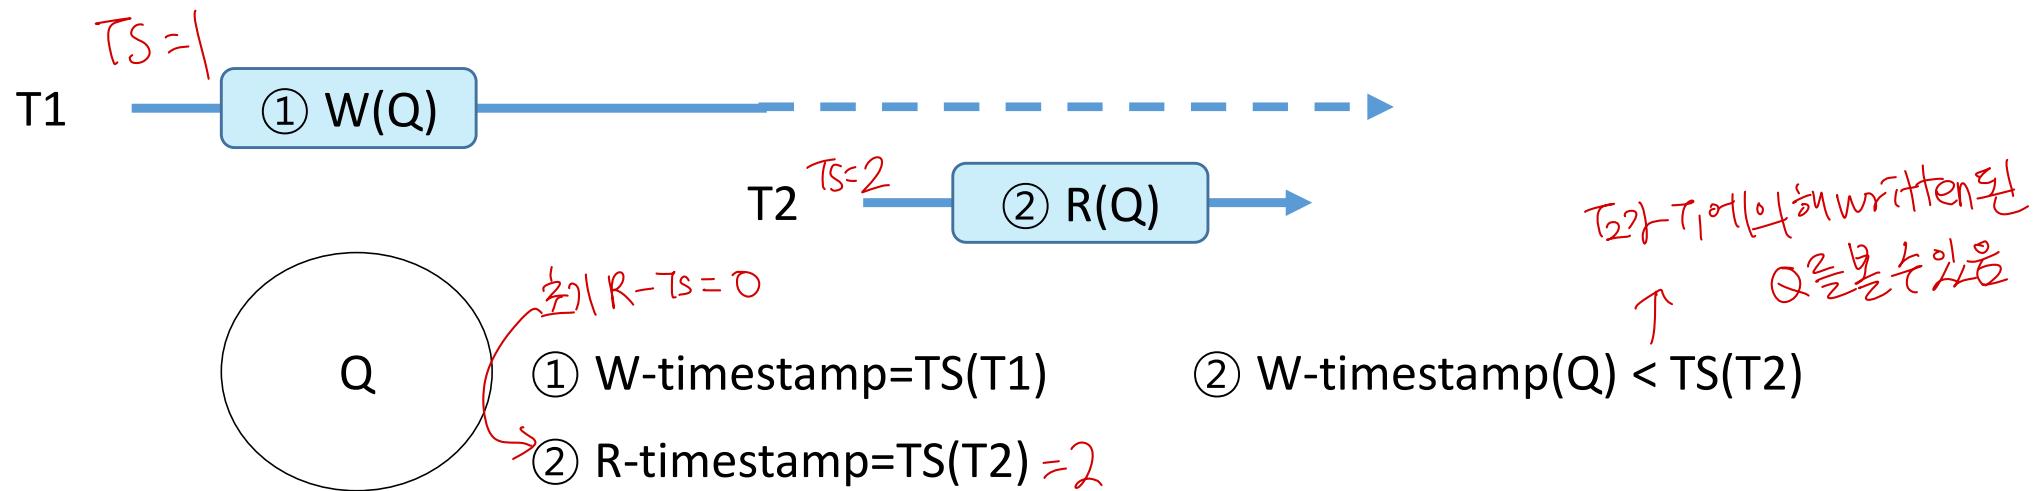
- If $W\text{-timestamp}(Q) \leq TS(T_i)$
 - the **read is allowed**, and
 - R-timestamp(Q) is set to $\max(R\text{-timestamp}(Q), TS(T_i))$.

T_i 가 old 일 때

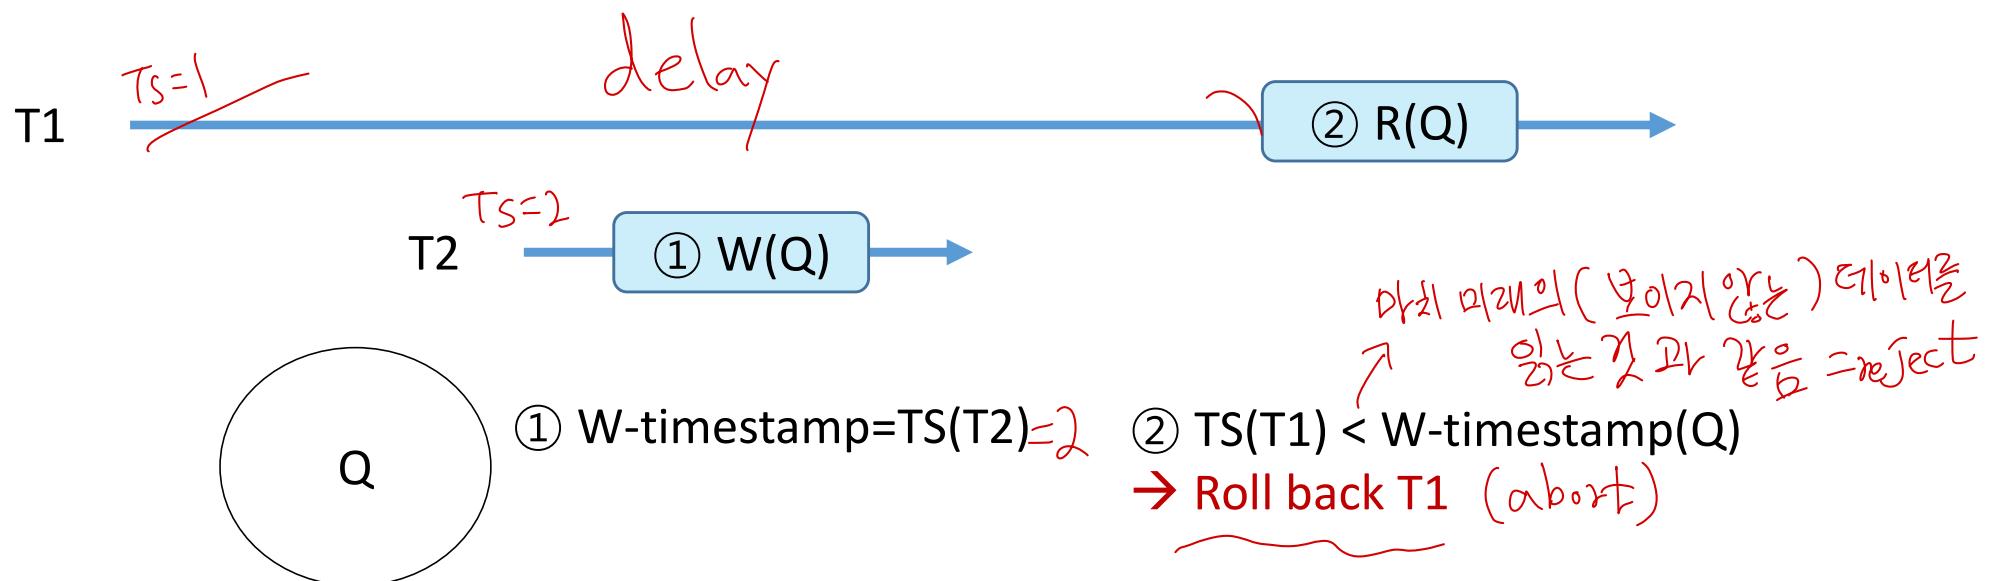
↳ 읽었으니 새로 갱신

- If $TS(T_i) \leq W\text{-timestamp}(Q)$
 - T_i needs to read Q that was already overwritten.
 - Reject read and rollback T_i**

Timestamp-Ordering Protocols (Cont.)



vs.



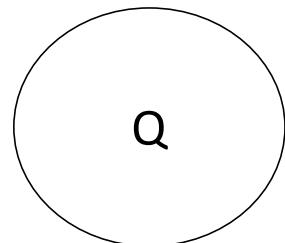
Timestamp-Ordering Protocols (Cont.)

- Suppose that transaction T_i issues **write(Q)**.

- $T_i = \text{old age}$

 - If $\text{TS}(T_i) < \text{R-timestamp}(Q)$ *T_{young} 은 Q 를 전부 T_i 가 write할 거라고 생각하는*
 - DBMS assumed that T_i would **never write Q**
 - Younger transaction already **read Q**
 - **Reject write and rollback T_i**
- If $\text{TS}(T_i) < \text{W-timestamp}(Q)$
 - DBMS assumed that T_i would **never write Q**
 - Younger transaction already **wrote Q** .
 - **Reject write and rollback T_i**
- Otherwise
 - Allow write**
 - Set W-timestamp(Q) to $\text{TS}(T_i)$

Timestamp-Ordering Protocols (Cont.)



① R-timestamp = TS(T2)

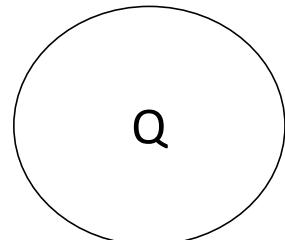
② TS(T1) < R-timestamp(Q)
→ Roll back T1 *{Simple Rule}*

T1

delayed write operation



Read \leq write \in
old tx \geq delayed tx
 \Rightarrow Roll back !!



① W-timestamp = TS(T2)

② TS(T1) < W-timestamp(Q)
→ Roll back T1

Example: Schedule Under TSO

- Is this schedule valid under TSO?

Assume $TS(T_{25}) = 25$ and
 $TS(T_{26}) = 26$

	R-TS	W-TS
A	0	0
B	0	0

T_{25}	T_{26}
read(B)	
$: RTS(B) \rightarrow 25$	$read(B) : RTS(B) \rightarrow 25$ $B := B - 50$
	$write(B) : WTS(B) \rightarrow 26$
read(A)	$: RTS(A) \rightarrow 25$ $\xrightarrow{\text{write}} R-TS, W-TS$
display($A + B$)	$read(A) : RTS(A) \rightarrow 26$ $\xrightarrow{\text{display}} R-TS, W-TS$ $A := A + 50$ $write(A)$ $display(A + B)$

Another Example Under TSO

T_1	T_2	T_3	T_4	T_5
<p>read (Y)</p> <p>read (X)</p>	<p>read (Y)</p> <p>read (Z)</p> <p>abort (delay S(M))</p>	<p>write (Y)</p> <p>write (Z)</p>		<p>read (X)</p> <p>read (Z)</p>

Handwritten annotations:

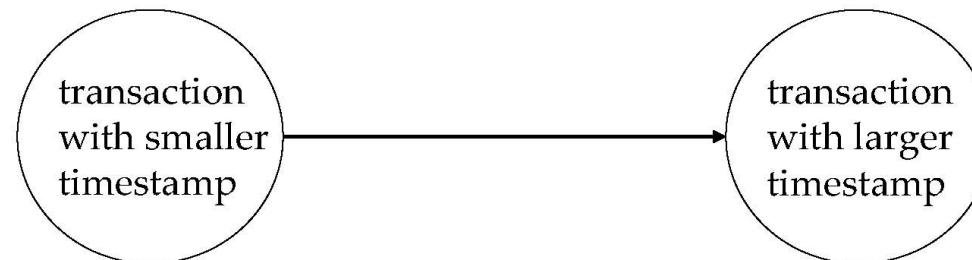
- T_1 : read (Y) circled in red.
- T_2 : read (X) circled in red.
- T_3 : write (Y)
- T_3 : write (Z)
- T_2 : abort (delay S(M)) in a red box.
- T_4 : read (W)
- T_4 : write (W)
- T_4 : abort (in a red box) with an arrow pointing to the previous write (W) in T_4 .
- T_4 : handwritten note: $WTS \leftarrow 4, RTS \leftarrow 4$.
- T_4 : handwritten note: $(\because T_4 \text{ is not committed})$ and $T_3 \text{ writes } W \text{ before } T_4 \text{ starts}$.
- T_5 : read (X) underlined in red.
- T_5 : read (Z)
- T_5 : write (Y)
- T_5 : write (Z)

	R-TS	W-TS
X	0	0
Y	0	0
Z	0	0
W	0	0

Correctness of Timestamp-Ordering Protocol

■ TSO guarantees serializability

- All edges in the precedence graph follow timestamp order
- No cycle in the graph



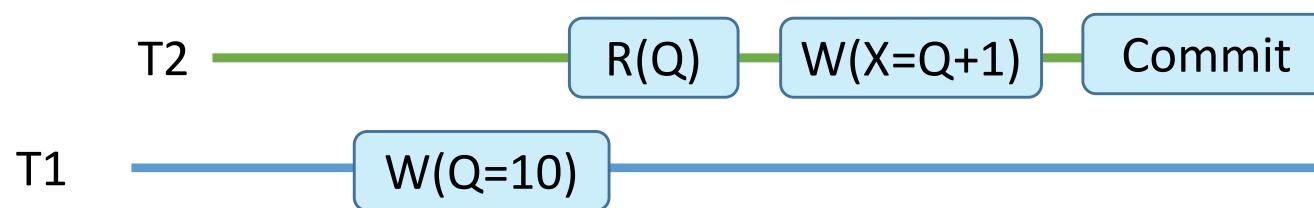
TX 간의 순서가 그들의 TS 값 순서에
따라 결정된다
⇒ serializability 를 보장한다!

■ TSO ensures freedom from deadlock

- no transaction ever waits

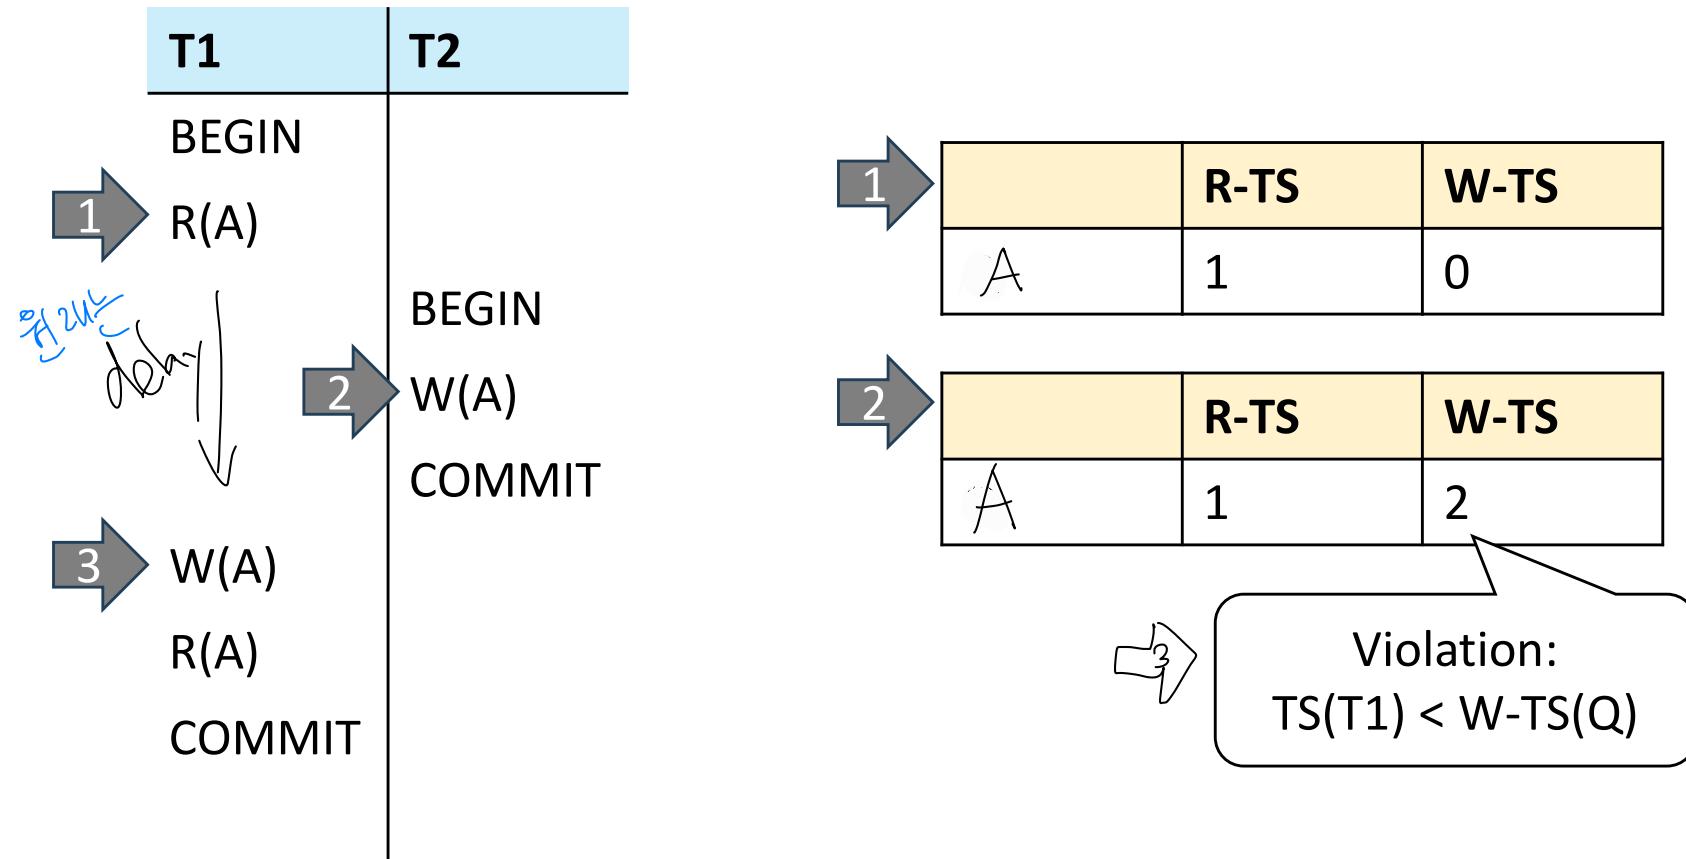
■ Limitations

- may not be cascade-free
- may not be recoverable
 - A schedule is recoverable if a transaction commits only after the transactions it depends on have committed.

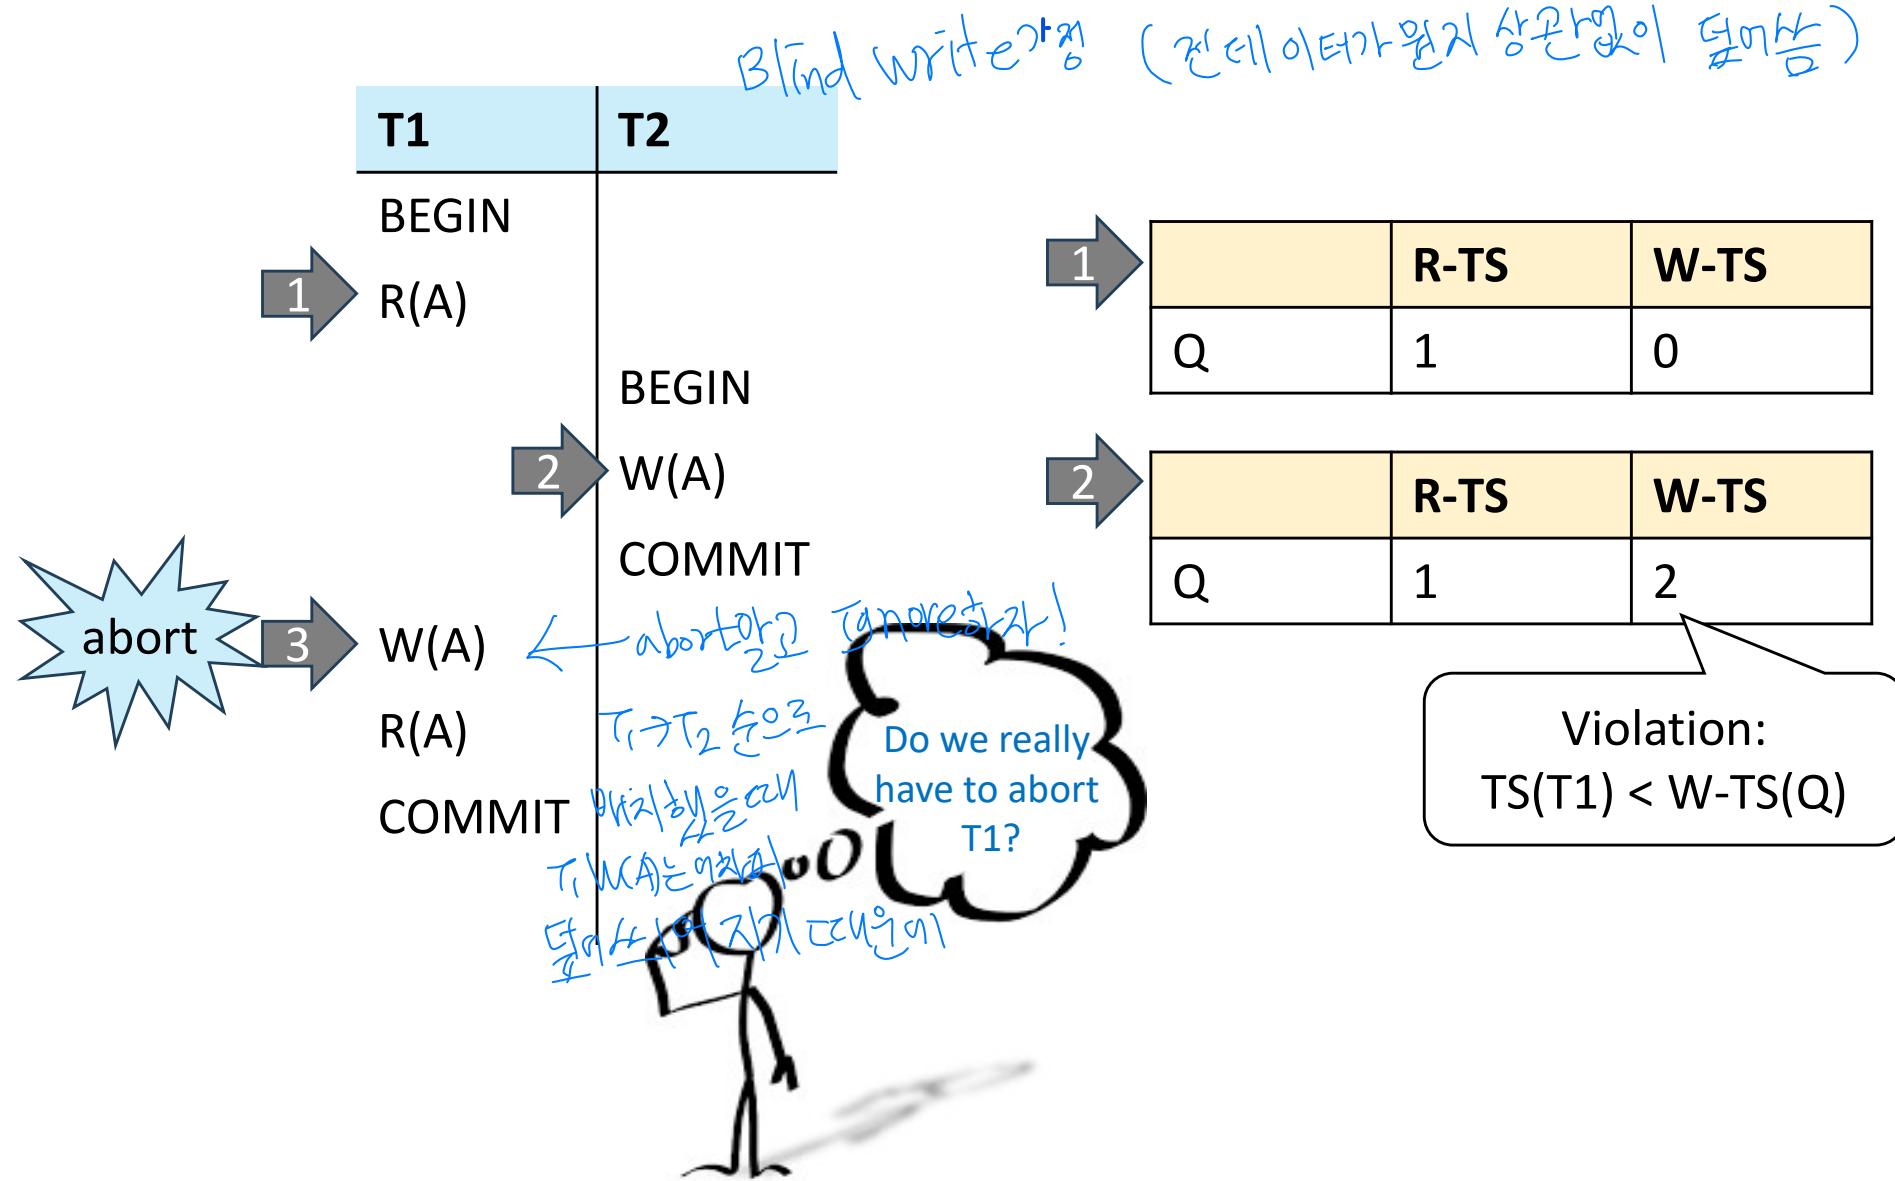


이상황이면 T2는 회복불가능
(Commit을 나중에
하거나 rollback을 해야하는 경우)
Cascade rollback
TSWE
T₂는 Abort되는 경우

Yet Another Example Under TSO



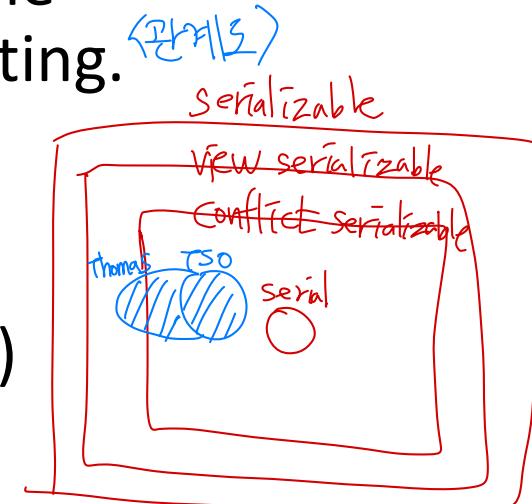
Yet Another Example Under TSO



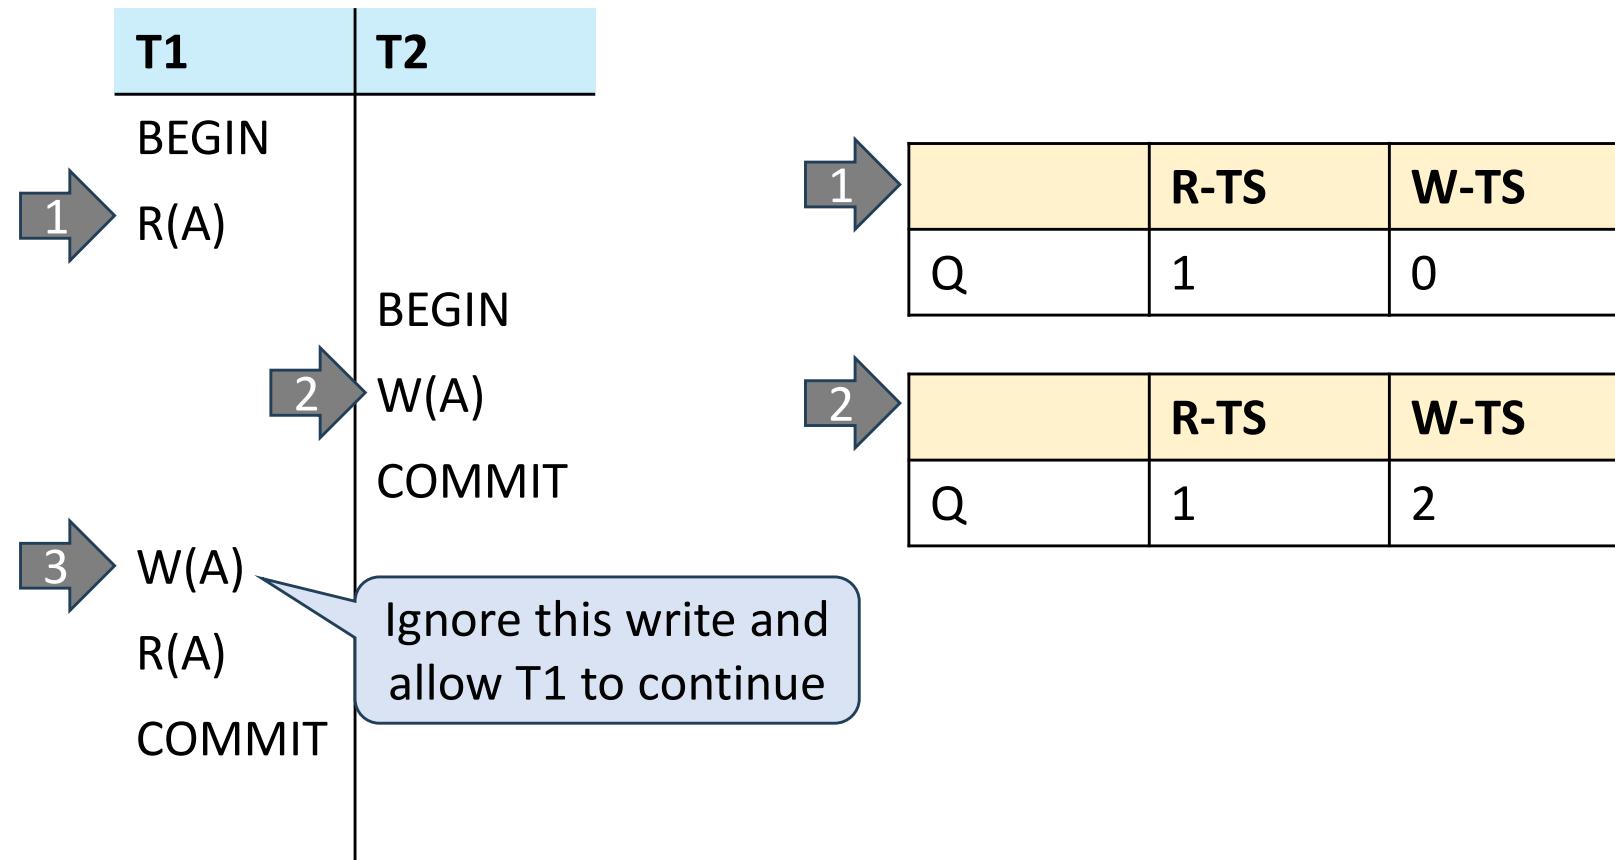
Thomas' Write Rule

- if $TS(T_i) < R\text{-timestamp}(Q)$
 - Abort and restart T_i
- if $TS(T_i) < W\text{-timestamp}(Q)$
 - **Thomas Write Rule:** Ignore the write to allow the transaction to continue executing without aborting.
하지만 abort 되지만, 수행됨
 - This violates timestamp order of T_i .
- Else
 - Allow T_i to write Q and update $W\text{-timestamp}(Q)$
- Thomas' Write Rule allows greater potential concurrency.
 - Allows some **view-serializable schedules** that are not conflict-serializable.

young tx의 write 가능하면
old tx의 write operation을 무시해도된다!



Thomas' Write Rule

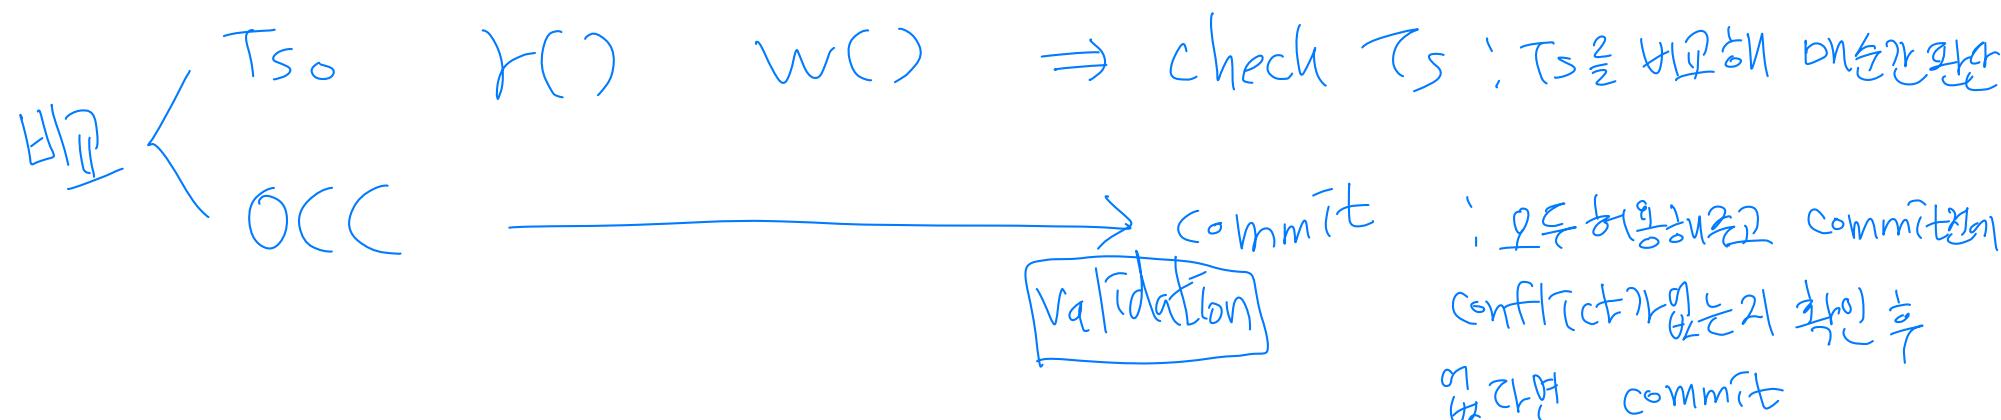


Validation-based Concurrency Control (Optimistic Concurrency Control - OCC)



Validation-Based Protocol (OCC)

- Idea: can we use **commit time** as serialization order?
(한국어)
- To do so:
 - Postpone writes to end of transaction
 - Keep track of data items read/written by transaction
 - Validation** performed at commit time, detect any out-of-serialization order reads/writes \Rightarrow conflict *commit 전에 validation*
- Also called as **Optimistic Concurrency Control (OCC)** since transaction executes fully in the hope that all will go well during validation



Validation-Based Protocol (OCC)

- Execution of transaction T_i is done in three phases.

1. Read and execution phase:

Transaction T_i writes only temporary local variables

2. Validation phase

Transaction T_i performs a “validation test” to determine if local variables can be written without violating serializability.

3. Write phase

If T_i is validated, the updates are applied to the database; otherwise, T_i is rolled back.
dirty page update

- We assume for simplicity that the validation and write phase occur together, atomically and serially
 - I.e., only one transaction executes validation/write at a time.
 - But, the three phases of concurrently executing transactions can be interleaved.

Validation-Based Protocol (OCC)

- Each transaction T_i has 3 timestamps
 - $\text{StartTS}(T_i)$: the time when T_i started its execution
 - $\text{ValidationTS}(T_i)$: the time when T_i entered its validation phase
 - $\text{FinishTS}(T_i)$: the time when T_i finished its write phase
- Validation tests use above timestamps and read/write sets to ensure that *serializability order is determined by validation time*
 - Thus, $\text{TS}(T_i) = \text{ValidationTS}(T_i)$
- Validation-based protocol gives greater degree of concurrency than locking/TSO if probability of conflicts is low.

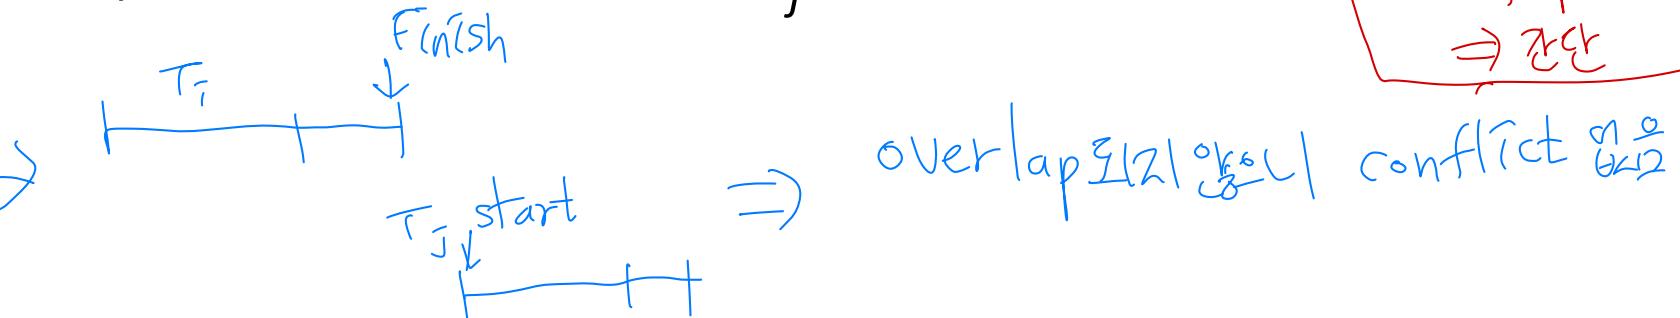
Validation Test for Transaction T_j

- If for all T_i with $\text{TS}(T_i) < \text{TS}(T_j)$ either one of the following condition holds:

- $\text{finishTS}(T_i) < \text{startTS}(T_j)$
 - execution is not concurrent
 - $\text{startTS}(T_j) < \text{finishTS}(T_i) < \text{validationTS}(T_j)$ and T_j does not read any data item written by T_i ,
 - there is no dependency (validation \rightarrow validation \rightarrow validation)
- ↑ young tx
Ol 2nd
장우자(여)
모두 Fail

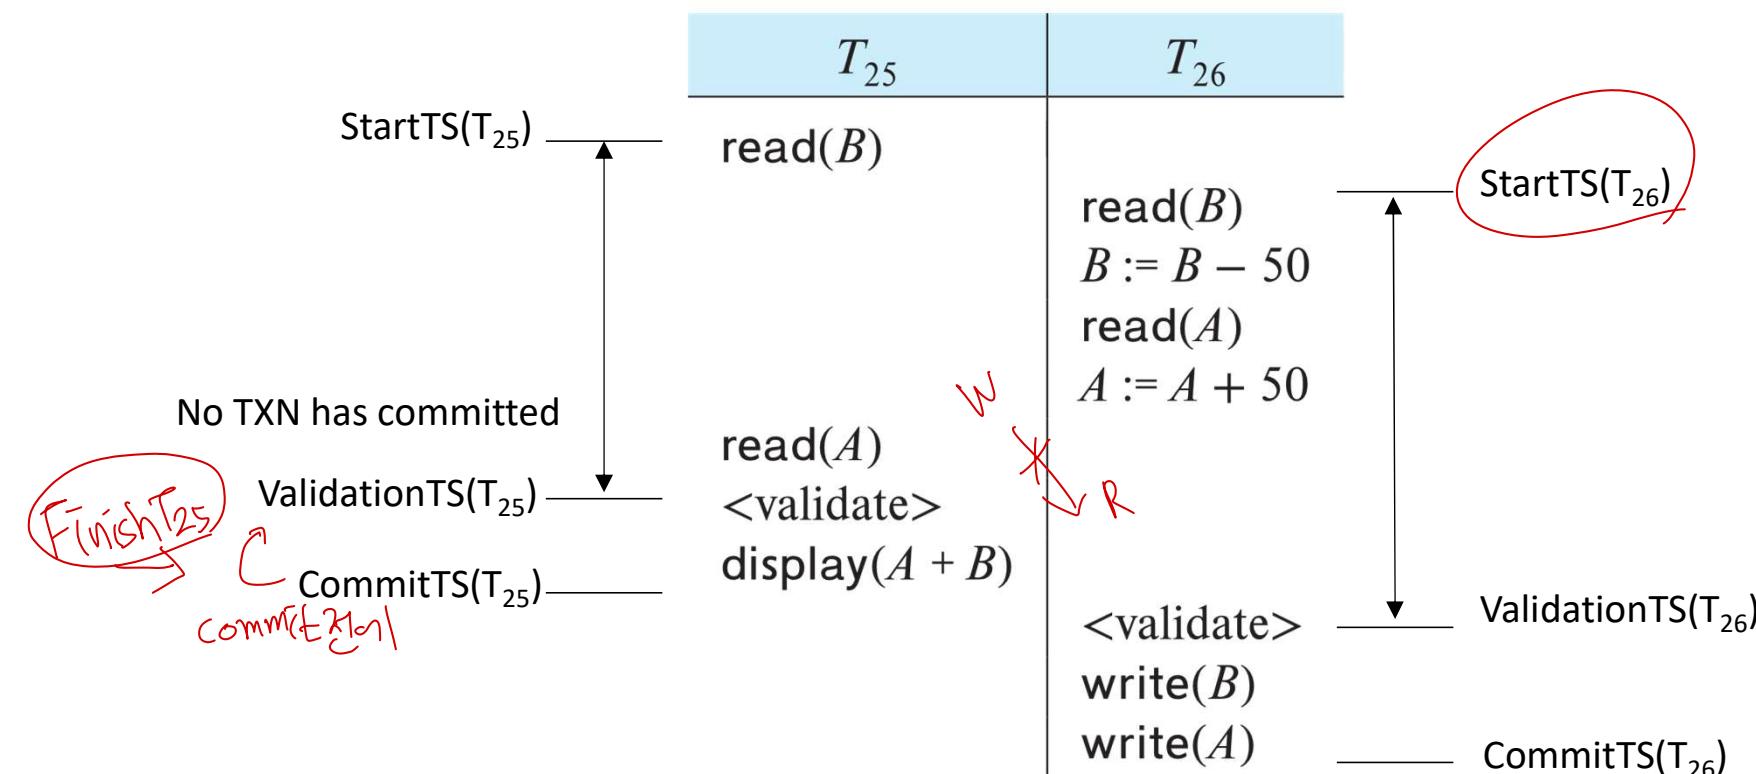
then validation succeeds and T_j can be committed.

- Otherwise, validation fails and T_j is aborted.



Schedule Produced by Validation

- Example of schedule produced using validation



startTS(T_{26}) < finishTS(T_{25}) < validationTS(T_{26})

T_{25} has committed, but it is read-only

Validation-Based Protocol (OCC) - Example

T1	T2
BEGIN	StartTS(T1)=1
Read-Exec Phase	
R(A)	BEGIN Read-Exec Phase
	R(A)
	Validate
	Write
	Commit
W(A)	
Validate	
Write	
Commit	

Database

Data	Value	W-TS
A	0	0
-	-	

Validation-Based Protocol (OCC) - Example

T1	T2
BEGIN	StartTS(T1)=1
Read-Exec Phase	
R(A)	BEGIN
	Read-Exec Phase
	R(A)
	Validate
	Write
	Commit
W(A)	
Validate	
Write	
Commit	

Database

Data	Value	W-TS
A	0	0
-	-	

Copy

T1's Workspace

Data	Value
A	0
-	-

Validation-Based Protocol (OCC) - Example

T1	T2
BEGIN	StartTS(T1)=1
Read-Exec Phase	
R(A)	BEGIN Read-Exec Phase
	StartTS(T2)=2
	R(A)
	Validate
	Write
	Commit
W(A)	
Validate	
Write	
Commit	

Database		
Data	Value	W-TS
A	0	0
-	-	

T1's Workspace

Data	Value
A	0
-	-

T2's Workspace

Data	Value
A	0
-	-

Validation-Based Protocol (OCC) - Example



Database		
Data	Value	W-TS
A	0	0
-	-	

T1's Workspace	
Data	Value
A	0
-	-

T2's Workspace	
Data	Value
A	0
-	-

Validation-Based Protocol (OCC) - Example



Database		
Data	Value	W-TS
A	0	0
-	-	

T1's Workspace		T2's Workspace	
Data	Value	Data	Value
A	0	A	0
-	-	-	-

Validation-Based Protocol (OCC) - Example



Database		
Data	Value	W-TS
A	0	0
-	-	

T1's Workspace	
Data	Value
A	1
-	-

Validation-Based Protocol (OCC) - Example



Database

Data	Value	W-TS
A	1	5
-	-	

T1's Workspace

Data	Value
A	1
-	-

$S_{T_1} \subset F_{T_2} \subset V_{T_1} + W-R$ 의 조건을 만족 \Rightarrow Valid

Validation-Based Protocol (OCC) – Another Example

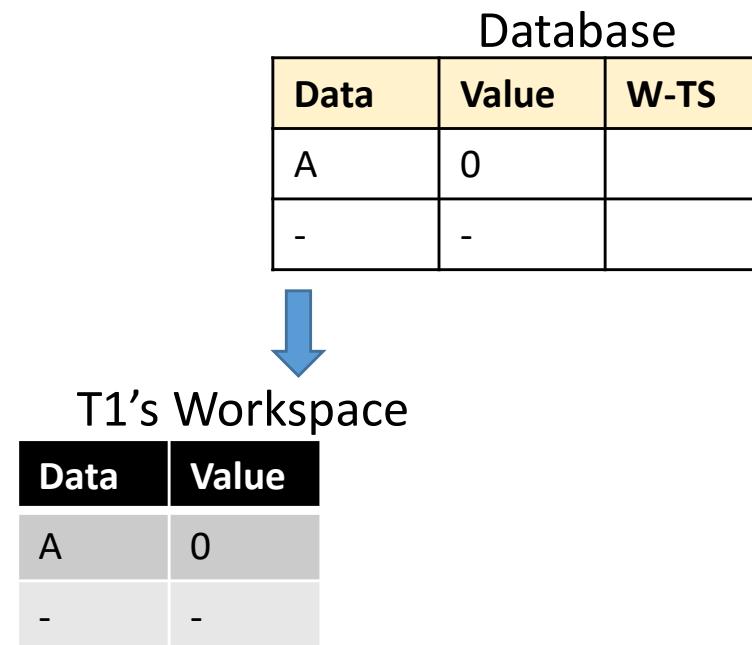
T1	T2
BEGIN	StartTS(T1)=1
Read-Exec Phase	
Read(A)	BEGIN Read-Exec Phase
	Read(A)
	A := A+1
	Validate
	Write(A)
	Commit
A := A+1	
Validate	
Write(A)	
Commit?	

Database

Data	Value	W-TS
A	0	
-	-	

Validation-Based Protocol (OCC) – Another Example

T1	T2
BEGIN	StartTS(T1)=1
Read-Exec Phase	
Read(A)	BEGIN
	Read-Exec Phase
	Read(A)
	A := A+1
	Validate
	Write(A)
	Commit
A := A+1	
Validate	
Write(A)	
Commit?	



Validation-Based Protocol (OCC) – Another Example

T1	T2
BEGIN	StartTS(T1)=1
Read-Exec Phase	
Read(A)	BEGIN Read-Exec Phase
	StartTS(T2)=2
	Read(A)
	A := A+1
	Validate
	Write(A)
	Commit
A := A+1	
Validate	
Write(A)	
Commit?	

Database		
Data	Value	W-TS
A	0	
-	-	

T1's Workspace

Data	Value
A	0
-	-

T2's Workspace

Data	Value
A	0
-	-

Validation-Based Protocol (OCC) – Another Example

T1	T2
BEGIN	StartTS(T1)=1
Read-Exec Phase	
Read(A)	BEGIN
	StartTS(T2)=2
	Read-Exec Phase
	Read(A)
	A := A+1
	Validate
	Write(A)
	Commit
A := A+1	
Validate	
Write(A)	
Commit?	

Handwritten annotations:

- T1**: FINISH, over H2H1, over H2H2, over H2H3, over H2H4, over H2H5, over H2H6, over H2H7, over H2H8.
- T2**: over H2H1, over H2H2, over H2H3, over H2H4, over H2H5, over H2H6, over H2H7, over H2H8.
- Write**: over H2H1, over H2H2, over H2H3, over H2H4, over H2H5, over H2H6, over H2H7, over H2H8.
- Valid**: over H2H1, over H2H2, over H2H3, over H2H4, over H2H5, over H2H6, over H2H7, over H2H8.

Database		
Data	Value	W-TS
A	0	
-	-	

T1's Workspace	
Data	Value
A	0
-	-

T2's Workspace	
Data	Value
A	1
-	-

Validation-Based Protocol (OCC) – Another Example

T1	T2
BEGIN	StartTS(T1)=1
Read-Exec Phase	
Read(A)	BEGIN
	StartTS(T2)=2
	Read-Exec Phase
	Read(A)
	A := A+1
	Validate
	ValTS(T2)=3
	Write(A)
	Commit
	FinTS(T2)=4
A := A+1	
Validate	
Write(A)	
Commit?	

Database		
Data	Value	W-TS
A	1	3
-	-	

T1's Workspace

Data	Value
A	0
-	-

T2's Workspace

Data	Value
A	1
-	-

Validation-Based Protocol (OCC) – Another Example



Database		
Data	Value	W-TS
A	1	3
-	-	

T1's Workspace

Data	Value
A	1
-	-

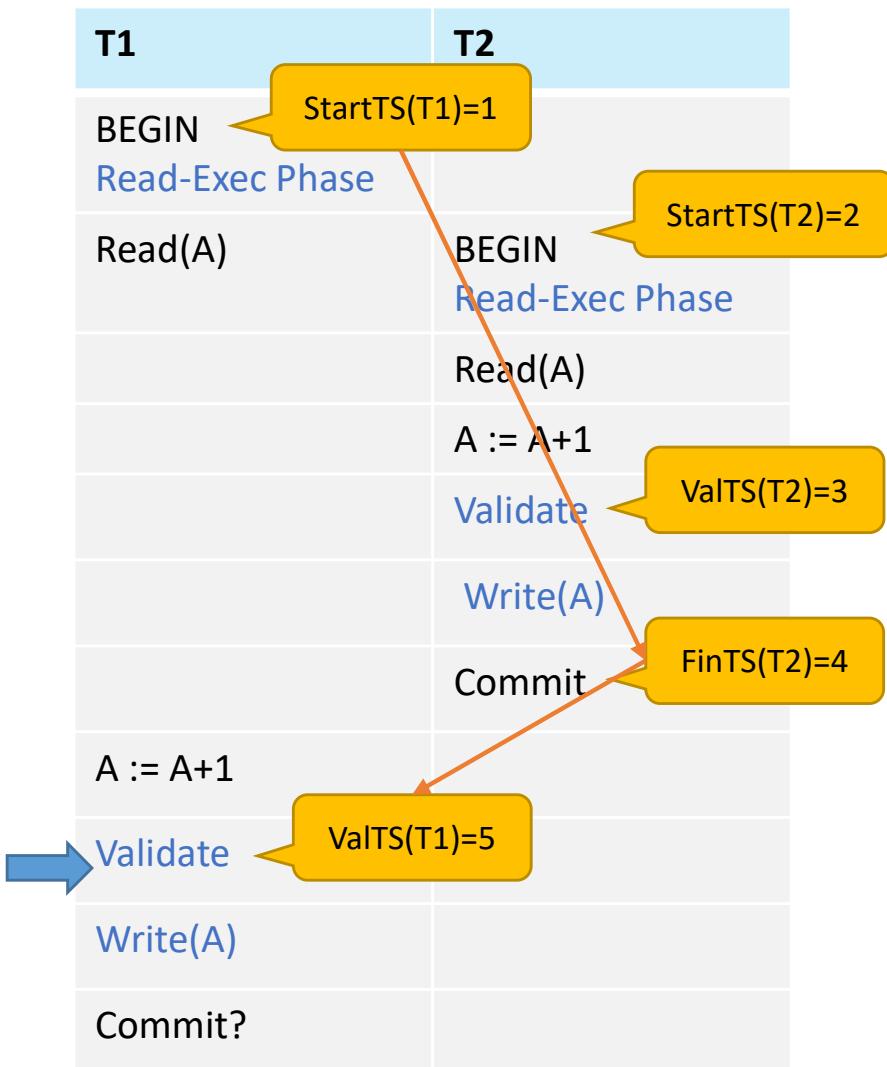
→ A := A+1

Validate

Write(A)

Commit?

Validation-Based Protocol (OCC) – Another Example



Validation-Based Protocol (OCC) – Another Example



Database		
Data	Value	W-TS
A	1	3
-	-	

T1's Workspace

Data	Value
A	1
-	-

W-R 의존성 존재

$\text{startTS}(T_j) < \text{finishTS}(T_i) < \text{validationTS}(T_j)$ and T_j has read a data item written by T_i ,
 → There is dependency!