



# Database Systems

## Lecture 20 – Chapter 18: Concurrency Control (Part III)

Beomseok Nam (남범석)  
[bnam@skku.edu](mailto:bnam@skku.edu)

# Multi-Version Concurrency Control

# Multi-Version Concurrency Control (MVCC)

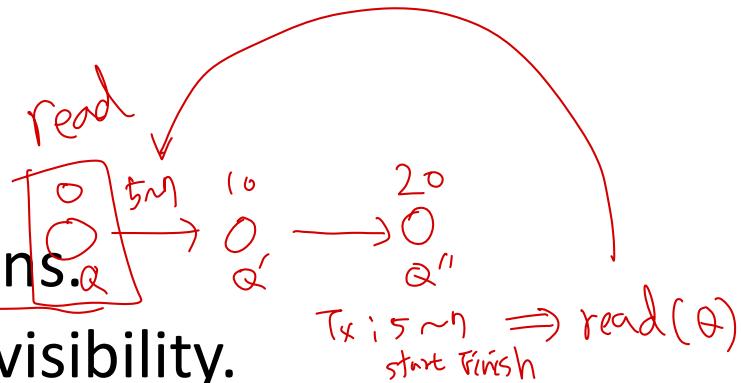
- The DBMS maintains ***multiple versions*** of an object:
  - When a transaction **writes** to an object, the DBMS **creates a new version** of that object.
  - When a transaction **reads** an object, it **reads the newest** version that existed when the transaction started.

~~Overwrite~~ Copy-on-write → young TXD  
latest version is 233  
version 221

# Multi-Version Concurrency Control (MVCC)

- Key ideas:

- Writers ~~do not block~~ readers.
- Readers ~~do not block~~ writers.
- Use timestamps to label versions.
- Use timestamps to determine visibility.
- Multi-versioning allows the DBMS to support time-travel queries.



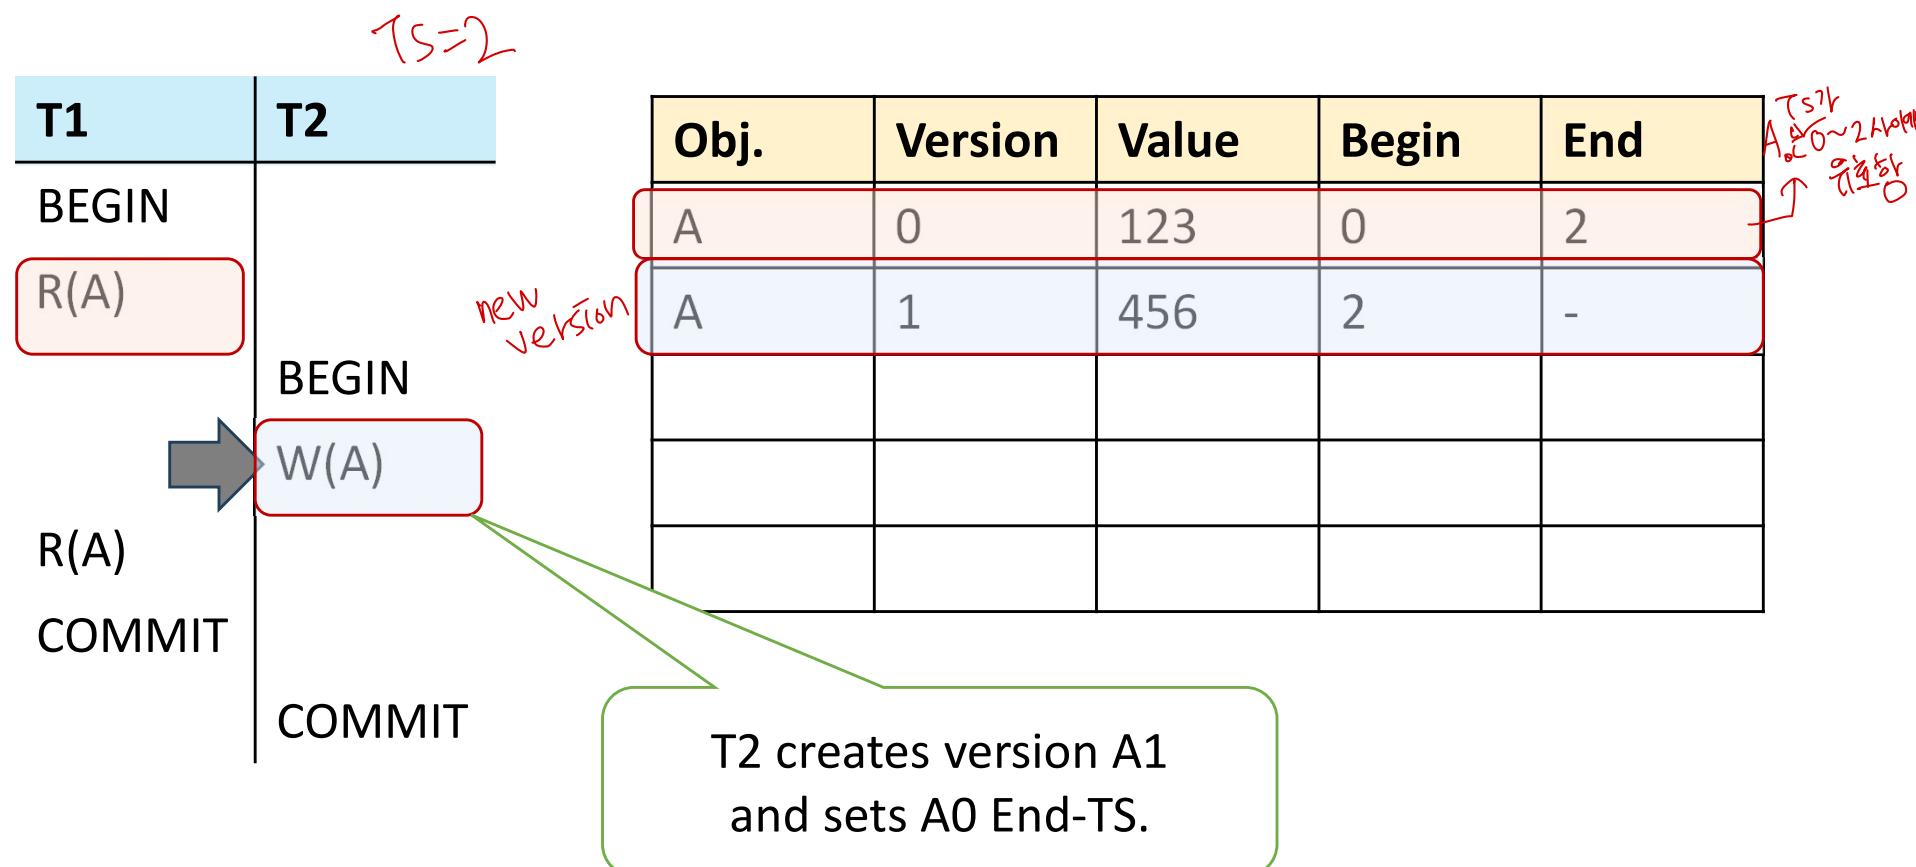
# MVCC – Example #1

T1	T2
BEGIN	
R(A)	
	BEGIN
	W(A)
R(A)	
COMMIT	
	COMMIT

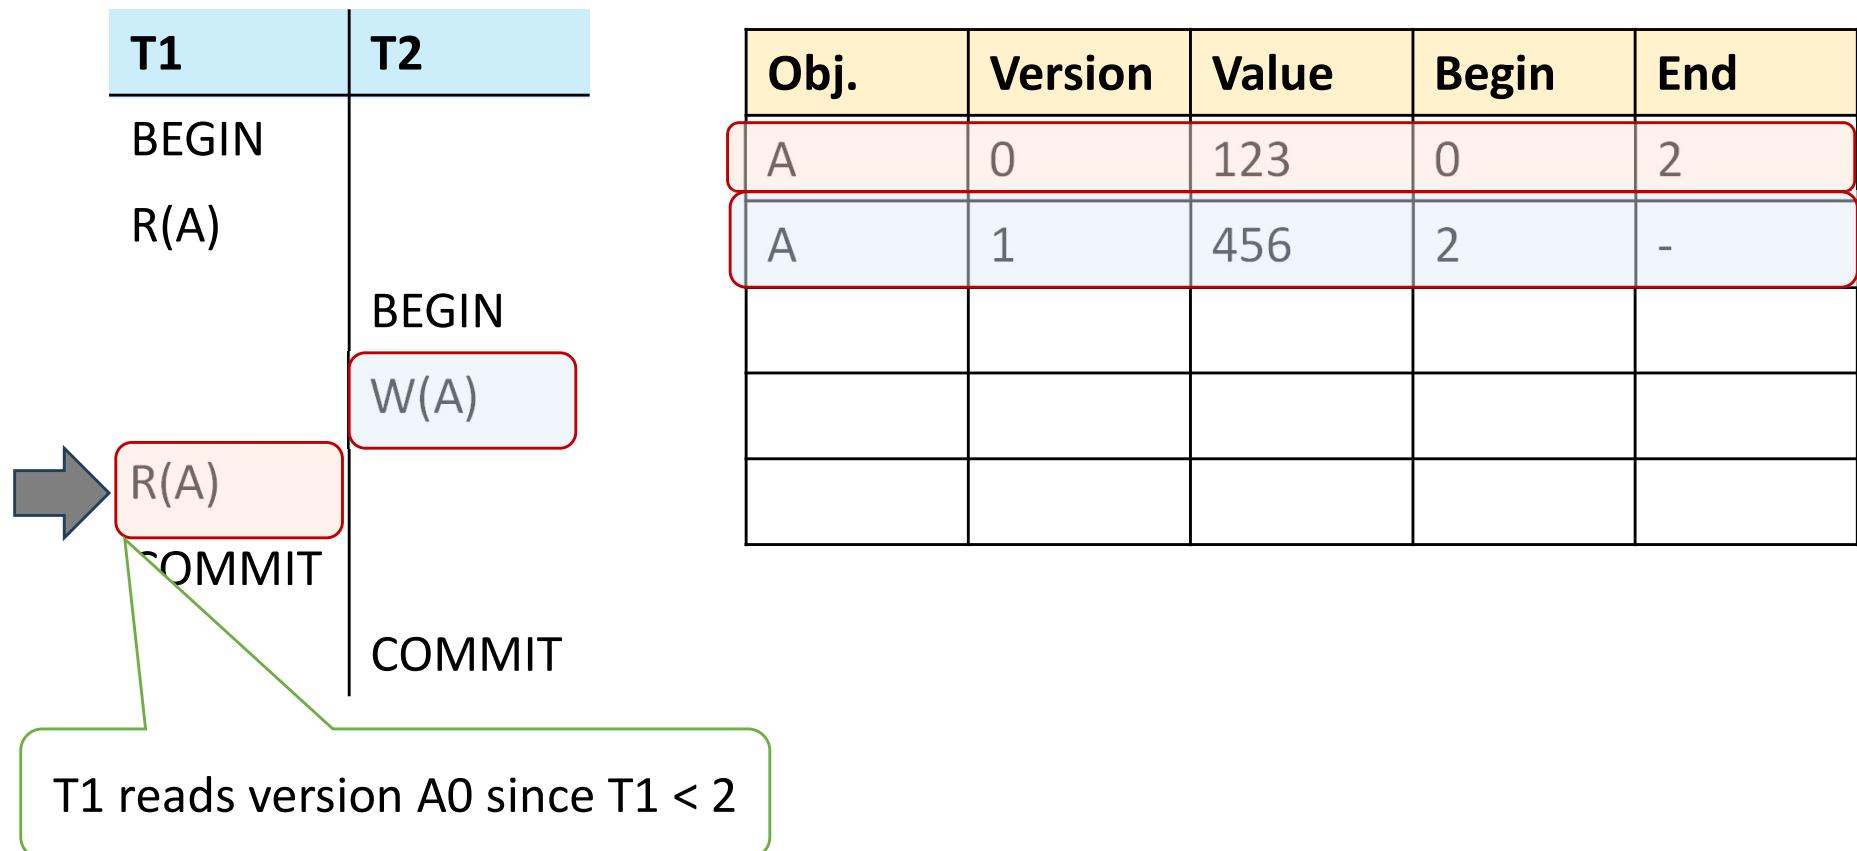
Obj.	Version	Value	Begin	End
A	0	123	0	-

New Version of  
A가 A의  
New Version

# MVCC – Example #1



# MVCC – Example #1



# MVCC – Example #2

T1	T2
BEGIN	
R(A)	
W(A)	BEGIN
	R(A)
	W(A)
R(A)	
COMMIT	

Obj.	Version	Value	Begin	End
A	0	123	0	-

# MVCC – Example #2

T1	T2
BEGIN	
R(A)	
W(A)	BEGIN
	R(A)
	W(A)
R(A)	
COMMIT	

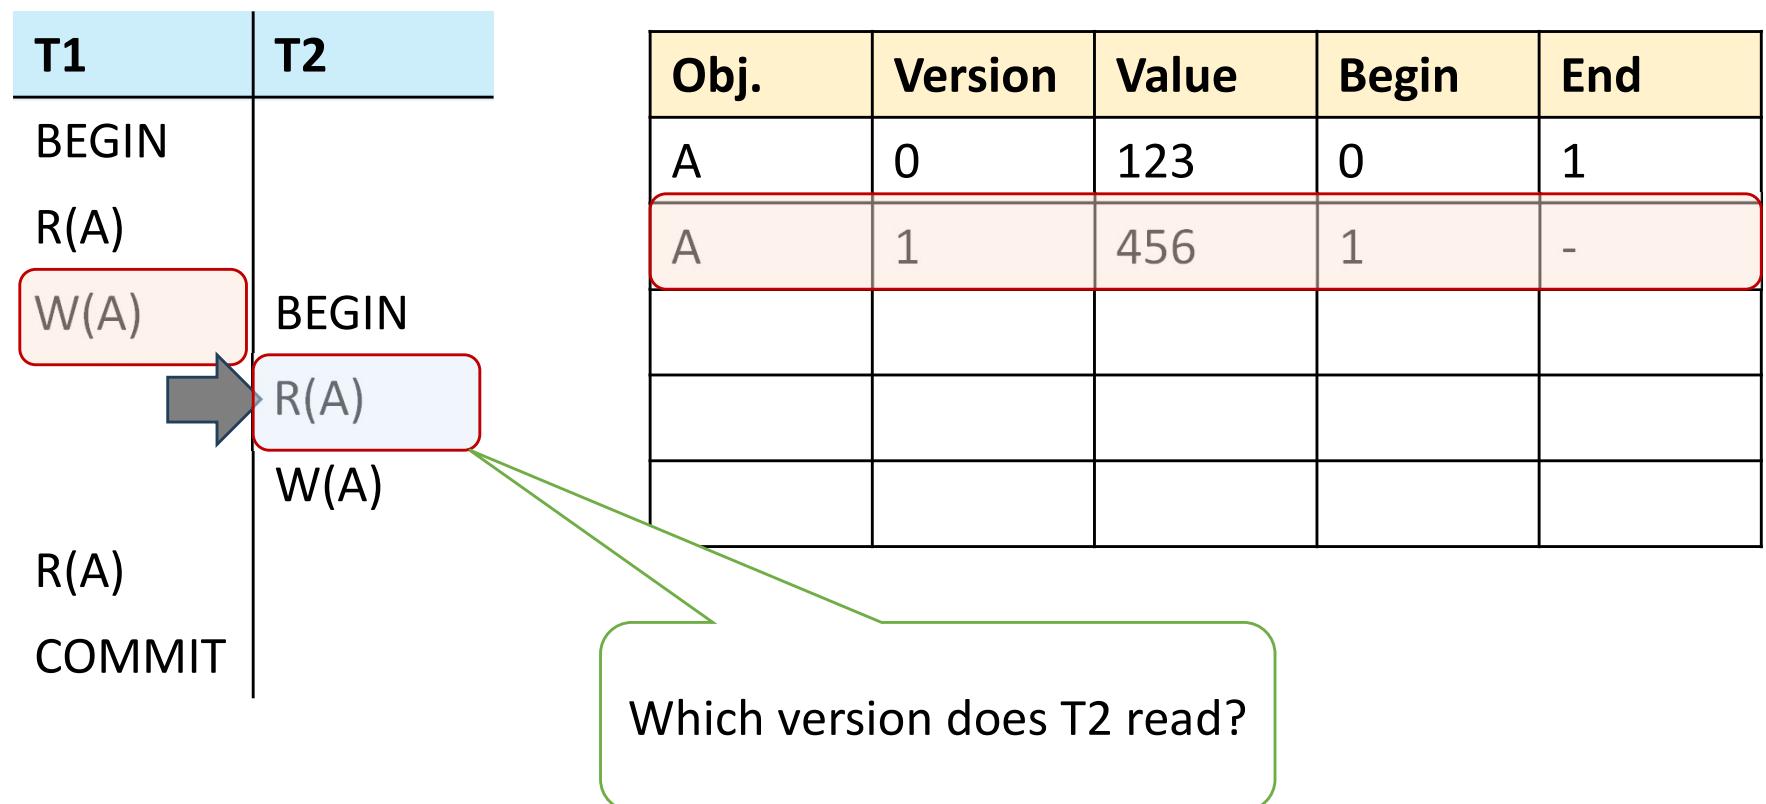
Obj.	Version	Value	Begin	End
A	0	123	0	1
A	1	456	1	-

# MVCC – Example #2

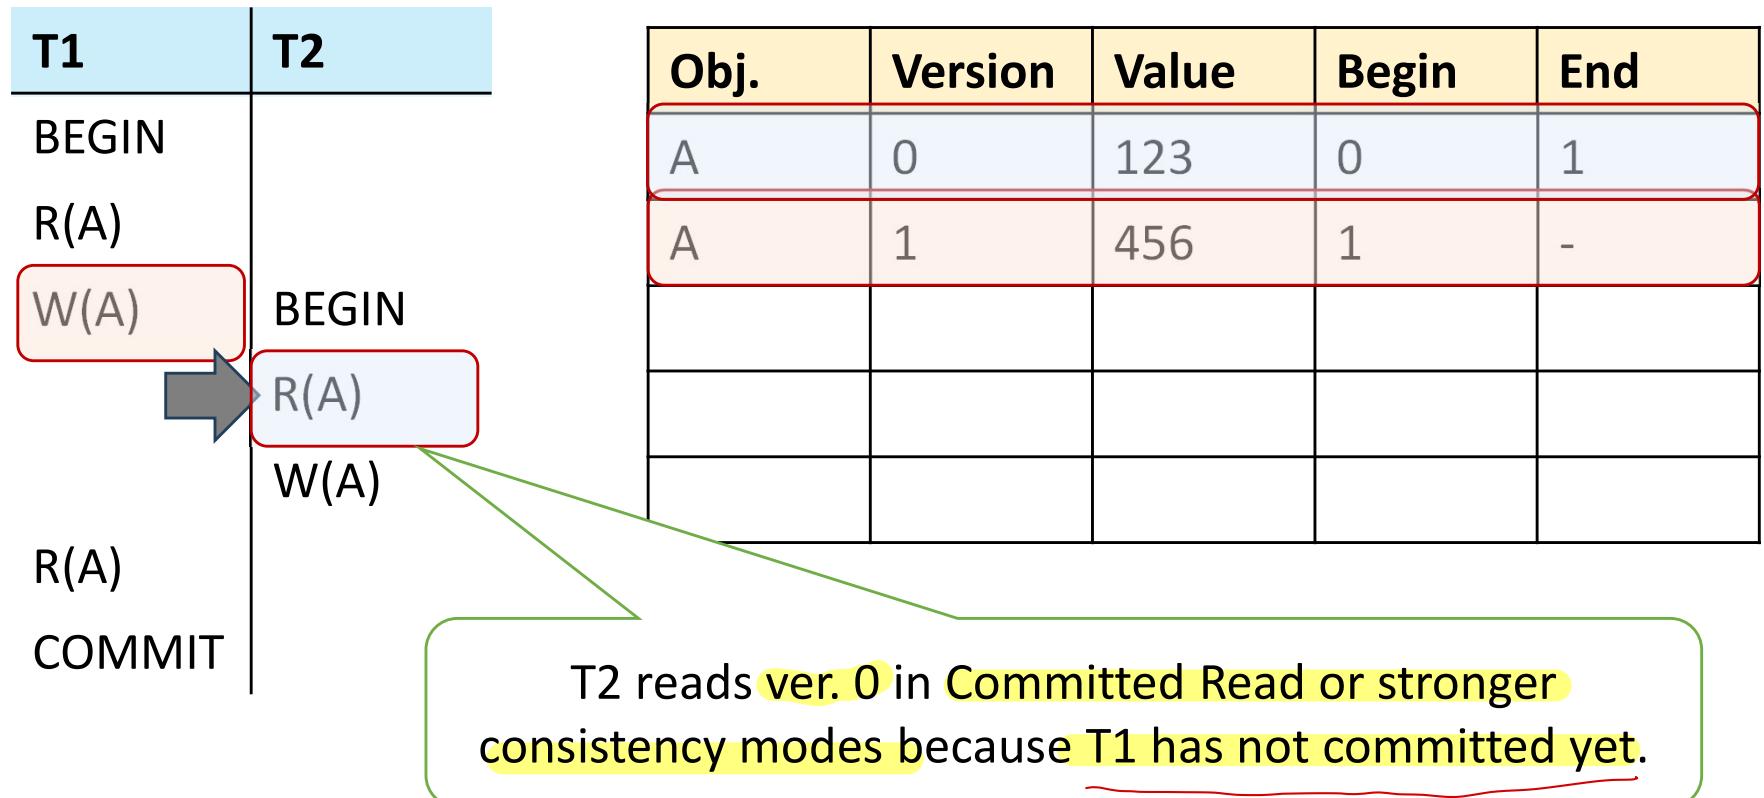
T1	T2
BEGIN	
R(A)	
W(A)	BEGIN
	R(A)
	W(A)
R(A)	
COMMIT	

Obj.	Version	Value	Begin	End
A	0	123	0	1
A	1	456	1	-

# MVCC – Example #2



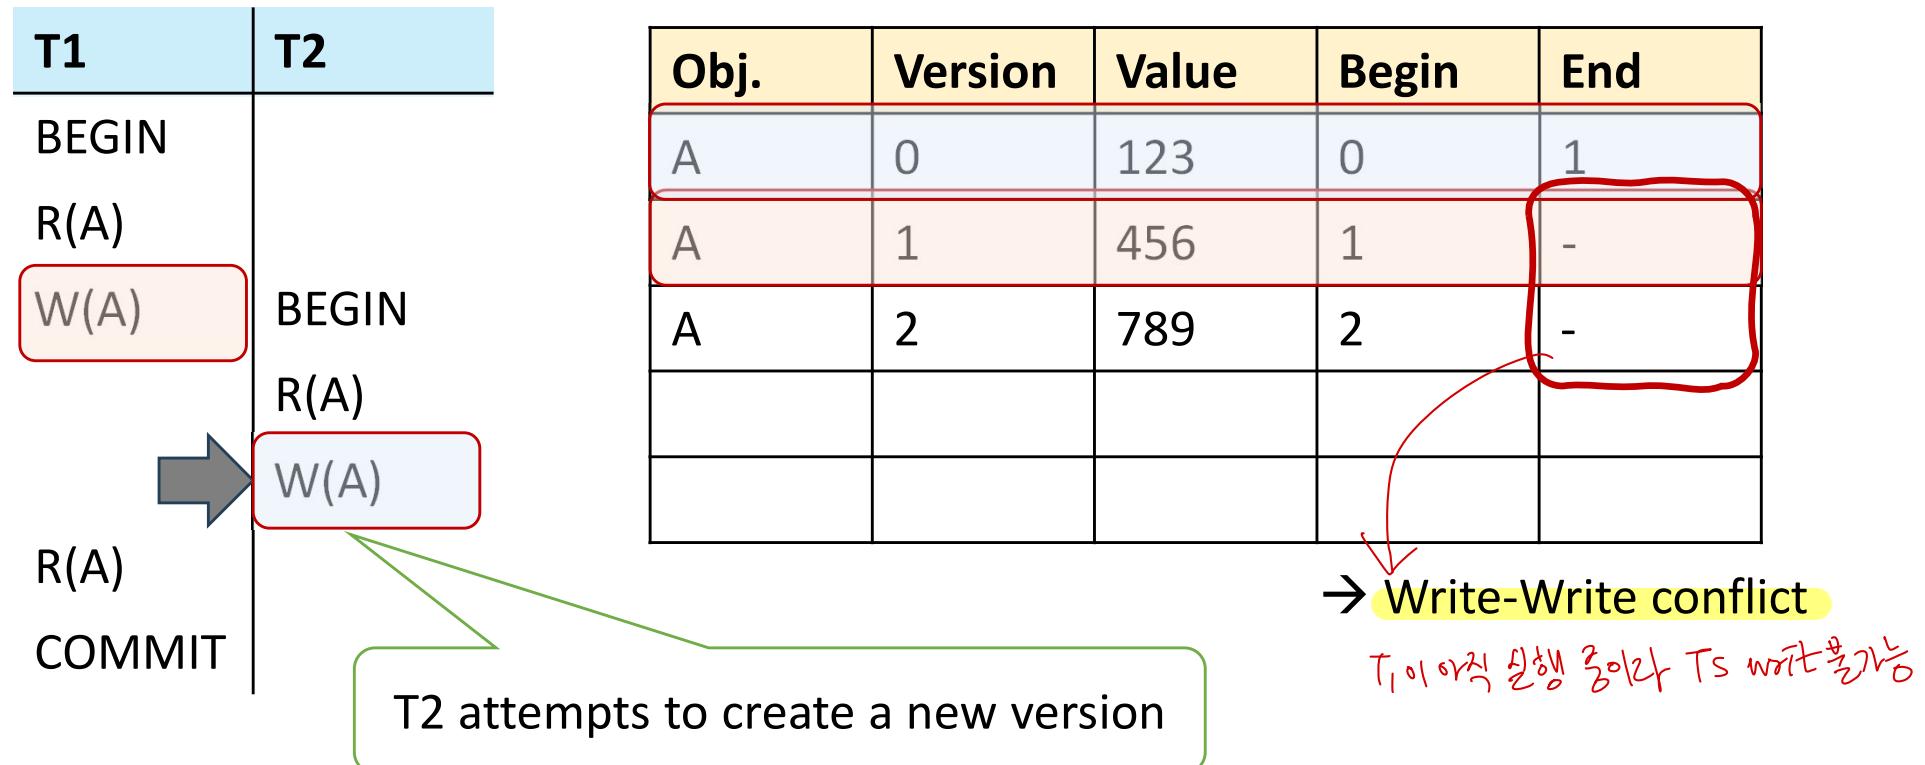
# MVCC – Example #2



T2 reads ver. 0 in Committed Read or stronger consistency modes because T1 has not committed yet.

⇒ T2는 dirty data에 접근하면 안됨  
⇒ old에 접근 (Version 0)

# MVCC – Example #2



# Multiversion Schemes

- Several variants to resolve conflicts:

- Multiversion Timestamp Ordering (MVTO) → TSO 이론
- Multiversion Two-Phase Locking (MV2PL) → 2PL 이론
- Snapshot isolation → first-committer-wins rule 이론

# Snapshot Isolation

# Snapshot Isolation

## Motivation

- OLAP (decision support query) transactions that read large amounts of data have concurrency conflicts with OLTP transactions that update a few rows

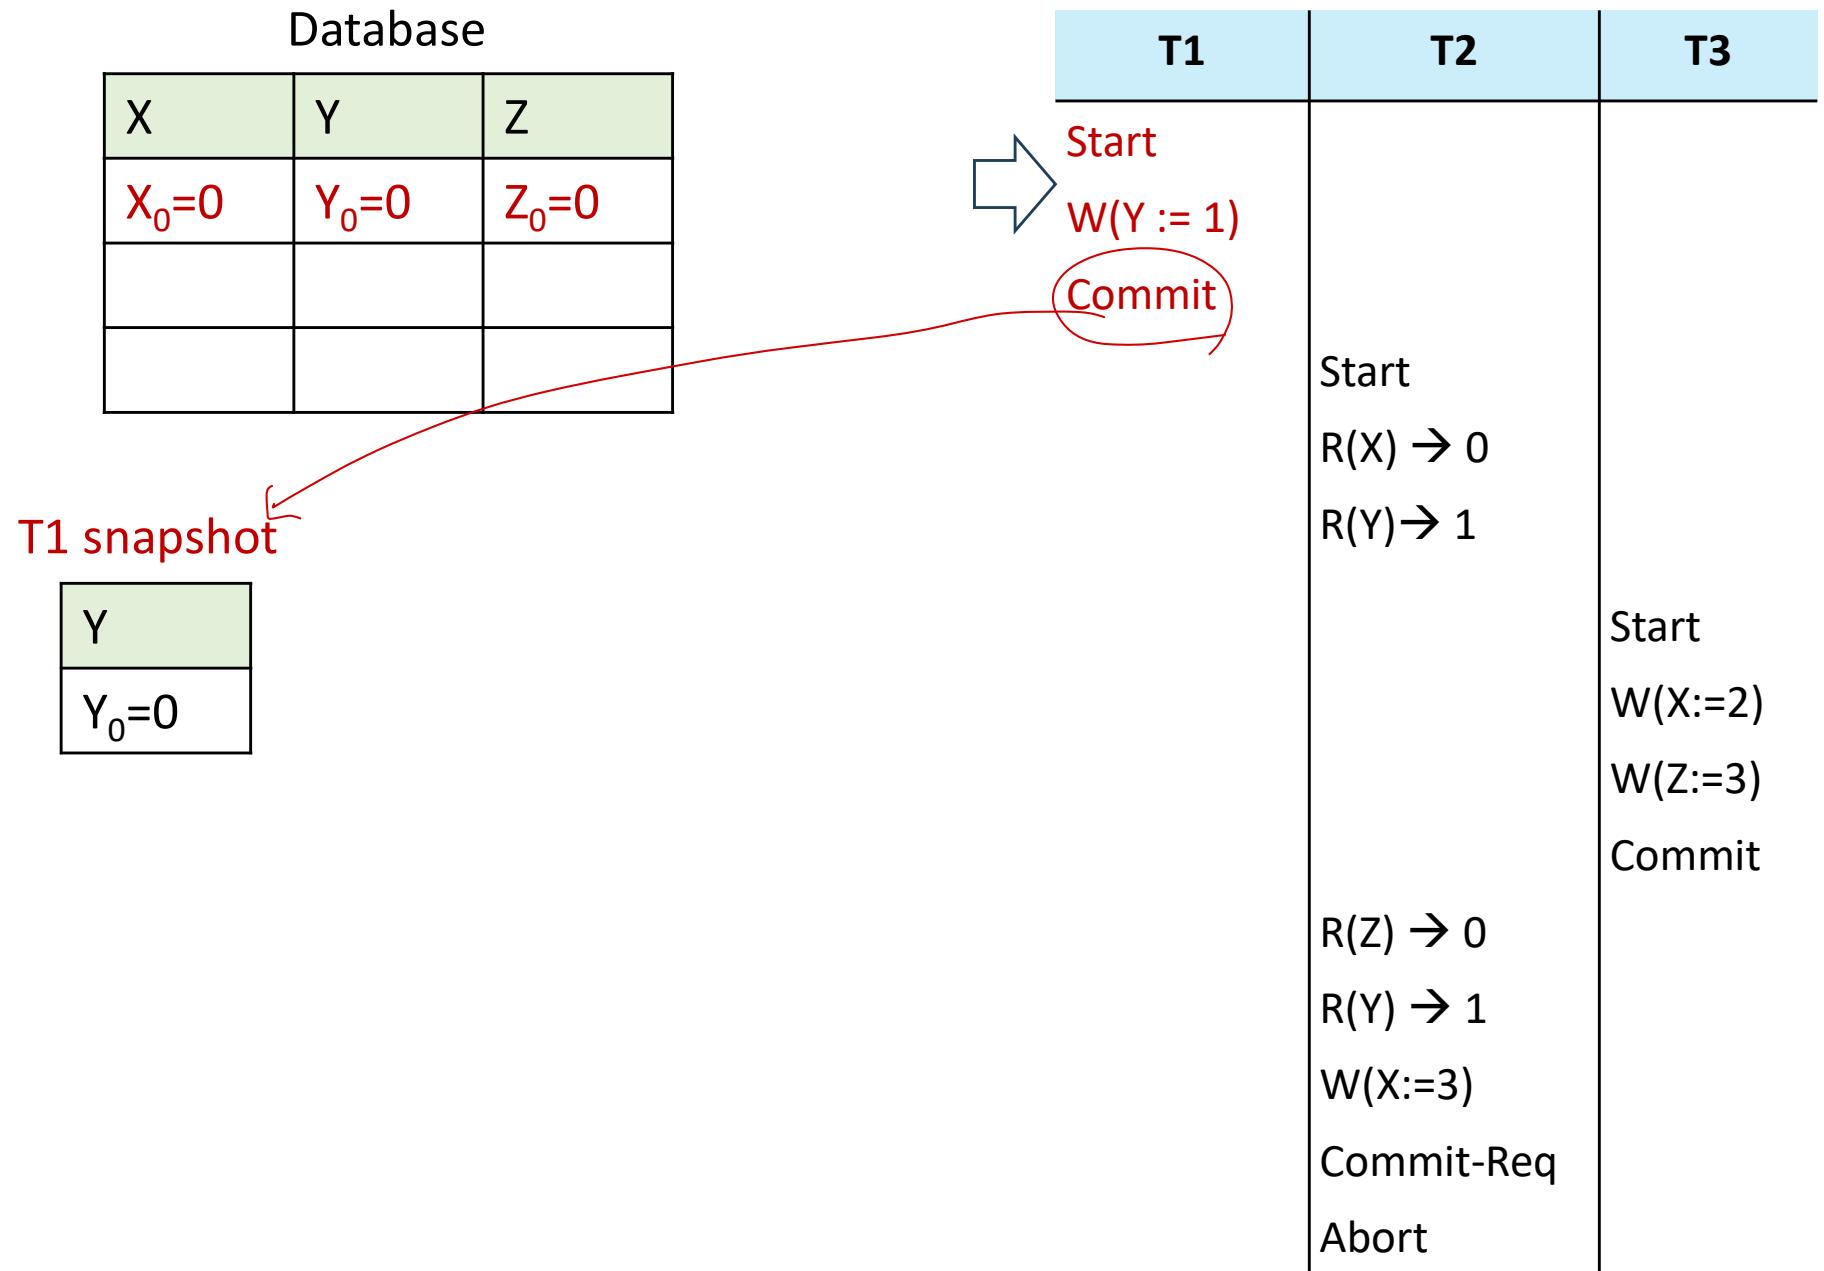
- Poor performance  
↳ DBMS가 느리고 흔히 snapshot 선택

OLTP ← → OLAP  
transaction 기반  
(계정내부 이체(화물교환) 등.)  
analytic 問

## Snapshot Isolation: Give snapshot of database to every transaction

- follows Optimistic Concurrency Control OCC 이용
- Takes snapshot of committed data at start → 트랜잭션이 시작될 때  
commit 된 데이터만 snapshot
- Always reads/modifies data in its own snapshot
- Updates of concurrent transactions are not visible
- First-committer-wins rule:
  - Commits only if no other concurrent transaction has already written data that the transaction intends to write.

# Snapshot Isolation



# Snapshot Isolation

Database

X	Y	Z
$X_0=0$	$Y_0=0$	$Z_0=0$

T1 snapshot

Y
$Y_1=1$

T1	T2	T3
Start		
W(Y := 1)		
Commit		
	Start	
	R(X) → 0	
	R(Y) → 1	
		Start
		W(X:=2)
		W(Z:=3)
		Commit
	R(Z) → 0	
	R(Y) → 1	
	W(X:=3)	
	Commit-Req	
	Abort	

# Snapshot Isolation

∈ MVCC

Database

X	Y	Z
$X_0=0$	$Y_0=0$	$Z_0=0$
	$Y_1=1$	

new  
version

T1	T2	T3
Start		
W(Y := 1)		
Commit		
	Start	
	R(X) → 0	
	R(Y) → 1	
		Start
		W(X:=2)
		W(Z:=3)
		Commit
		R(Z) → 0
		R(Y) → 1
		W(X:=3)
		Commit-Req
		Abort

# Snapshot Isolation

Database

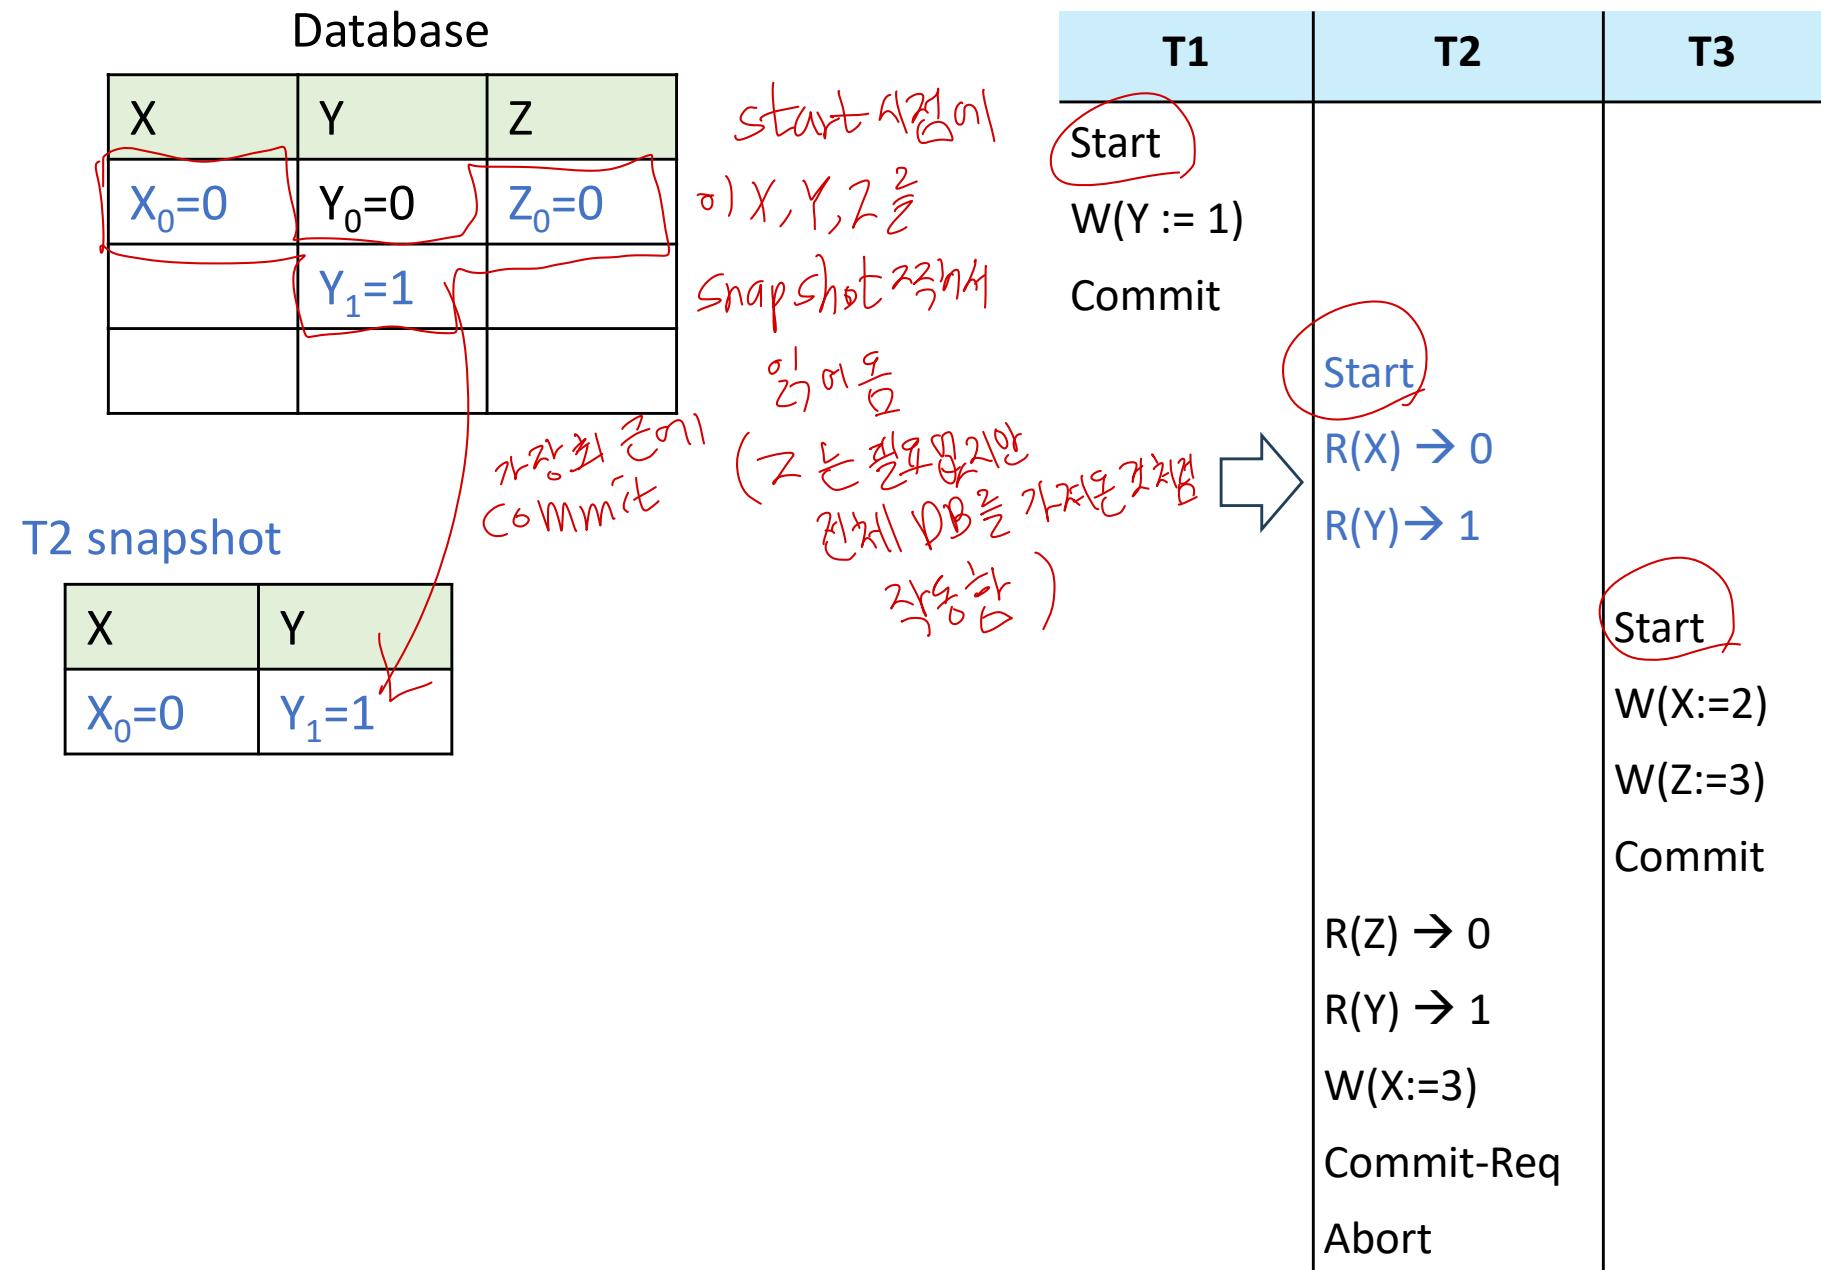
X	Y	Z
$X_0=0$	$Y_0=0$	$Z_0=0$
	$Y_1=1$	

T2 snapshot

X
$X_0=0$

T1	T2	T3
Start		
$W(Y := 1)$		
Commit		
	Start	
	$R(X) \rightarrow 0$	
	$R(Y) \rightarrow 1$	
		Start
		$W(X:=2)$
		$W(Z:=3)$
		Commit
		$R(Z) \rightarrow 0$
		$R(Y) \rightarrow 1$
		$W(X:=3)$
		Commit-Req
		Abort

# Snapshot Isolation



# Snapshot Isolation

Database

X	Y	Z
$X_0=0$	$Y_0=0$	$Z_0=0$
	$Y_1=1$	

T2 snapshot

X	Y
$X_0=0$	$Y_1=1$

T3 snapshot

X
$X_0=0$

T1	T2	T3
Start		
$W(Y := 1)$		
Commit		
	Start	
	$R(X) \rightarrow 0$	
	$R(Y) \rightarrow 1$	
		Start
		$W(X:=2)$
		$W(Z:=3)$
		Commit
	$R(Z) \rightarrow 0$	
	$R(Y) \rightarrow 1$	
		$W(X:=3)$
		Commit-Req
		Abort

# Snapshot Isolation

Database

X	Y	Z
$X_0=0$	$Y_0=0$	$Z_0=0$
	$Y_1=1$	

T2 snapshot

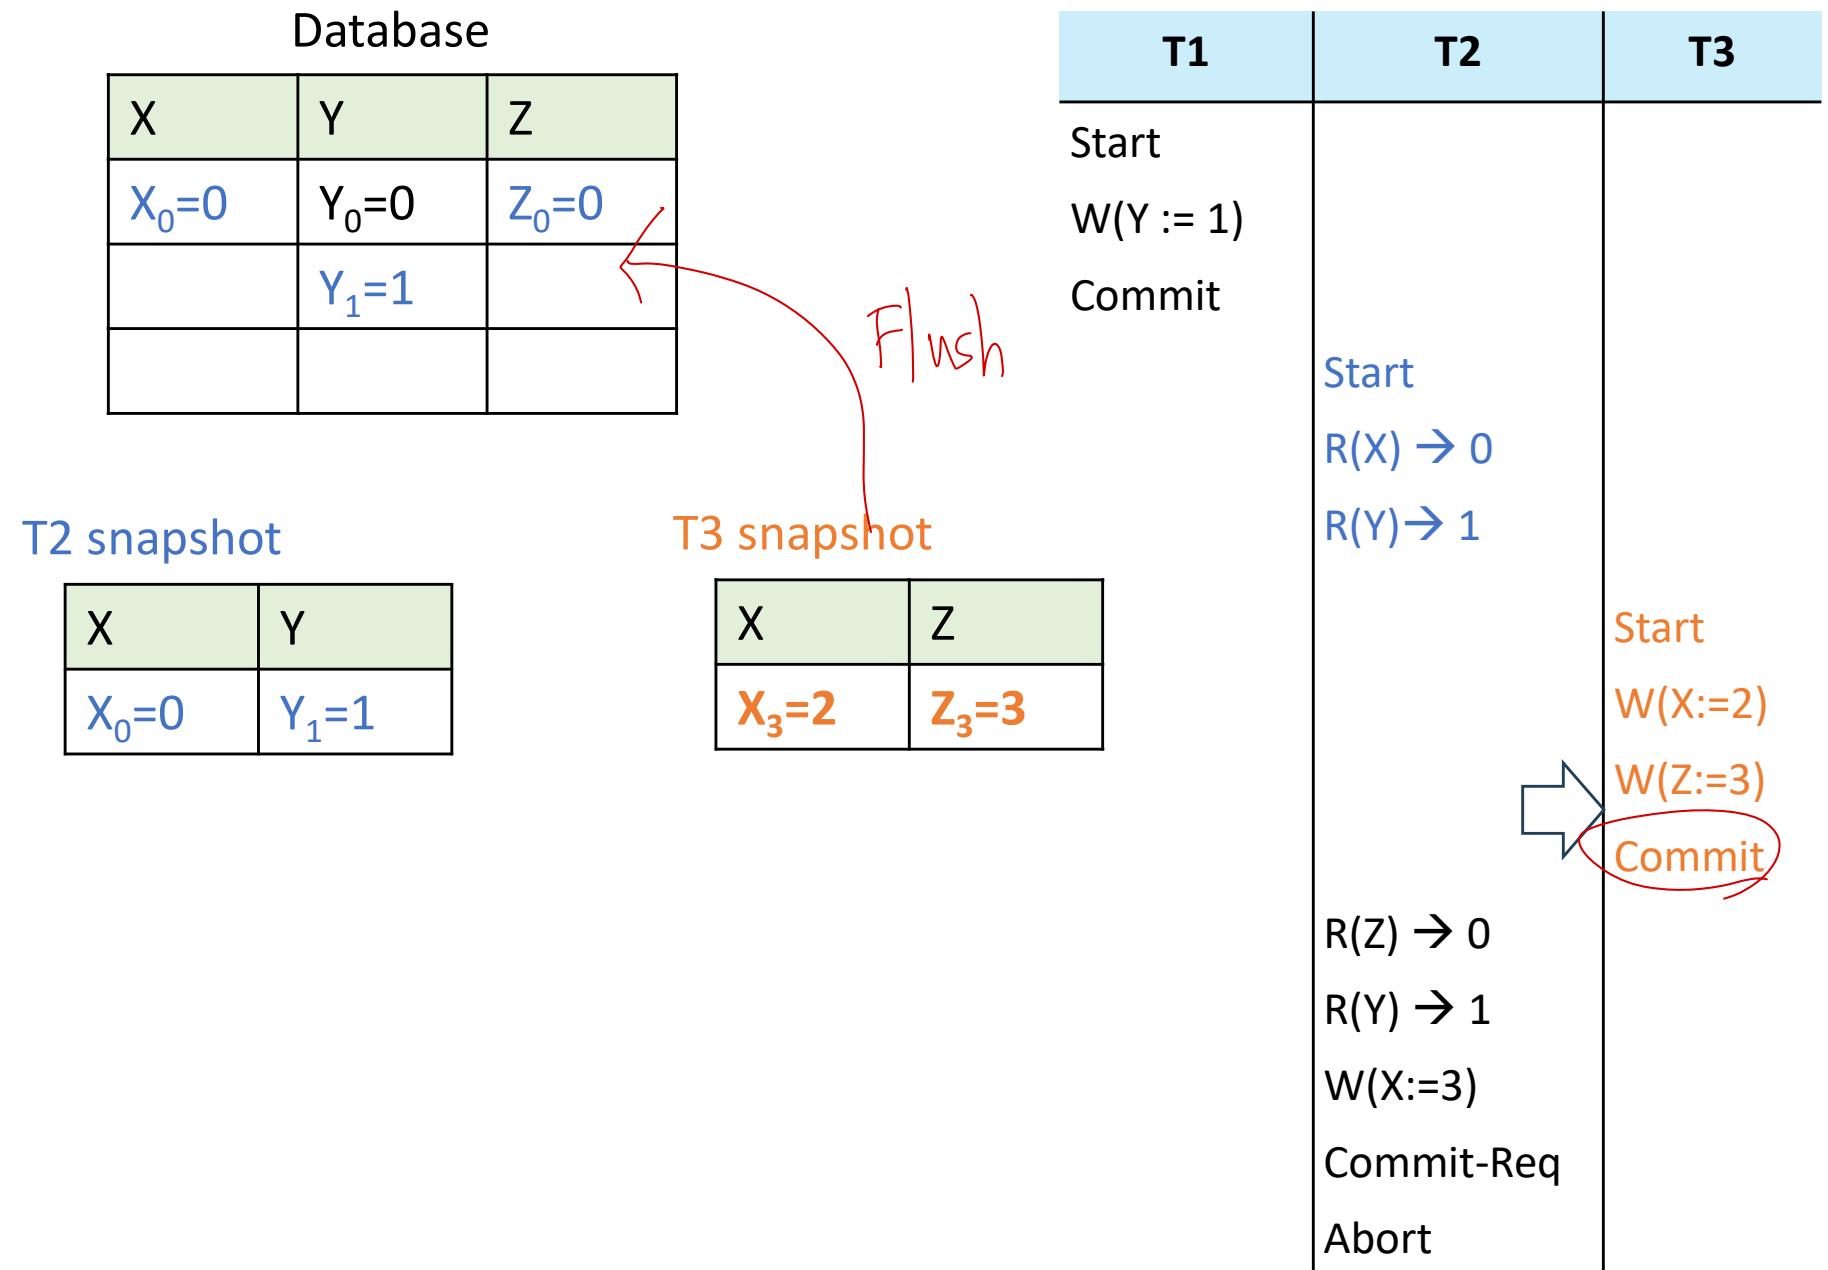
X	Y
$X_0=0$	$Y_1=1$

T3 snapshot

X	Z
$X_3=2$	$Z_0=0$

T1	T2	T3
Start		
$W(Y := 1)$		
Commit		
	Start	
	$R(X) \rightarrow 0$	
	$R(Y) \rightarrow 1$	
		Start
		$W(X:=2)$
		$W(Z:=3)$
		Commit
	$R(Z) \rightarrow 0$	
	$R(Y) \rightarrow 1$	
	$W(X:=3)$	
	Commit-Req	
		Abort

# Snapshot Isolation



# Snapshot Isolation

Database

X	Y	Z
$X_0=0$	$Y_0=0$	$Z_0=0$
$X_3=2$	$Y_1=1$	$Z_3=3$

T2 snapshot

X	Y	Z
$X_0=0$	$Y_1=1$	$Z_0=0$

T1	T2	T3
Start		
$W(Y := 1)$		
Commit		
	Start	
	$R(X) \rightarrow 0$	
	$R(Y) \rightarrow 1$	
		Start
		$W(X:=2)$
		$W(Z:=3)$
		Commit
		
	$R(Z) \rightarrow 0$	
	$R(Y) \rightarrow 1$	
	$W(X:=3)$	
		Commit-Req
		Abort

# Snapshot Isolation

Database

X	Y	Z
$X_0=0$	$Y_0=0$	$Z_0=0$
$X_3=2$	$Y_1=1$	$Z_3=3$

T2 snapshot

X	Y	Z
$X_0=0$	$Y_1=1$	$Z_0=0$

②  $\because 2 < 3 \circ 1 \leq 3$  old(Z) ~~2, 0~~ ~~2, 1~~

T1	T2	T3
Start		
$W(Y := 1)$		
Commit		
	Start	
	$R(X) \rightarrow 0$	
	$R(Y) \rightarrow 1$	
		Start
		$W(X:=2)$
		$W(Z:=3)$
		Commit
	$R(Z) \rightarrow 0$	
	$R(Y) \rightarrow 1$	
	$W(X:=3)$	
	Commit-Req	
		Abort

# Snapshot Isolation

Database

X	Y	Z
$X_0=0$	$Y_0=0$	$Z_0=0$
$X_3=2$	$Y_1=1$	$Z_3=3$

T2 snapshot

X	Y	Z
$X_2=3$	$Y_1=1$	$Z_0=0$

T1	T2	T3
Start		
$W(Y := 1)$		
Commit		
	Start	
	$R(X) \rightarrow 0$	
	$R(Y) \rightarrow 1$	
		Start
		$W(X:=2)$
		$W(Z:=3)$
		Commit
		$R(Z) \rightarrow 0$
		$R(Y) \rightarrow 1$
		$W(X:=3)$
		Commit-Req
		Abort

# Snapshot Isolation

Database

X	Y	Z
$X_0=0$	$Y_0=0$	$Z_0=0$
$X_3=2$	$Y_1=1$	$Z_3=3$

T2 snapshot

X	Y	Z
$X_2=3$	$Y_1=1$	$Z_0=0$

T1	T2	T3
Start		
$W(Y := 1)$		
Commit		
	Start	
	$R(X) \rightarrow 0$	
	$R(Y) \rightarrow 1$	
		<i>F1rst-commit</i>
	Start	
	$W(X:=2)$	
	$W(Z:=3)$	
	Commit	
	$R(Z) \rightarrow 0$	
	$R(Y) \rightarrow 1$	
	$W(X:=3)$	
		<i>First-committter</i>
		<i>2차 커미터 의해 reject</i>
	Commit-Req	
	Abort	
		<i>conflict</i>

Concurrent updates not visible

Not first-committer of X

Serializability error, T2 is rolled back

# Snapshot Isolation – another example (1)

Database

X	Y	Z
$X_0=0$	$Y_0=0$	$Z_0=0$

T1 snapshot

Y
$Y_1=2$

T2 snapshot

X	Y
$X_0=0$	$Y_0=0$

T1	T2
Start	
$W(Y := 2)$	
	Start
	$R(X) \rightarrow 0$
	$R(Y) \rightarrow 0$
	$W(X:=1)$
	$W(Z:=3)$
	Commit
	$R(Z) \rightarrow 0$
	$R(Y) \rightarrow 0$
	$W(X:=3)$
	Commit-Req
	Abort

# Snapshot Isolation – another example (2)

Database

X	Y	Z
$X_0=0$	$Y_0=0$	$Z_0=0$

T1 snapshot

X	Y	Z
$X_1=1$	$Y_1=2$	$Z_1=3$

T2 snapshot

X	Y
$X_0=0$	$Y_0=0$

T1	T2
Start	
$W(Y := 2)$	
	Start
	$R(X) \rightarrow 0$
	$R(Y) \rightarrow 0$
	$W(X:=1)$
	 $W(Z:=3)$
	Commit
	$R(Z) \rightarrow 0$
	$R(Y) \rightarrow 0$
	$W(X:=3)$
	Commit-Req
	Abort

# Snapshot Isolation – another example (3)

Database

X	Y	Z
$X_0=0$	$Y_0=0$	$Z_0=0$
$X_1=1$	$Y_1=2$	$Z_1=3$

T2 snapshot

X	Y	Z
$X_0=0$	$Y_0=0$	$Z_0=0$

T1	T2
Start	
$W(Y := 2)$	
	Start
	$R(X) \rightarrow 0$
	$R(Y) \rightarrow 0$
	<u>Commit</u> $\xrightarrow{X=1, Y=2}$
	$W(X:=1)$
	$W(Z:=3)$
Commit	 older
	$R(Z) \rightarrow 0$
	$R(Y) \rightarrow 0$
	$W(X:=3)$
	Commit-Req
	Abort

$T_2$  starts with data  $(0, 0, 0)$   
 $(1, 2, 3)$  in old old old

# Snapshot Isolation – another example (3)

Database

X	Y	Z
$X_0=0$	$Y_0=0$	$Z_0=0$
$X_1=1$	$Y_1=2$	$Z_1=3$

T2 snapshot

X	Y	Z
$X_2=3$	$Y_0=0$	$Z_0=0$

T1	T2
Start	
$W(Y := 2)$	
	Start
	$R(X) \rightarrow 0$
	$R(Y) \rightarrow 0$
	$W(X:=1)$
	$W(Z:=3)$
	Commit
	$R(Z) \rightarrow 0$
	$R(Y) \rightarrow 0$
	 $W(X:=3)$
	Commit-Req
	Abort

# Snapshot Isolation – another example (4)

Database

X	Y	Z
$X_0=0$	$Y_0=0$	$Z_0=0$
$X_1=2$	$Y_1=1$	$Z_1=3$

T2 snapshot

X	Y	Z
$X_2=3$	$Y_0=0$	$Z_0=0$

Concurrent updates not visible

Not first-committer of X  
Serialization error, T2 is rolled back

T1	T2
Start	
$W(Y := 2)$	
	Start
	$R(X) \rightarrow 0$
	$R(Y) \rightarrow 0$
	$W(X:=1)$
	$W(Z:=3)$
Commit	
	$R(Z) \rightarrow 0$
	$R(Y) \rightarrow 0$
	$W(X:=3)$
	Commit-Req
	Abort

conflict

# Benefits of Snapshot Isolation

- Reads are never blocked (Good for OLAP queries) → 매우 빠르게 실행
- Avoids several anomalies
  - No dirty read
    - no read of uncommitted data
  - No lost update  $\leftarrow$  ignore
    - Lost update is the update overwritten by another transaction that did not see the update
  - No non-repeatable read
    - I.e., if read is executed again, it will see the same value

# Snapshot Isolation Anomaly

- '온제점', skew Read
- Snapshot Isolation does not always give serializable executions
    - Serializable: among two concurrent txns, one sees the effects of the other
    - In SI: neither sees the effects of the other
  - Result: Integrity constraints can be violated

# Example: Skew Write

- Initially  $A = 3$  and  $B = 17$

- Serial execution:  $A = ??$ ,  $B = ??$
- if both transactions start at the same time, with snapshot isolation:  $A = ??$ ,  $B = ??$

storage	
A	B
$A_0=3$	$B_0=17$

T1 snapshot

$A=3$   
 $B=17$

T1	T2
Start	
Read(A)	
Read(B)	
	Start
	Read(A)
	Read(B)
A=B	
	B=A
Write(A)	
Commit	
	Write(B)
	Commit

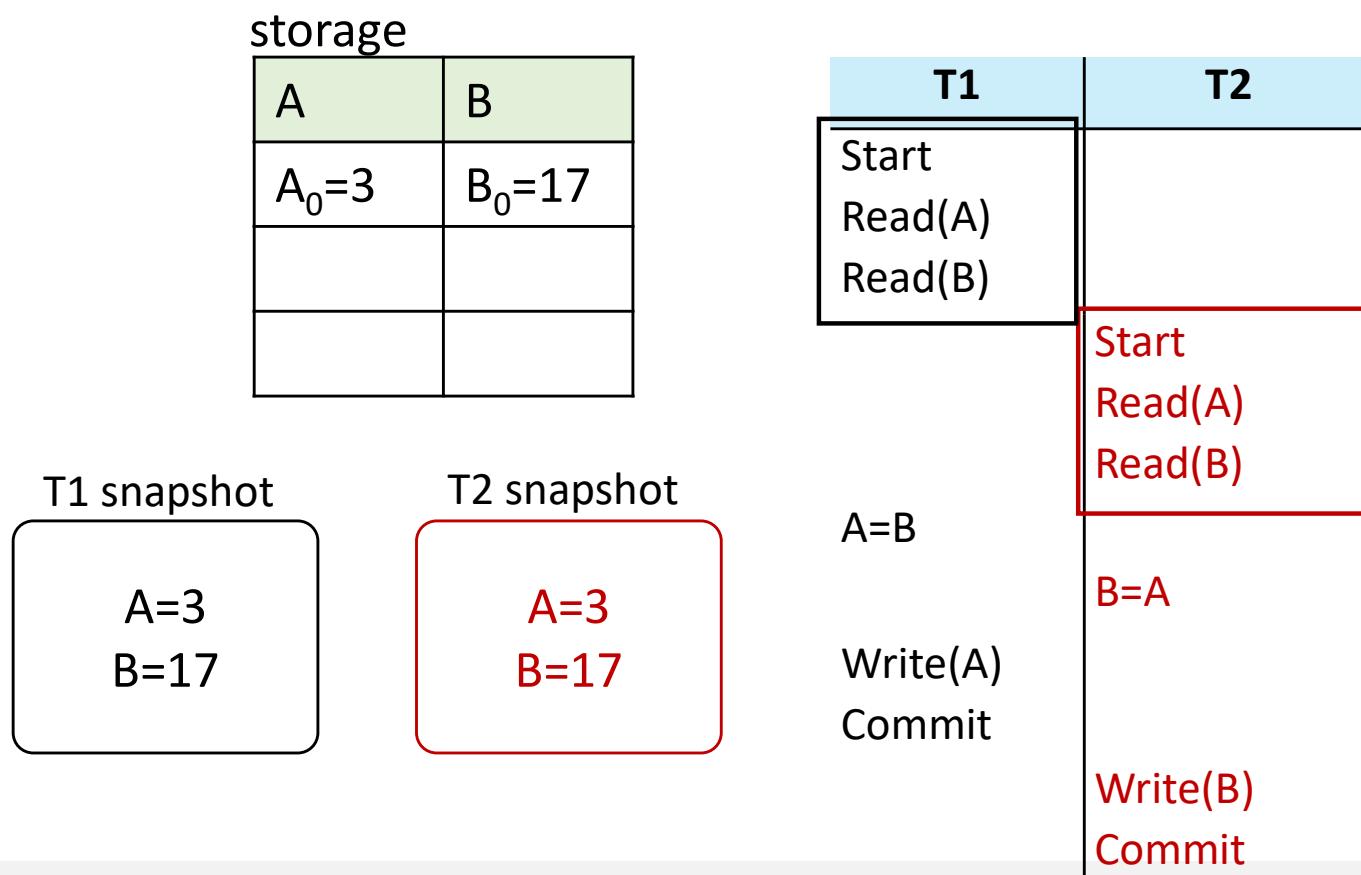
$A, B(3, 17)$ 를 serial로 실행해보자.  
 $T_1 \rightarrow T_2$  상각해보면  
 $(17, 17) \rightarrow (17, 17)$

$T_2 \rightarrow T_1$  상각해보면  
 $(3, 3) \rightarrow (3, 3)$

둘 다 가능함

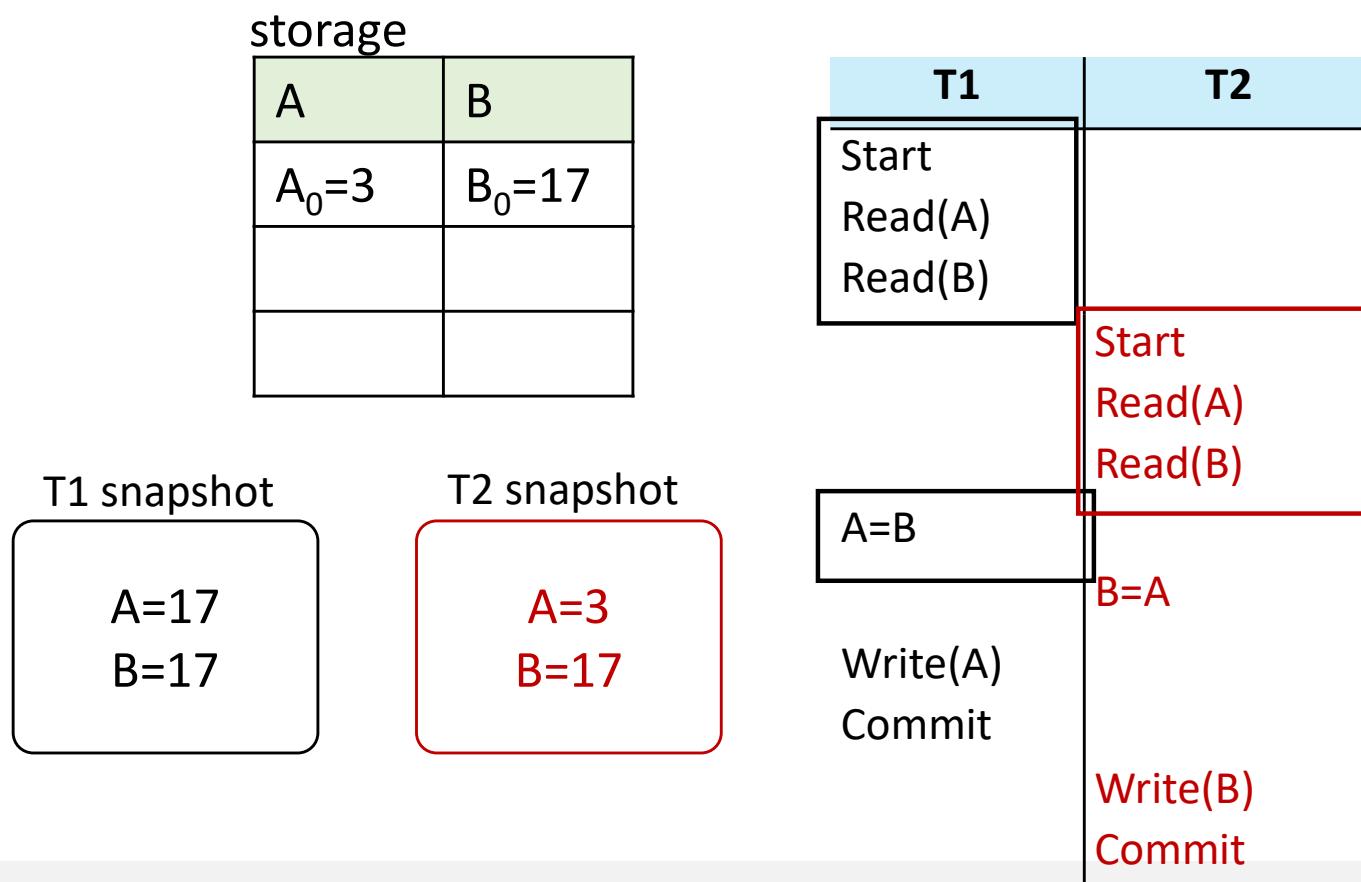
# Example: Skew Write

- Initially  $A = 3$  and  $B = 17$ 
  - Serial execution:  $A = ??$ ,  $B = ??$
  - if both transactions start at the same time, with snapshot isolation:  $A = ??$ ,  $B = ??$



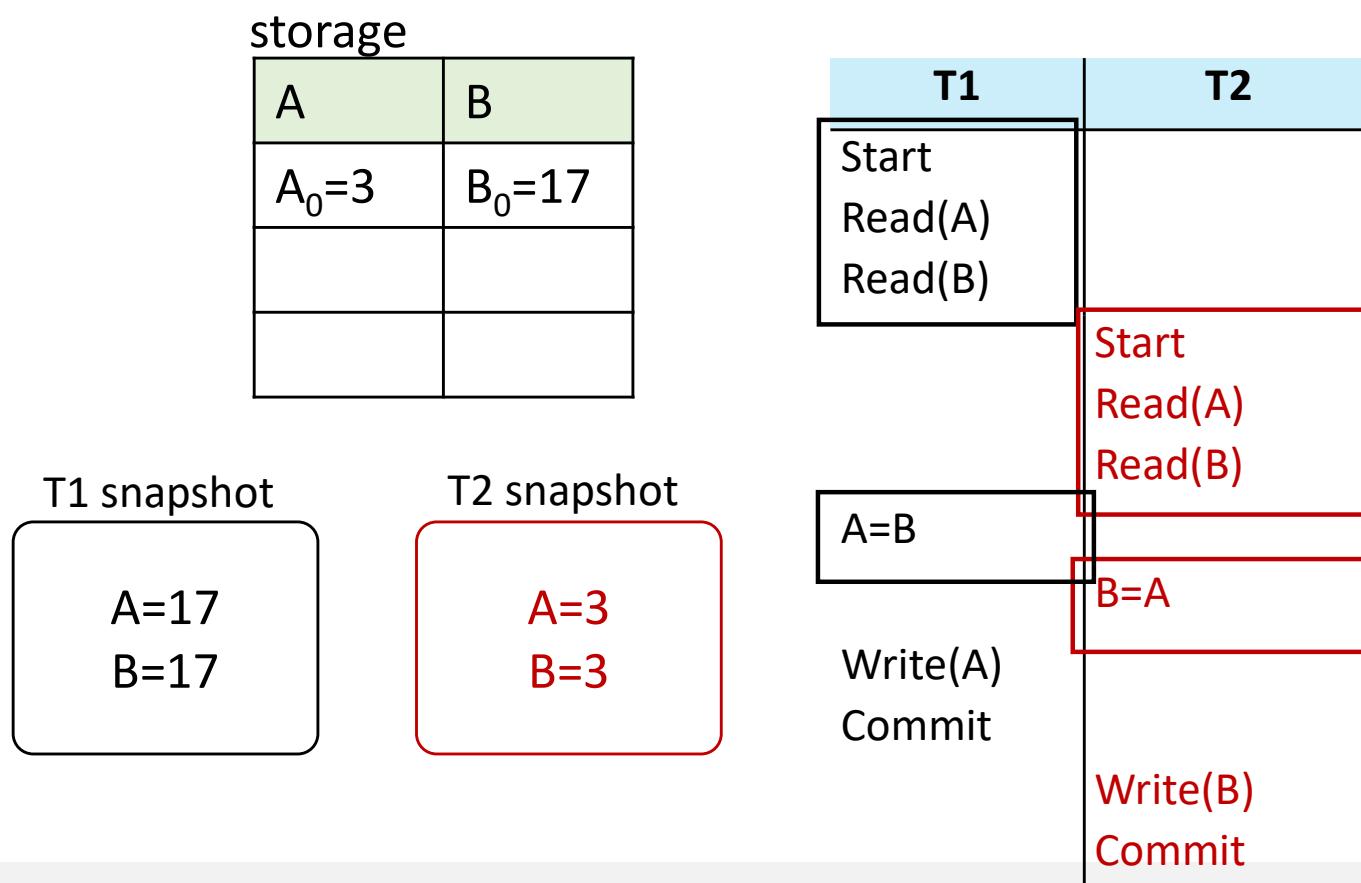
# Example: Skew Write

- Initially  $A = 3$  and  $B = 17$ 
  - Serial execution:  $A = ??$ ,  $B = ??$
  - if both transactions start at the same time, with snapshot isolation:  $A = ??$ ,  $B = ??$



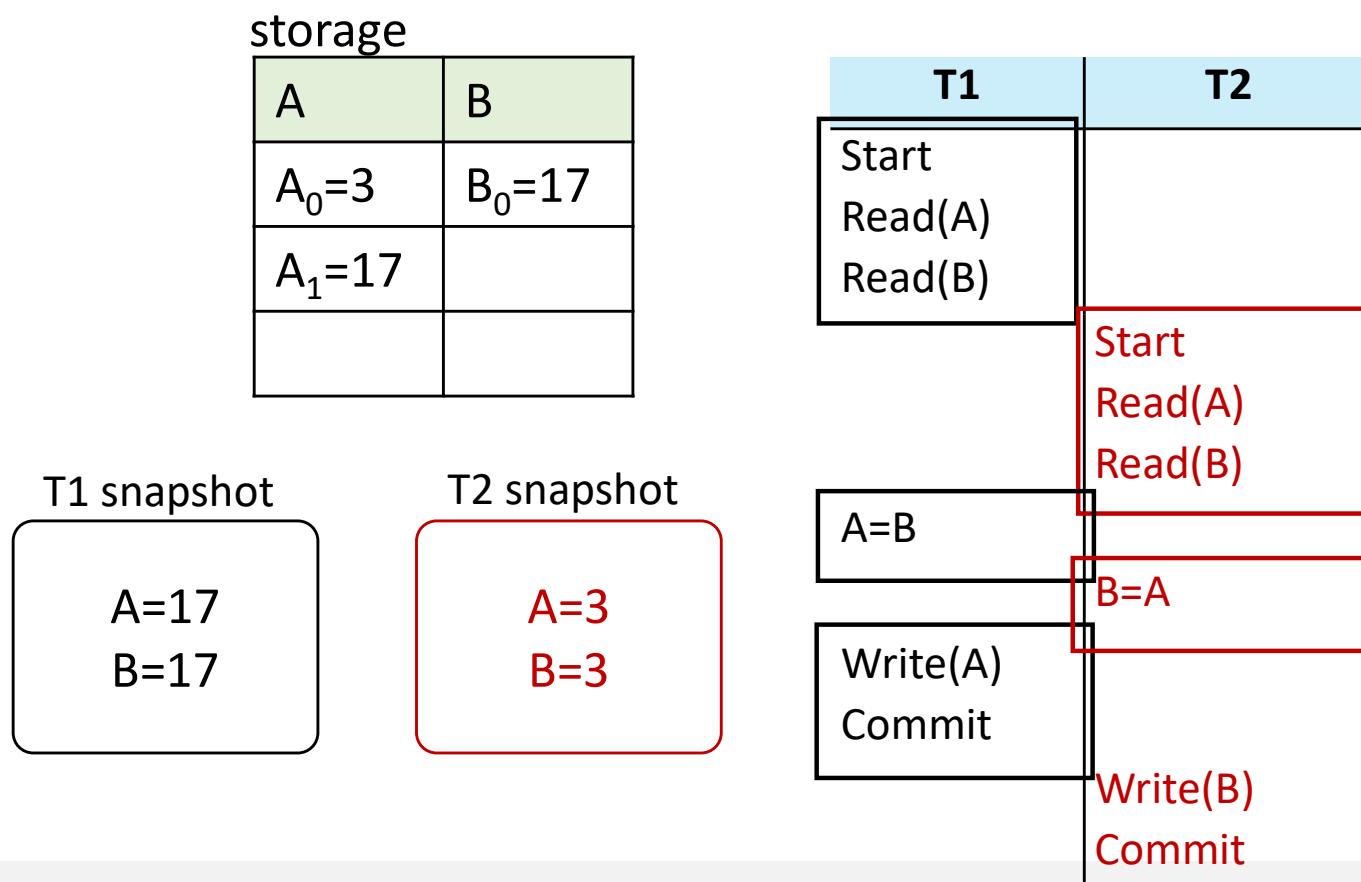
# Example: Skew Write

- Initially  $A = 3$  and  $B = 17$ 
  - Serial execution:  $A = ??$ ,  $B = ??$
  - if both transactions start at the same time, with snapshot isolation:  $A = ??$ ,  $B = ??$



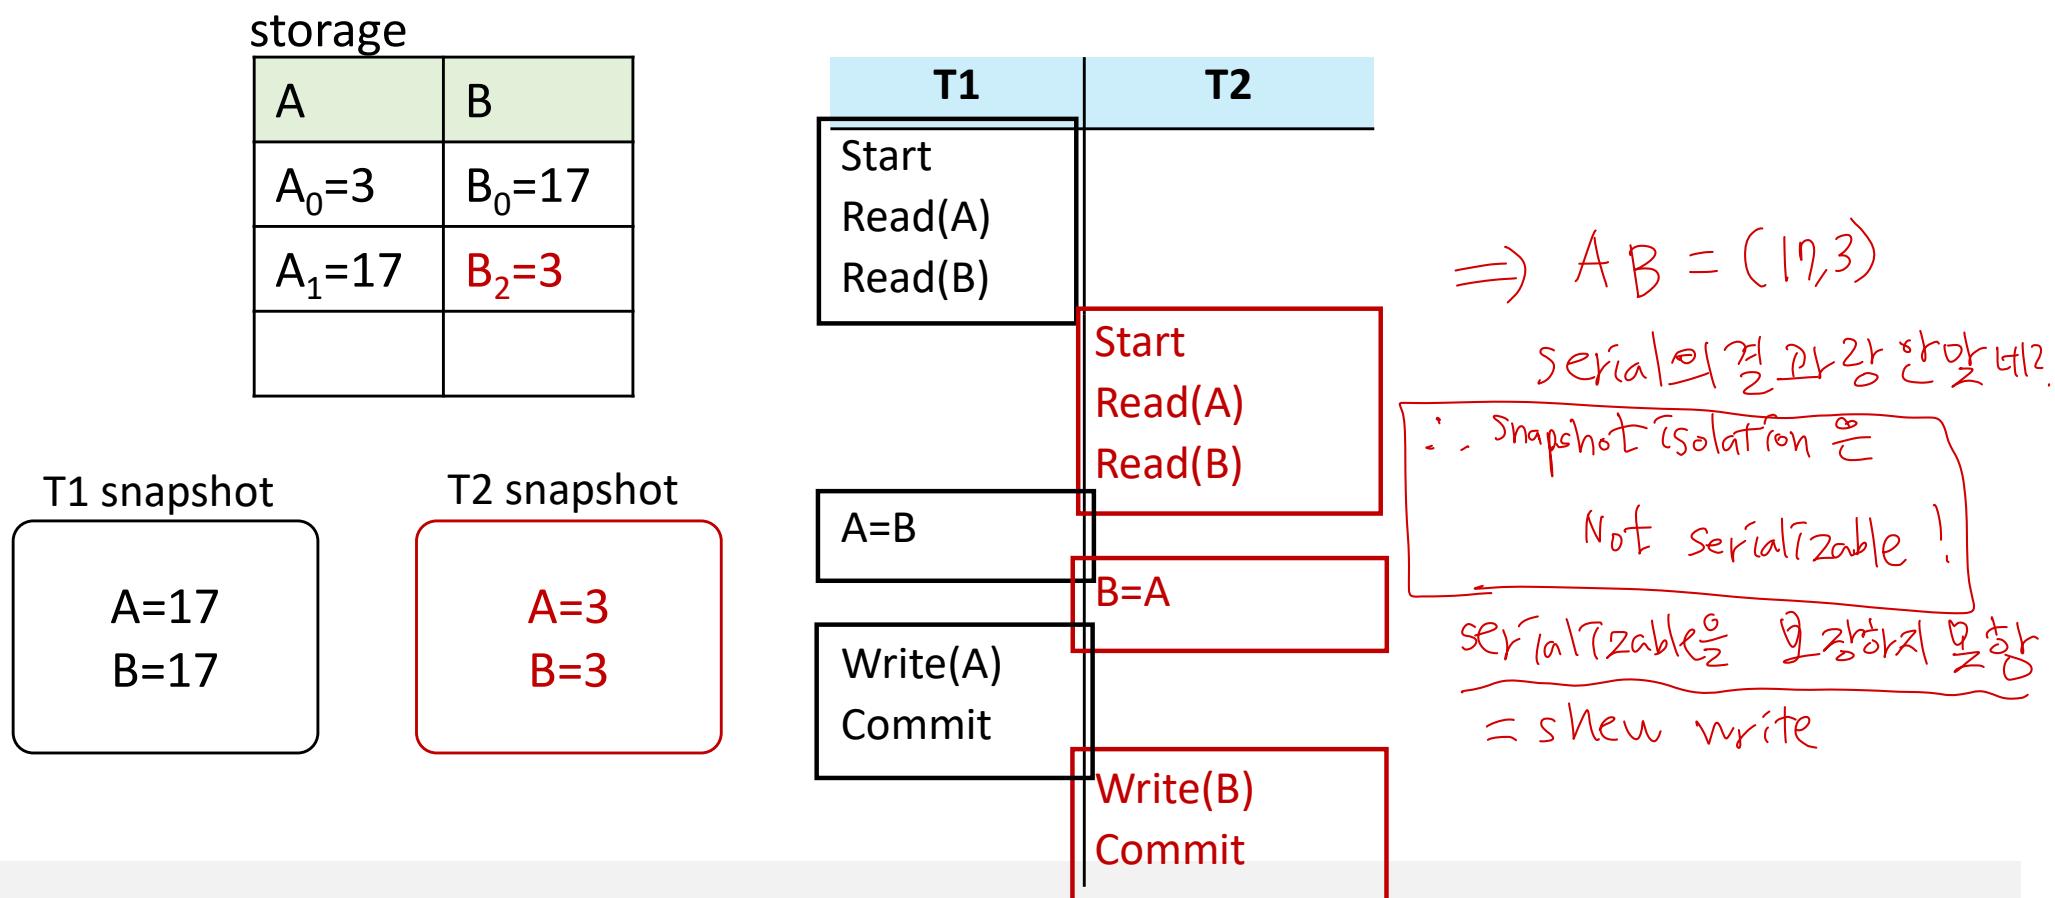
# Example: Skew Write

- Initially  $A = 3$  and  $B = 17$ 
  - Serial execution:  $A = ??$ ,  $B = ??$
  - if both transactions start at the same time, with snapshot isolation:  $A = ??$ ,  $B = ??$



# Example: Skew Write

- Initially  $A = 3$  and  $B = 17$ 
  - Serial execution:  $A = ??$ ,  $B = ??$
  - if both transactions start at the same time, with snapshot isolation:  $A = ??$ ,  $B = ??$



# Skew Write

- Skew Write is not very common in practice
- Skew also occurs with inserts, e.g.,
  - Find max order number among all orders
  - Create a new order with order number = previous max + 1
  - Two transaction can both create order with same number

*snapshot*

$T_i$	$T_j$
read(A)	
read(B)	
	read(A) read(B)
A=B	
	B=A
write(A)	
	write(B)

A	B
$A_0=3$	$B_0=17$
$A_1=17$	$B_2=3$

*serial*

T1	T2
Read(A)	
Read(B)	
A=B	
Write(A)	
	Read(A) Read(B) B=A Write(B)

A	B
$A_0=3$	$B_0=17$
$A_1=17$	$B_2=17$

*serial*

T1	T2
Read(A)	
Read(B)	
B=A	
Write(B)	
	Read(A) Read(B) A=B Write(A)

A	B
$A_0=3$	$B_0=17$
$A_1=3$	$B_2=3$

# Working Around Snapshot Isolation Anomalies

- Can work around SI anomalies for specific queries by using **select .. for update** (supported in Oracle)
  - Example
    - **select max(orderno) from orders for update**
    - read value into local variable maxorder
    - insert into orders (maxorder+1, ...)
- **select for update (SFU) clause** treats all data read by the query as if it were also updated, preventing concurrent updates
- Can be added to queries to ensure serializability in many applications
  - Does not handle phantom phenomenon/predicate reads though

*skew writes for update를 통해 피할 수 있다*

# Multi-Version Timestamp Ordering (MVTO)

# Multiversion Timestamp Ordering (MVTO)

- Each data item  $Q$  has a sequence of versions  $\langle Q_1, Q_2, \dots, Q_m \rangle$ .  
*TS 순서  
각 버전은 Version의 시퀀스를 지정*
- Each version  $Q_k$  contains three data fields:
  - **Content**
    - the value of version  $Q_k$ .
  - **W-timestamp( $Q_k$ )**
    - timestamp of the transaction that created version  $Q_k$
  - **R-timestamp( $Q_k$ )**
    - largest timestamp of a transaction that successfully read version  $Q_k$

# Multiversion Timestamp Ordering (Cont)

- Let  $Q_k$  denote the version of  $Q$  whose write timestamp is the largest write timestamp less than or equal to  $TS(T_i)$ .

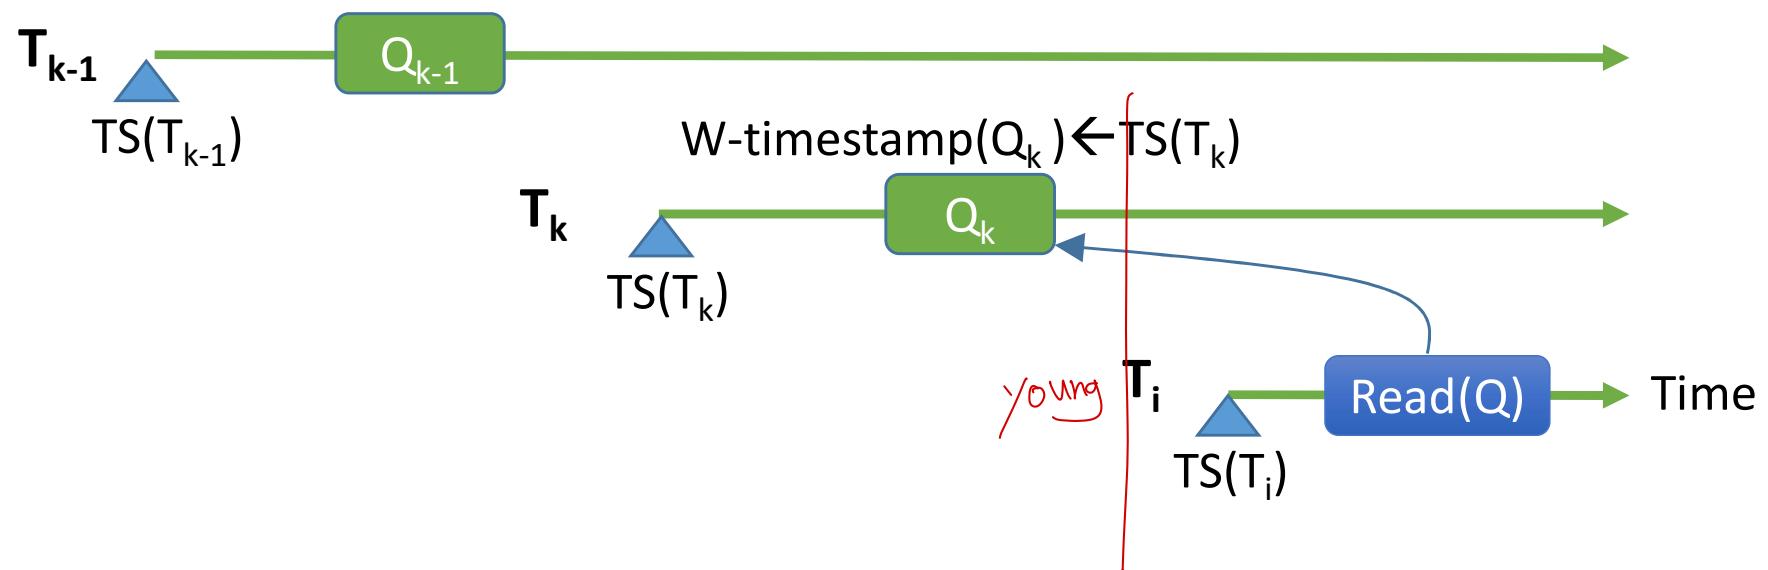
1. If transaction  $T_i$  issues a **read**( $Q$ ), then

- return the content of version  $Q_k$

- If  $R\text{-timestamp}(Q_k) < TS(T_i)$ ,

- $\text{set } R\text{-timestamp}(Q_k) \leftarrow TS(T_i),$

$W\text{-}TS \leq TS(T_i)$  인 애들 중에  
가장 큰 값으로!



# Multiversion Timestamp Ordering (Cont)

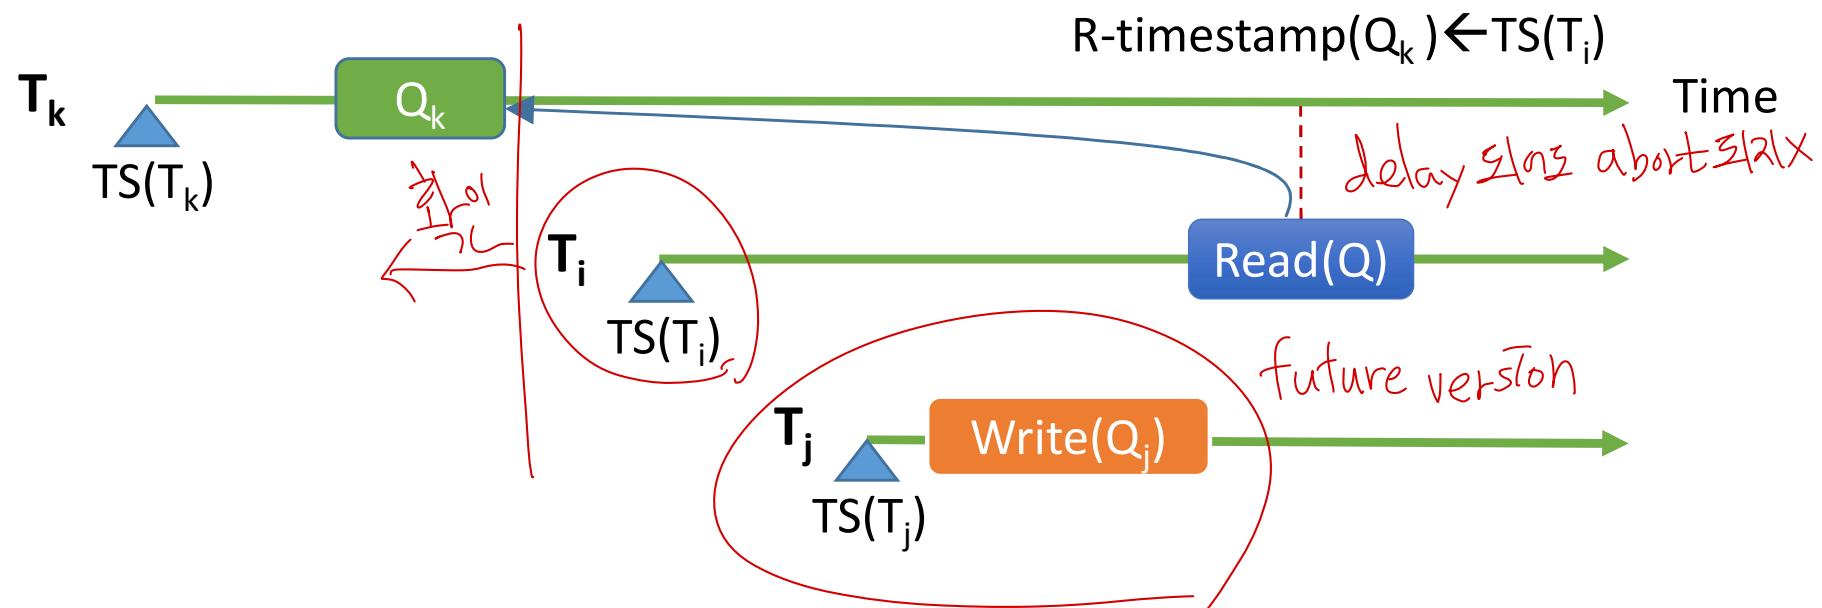
- Let  $Q_k$  denote the version of  $Q$  whose write timestamp is the largest write timestamp less than or equal to  $\text{TS}(T_i)$ .

1. If transaction  $T_i$  issues a **read**( $Q$ ), then

- return the content of version  $Q_k$

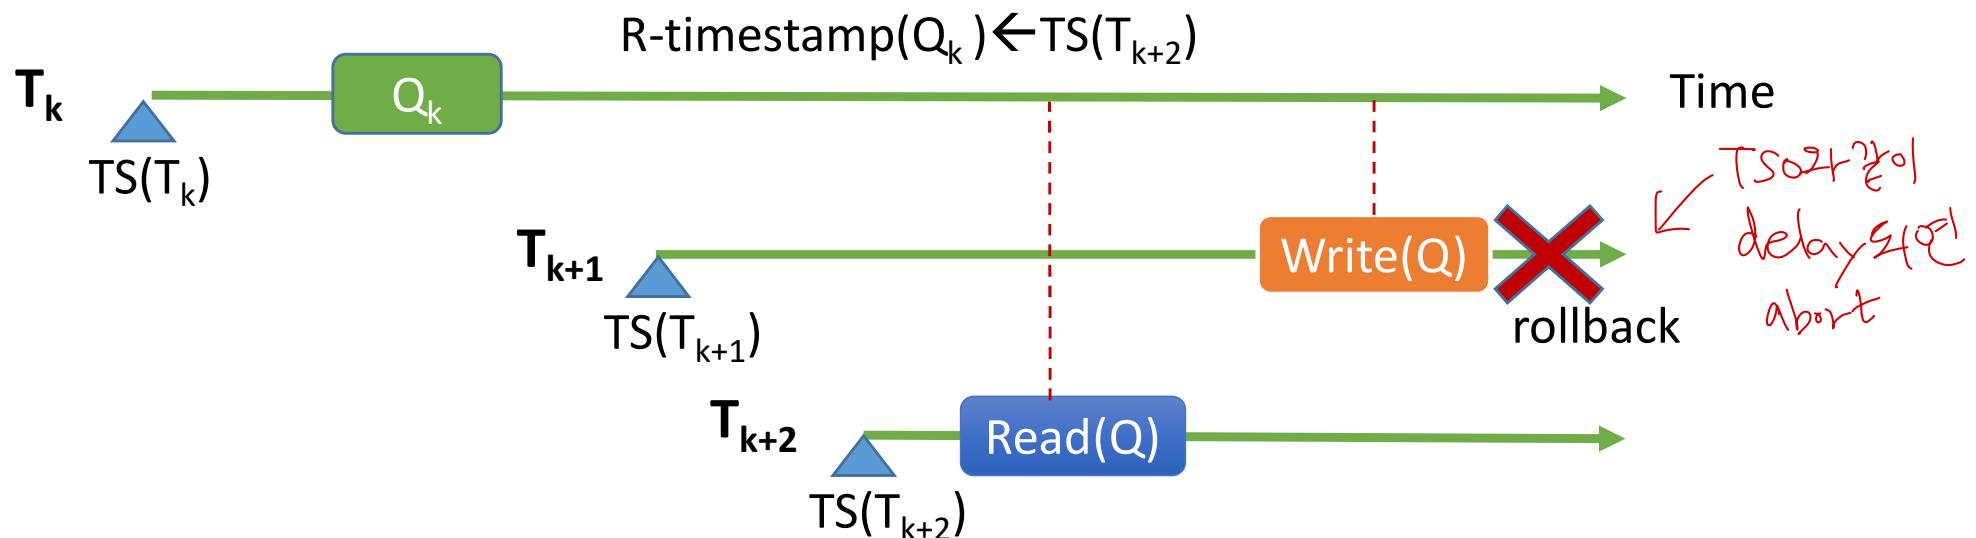
- If  $\text{R-timestamp}(Q_k) < \text{TS}(T_i)$ ,

- set  $\text{R-timestamp}(Q_k) \leftarrow \text{TS}(T_i)$ ,



# Multiversion Timestamp Ordering (Cont)

2. If transaction  $T_i$  issues a **write**( $Q$ )
1. if  $TS(T_i) < R\text{-timestamp}(Q_k)$ , roll back transaction  $T_i$ .
  2. if  $TS(T_i) = W\text{-timestamp}(Q_k)$ , overwrite  $Q_k$ .
  3. Otherwise, create a new version  $Q_i$ 
    - Set  $W\text{-timestamp}(Q_i)$  and  $R\text{-timestamp}(Q_i)$  to  $TS(T_i)$ .



# Multiversion Timestamp Ordering (Cont)

## ■ Observations

- Reads always succeed → historical data에 접근하기(n)
- Writes may be rejected

## ■ MVTO guarantees serializability

---

## Multi-Version Two-Phase Locking (MV2PL)

각도별식 2f 과 비슷 (가벼운 lock 방식 채용)

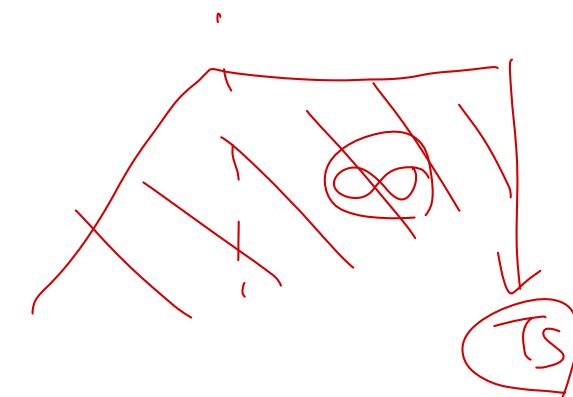
# Multiversion Two-Phase Locking (MV2PL)

- Main Goal: Differentiate read-only transactions from write transactions

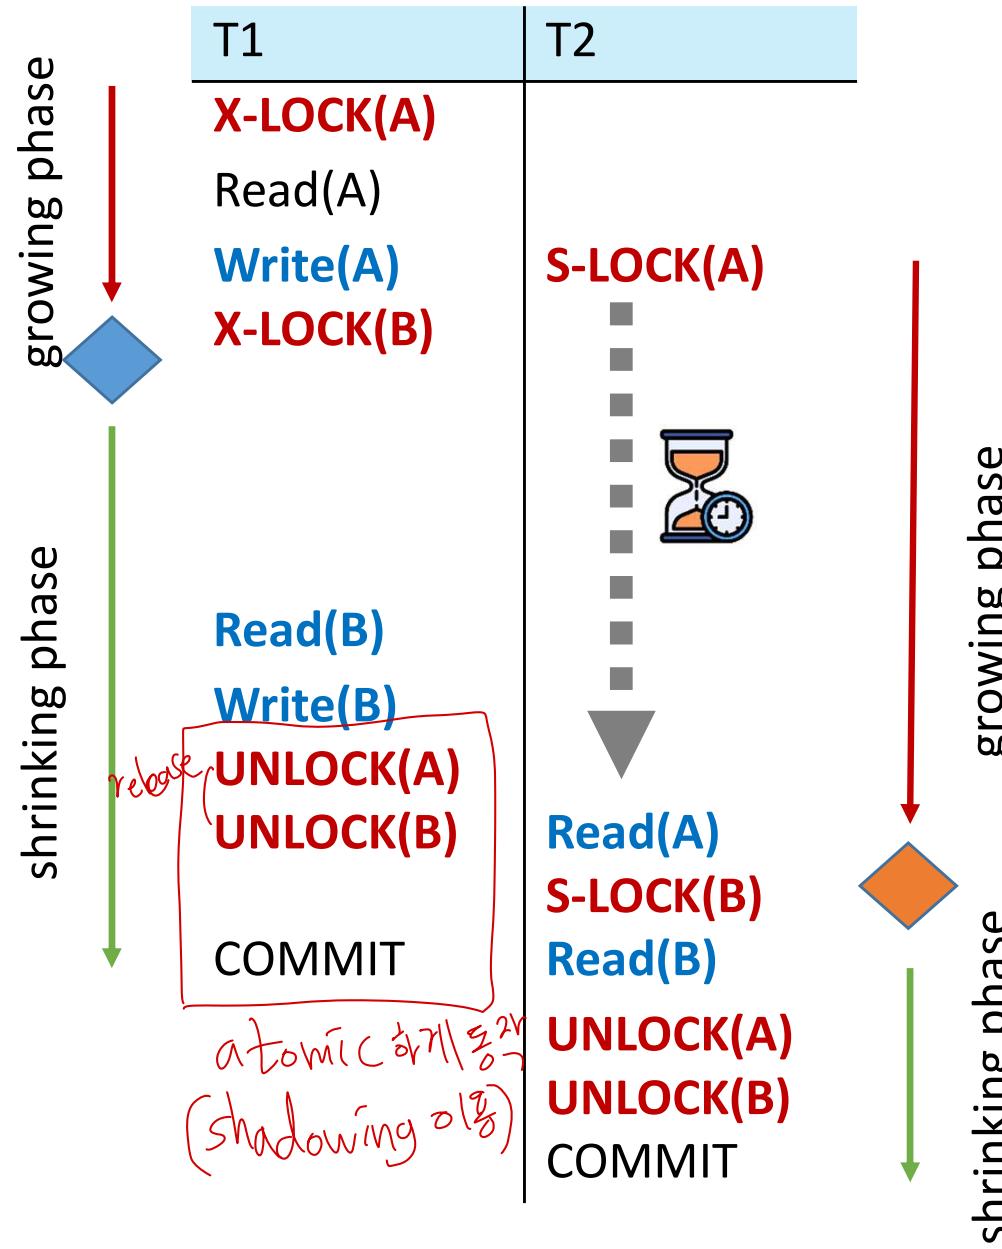
동시성  
Concurrency 높이↓  
→ commit 할 준비가 되면,  
모든 lock 해제

- Update transactions use strict two-phase locking.

- Read(Q) returns the latest version of the item
- The first **write** of Q by  $T_i$  creates a new version  $Q_i$ 
  - W-timestamp( $Q_i$ ) set to  $\infty$  initially → Read 가 W-TS 를 읽을 때  
아직 commit 되지 않았음을 표시해(위해)
- When  $T_i$  commits:
  - Set  $TS(T_i) = \text{ts-counter} + 1$
  - Set  $W\text{-timestamp}(Q_i) = TS(T_i)$
  - **ts-counter = ts-counter + 1**



# Strict Two-Phase Locking (2PL) - Review



# Implementation of MV2PL

2PL  $\leftrightarrow$  MV2PL의 차이

T10	T11
S-Lock(A)	Read CNT 1 증가
Read(A)	
X-Lock(B)	
Write(B)	

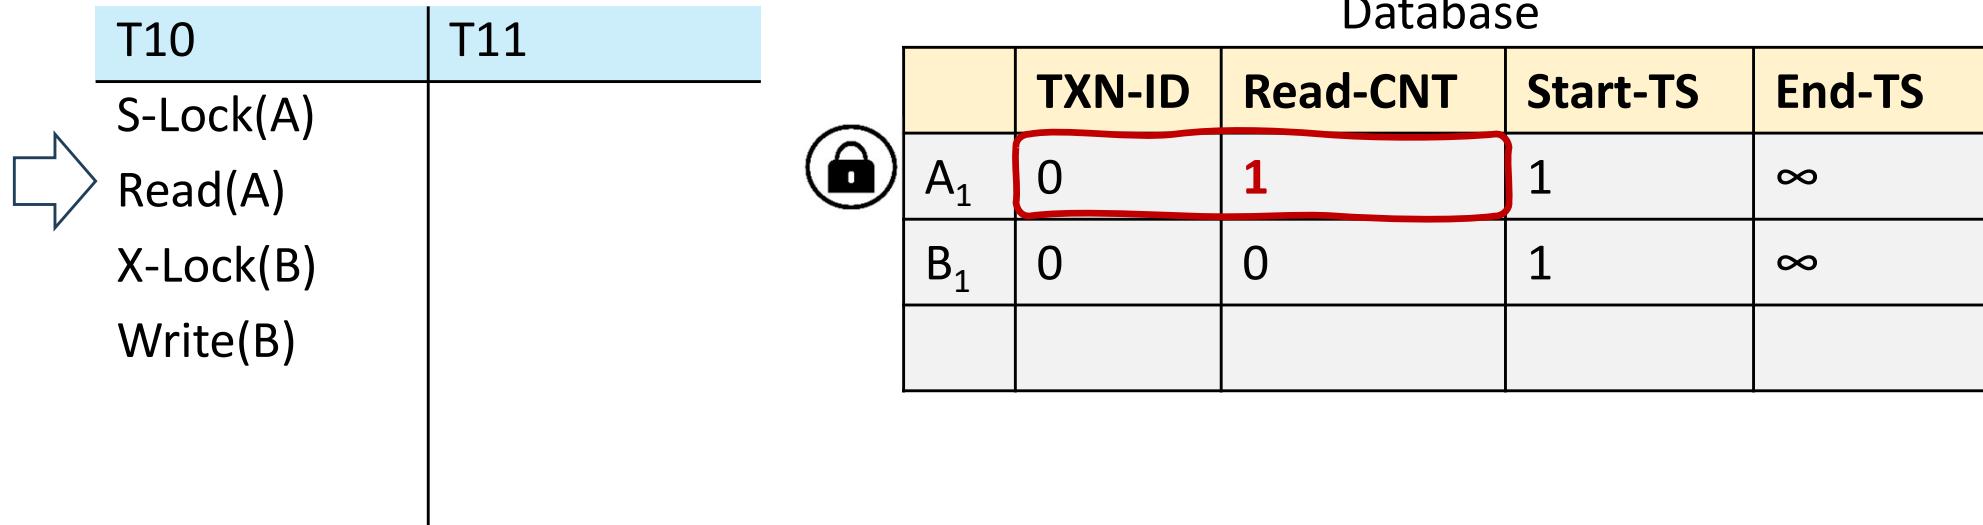
T10과 T11의 차이점

	TXN-ID	Read-CNT	Start-TS	End-TS
A <sub>1</sub>	0	0	1	$\infty$
B <sub>1</sub>	0	0	1	$\infty$

현재 다른 tx가 읽고 있는지 같다  
(현재 블록의 tx가 있고 있는지)

- Txns use the tuple's Read-CNT field as SHARED lock.
- If TXN-ID is zero, then the transaction acquires the SHARED lock by incrementing the Read-CNT field.

# Implementation of MV2PL



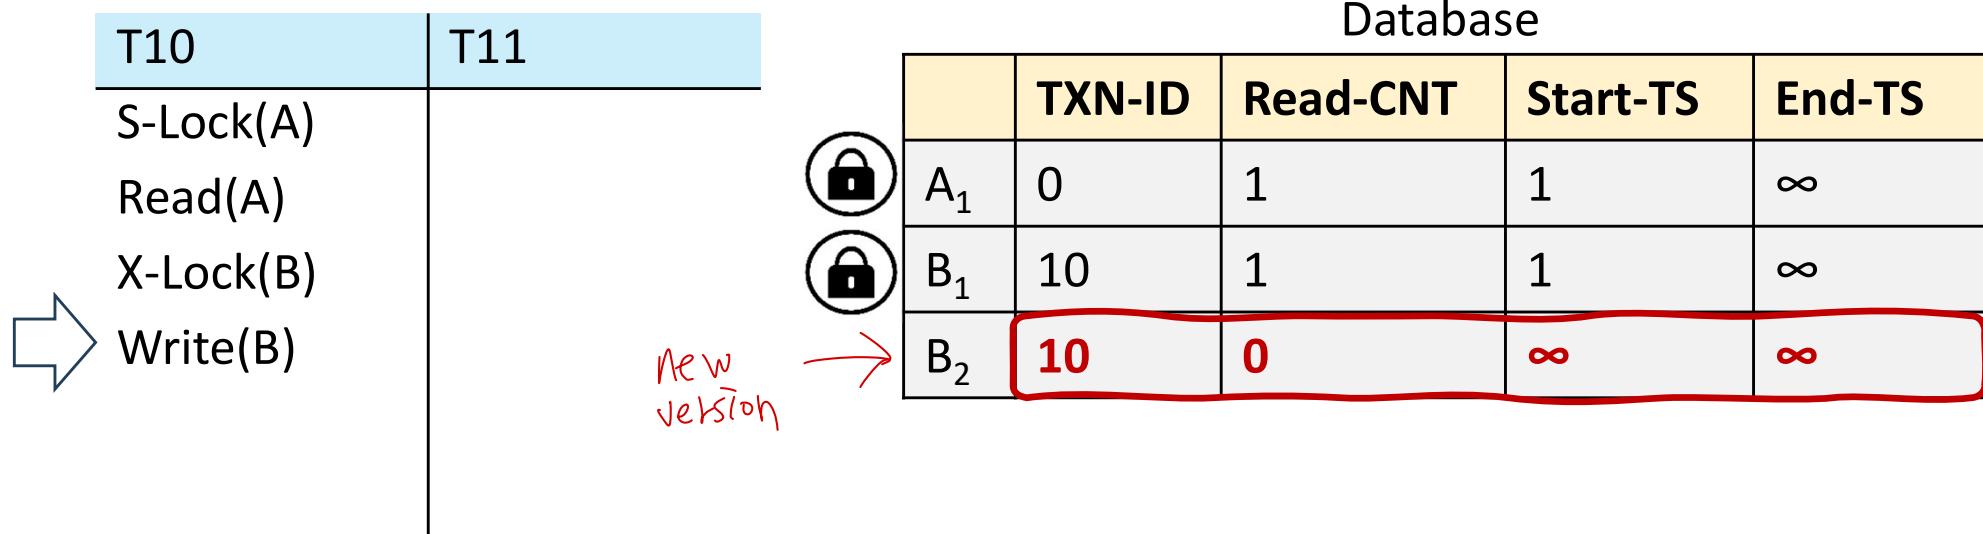
- Txns use the tuple's Read-CNT field as SHARED lock.
- If TXN-ID is zero, then the transaction acquires the SHARED lock by incrementing the Read-CNT field.

# Implementation of MV2PL

T10	T11	Database				
		TXN-ID	Read-CNT	Start-TS	End-TS	
S-Lock(A)		A <sub>1</sub>	0	1	1	$\infty$
Read(A)		B <sub>1</sub>	10	1	1	$\infty$
X-Lock(B)						
Write(B)						

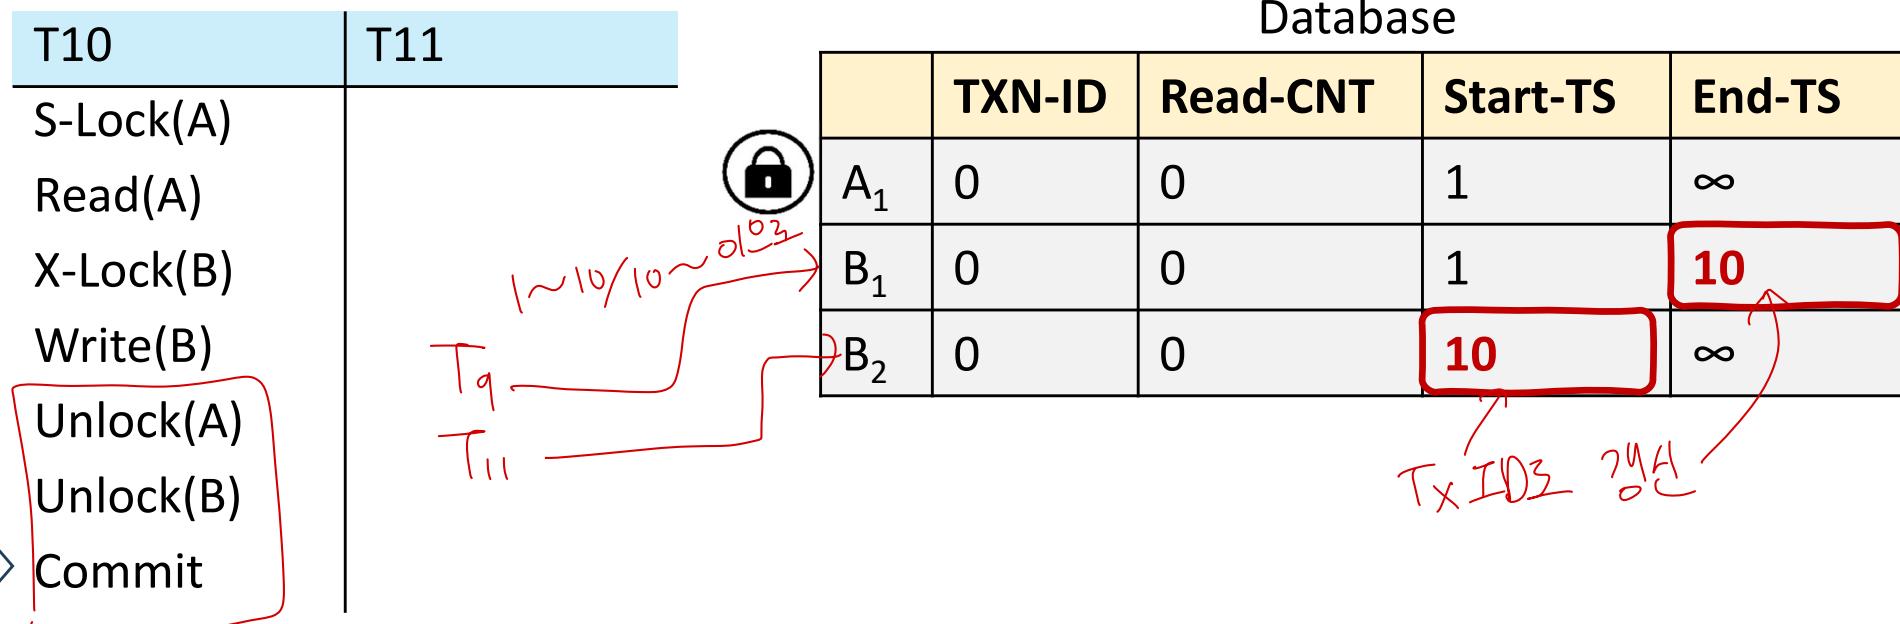
- Use TXN-ID and Read-CNT together as EXCLUSIVE lock.
- If both TXN-ID and Read-CNT are zero, then transaction acquires the EXCLUSIVE lock by setting both of them.

# Implementation of MV2PL



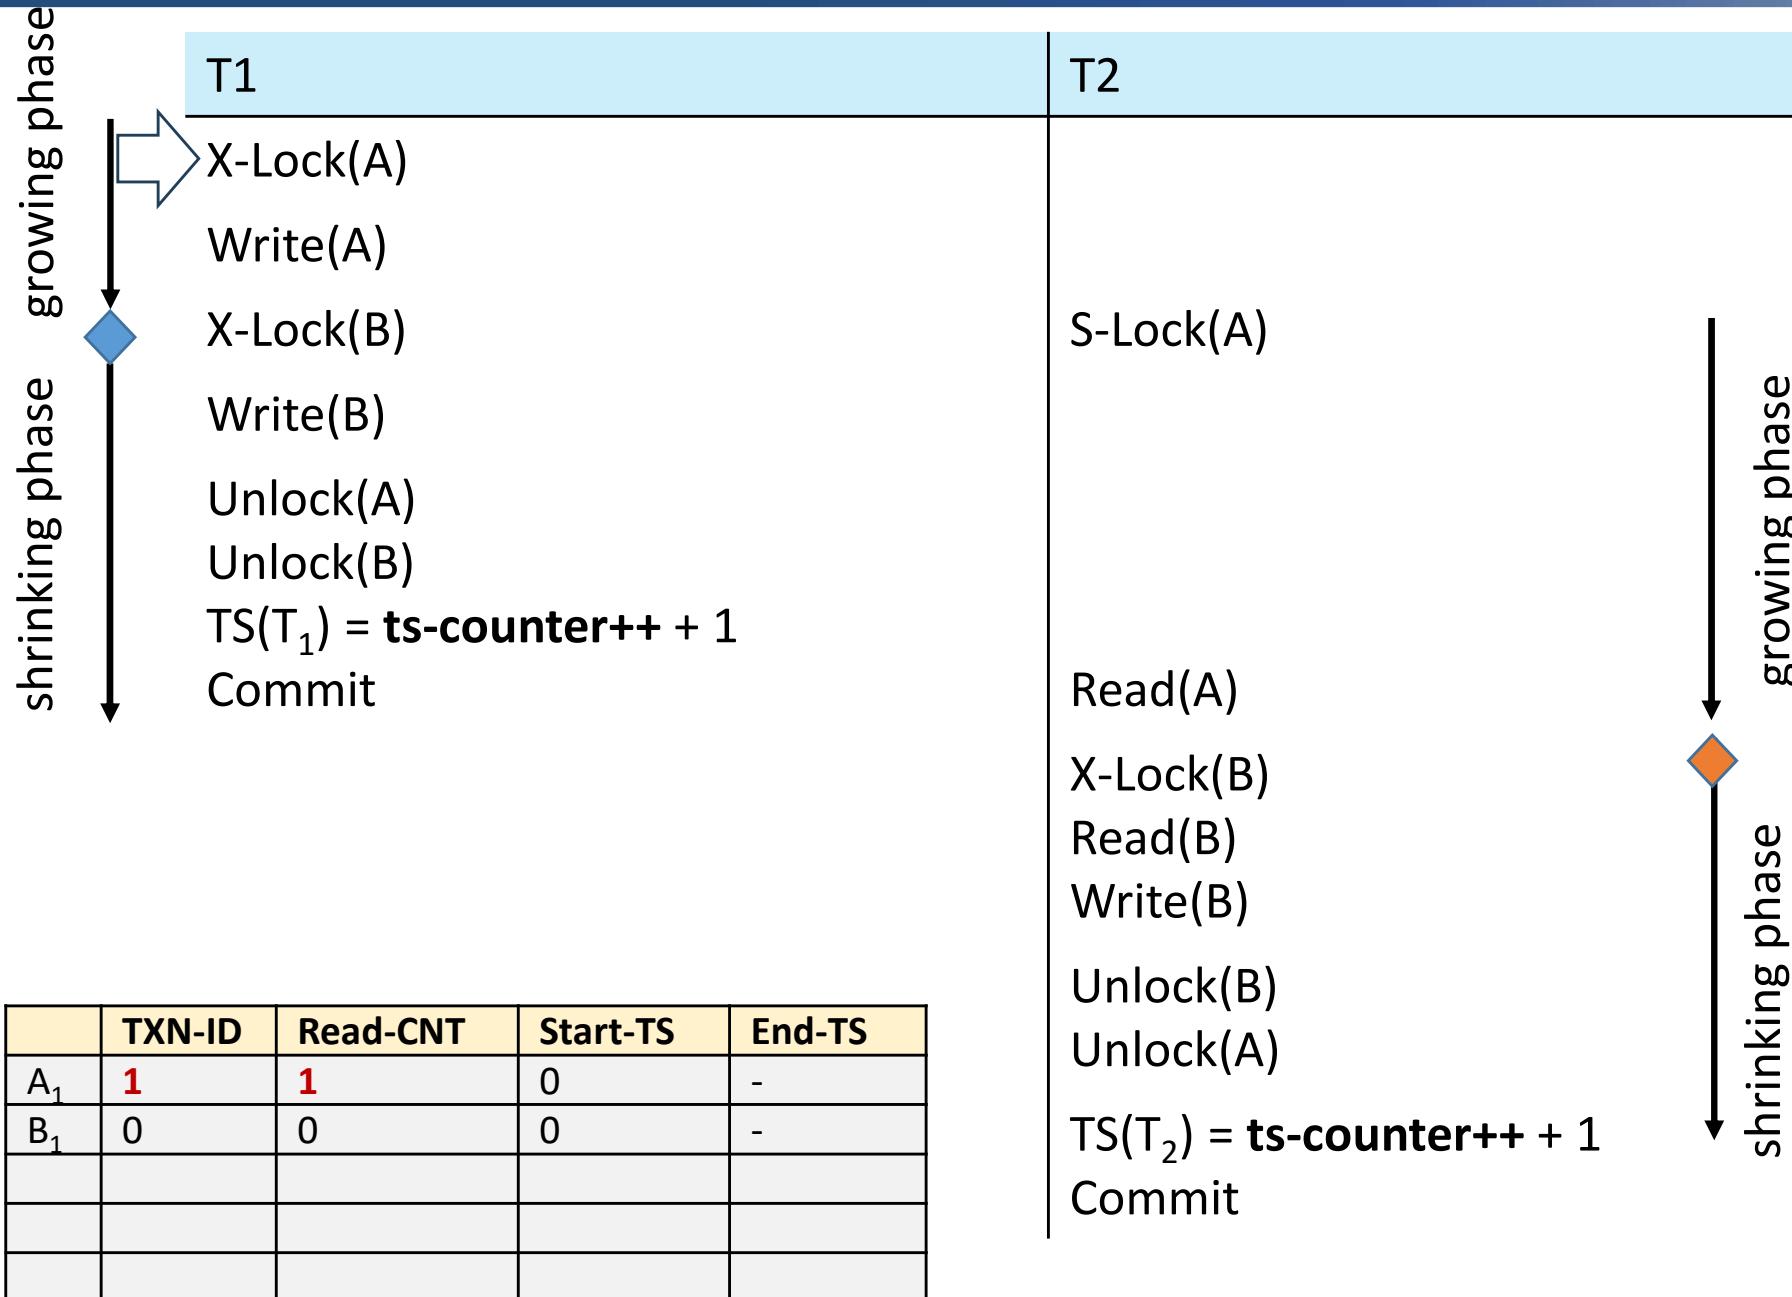
- Use TXN-ID and Read-CNT together as EXCLUSIVE lock.
- If both TXN-ID and Read-CNT are zero, then transaction acquires the EXCLUSIVE lock by setting both of them.

# Implementation of MV2PL

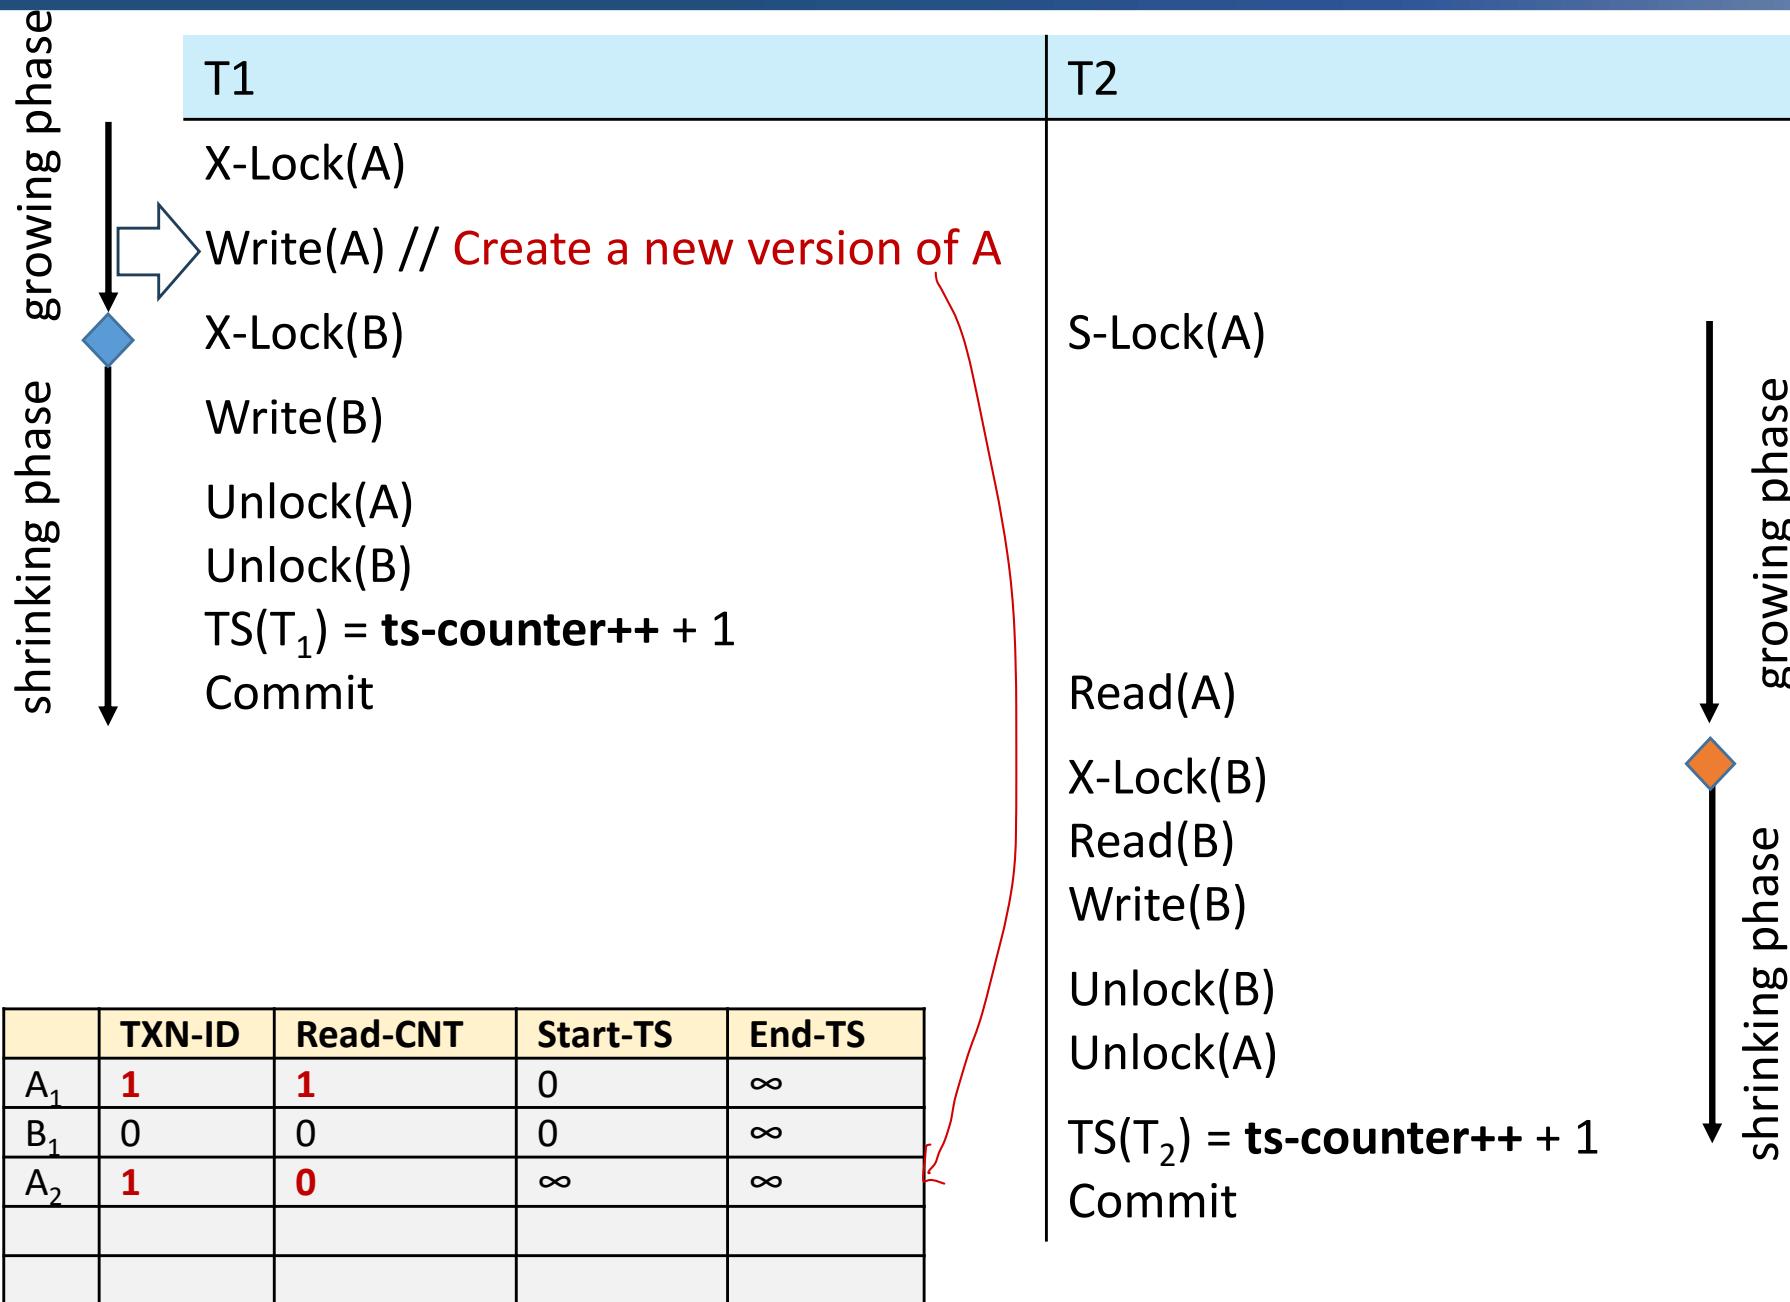


- Use TXN-ID and Read-CNT together as EXCLUSIVE lock.
- If both TXN-ID and Read-CNT are zero, then transaction acquires the EXCLUSIVE lock by setting both of them.

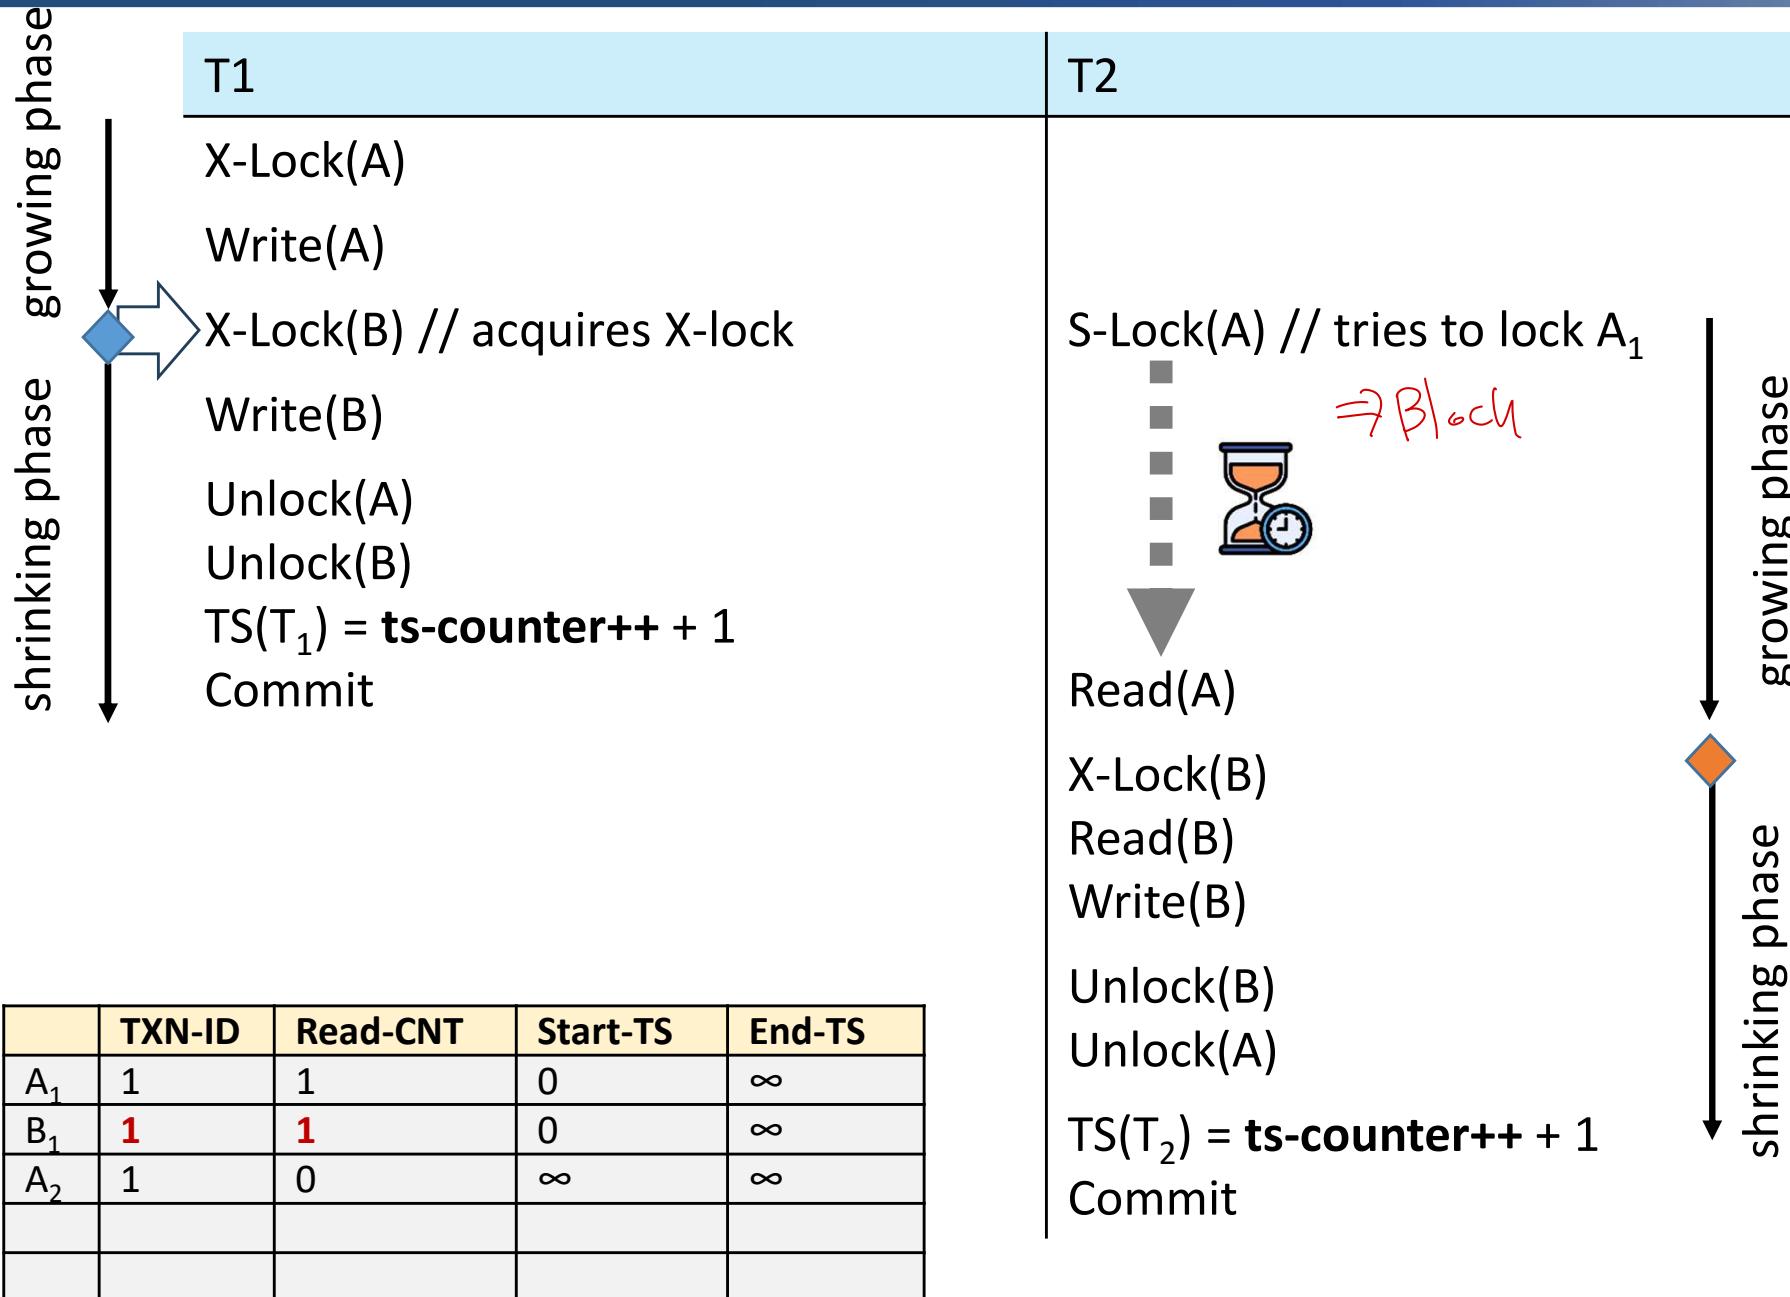
# Multi-Version Two-Phase Locking Protocol (MV2PL) - Example



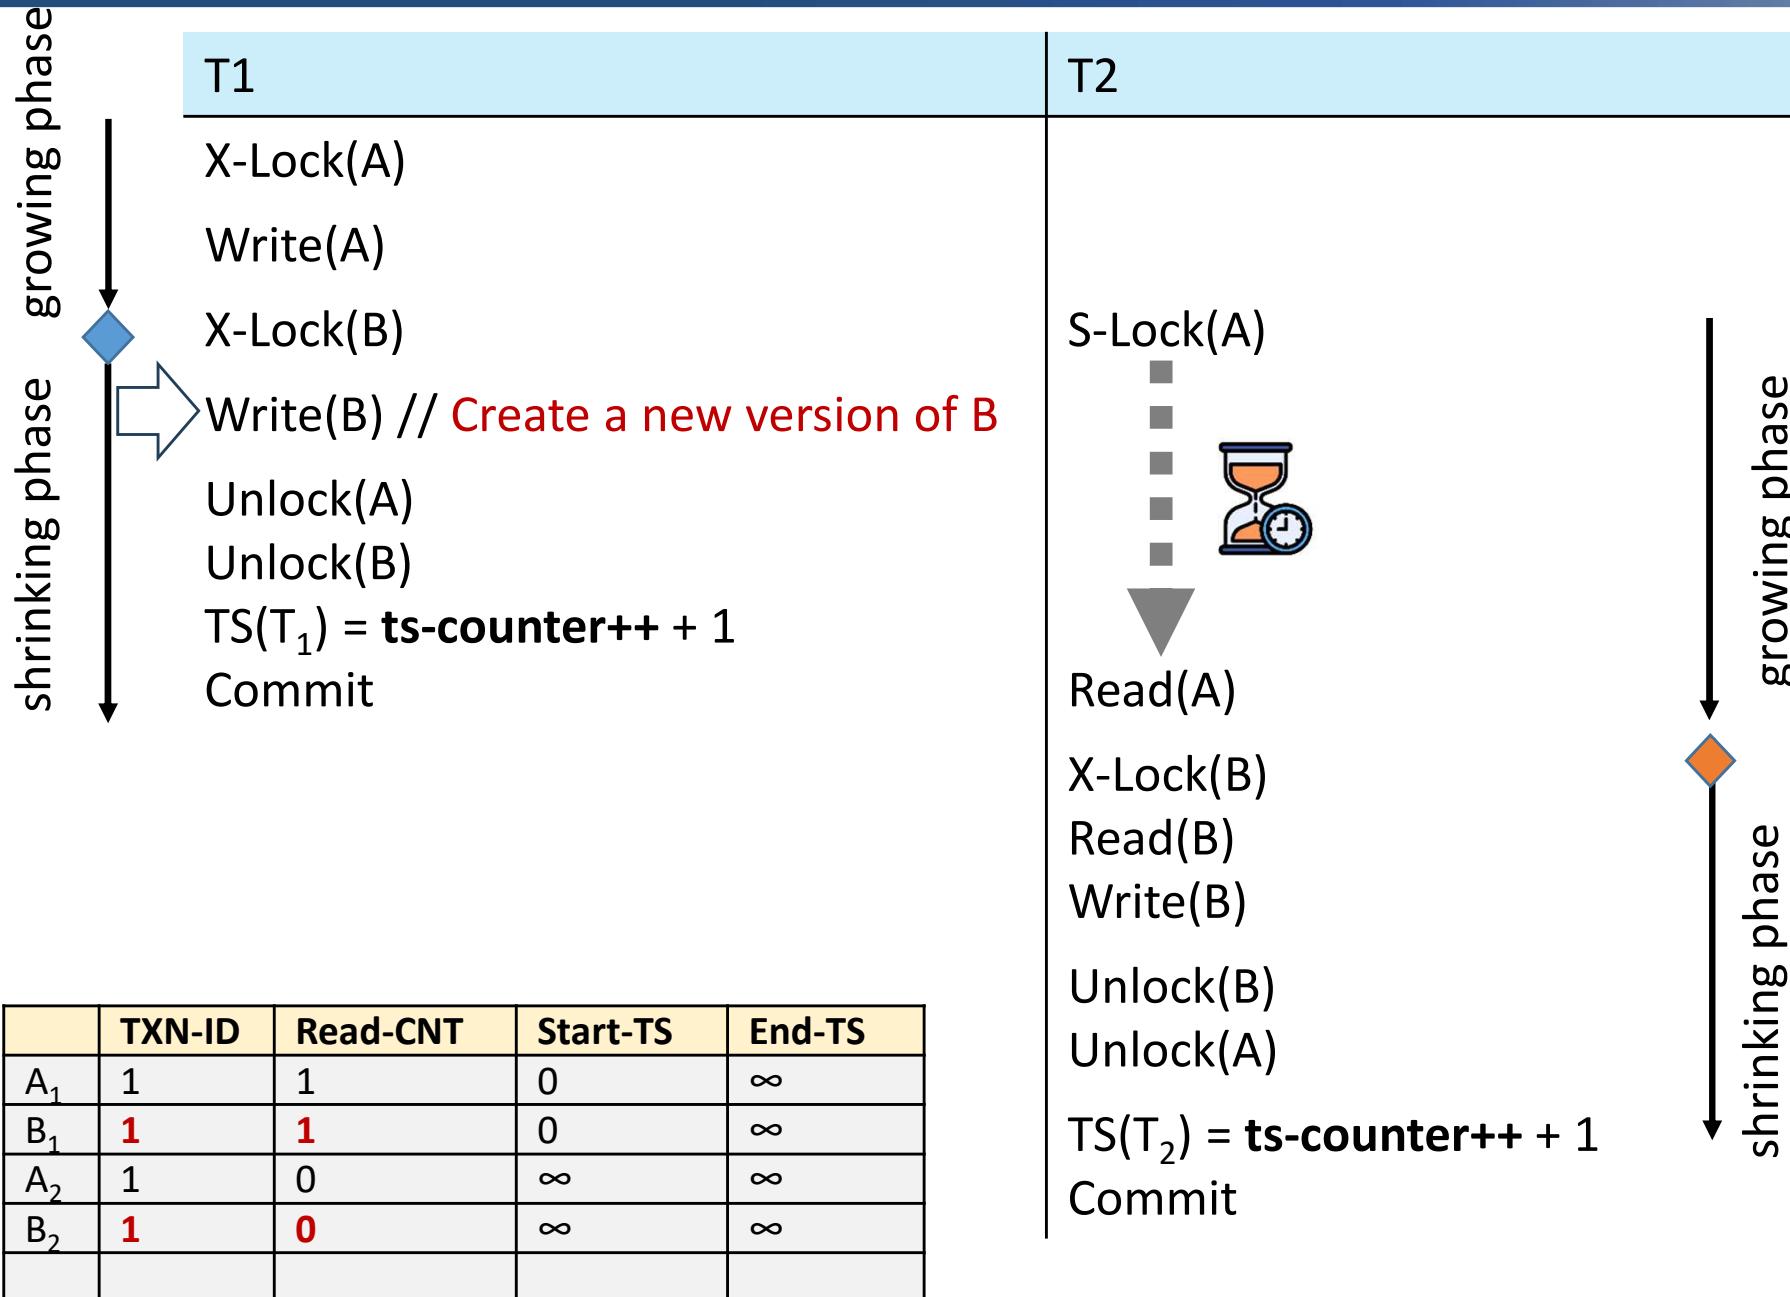
## Multi-Version Two-Phase Locking Protocol (MV2PL) - Example



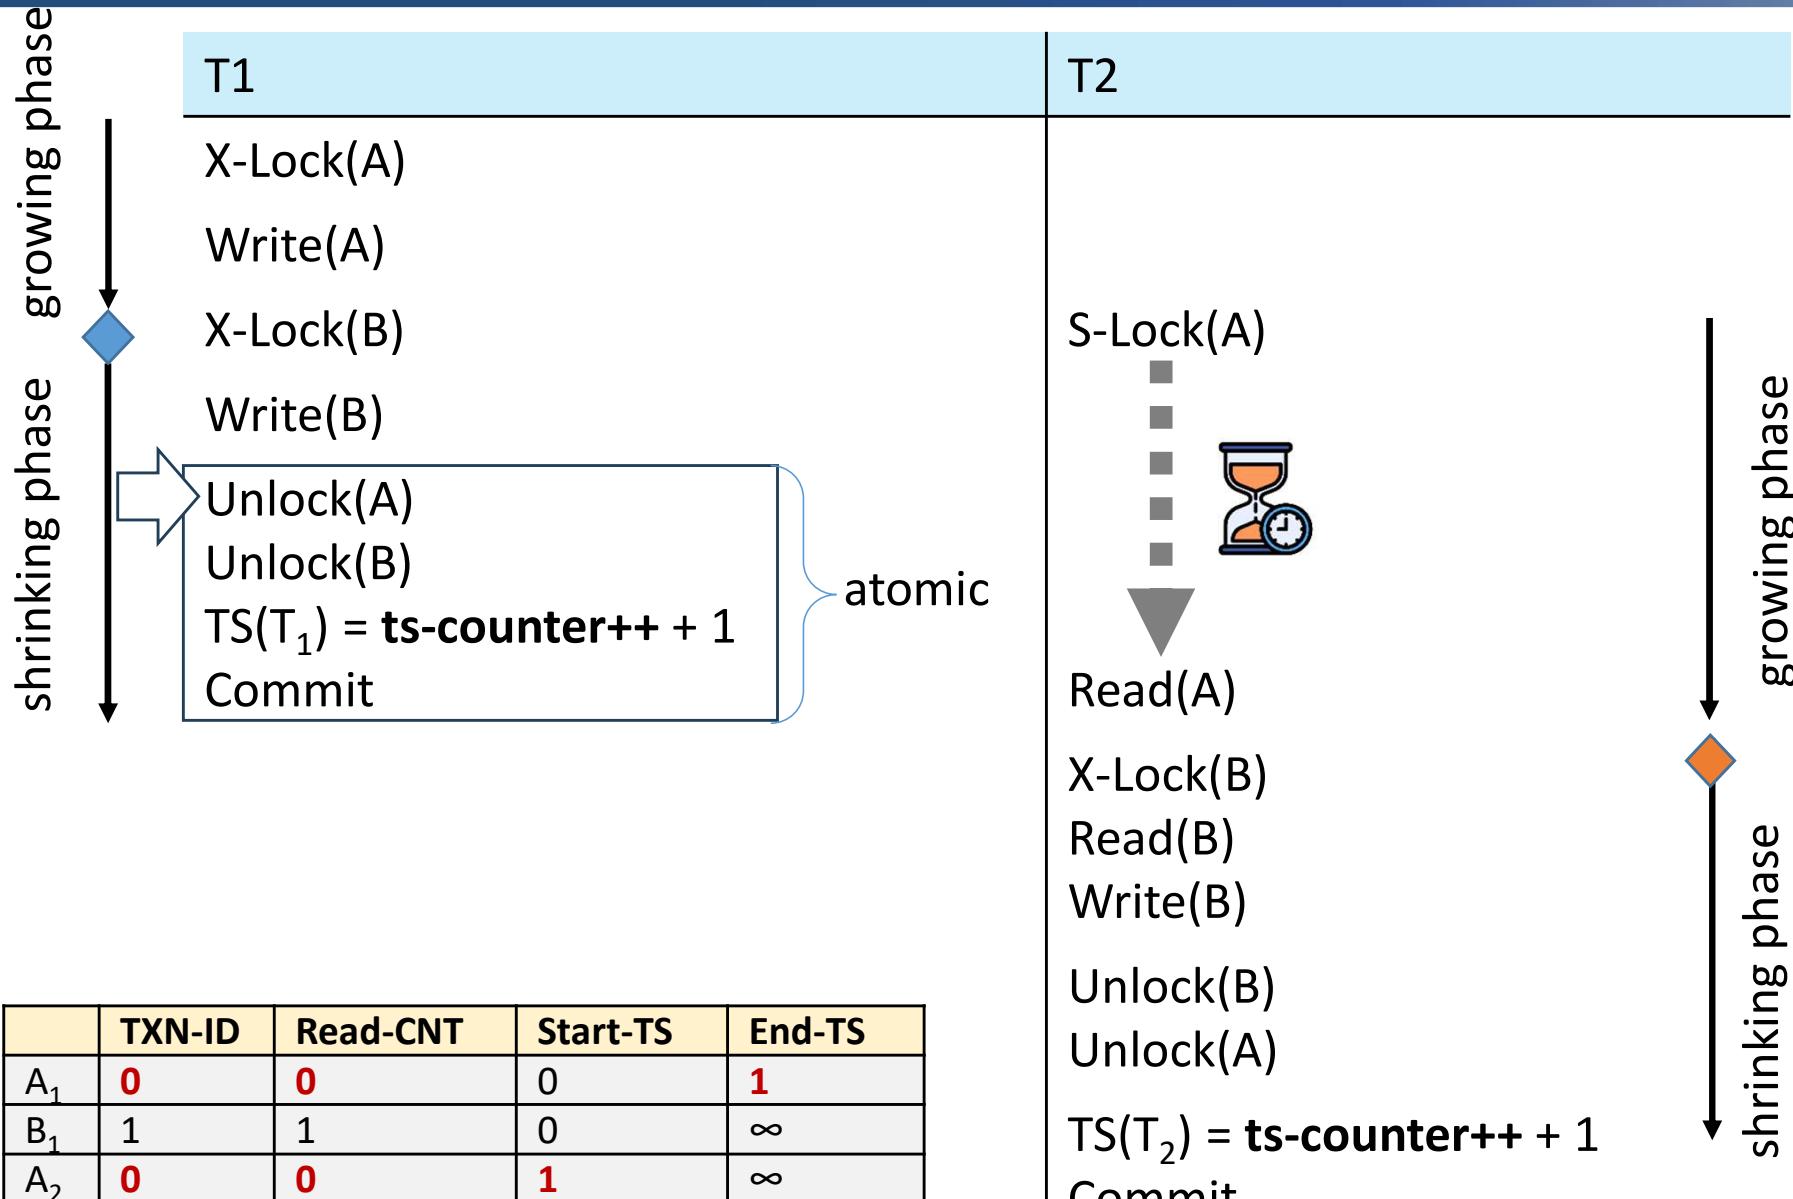
# Multi-Version Two-Phase Locking Protocol (MV2PL) - Example



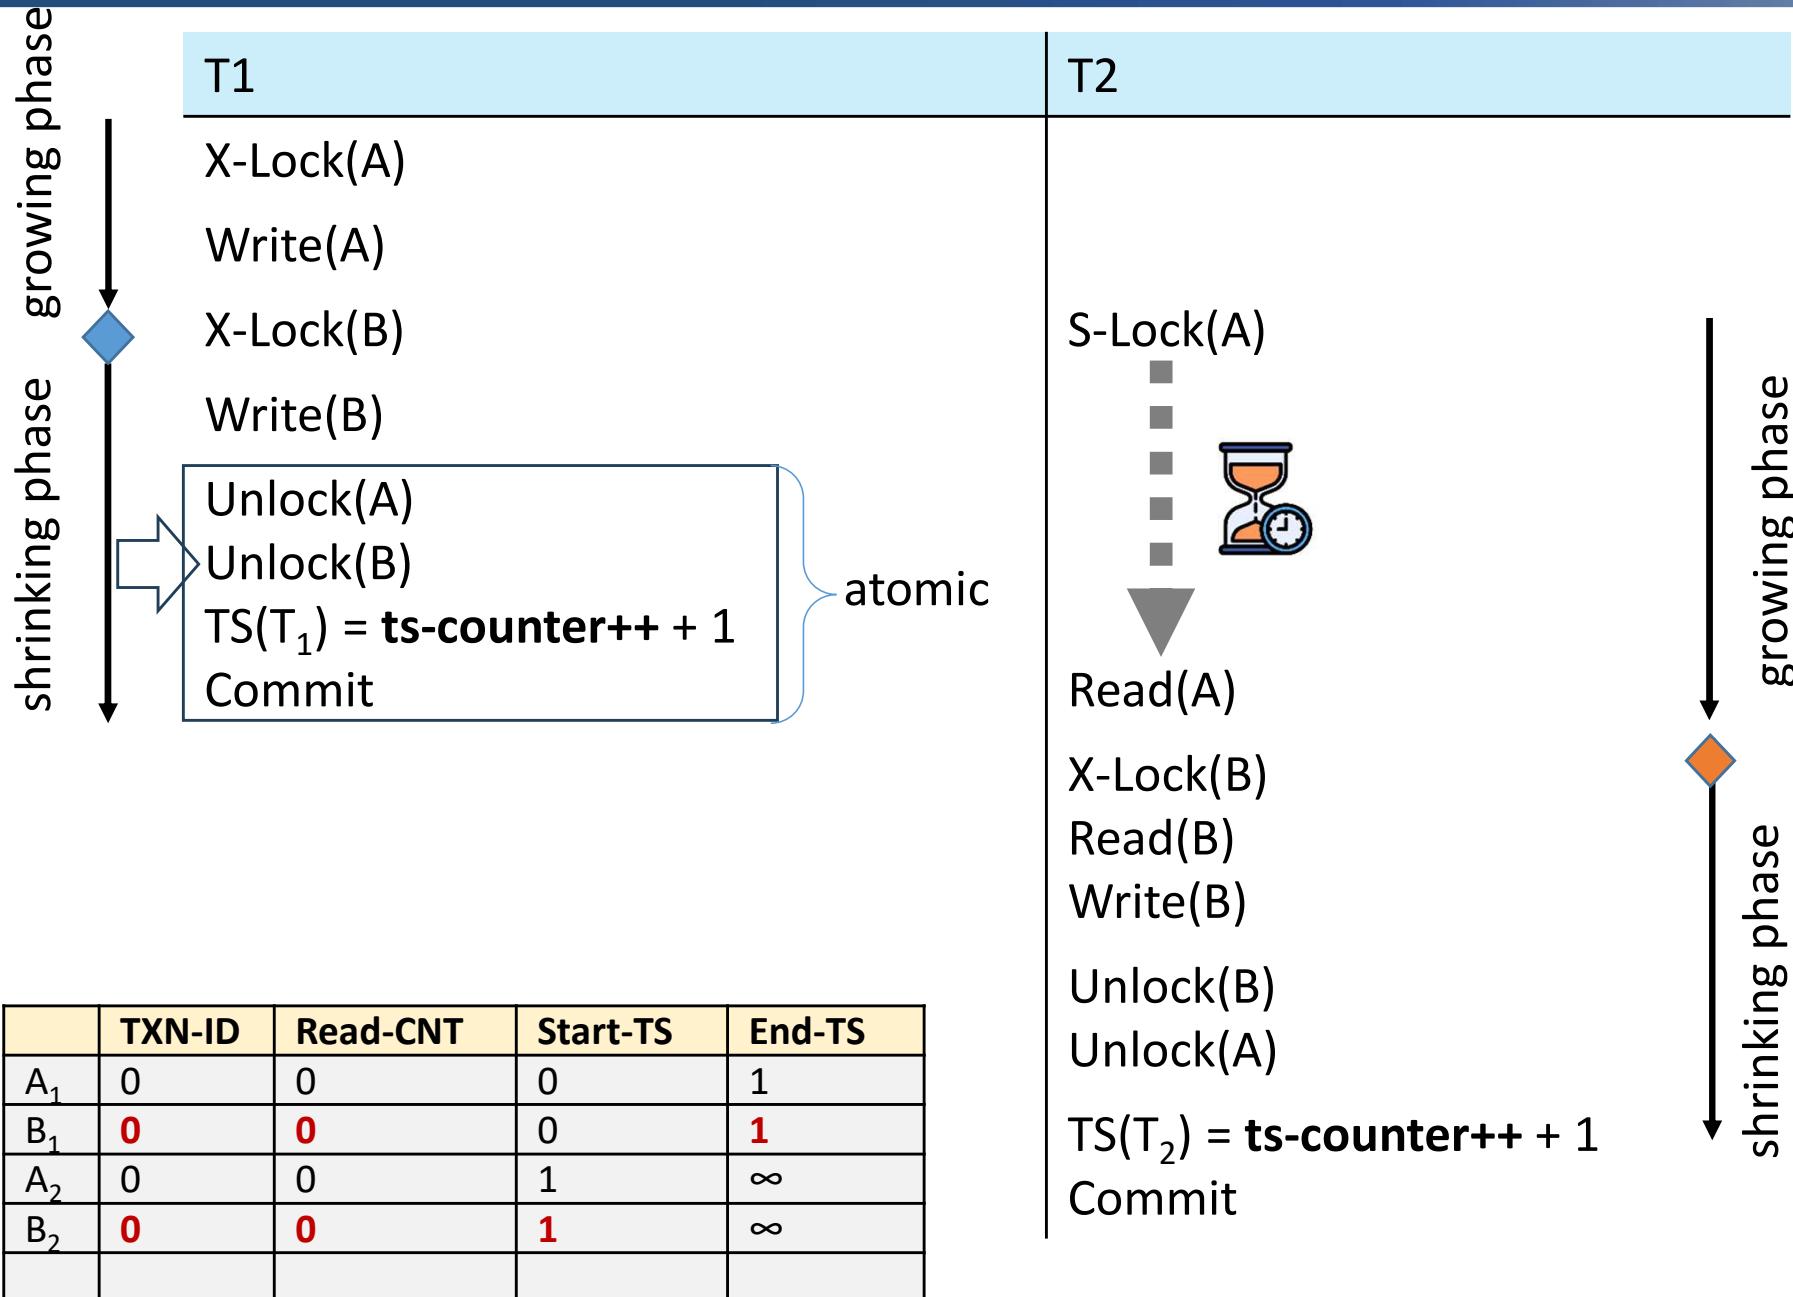
# Multi-Version Two-Phase Locking Protocol (MV2PL) - Example



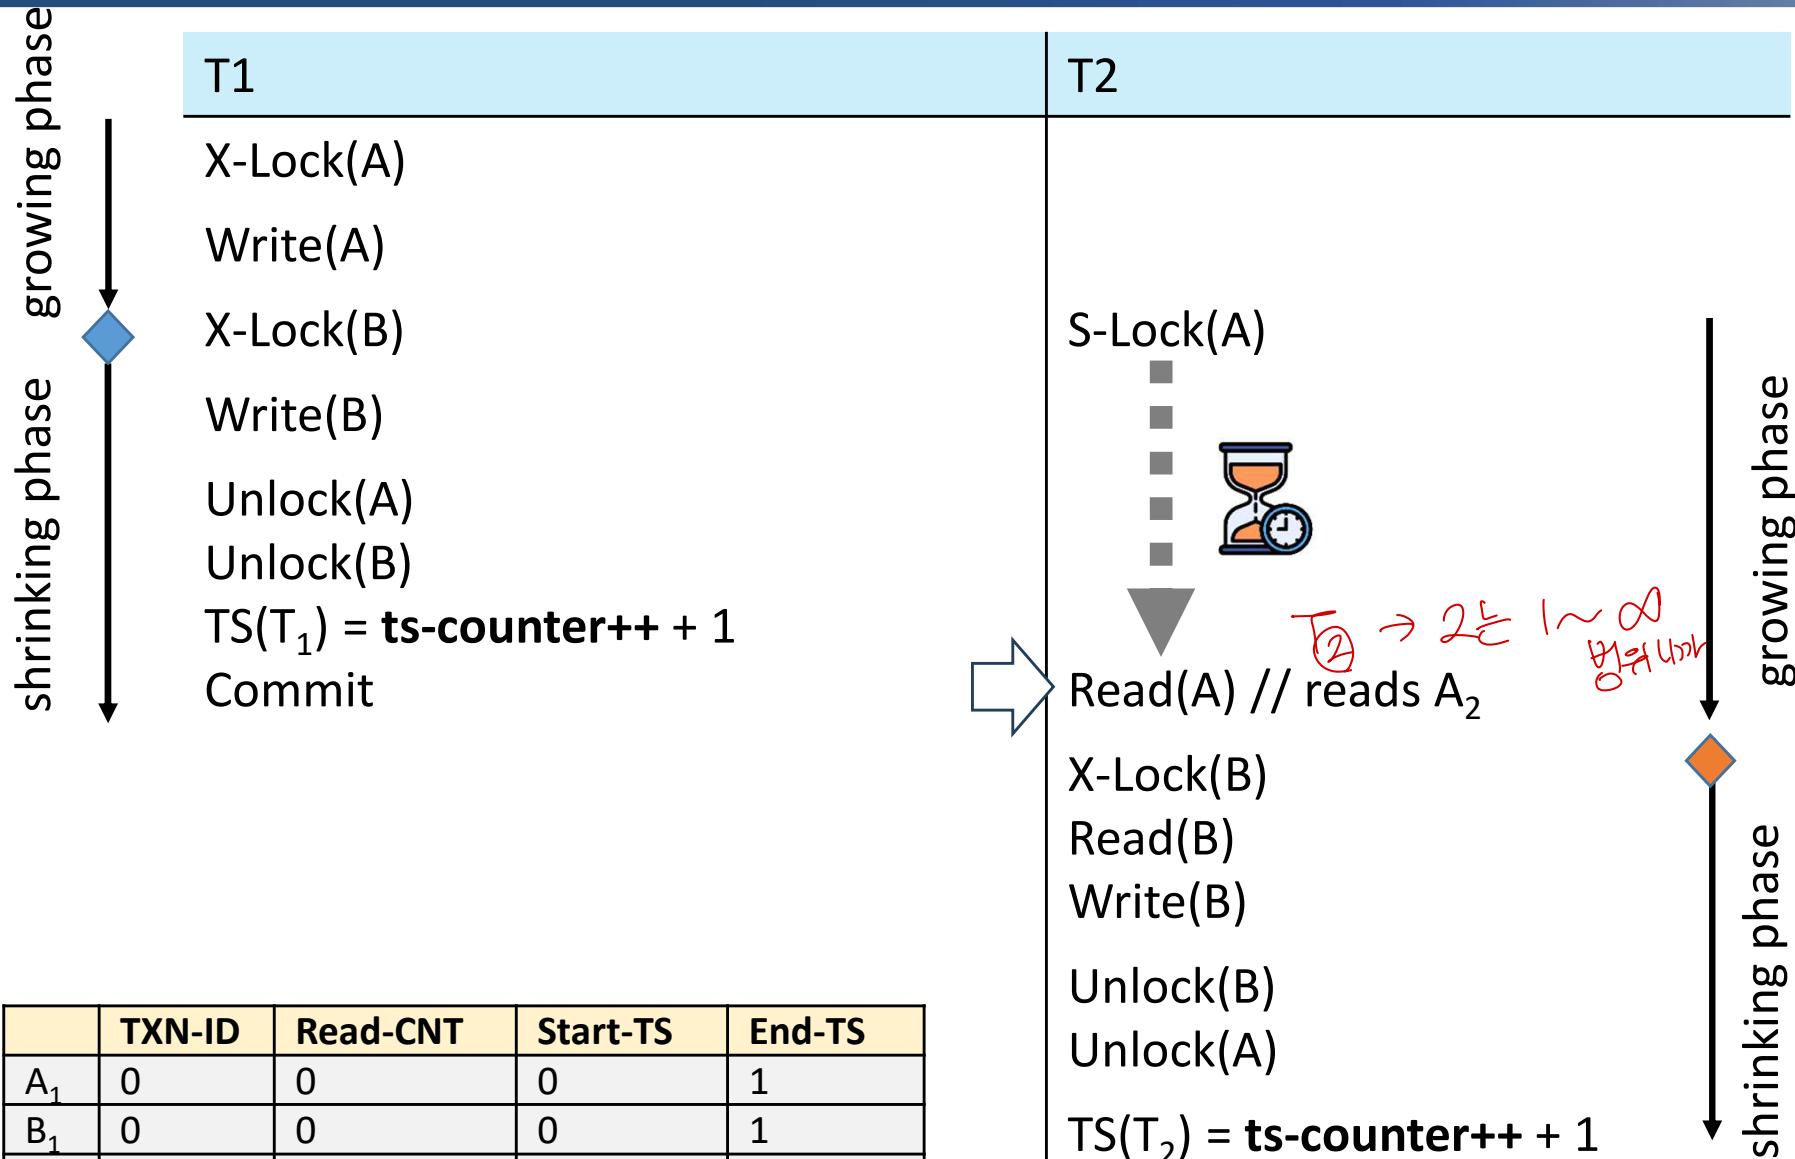
# Multi-Version Two-Phase Locking Protocol (MV2PL) - Example



# Multi-Version Two-Phase Locking Protocol (MV2PL) - Example

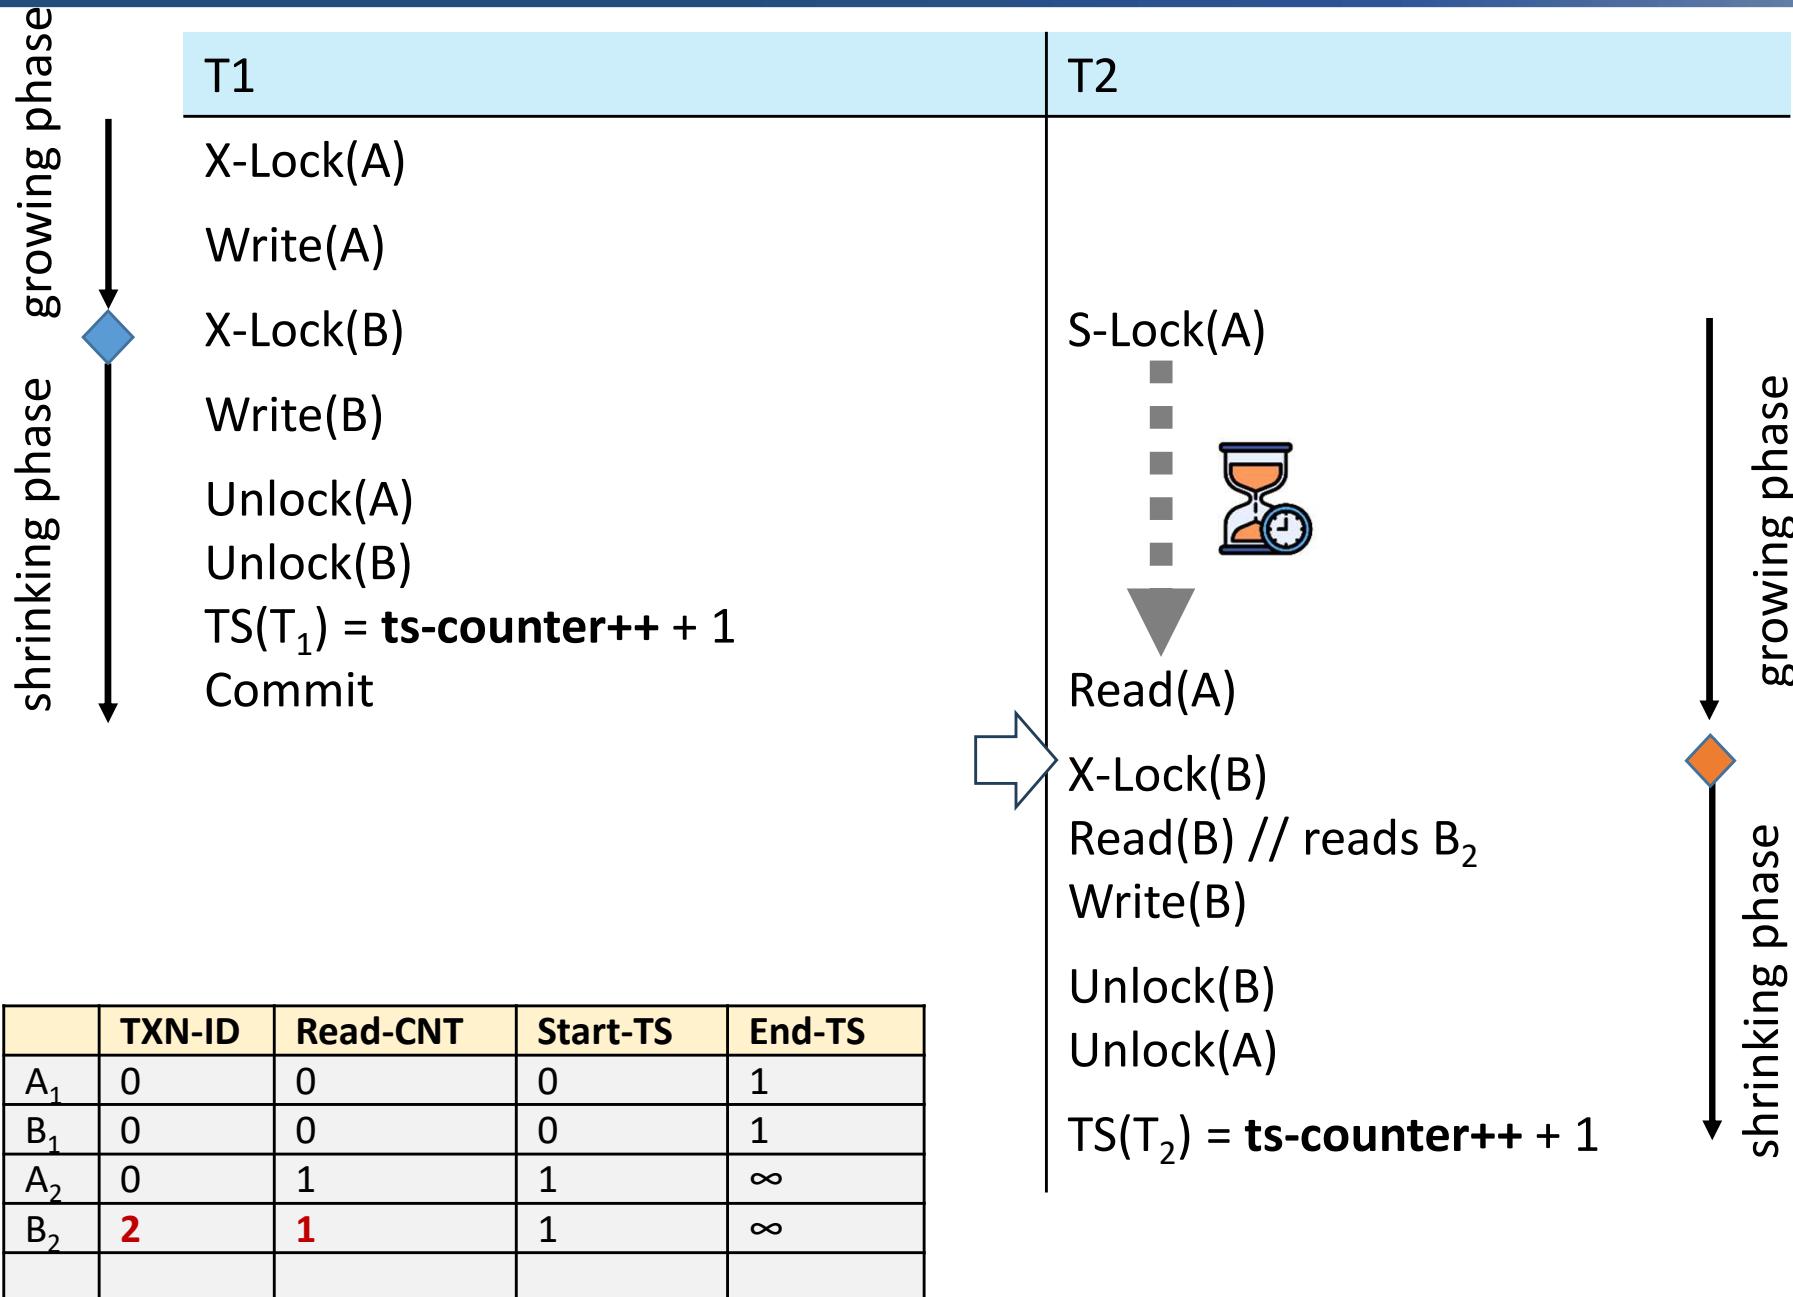


# Multi-Version Two-Phase Locking Protocol (MV2PL) - Example

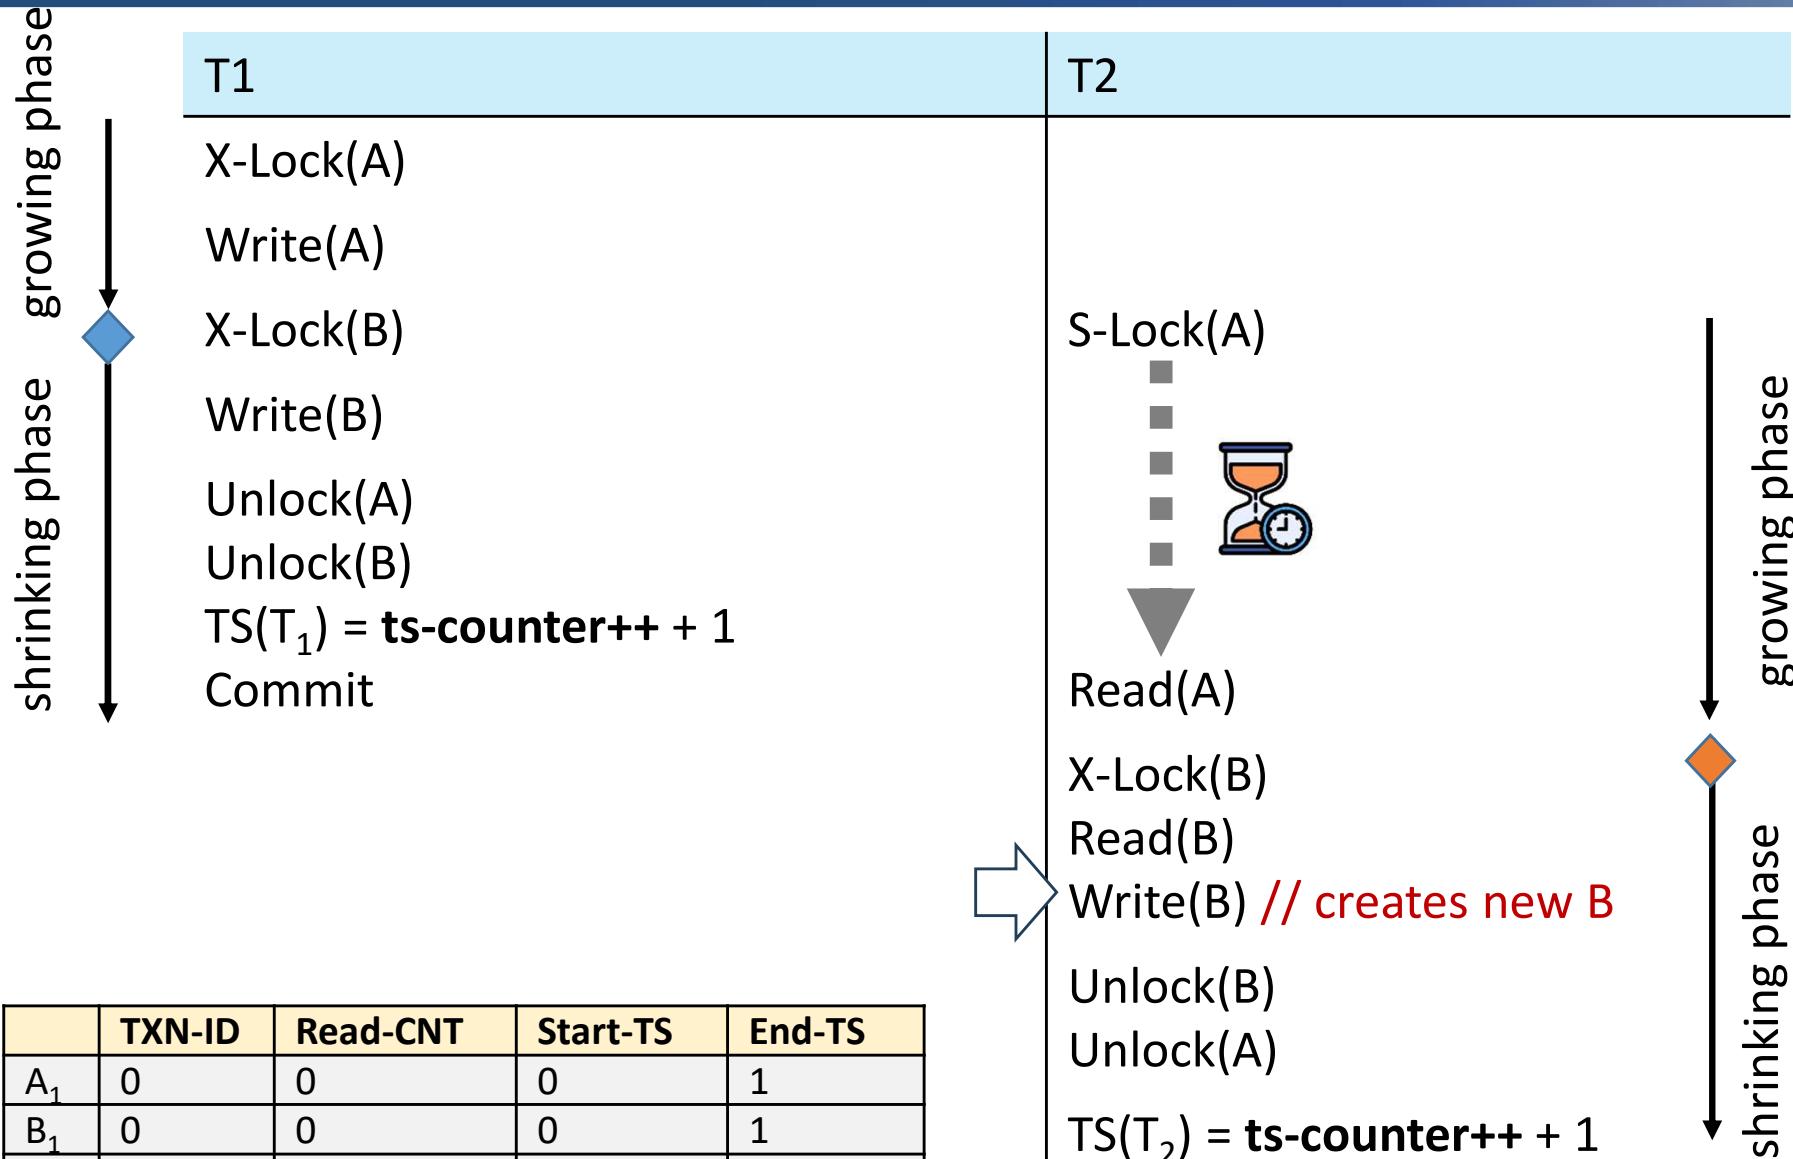


	TXN-ID	Read-CNT	Start-TS	End-TS
A <sub>1</sub>	0	0	0	1
B <sub>1</sub>	0	0	0	1
A <sub>2</sub>	0	1	1	$\infty$
B <sub>2</sub>	0	0	1	$\infty$

# Multi-Version Two-Phase Locking Protocol (MV2PL) - Example

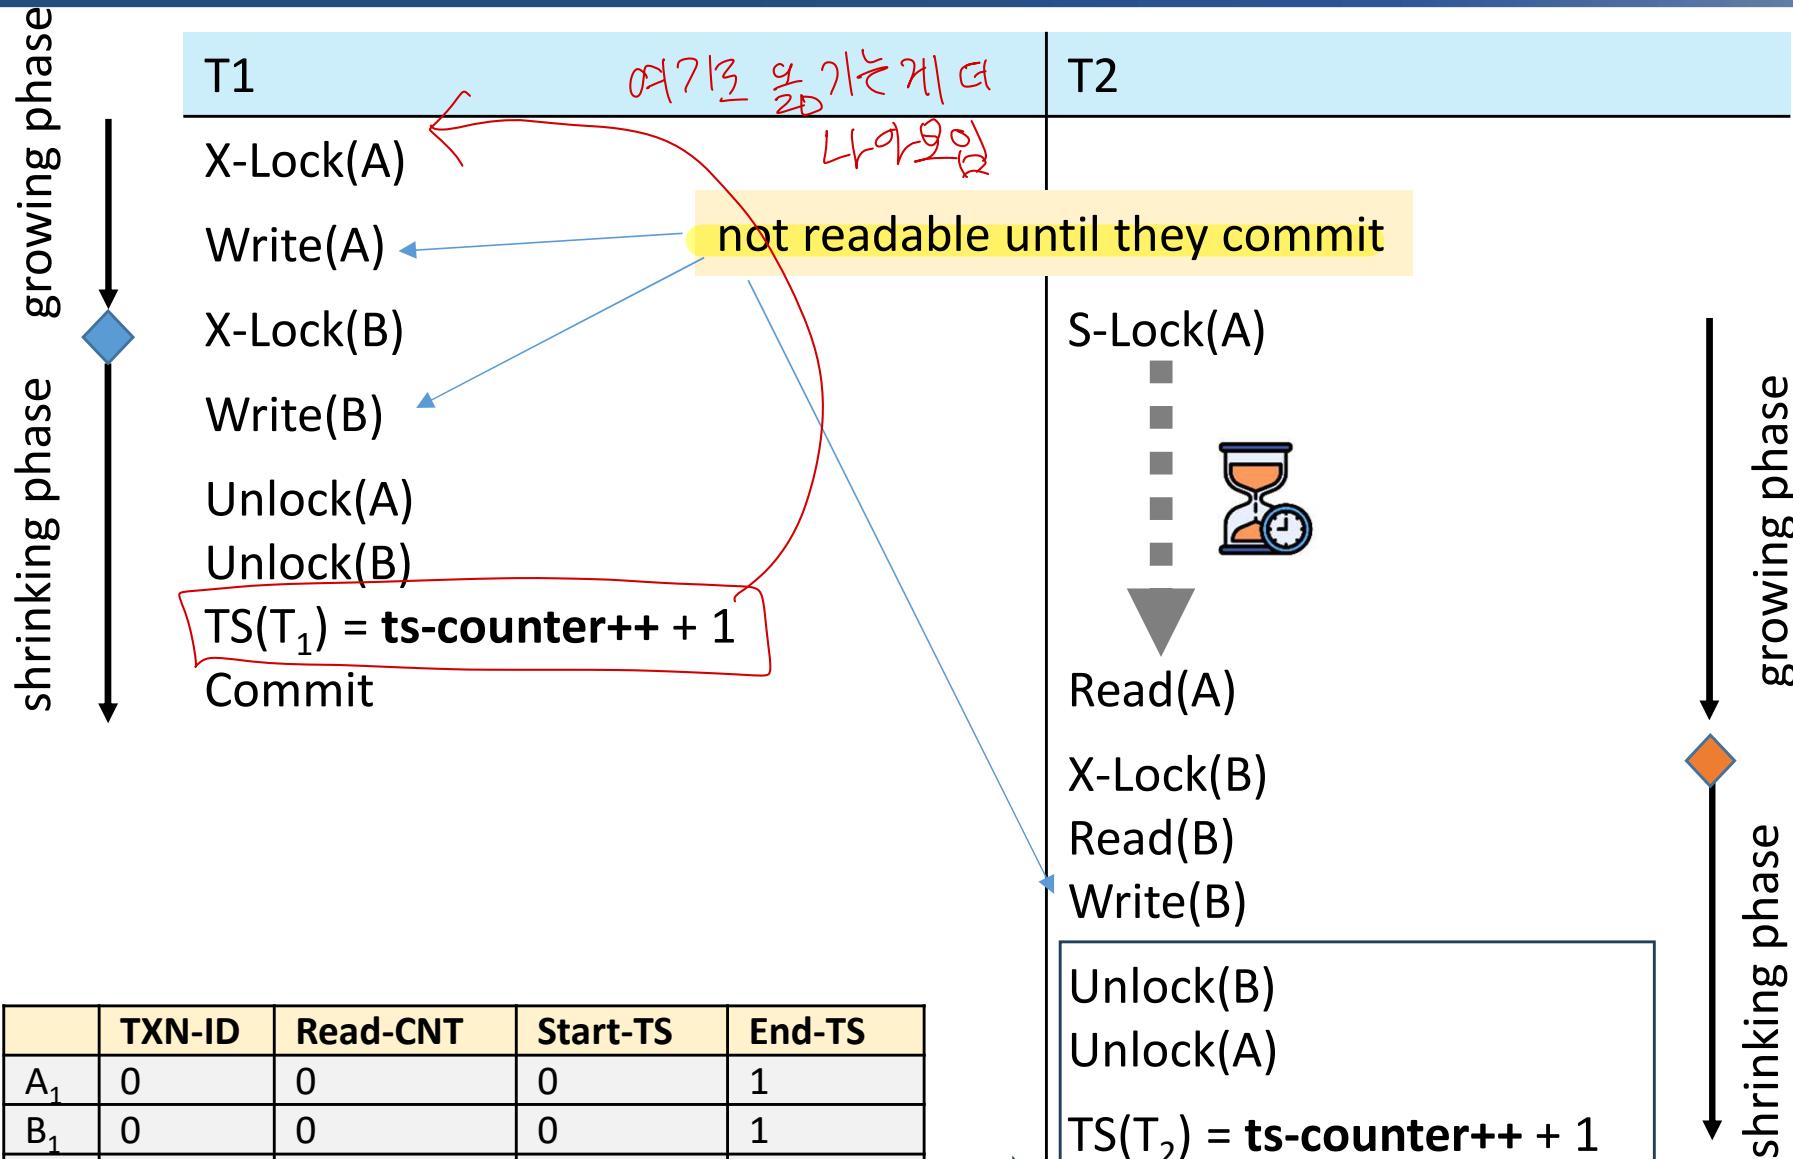


# Multi-Version Two-Phase Locking Protocol (MV2PL) - Example



	TXN-ID	Read-CNT	Start-TS	End-TS
A <sub>1</sub>	0	0	0	1
B <sub>1</sub>	0	0	0	1
A <sub>2</sub>	0	1	1	$\infty$
B <sub>2</sub>	2	1	1	$\infty$
B <sub>3</sub>	<b>2</b>	<b>0</b>	$\infty$	$\infty$

# Multi-Version Two-Phase Locking Protocol (MV2PL) - Example



	TXN-ID	Read-CNT	Start-TS	End-TS
A <sub>1</sub>	0	0	0	1
B <sub>1</sub>	0	0	0	1
A <sub>2</sub>	0	1	1	$\infty$
B <sub>2</sub>	2	1	1	2
B <sub>3</sub>	2	0	2	$\infty$

# MVCC: Implementation Issues

- Creation of multiple versions increases storage overhead
  - Extra tuples
  - Extra space in each tuple for storing version information
- Versions need to be *garbage collected*
  - E.g., if Q has two versions Q5 and Q9, and the oldest active transaction has timestamp > 5, Q5 will never be required again

<기존 방식과의 차이점>

multiversion → historical data 를 잡아서 기존에 abort되는 경우를  
해결하겠다.