

Database Systems

Lecture22 – NoSQL

Not only SQL

ACID

기존 RDBMS와 strong consistency 를
회피한 프레임워크

Beomseok Nam (남범석)
bnam@skku.edu

Ch.8 Complex Data Types

Ch 8. Complex Data Types

- **Semi-Structured Data**
- Object Orientation
- Textual Data
- Spatial Data

Structured Data

- Data organized in a predefined format
 - E.g., database tables

instructor

ID	name	dept_name	salary
1	Alice	Comp. Sci.	20,000
2	Bob	Biology	10,000

department

dept_name	building
Comp. Sci.	Taylor
Biology	Taylor
Elec. Eng.	Watson

- First Normal Form (1NF)
• Each column in a table contains atomic values

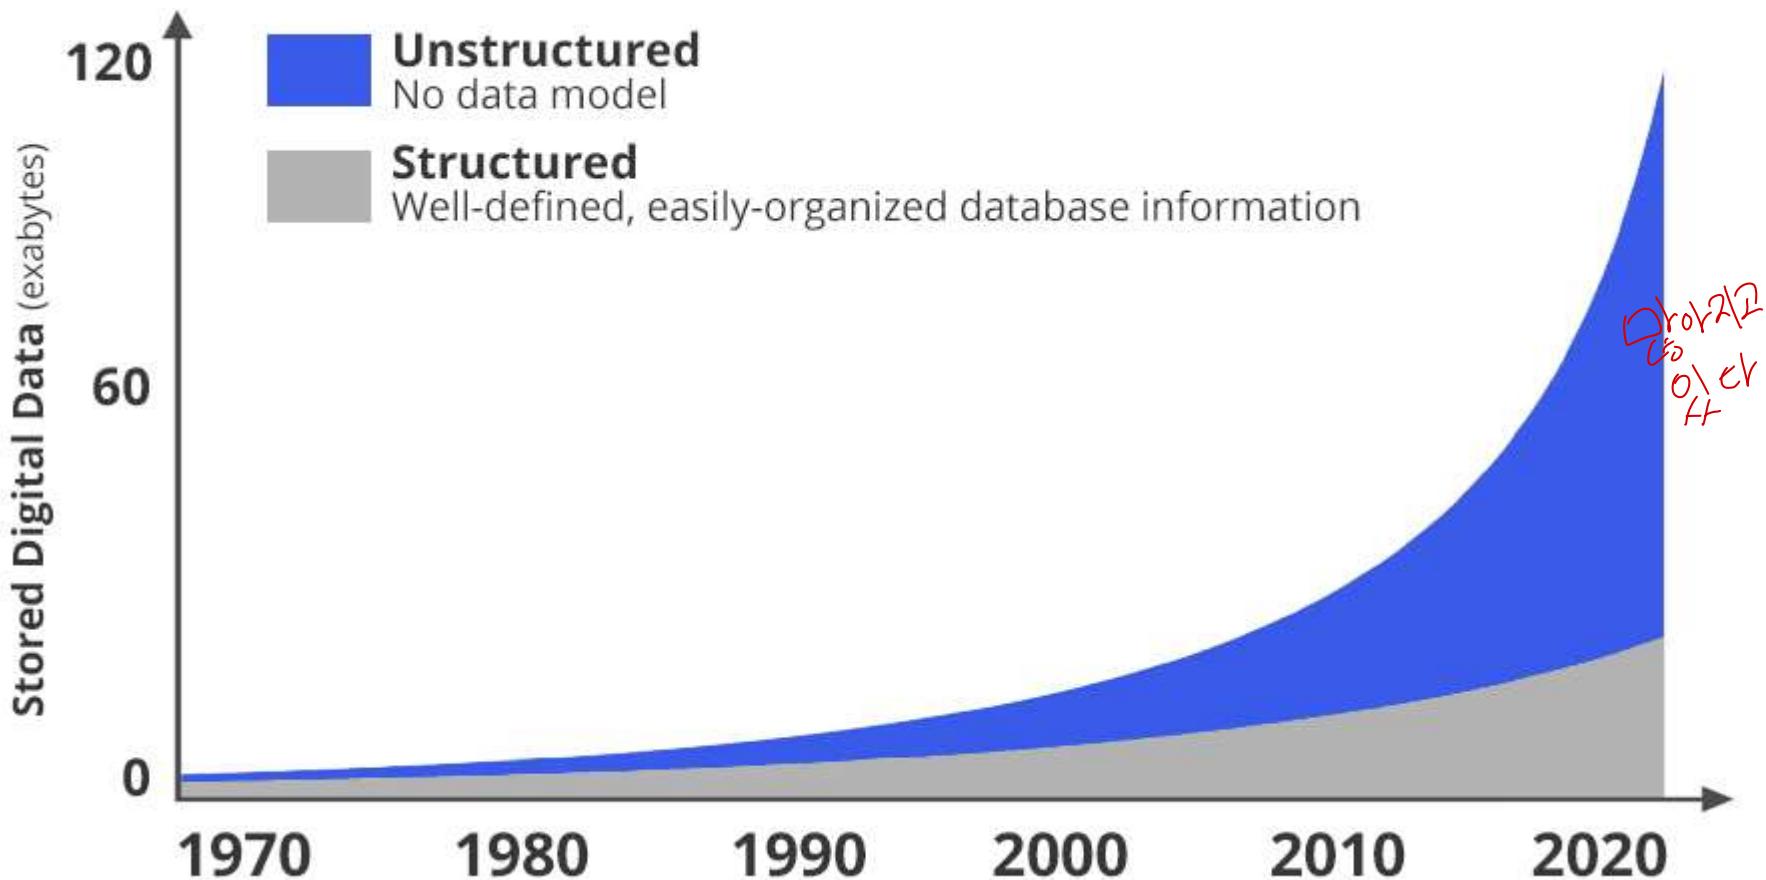
multiple value \rightarrow \nrightarrow value

Unstructured Data

- Data that doesn't have a predefined structure or format
- **Text-based:** Emails, news articles, SNS posts, text files.
- **Non-textual:** Images, audio, video.
- **Challenges:** Lack of structure can make it difficult to search, sort, and analyze.
- **Benefits:** Provides a rich source of information and can be used for various applications like sentiment analysis, customer behavior analysis, and predictive analytics.

Explosion of Unstructured/Semi-Structured Data

- A poor fit for the legacy RDBMS



Graph Source: IDC

Semi-Structured Data </>

- Data that does not conform to a strict, fixed schema, but contains **tags** or **markers** to separate semantic elements and enforce hierarchies of records and fields.
- **Characteristics:**
 - **Self-describing:** contains metadata about its own structure.
 - **Flexible Schema:** Allows for variations in structure, new fields can be added easily.
 - **Hierarchy:** Can represent hierarchical relationships between data elements.
 - **Irregular or Incomplete:** Not all data elements may be present for every record.

Why is Semi-Structured Data Important?

- **Web Data:** Dominant format for data exchange on the internet.
- **Big Data:** Handles the variability and volume of modern data sources.

Patton Oswalt

@pattonoswalt

Follow

My dong is super-friendly and loves getting rubbed by children. #CareerEndingTwitterTypos

12:39 PM - 4 Nov 2013

462 RETWEETS 639 FAVORITES

{ key-value pair

```
"tweet_id": "RF_pattonoswalt_201311041239",
"user": { "username": "pattonoswalt", "id": "pattonoswalt" },
"text": "My dong is super-friendly and loves getting rubbed by children.",
"created_at": "Mon Nov 04 12:39:00 +0000 2013",
"hashtags": [ "CareerEndingTwitterTypos" ]
```

}

값과 키로
구조를 갖게
변환

XML

서버 구조 데이터의 어떤 파일로 맷 ①

- XML uses tags to mark up text

- E.g.

<course>

any name이든 가능
↓

<course id> CS-101 </course id>

<title> Intro. to Computer Science </title>

<dept name> Comp. Sci. </dept name>

<credits> 4 </credits>

</course>

- Tags make the data self-describing

- Tags can be hierarchical

nested 가능

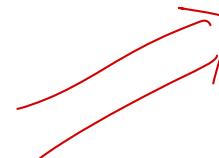
JSON (JavaScript Object Notation)

메인 파일 티켓 ②

- Example 1:

```
{  
    "ID": "1111",  
    "name": {  
        "firstname": "Albert",  
        "lastname": "Einstein"  
    },  
    "deptname": "Physics",  
    "children": [  
        {"firstname": "Hans", "lastname": "Einstein"},  
        {"firstname": "Eduard", "lastname": "Einstein"}  
    ]  
}
```

No Fixed Schema!!



- Example 2:

```
{  
    "ID": "22222",  
    "name": {  
        "Beomseok Nam"  
    },  
    "deptname": "Computer Science",  
    "e-mail": "bnam@skku.edu"  
}
```

Features of Semi-Structured Data Models

■ Multivalued data types

JSON: INF를 만들 수 있음 → 멀티밸류 사용 가능

- **Sets, multisets**

- E.g.,: set of interests {'soccer', 'cooking', 'game', 'rock'}

- **Key-value map**

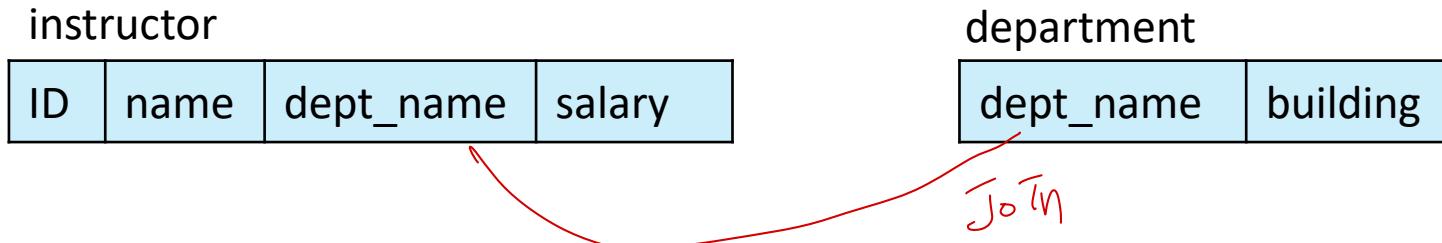
- Store a set of key-value pairs
 - E.g., {(brand, Apple), (ID, MacBook), (size, 13), (color, silver)}
 - Operations on maps: *put(key, value)*, *get(key)*, *delete(key)*

- **Arrays**

- Widely used for scientific and monitoring applications

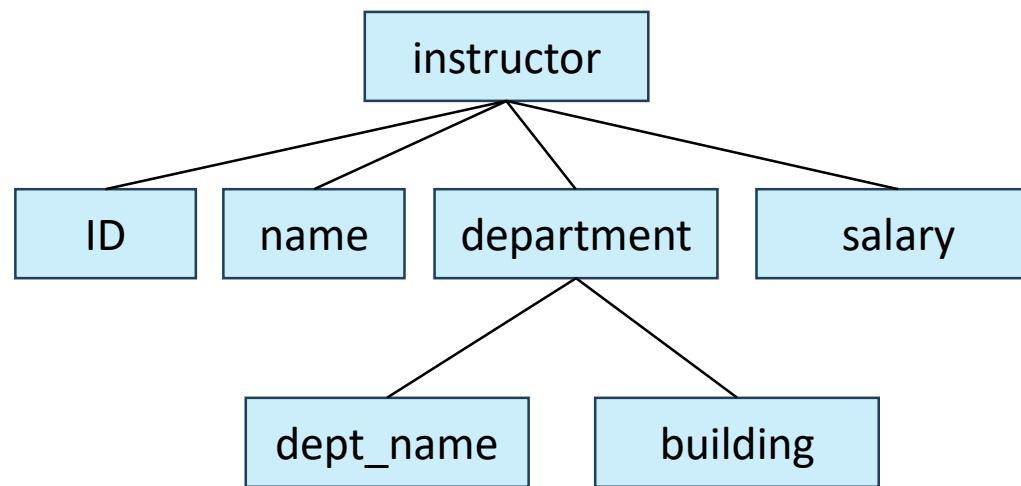
1st Normal Form vs. Nested Data

- 1st Normal Form *RDBMS*



Vs.

- Nested Data (JSON)



Data Model for Semi-Structured Data

- ## ▪ Structured Table for Semi-Structured Data?

- Does not allow schema changes
 - Too many columns
 - Sparse tables



The figure displays a sparse distribution of green squares on a white grid. The grid is composed of small, equal-sized squares. The green squares are concentrated in several vertical columns, with varying heights. A notable feature is a single green square located in the bottom-left corner of the grid. The grid extends from approximately x=50 to x=950 and y=50 to y=1000.

Data Model for Semi-Structured Data

- Structured Table for Semi-Structured Data?
 - Does not allow schema changes
 - Too many columns
 - Sparse tables

Structured Table (Schema)

ID	name	dept_name	salary
22222	Einstein	Physics	95000
12121	Wu	Finance	90000
32343	El Said	History	60000
45565	Katz	Comp. Sci.	75000
98345	Kim	Elec. Eng.	80000
76766	Crick	Biology	7
10101	Srinivasan	Comp. Sci.	65000
58583	Califieri	History	62000
83821	Brandt	Comp. Sci.	92000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
76543	Singh	Finance	80000

열 우선
Column-wise
↓ 방식
row

Key-Value Stores (Schema-less)

KeySpace

column family
=Table → el 테이블

settings

column *EA*
NAME (KEY) VALUE TIMESTAMP

열 우선
Column-wise
↓ 방식
row

Key-Value Store (KVStore)

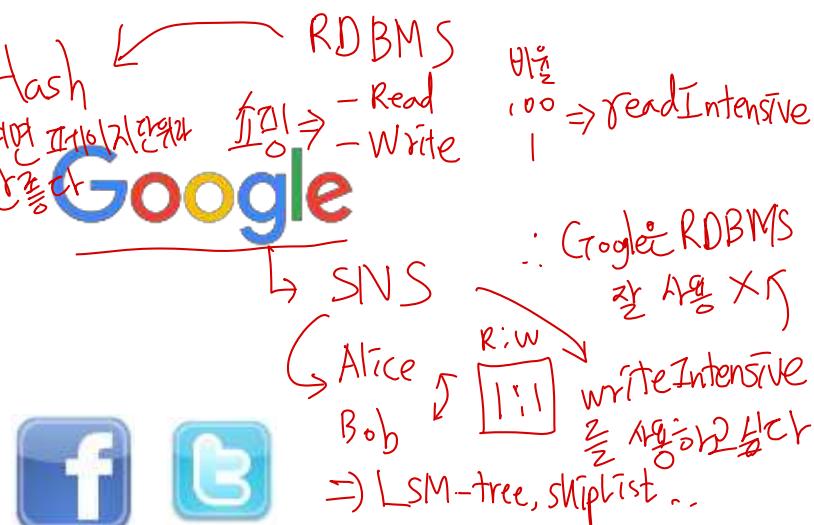
- KV-Stores seem very simple indeed
 - They are nothing but multi-threaded indexing systems ⇒ 일상적인 indexing 시스템
 - A simpler and more scalable “database”

≈ BST, B+tree, Hash

- Interface
 - `put(key, value);` // insert “value” associated with “key”
 - `value = get(key);` // read data associated with “key”

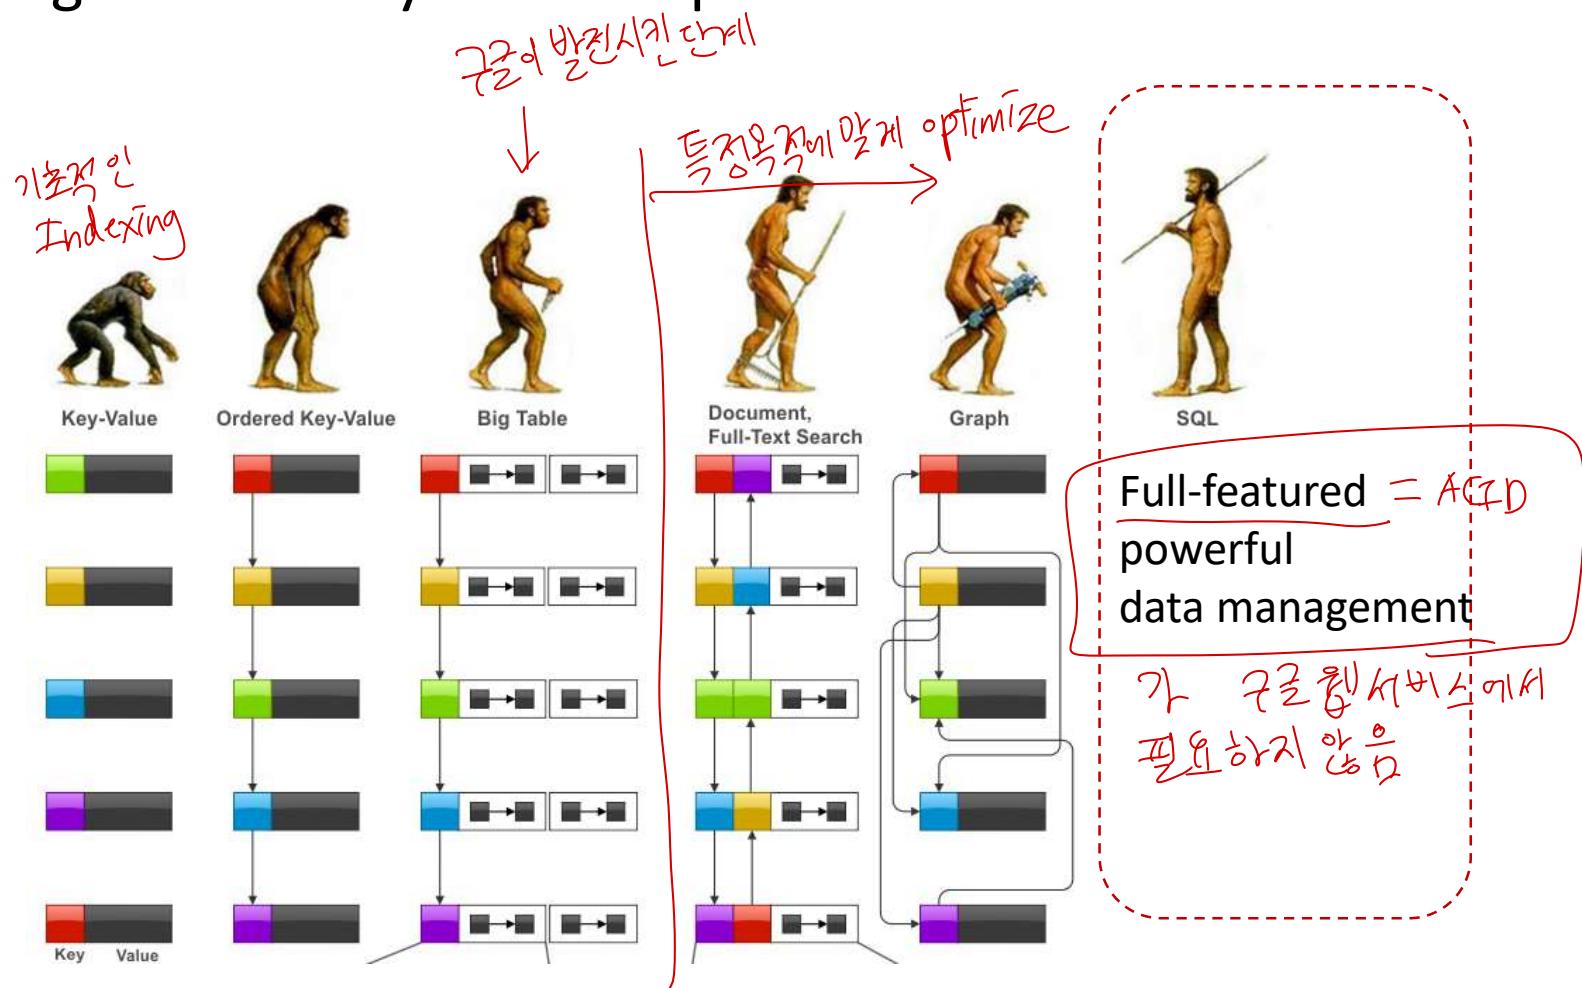
- Examples
 - Google BigTable & internal codes:
– Key: hangoutID
– Value: Hangout conversations

- Facebook, Twitter:
 - Key: UserID
 - Value: user profile (e.g., posting history, photos, friends, ...)



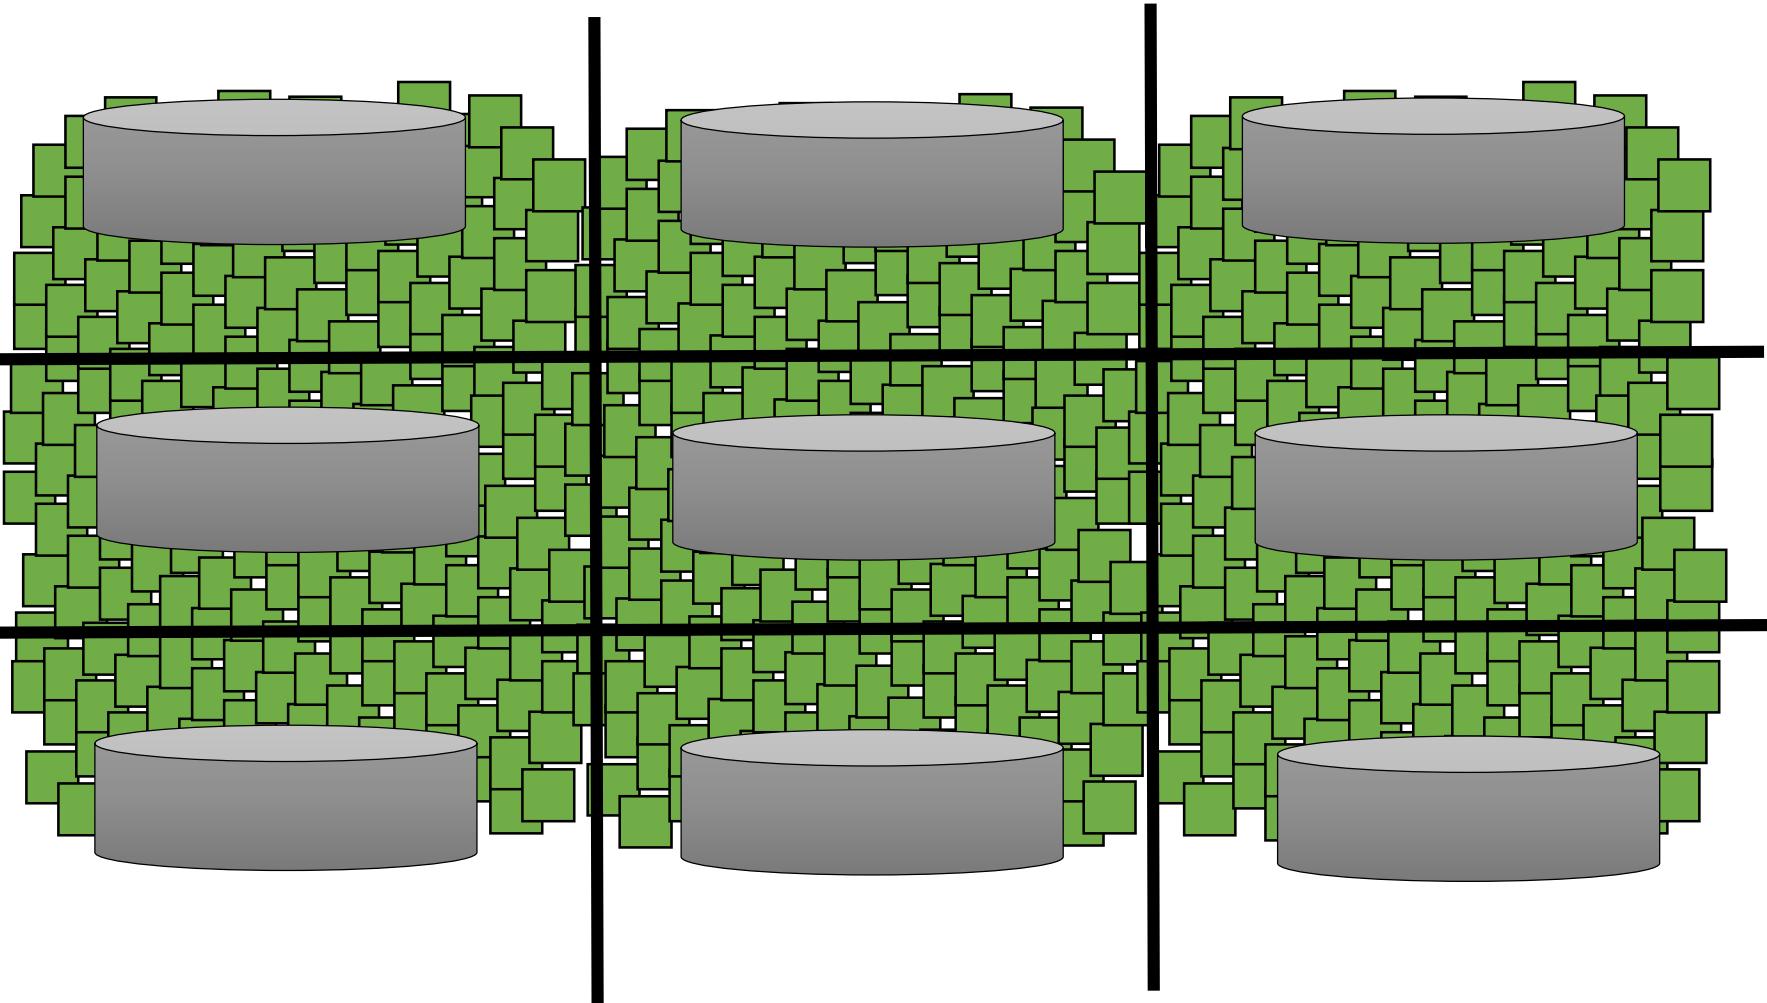
When to use Key-Value Stores over RDBMS?

- If workload is write-intensive
- If strong consistency is not required



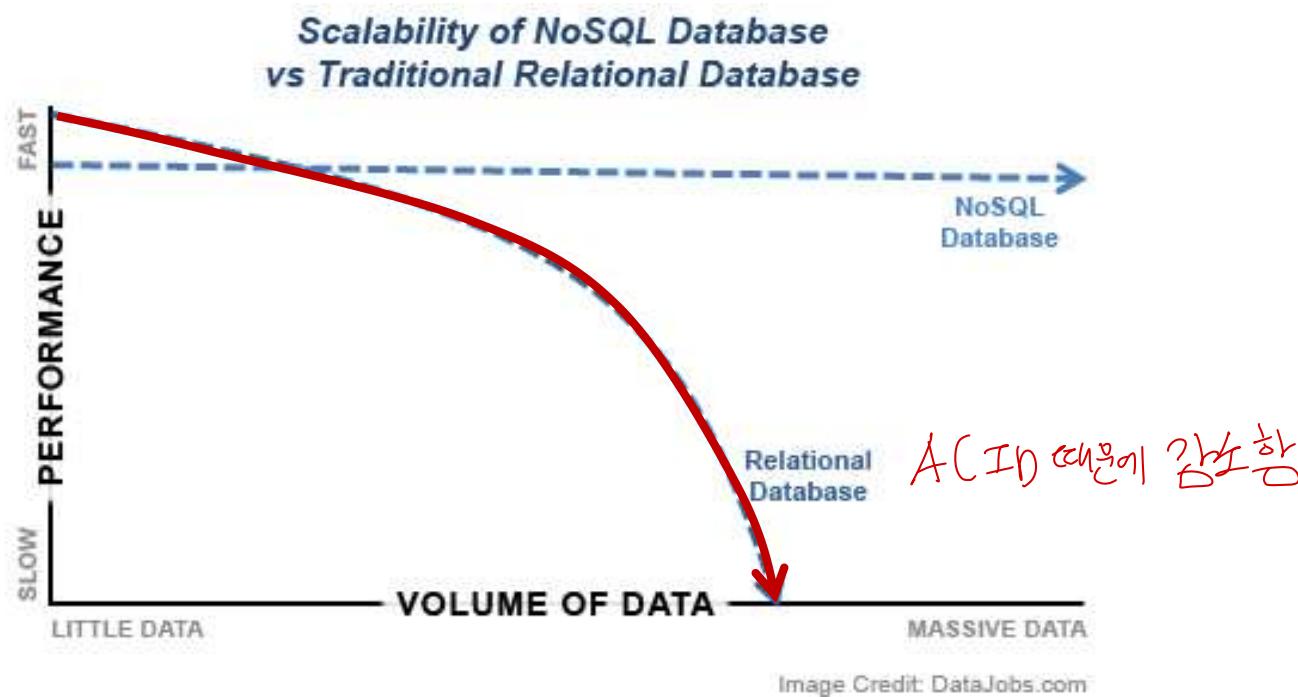
How to Manage Huge Datasets?

- Divide and Conquer! (하나의 단위인 대용량 데이터)

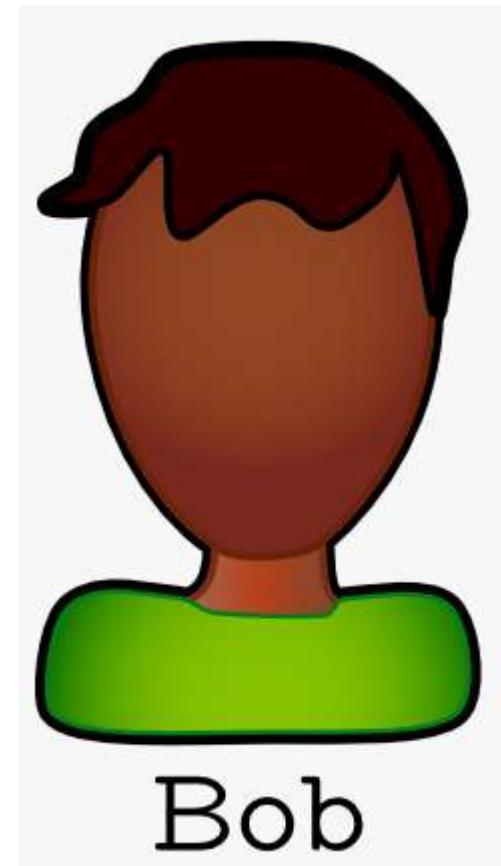
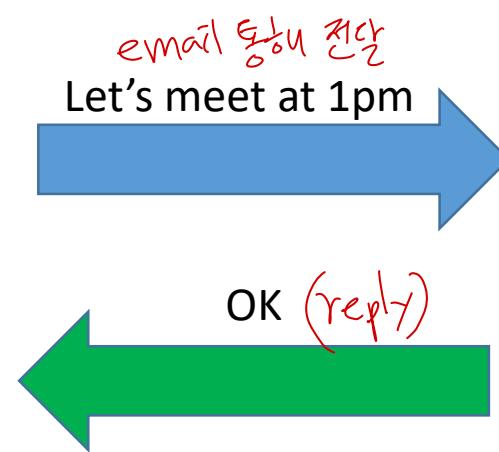
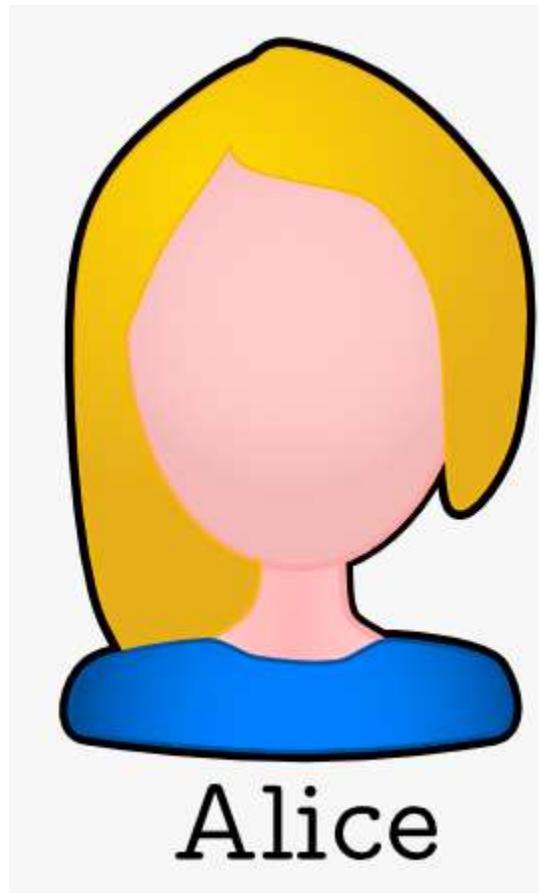


Distributed and Parallel DBMS

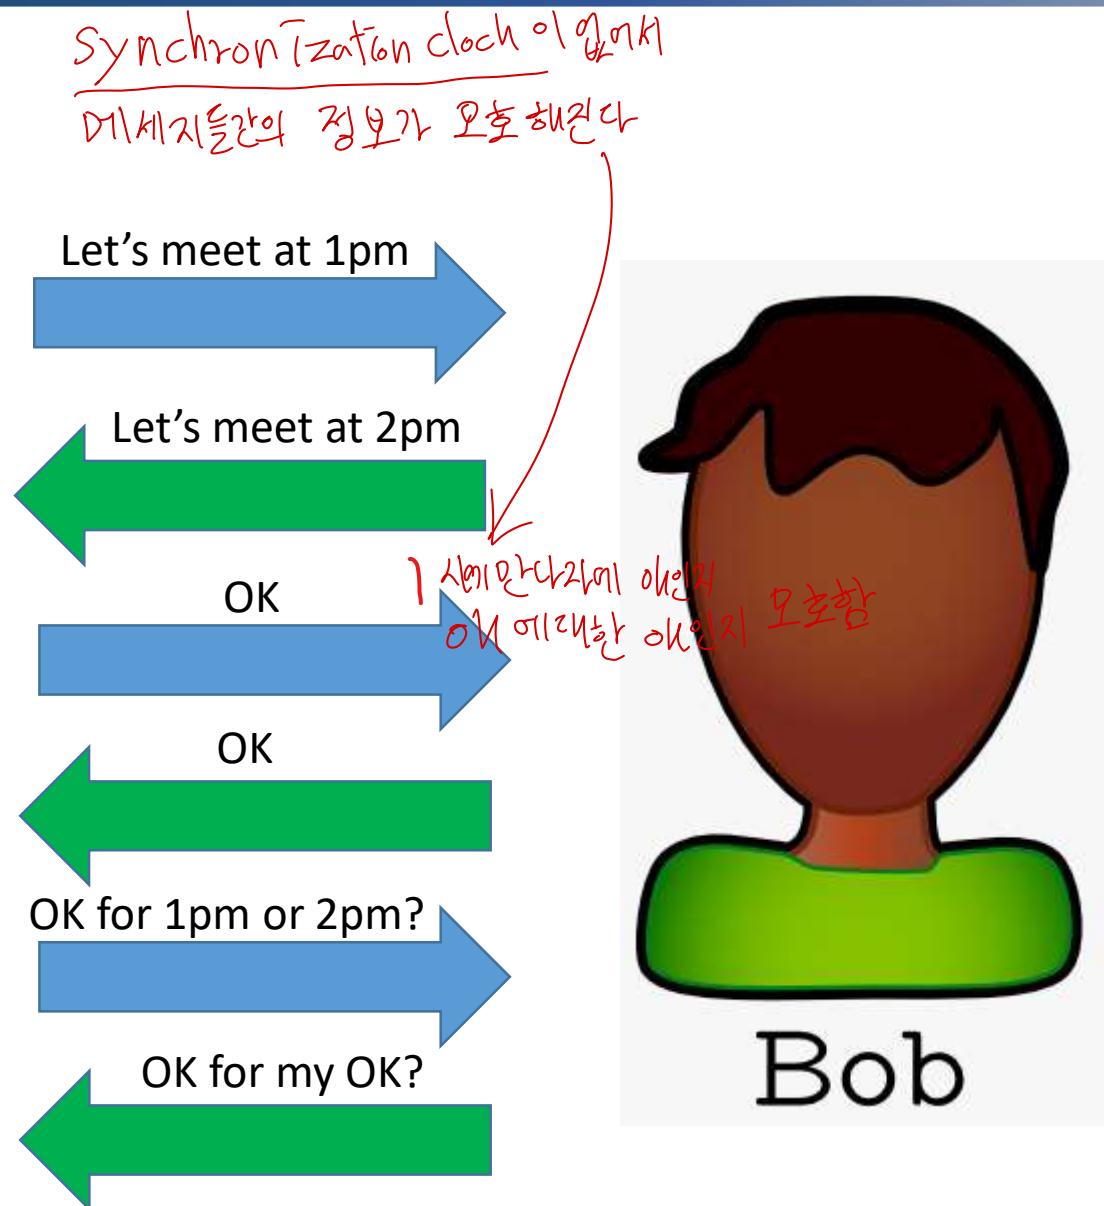
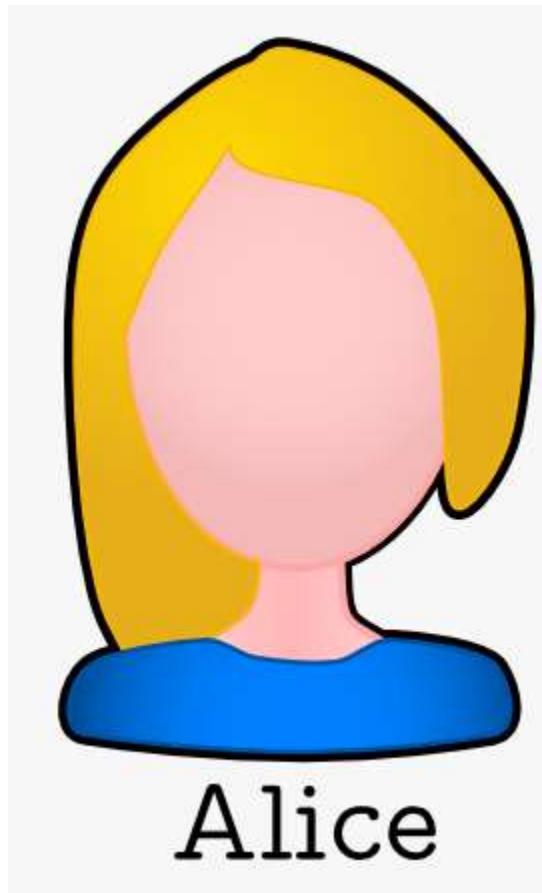
- Not scalable because “consistency” is hard to achieve in large scale systems



Distributed Consensus is a Hard Problem



Distributed Consensus is a Hard Problem

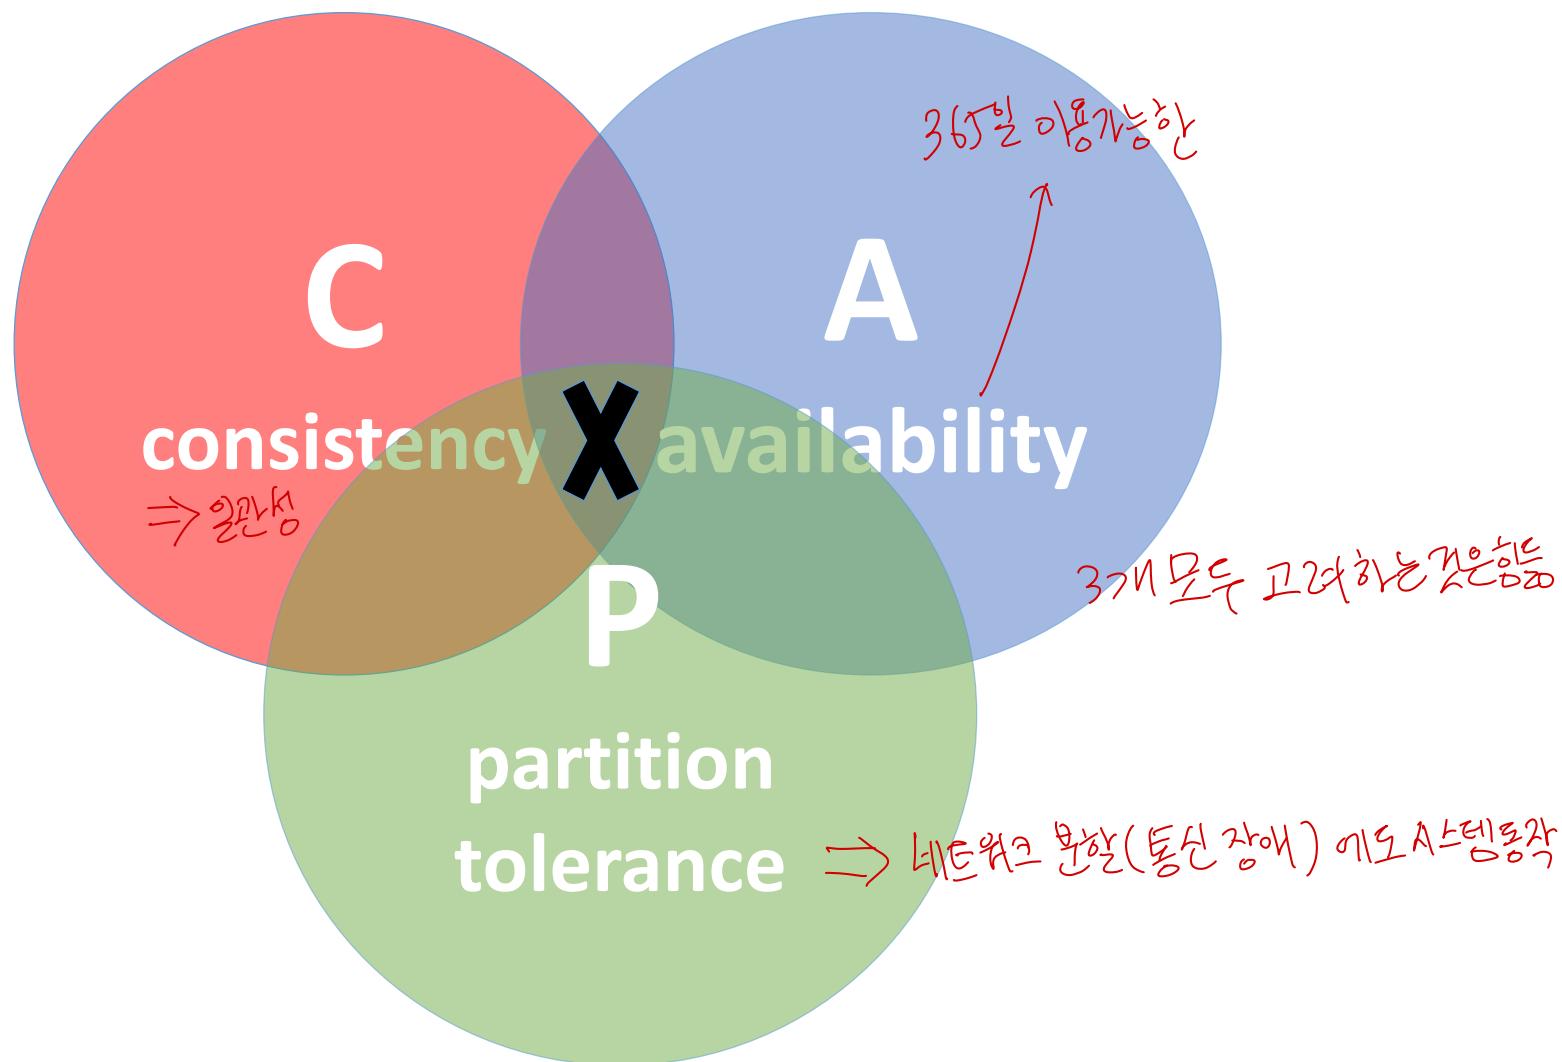


Agreement with Unreliable Message Channel

↑ *가장 같은 문제 예제*

- In distributed systems, network is unreliable
- Then, distributed consensus becomes even harder.
 - Alice → Bob : Let's meet at noon on the 2nd floor
 - Alice ← Bob : Ok!
 - Alice (What if Bob doesn't know that I received his message?)
 - Alice → Bob : I received your message, so it's ok.
 - Bob (What if Alice doesn't know that I received her message?)
 - ...

CAP Theorem



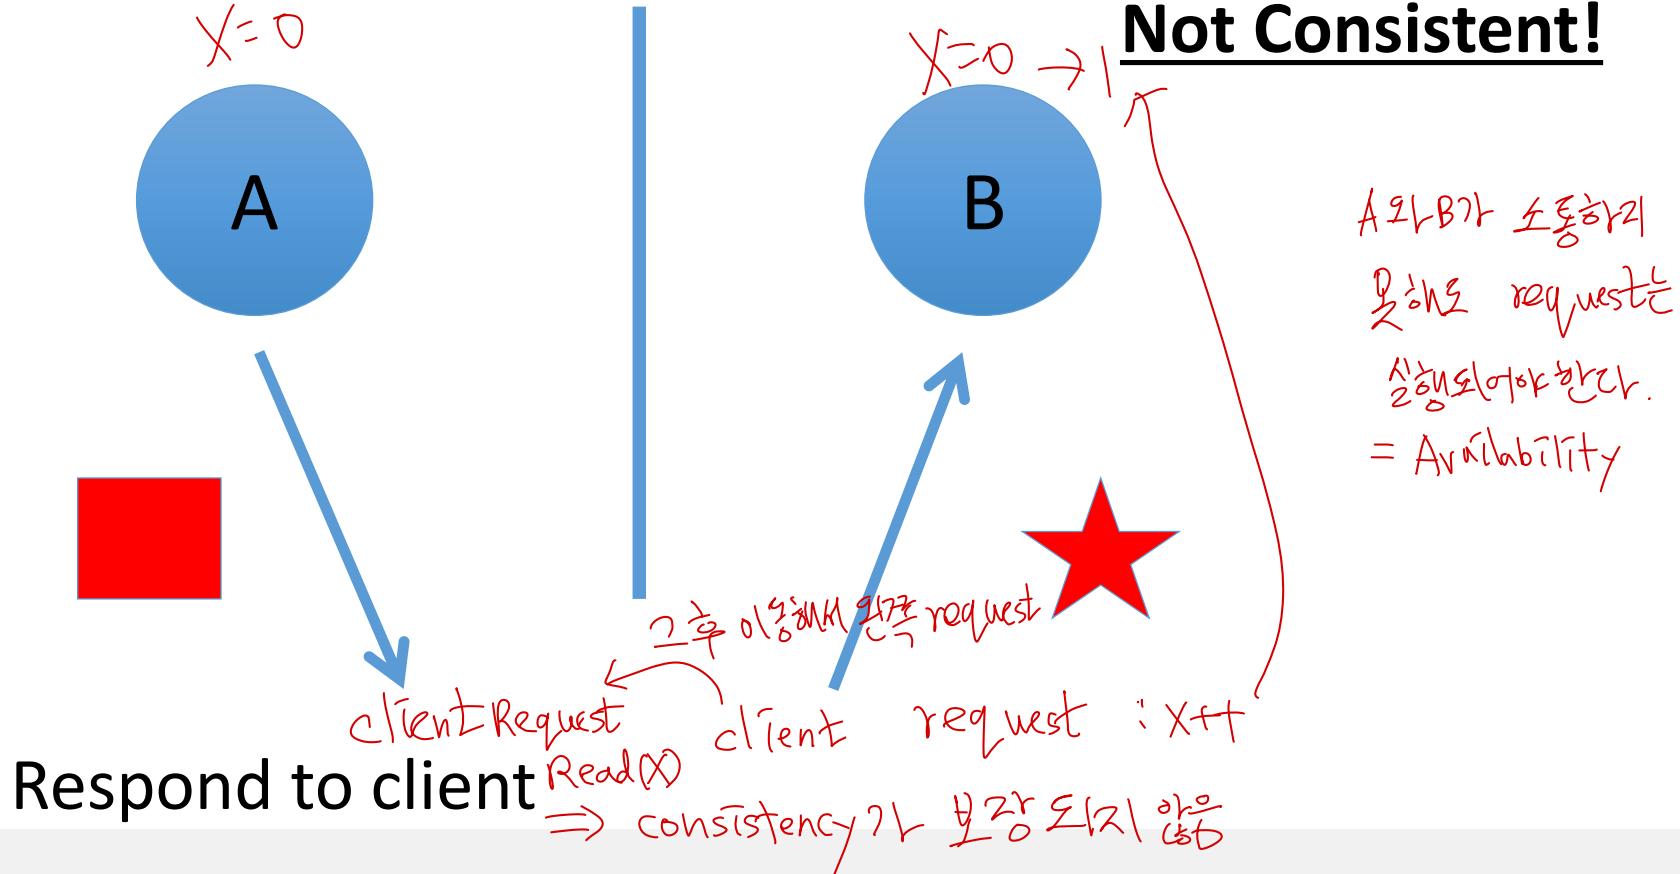
CAP Theorem

- Consistency:
 - All nodes should see the same data at the same time
- Availability: *각 request를 반드시 실행 해야 함*
 - Node failures do not prevent survivors from continuing to operate
≈ Availability 보장
- Partition-tolerance:
 - The system continues to operate despite network partitions
- A distributed system can satisfy any two of these guarantees at the same time **but not all three**

CAP Theorem: AP Systems sacrifice C

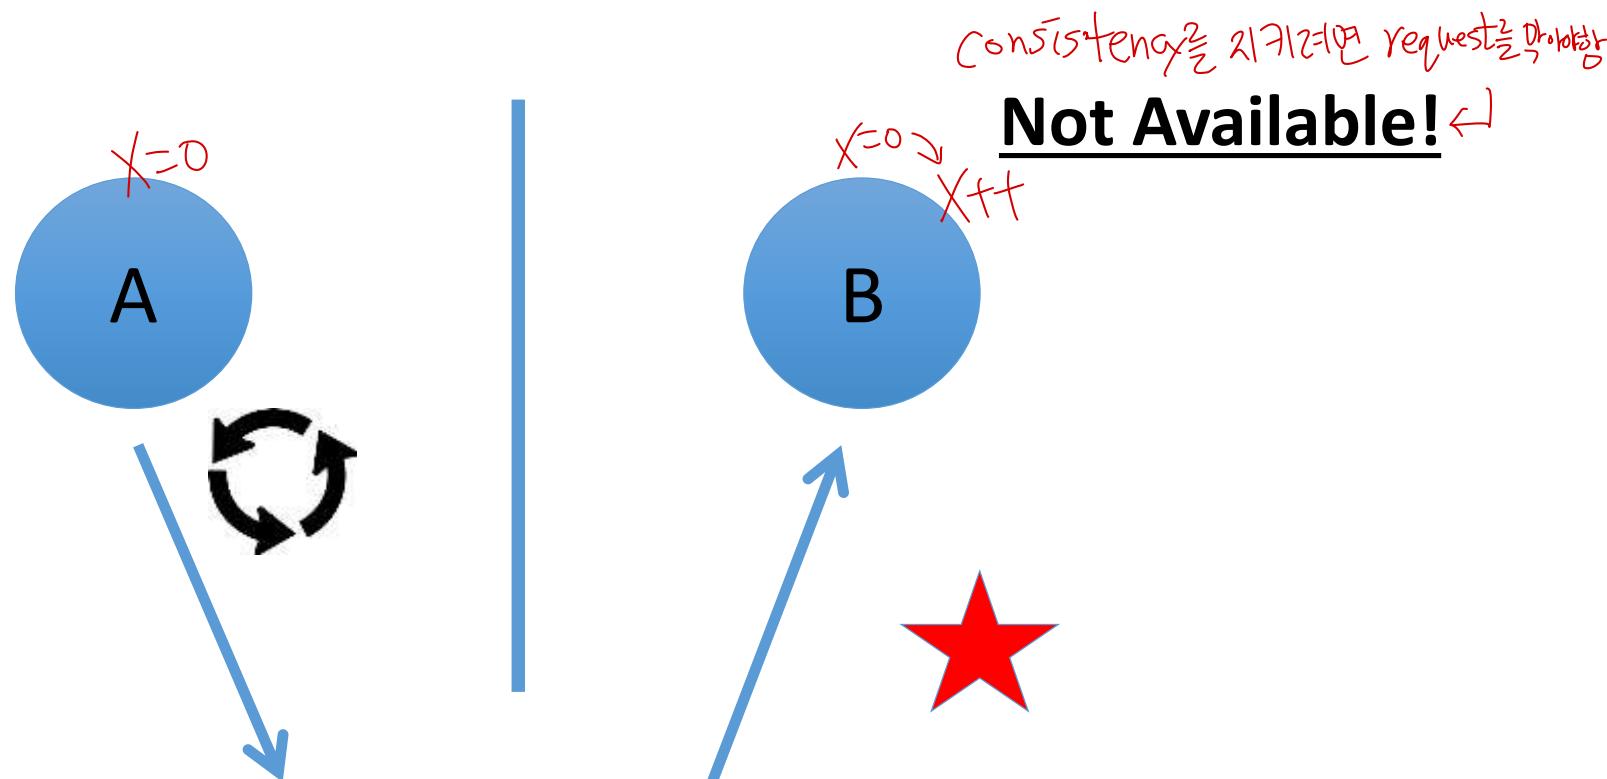
- To make a distributed system **available** when a network **partition occurs**, **consistency** must be given up.

partition은 전제조건



CAP Theorem: CP Systems sacrifice A

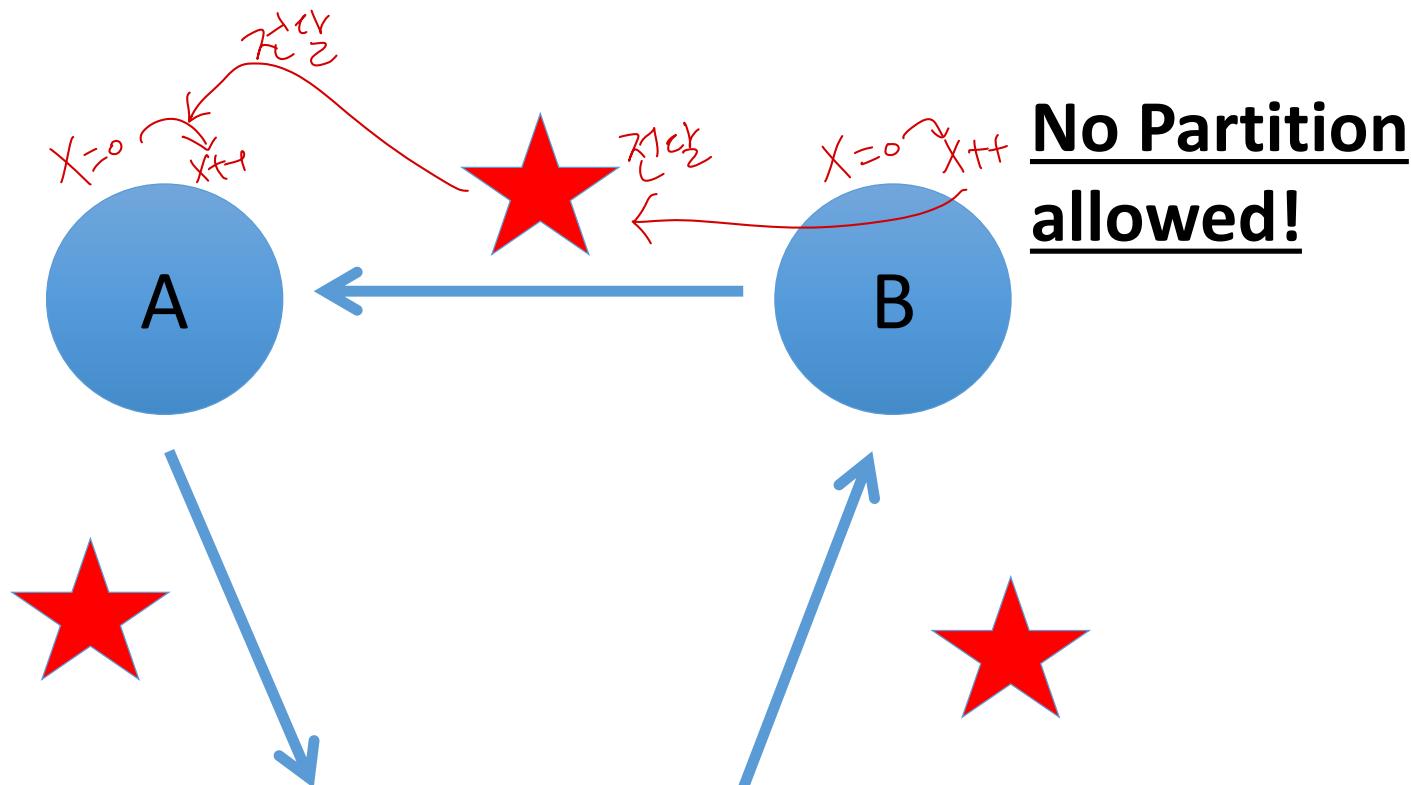
- To make a distributed system **consistent** when a network **partition occurs**, one of the partitions must not be **available**.



Wait to be updated

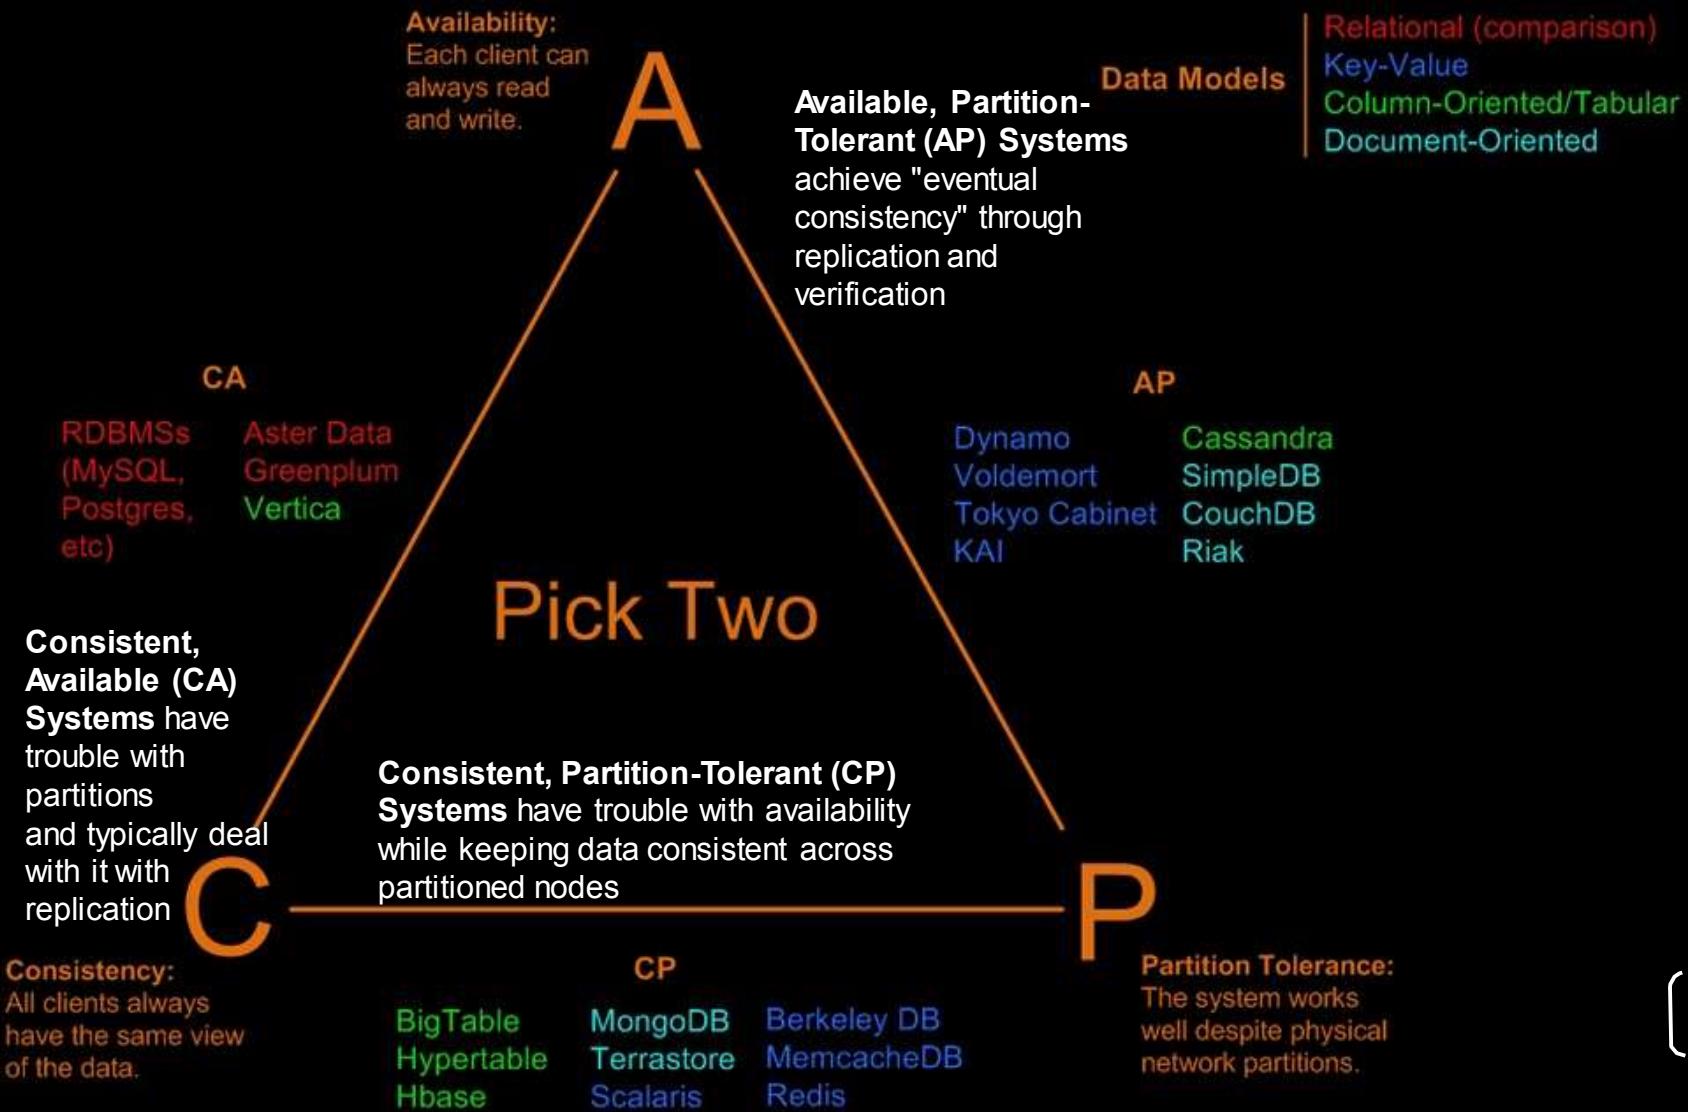
CAP Theorem: CA Systems sacrifice P

- To make a distributed system *consistent* and *available* there must not be partition.



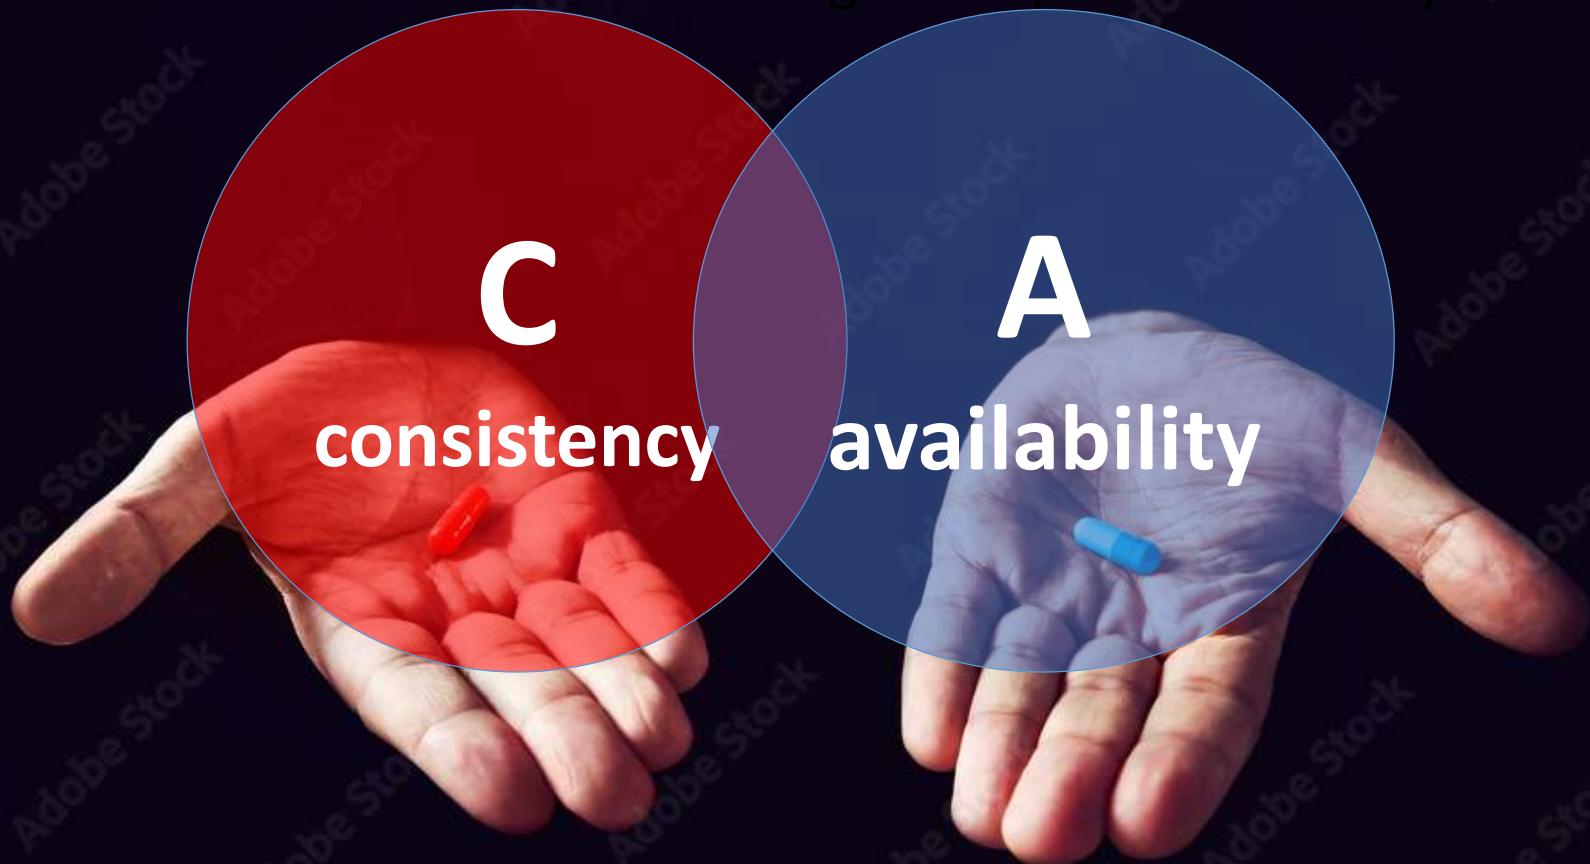
A gets updated from B

Visual Guide to NoSQL Systems



Which one would you sacrifice?

Partition Tolerance cannot be given up in scalable systems



Sacrificing Availability → Financial Penalty

삼성화재

시스템 성능 개선을 위해
홈페이지 서비스가 일시 중단 중입니다.
10월 10일(화) 낮 12시에 다시 오픈할 예정입니다.
이용에 불편을 드려 대단히 죄송합니다.

시스템 중단기간 중, 아래의 자동차보험 업무에 한하여
제한적으로 서비스가 제공되며, 그 외의 경우는 홈페이지 오픈 후
다시 방문하셔서 이용해 주세요.

중단기간 2017년 9월 29일(금) 밤 10시 ~ 10월 10일(화) 낮 12시



카카오뱅크는 점검 중

더욱 안정적인 서비스를 제공하기 위해
시스템 점검을 진행하고 있습니다.
점검이 완료된 후 다시 접속해주세요.

- 점검대상 카카오뱅크 전체 서비스
- 점검시간 2020년 6월 28일(일) 02:00 ~ 09:00 (7시간)



Consistency 중지!

제주특별자치도 홈페이지 및 패밀리사이트

서비스 일시 중단 안내



서비스 점검 사전 안내

안정적인 서비스 제공을 위해 시스템 점검 예정입니다.
점검 이후 자동 로그인이 해제되니 참고해 주세요.

계정 확인 : 프로필 사진 → MY 원티드

10월 17일 오전 1시 – 5시 (약 4시간 소요 예상)

계정 확인하기

NAVER

잠시 후 다시 확인해주세요!

지금 이 서비스와 연결할 수 없습니다.
문제를 해결하기 위해 열심히 노력하고 있습니다.
잠시 후 다시 확인해주세요.

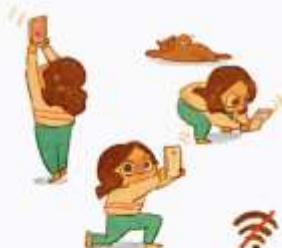


Illustration by 마법

이전 페이지

네이버쇼핑 홈

Eventual Consistency : B2C (Business to Consumer)

수집 정보
업데이트 정보



Availability 증시! (consistency 가 학생)

- Facebook example:
 - Bob finds an interesting story and shares with Alice by posting on her Facebook wall
 - Bob asks Alice to check it out
 - Alice logs in her account, checks her Facebook wall but finds nothing is there.

- What can they do?

- Nothing. They use the SNS for free

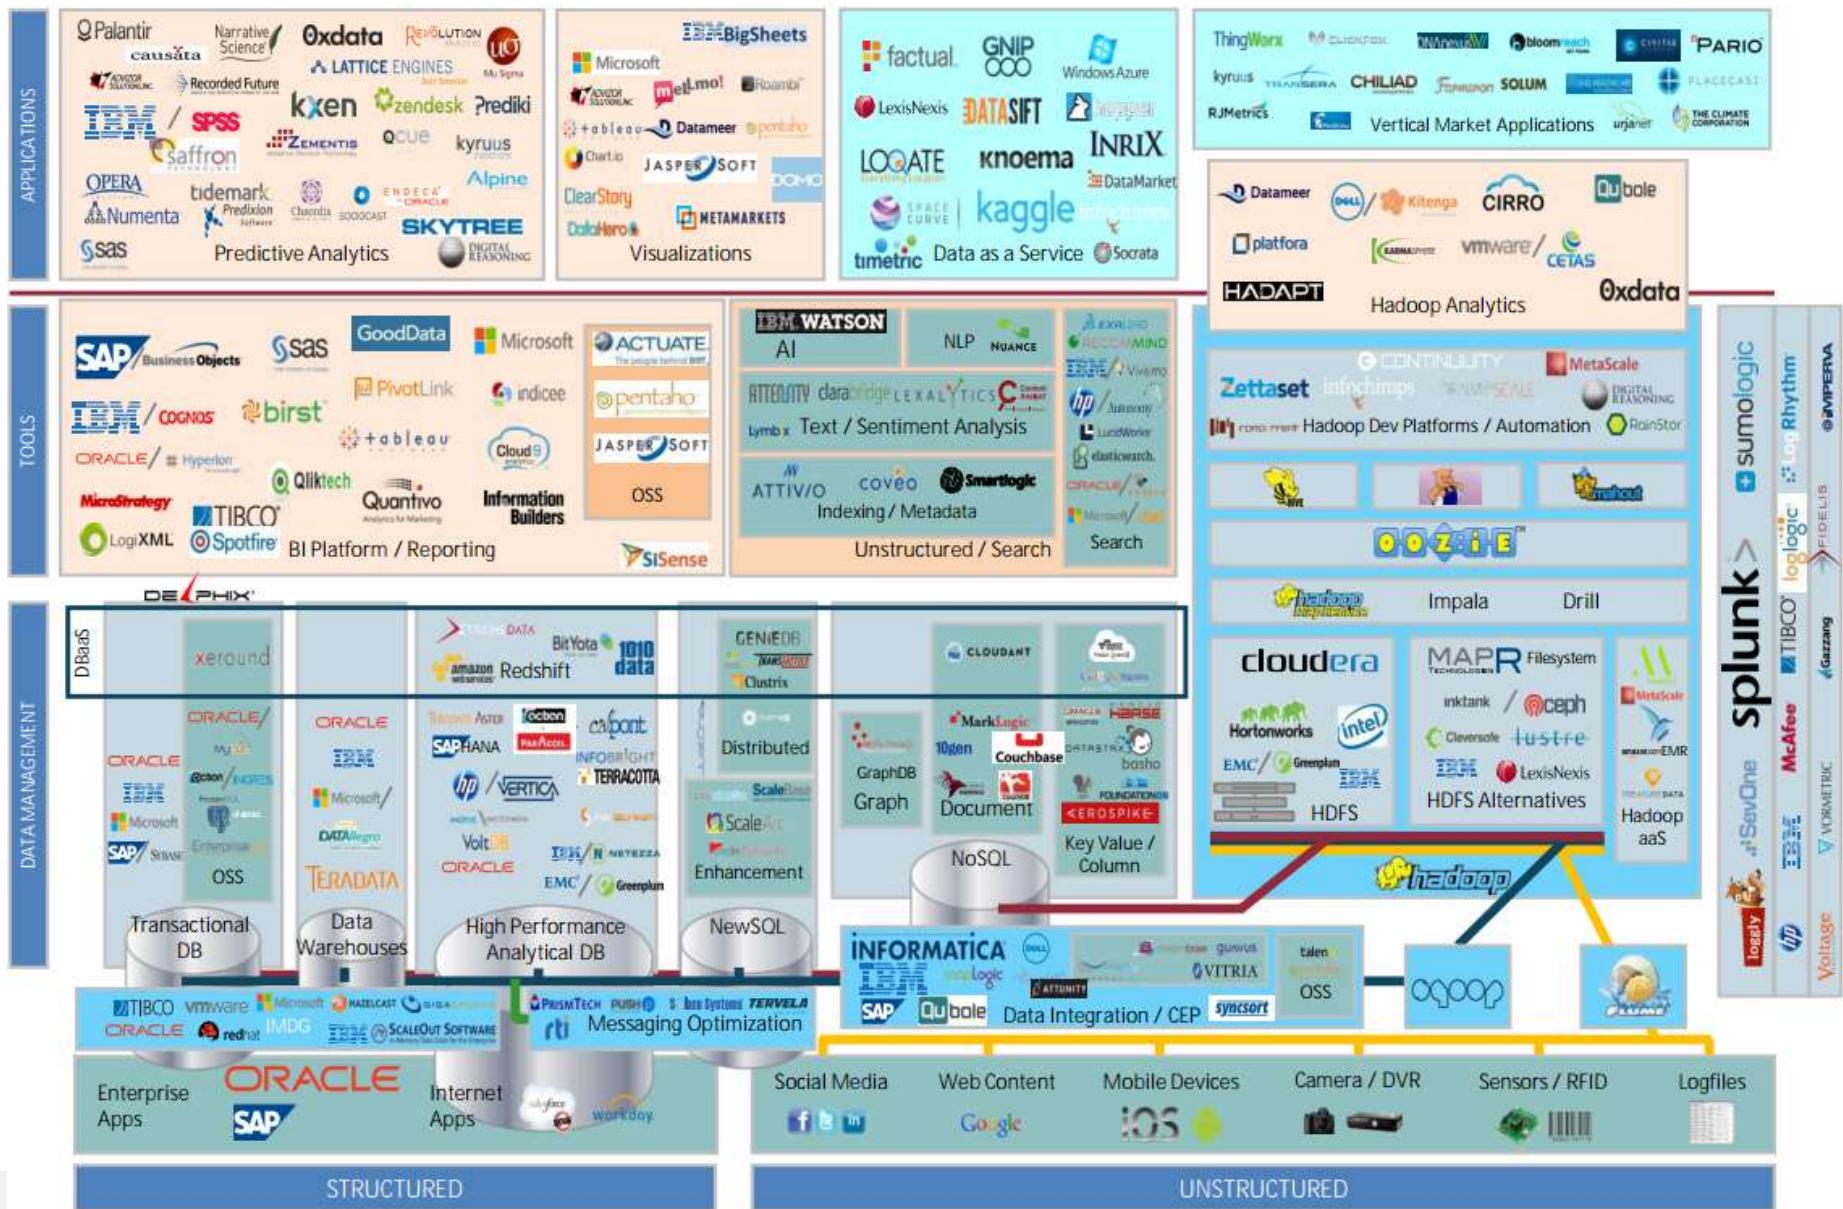
무료니까 기다릴 수밖에 없고 결국 일관되지 않음

- Eventual consistency is not for B2B

consistency 중요 (RDBMS 사용..)



So many data processing engines out there...



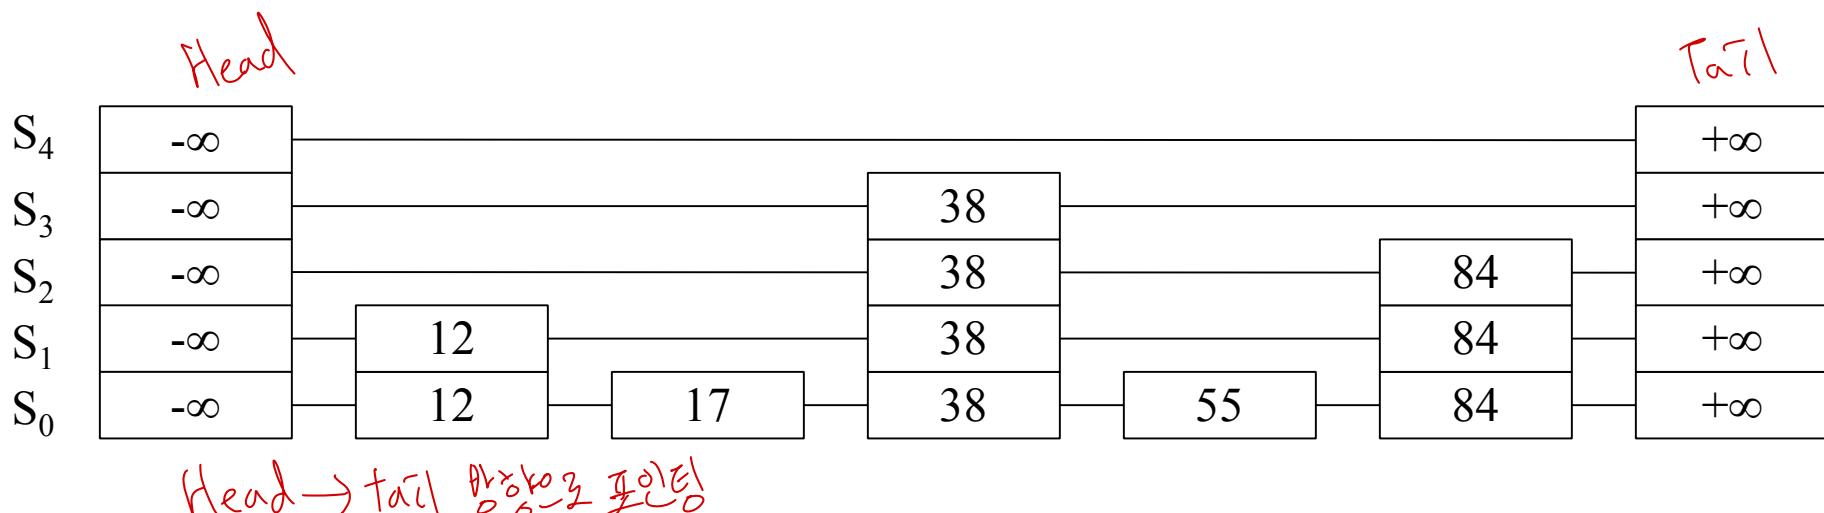
Skip Lists

Skip Lists (a.k.a Skiplist)

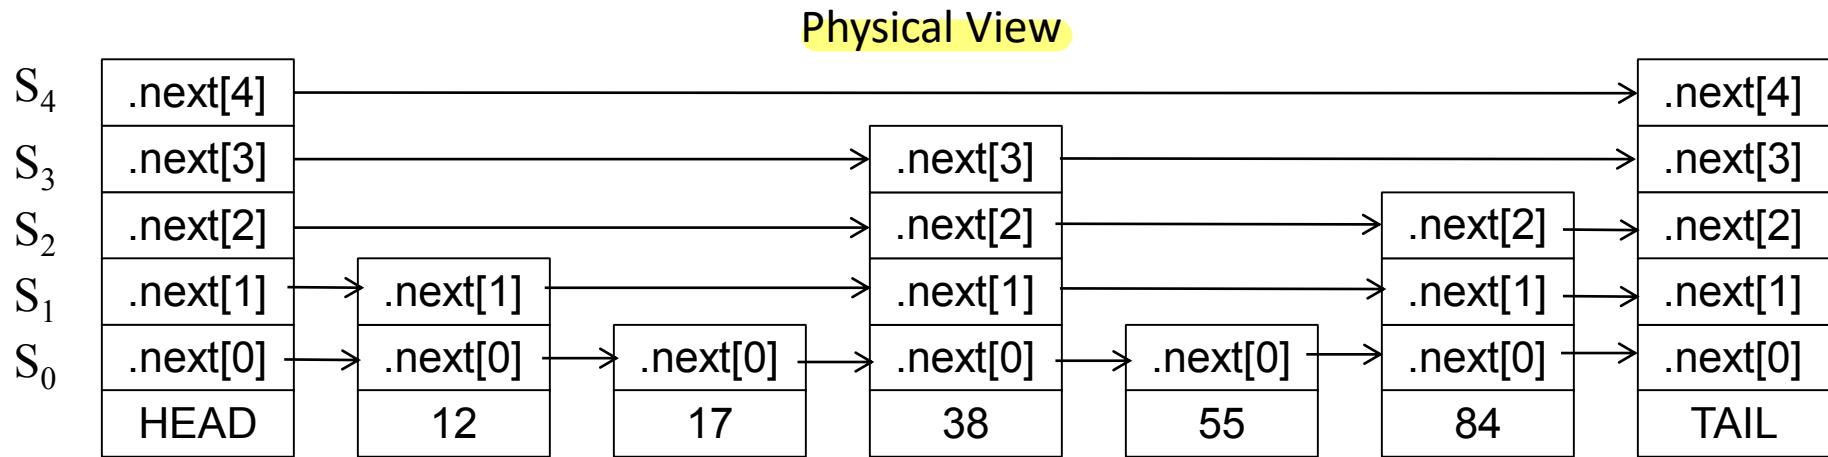
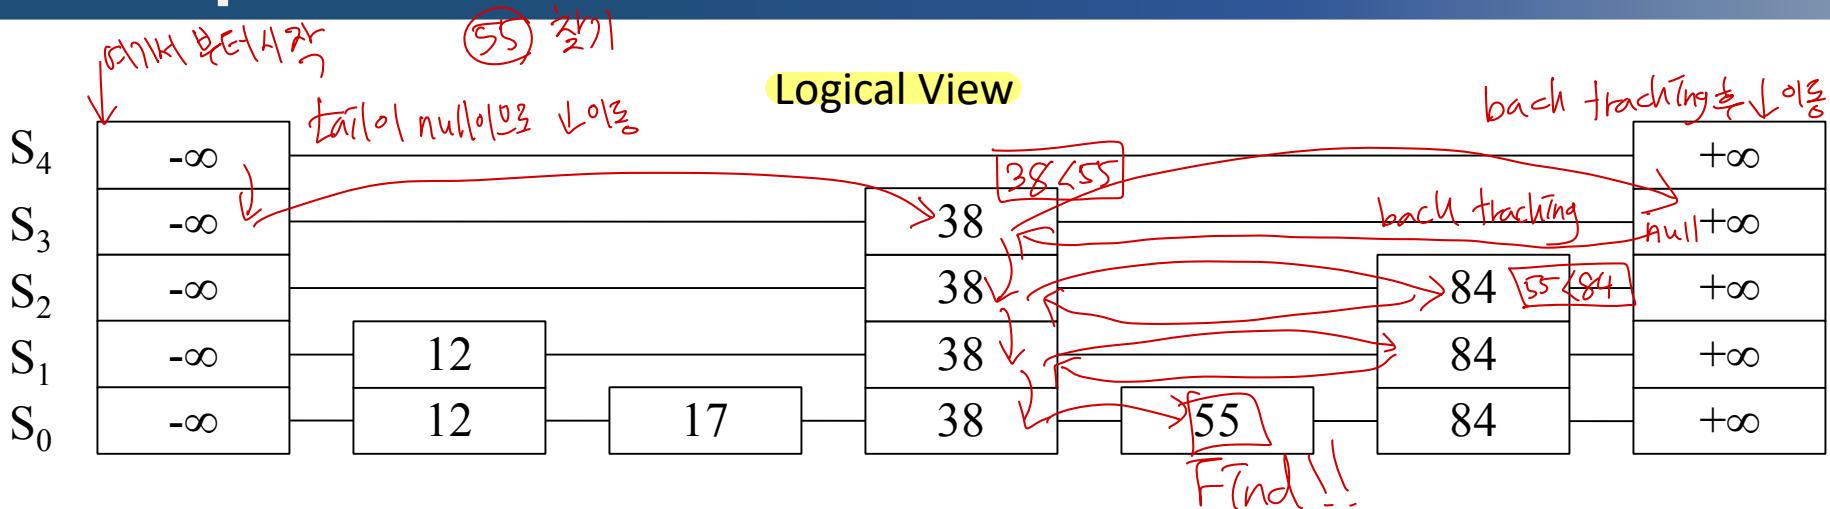
- Binary Search Trees: $O(\log n)$
 - If, and only if, the tree is “balanced”
SkipList 가 BST보다 나음
- Insertion of partially sorted data is optimally bad for binary search trees.
- Skip Lists make use of some randomization to maintain $O(\log n)$ search and update times
 - These are on average figures
- Analysis of skip lists is interesting
 - Worst case is potentially very bad
 - We will not cover the analysis here.

Skip Lists (Skip List)

- A Skip List is a series of lists: $\{S_0, S_1, \dots, S_h\}$
 - S_0 contains all the elements
 - S_i contains a randomly chosen subset of elements from S_{i-1}
 - Each list S_i is sorted
- We introduce two special elements: $-\infty$ & $+\infty$
 - Each list S_i contains these, S_h contains only these.



Skip Lists



Skip List - Informal

- We set the lists up so that:
 - An item that is in S_i is in S_{i+1} with probability $\frac{1}{2}$ (or $1/4, 1/8$, etc)
- So S_1 has about $n/2$ items in it
 - S_2 has about $n/4$ items
 - S_i has about $n/(2^i)$ items
- The height of the skip list is about $\log_2(n)$
- 초상위 레벨에서 모든 노드를 skip 했을 때 노드 개수는 $k \times 2^{\log_2 n}$ 개를 skip한 것과

S_{i+1}/S_i $\frac{1}{2}$
즉, 1보다 더 작은 분수임

$$\frac{n}{2^H} = 1 \text{ 일 때 } H \text{ 를 찾기 } \Rightarrow H = \log_2 n$$

\therefore 시간복잡도 : $O(\log n)$

NOTE:

- Unlike a binary search tree, this “halving” property is not enforced
 - It is approximate

Skip List Operations

- We use a positional abstraction
- Operations
- After(p) – position after p on the same level
- Before(p) – position before p on the same level
- Below(p) – position below p in the same tower
- Above(p) – position above p in the same tower

Searching

Simple: binary search (modified)

- The upper level lists provide something like a “thumb index”
- Algorithm for searching for the key k in a skip list
 - Returns p = reference to the key in the list which is equal or larger than k

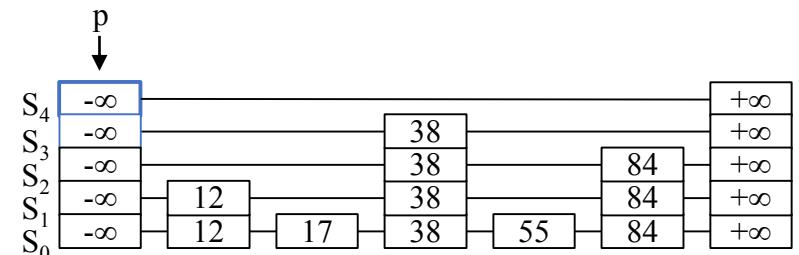
$p \leftarrow$ first element of S_{\max_height}

while $\text{below}(p) \neq \text{null}$ **do** ↓ 방향 이동

$p \leftarrow \text{below}(p)$ {drop down}

while $\text{key}(\text{after}(p)) \leq k$ **do**

$p \leftarrow \text{after}(p)$ {scan forward}



Inserting

- Uses insertAfterAbove(p,q,(k,e)) which inserts (k,e) after p and above q.

$p \leftarrow \text{SkipSearch}(k)$

$0 \leq \text{random}() \leq 1$

$q \leftarrow \text{insertAfterAbove}(p, \text{null}, (k, e))$

50% 확률로 자기 위에 노드 생성(반복)

while $\text{random}() < \frac{1}{2}$ **do**

while $\text{above}(p) = \text{null}$ **do**

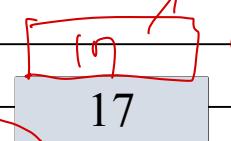
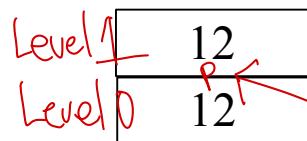
해당 레벨의 predecessor (자신보다 작은
가장 큰 값) 찾기

$p \leftarrow \text{before}(p)$ {scan backwards}

$p \leftarrow \text{above}(p)$ {jump up}

$q \leftarrow \text{insertAfterAbove}(p, q, (k, e))$

- E.g. $\text{above}(\text{before}(p))$

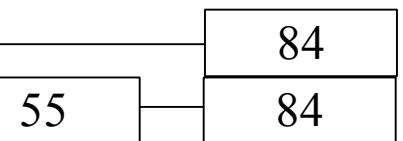
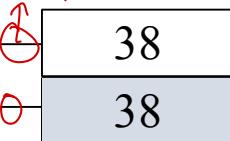


50% 확률로 sum

25%

50%

100%



p
전입자

q
(k, e) Insert

Deleting

- Removal of an item is easier than insertion
- We simply find the item ($\text{SkipSearch}(k)$)
 - Then we remove this from the list S_0 *지우고 연결 시로 (링크드리스트 외에는)*
 - We recursively remove it from the lists S_i above this
 - This is simply deletion from a linked list, which is simple...

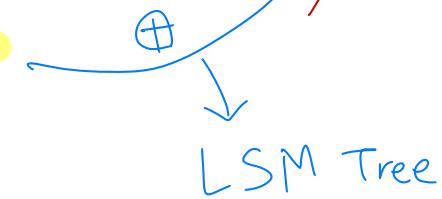
note

- The *above* and *before* methods can be eliminated
 - We can always insert, search and delete using a simple scan-forward-and-down regime
- Storage overhead is therefore just a (key,next,down) triple for each tower element...

Chap. 24 Log-Structured Merge Tree

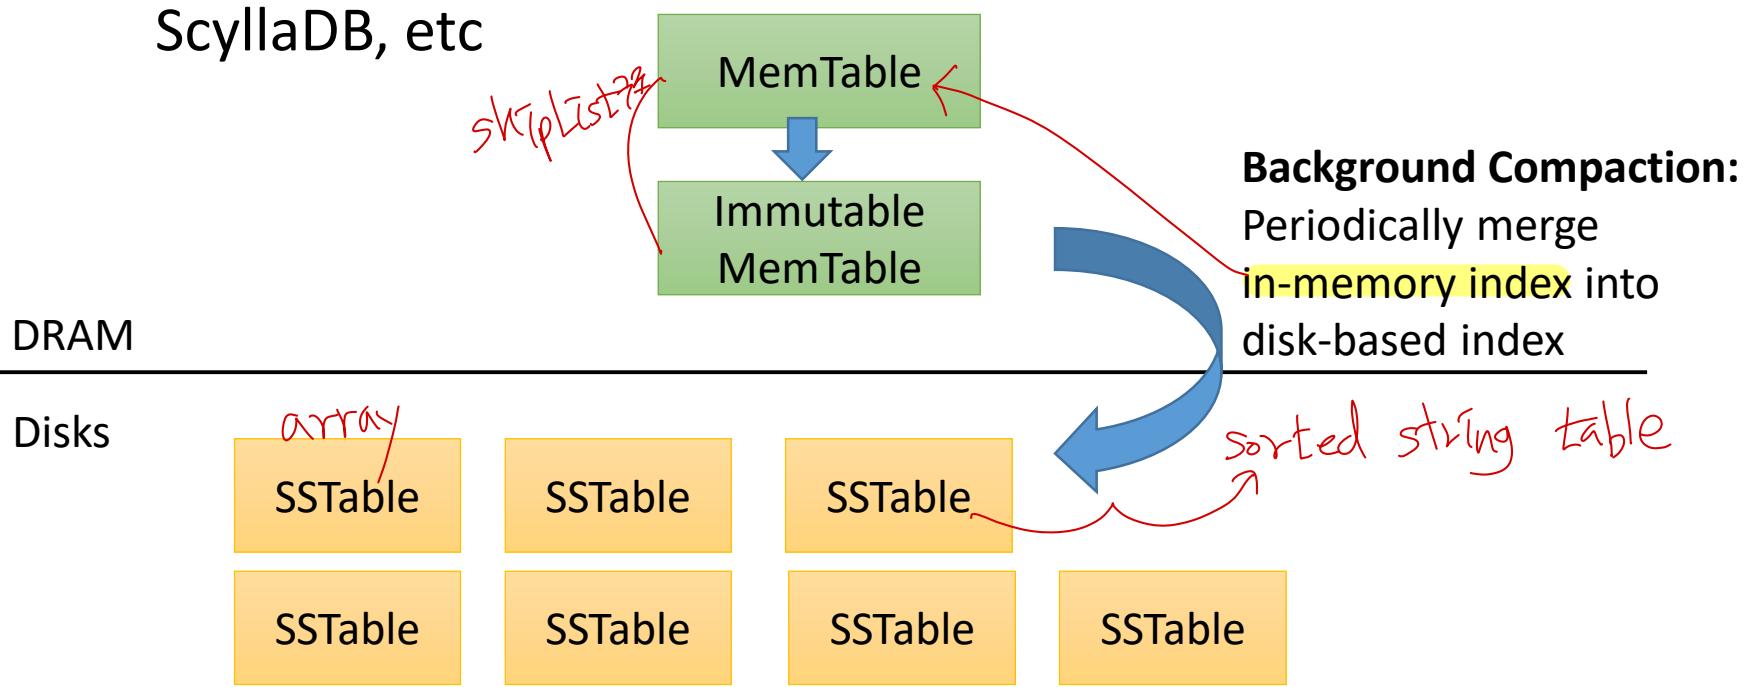
B+tree : disks → sequential or 성능이 좋음 → array

BST : DRAM → NoSQL은 skip list

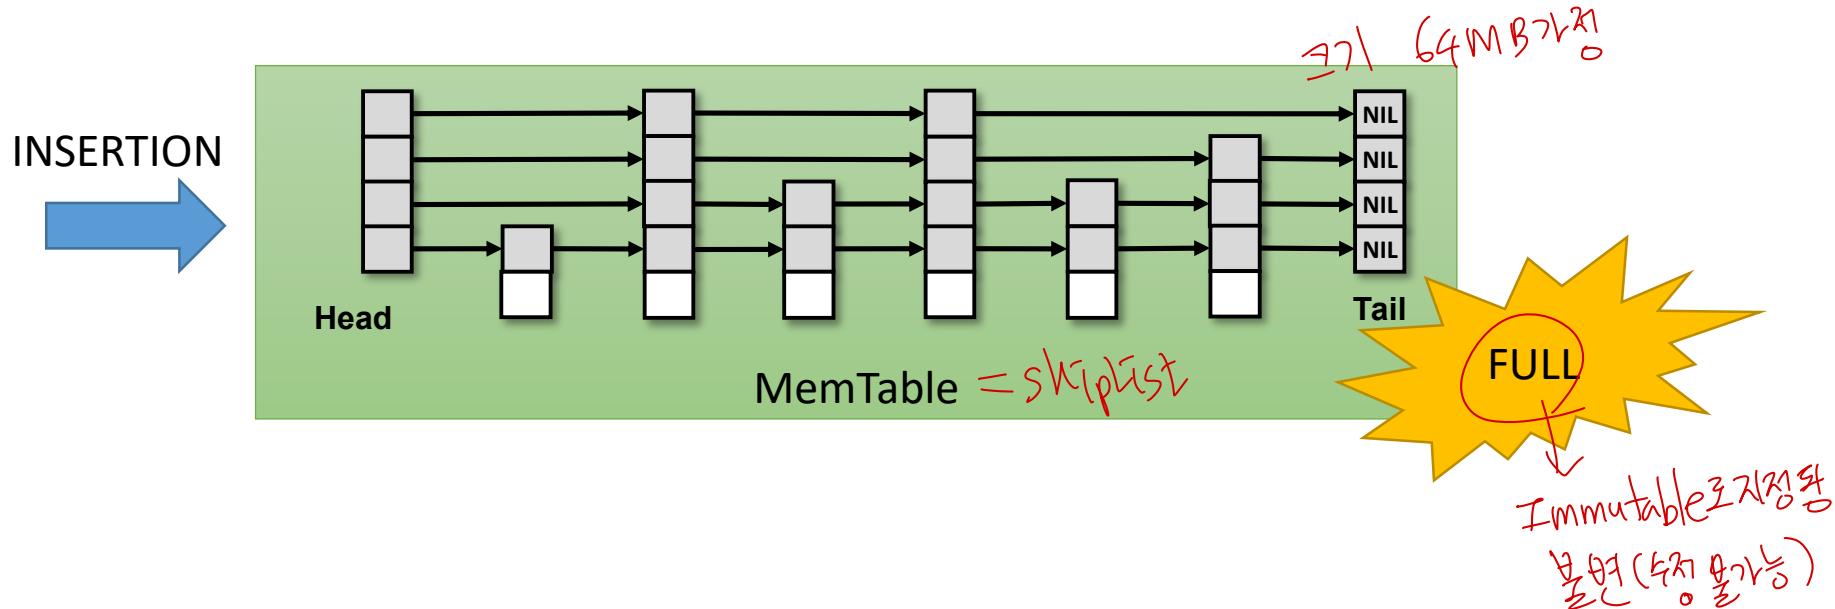


Log-Structured Merge-Tree LSM-Tree

- Designed for Write-intensive workloads
 - Fast insertion *이거 빠르면 write에 좋다!*
 - Moderate search performance *검색 느려도*
 - Widely used in various key-value stores
 - BigTable, Hbase, Cassandra, RocksDB, WiredTiger, InfluxDB, ScyllaDB, etc



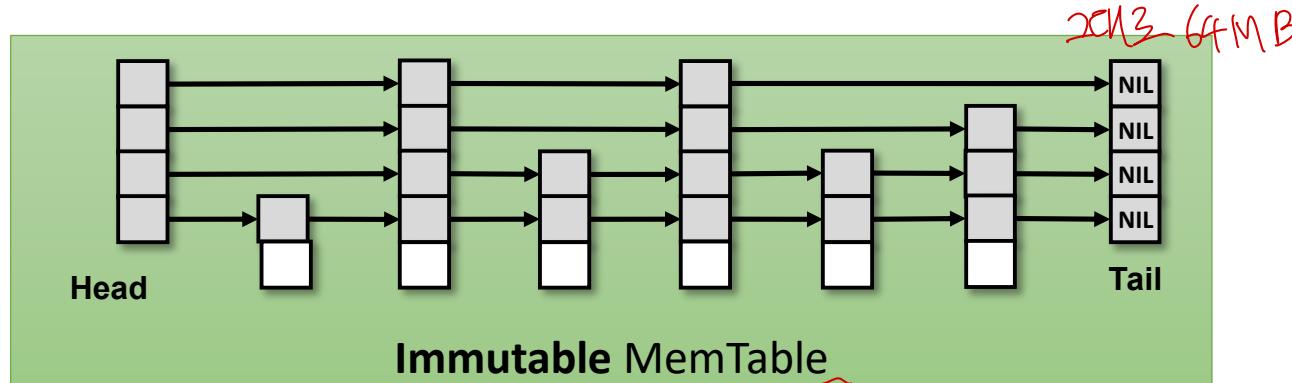
Log-Structured Merge-Tree



DRAM

Disks

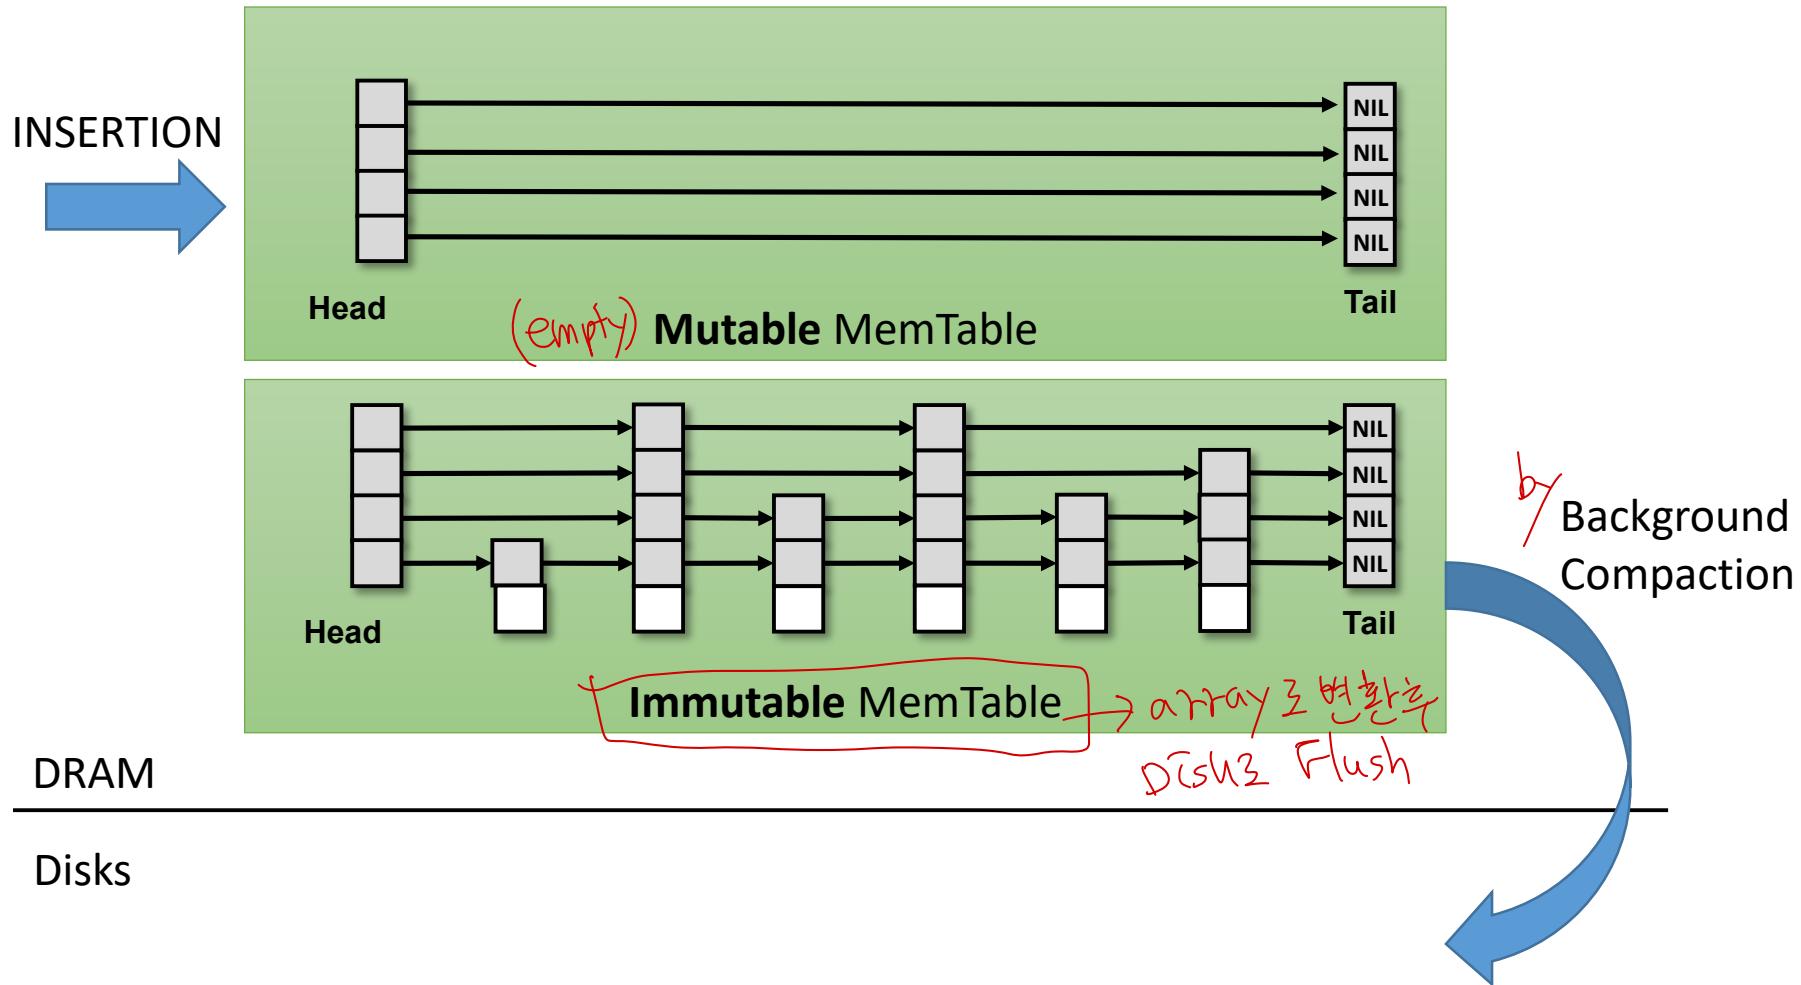
Log-Structured Merge-Tree



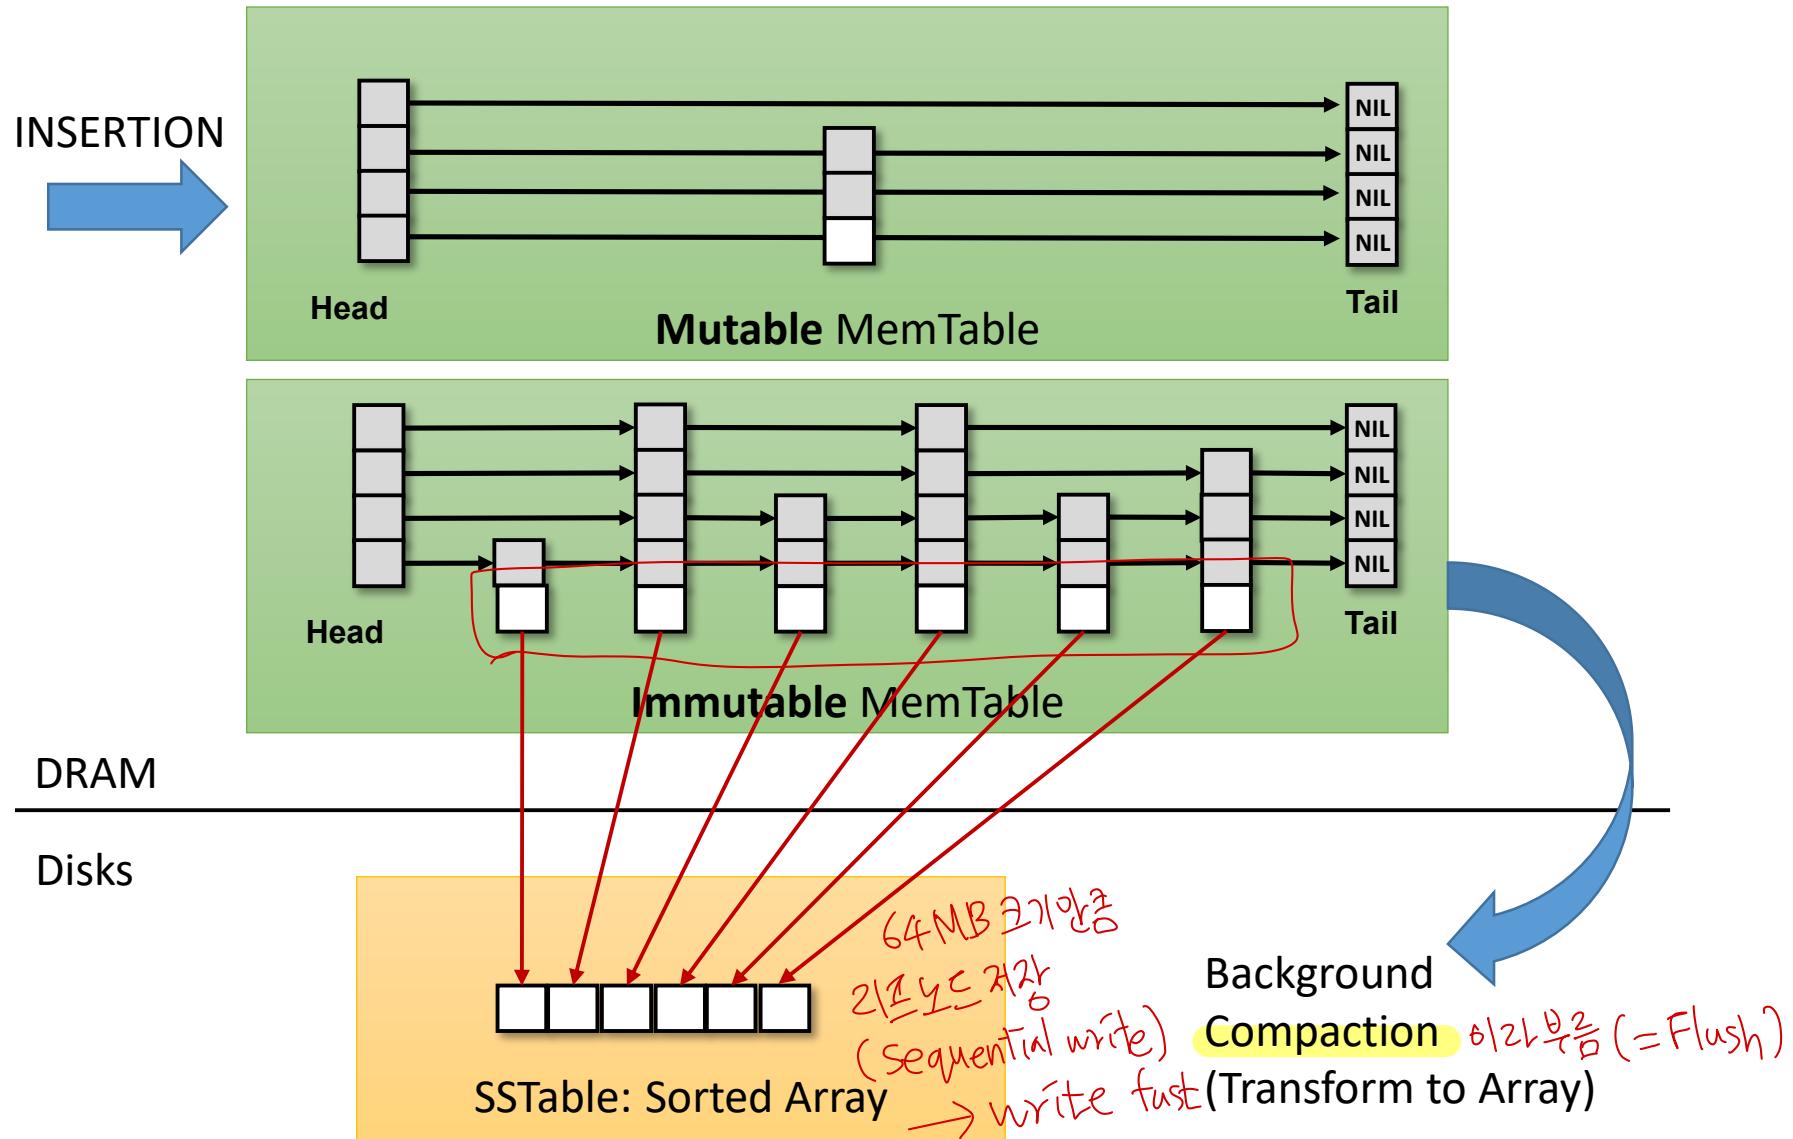
DRAM

Disks

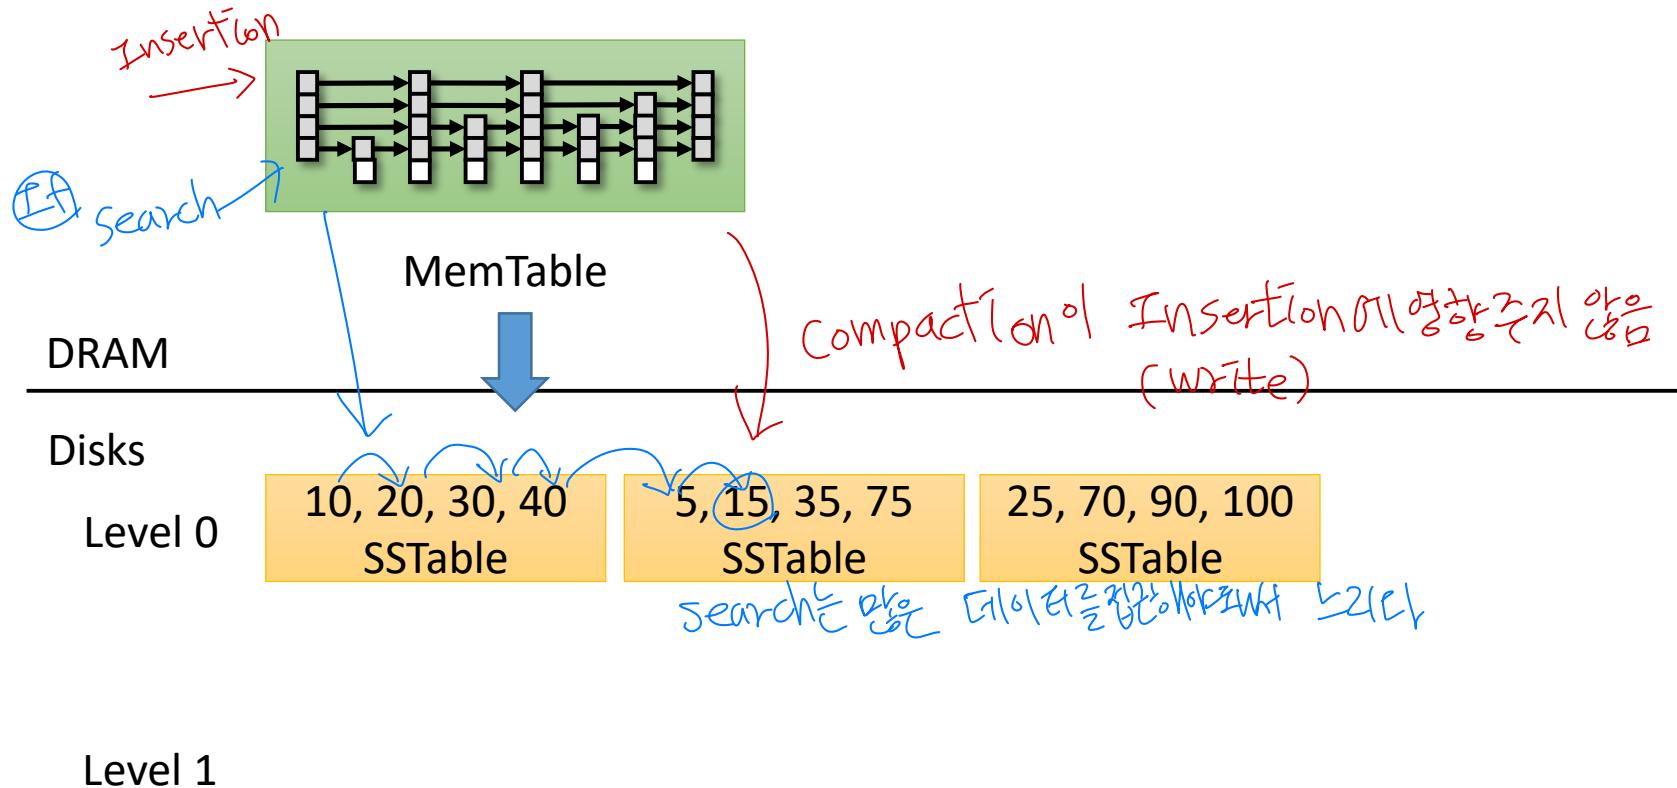
Log-Structured Merge-Tree



Log-Structured Merge-Tree



Log-Structured Merge-Tree



Log-Structured Merge-Tree

DRAM

Disks

Level 0

10, 20, 30, 40
SSTable

5, 15, 35, 75
SSTable

25, 70, 90, 100
SSTable

Level 1

5, 10, 15, 20
SSTable

25, 30, 35, 40
SSTable

70, 75, 90, 100
SSTable

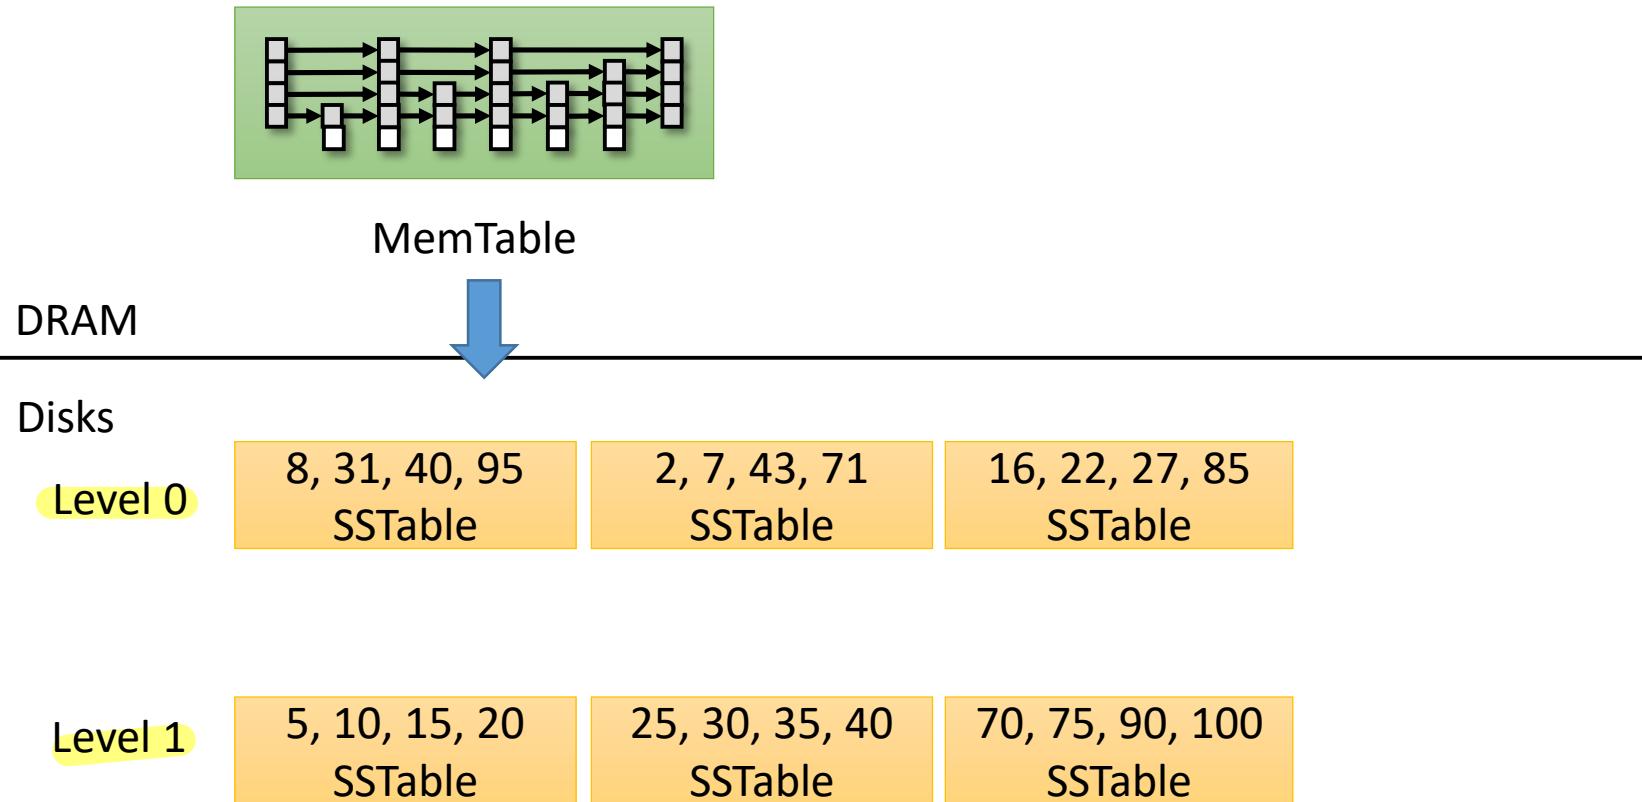
5 - 20 25 - 40 70 - 100 → 범위(조각) 병합(слияние) 가능합니다.



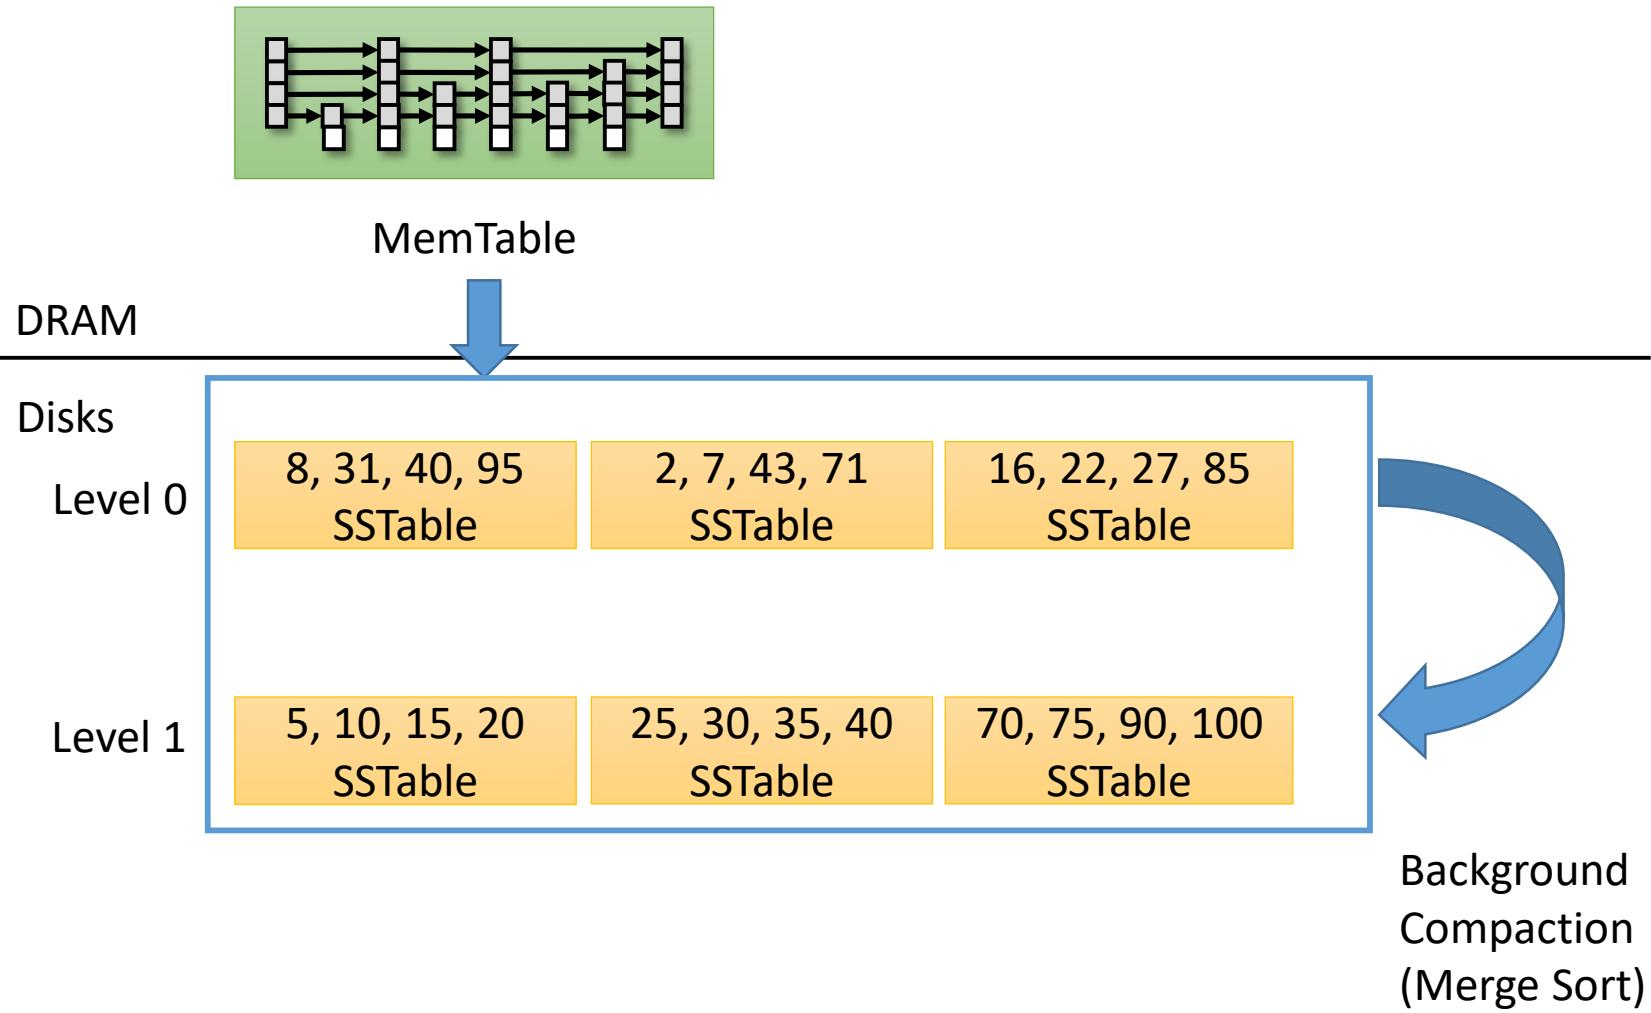
Background
Compaction
(Merge Sort)

Search를 향상시키는 방식

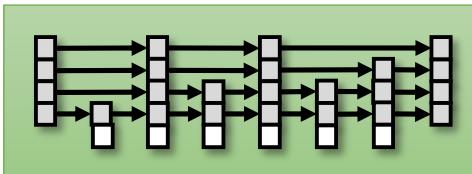
Log-Structured Merge-Tree



Log-Structured Merge-Tree



Log-Structured Merge-Tree



MemTable

DRAM



Disks

Level 0

Level 1

2, 5, 7, 8
SSTable

10,15,16,20
SSTable

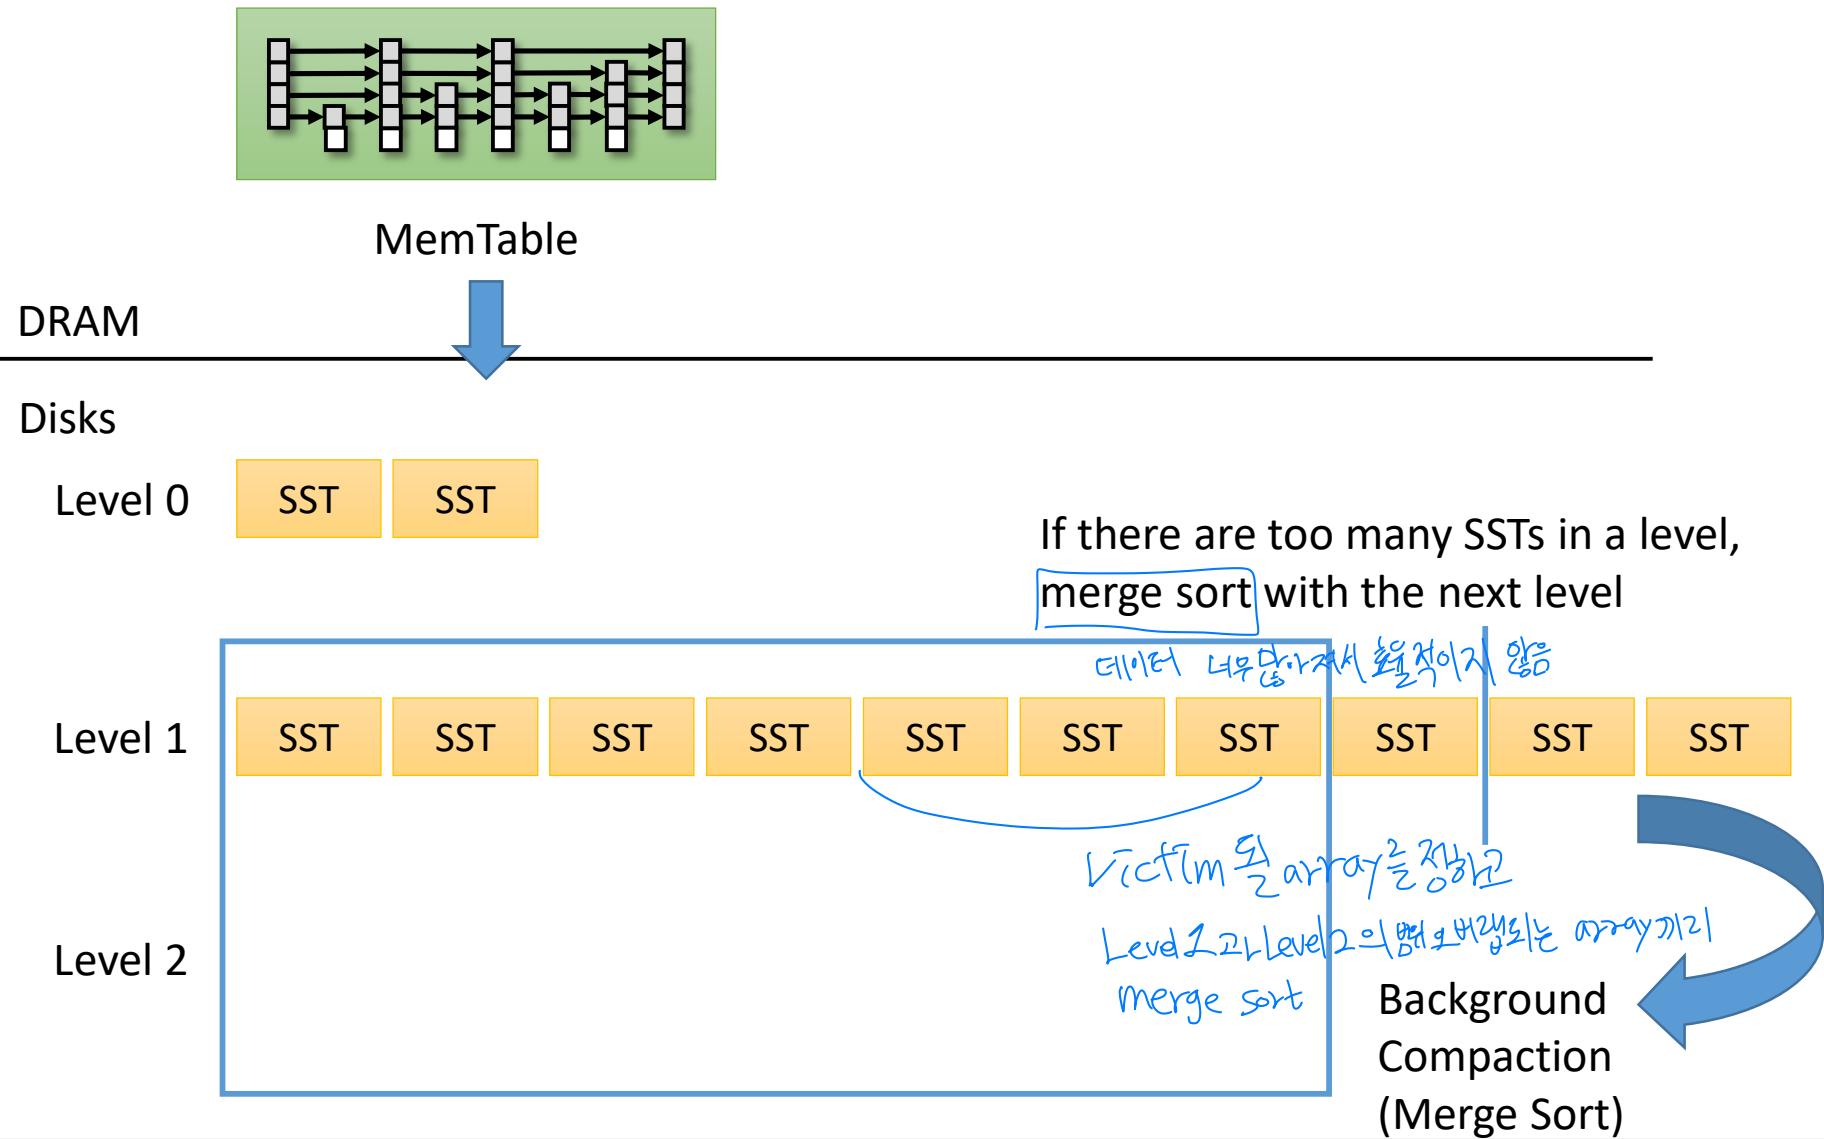
22,25,27,30
SSTable

31,35,40,43
SSTable

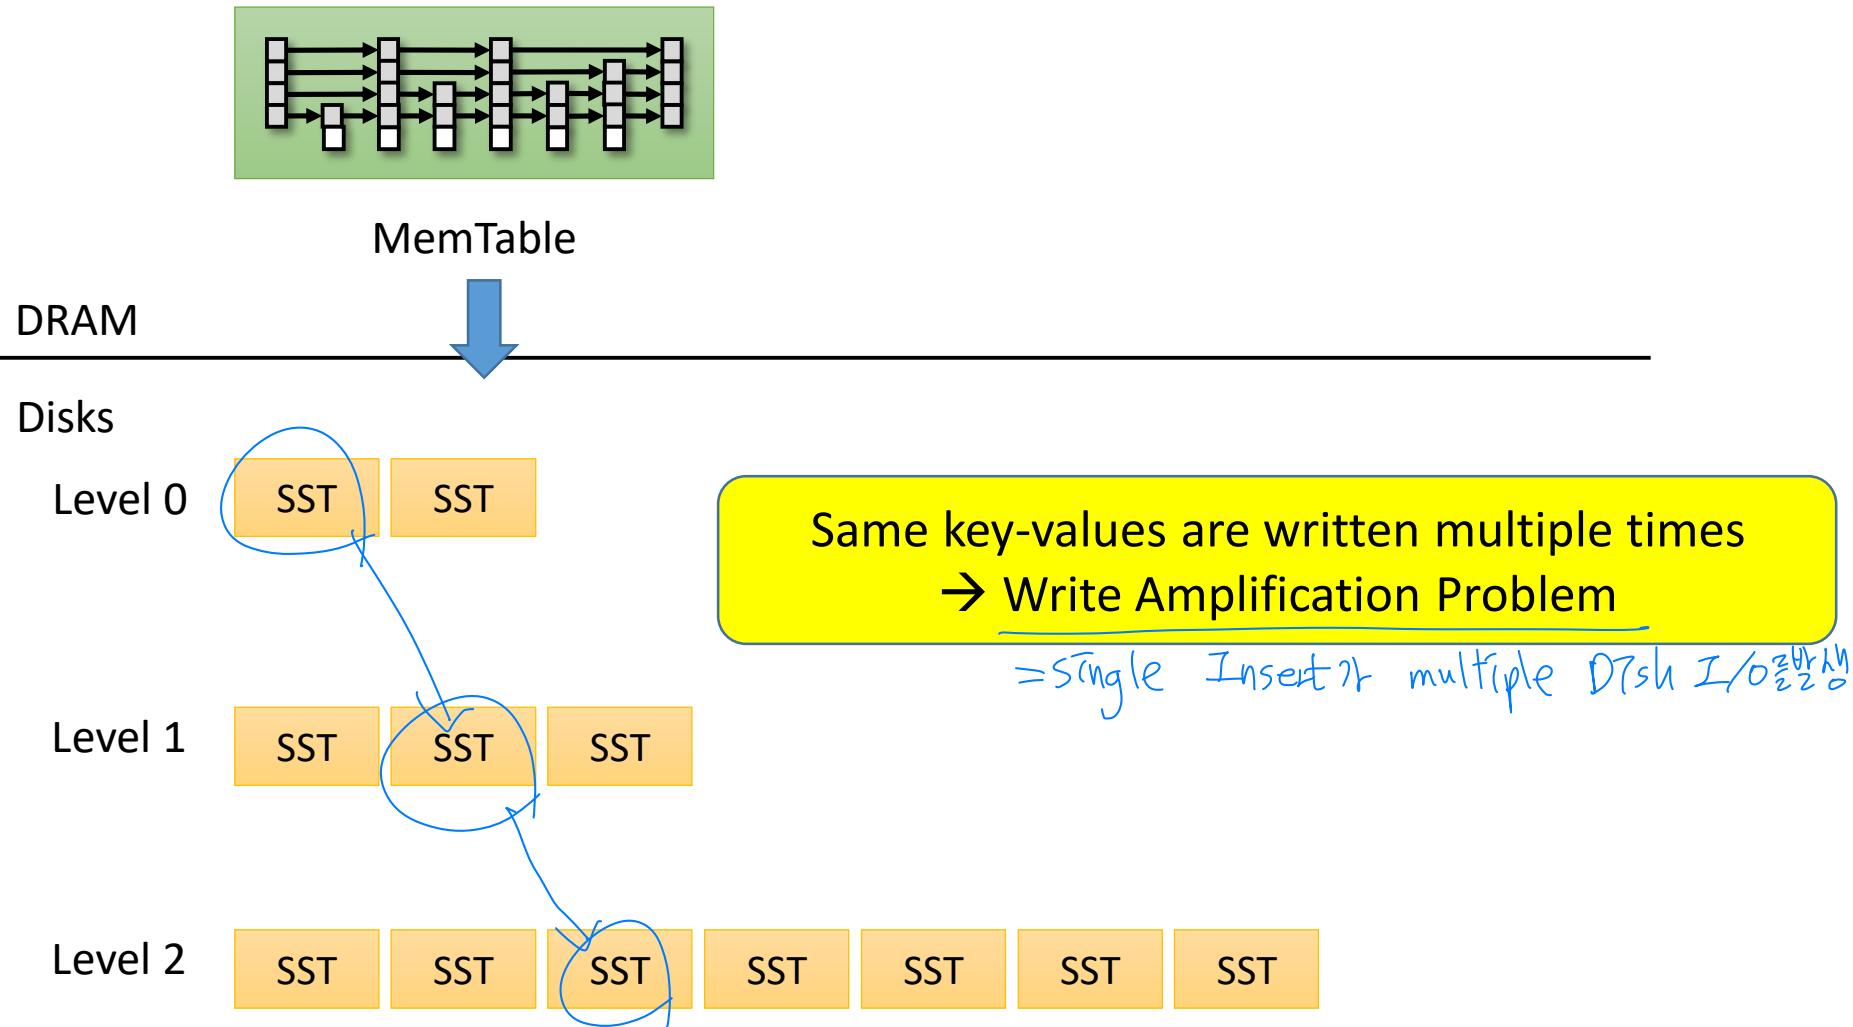
70,71,75,85
SSTable

90,95,100
SSTable

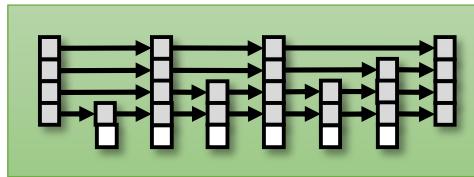
Log-Structured Merge-Tree



Log-Structured Merge-Tree

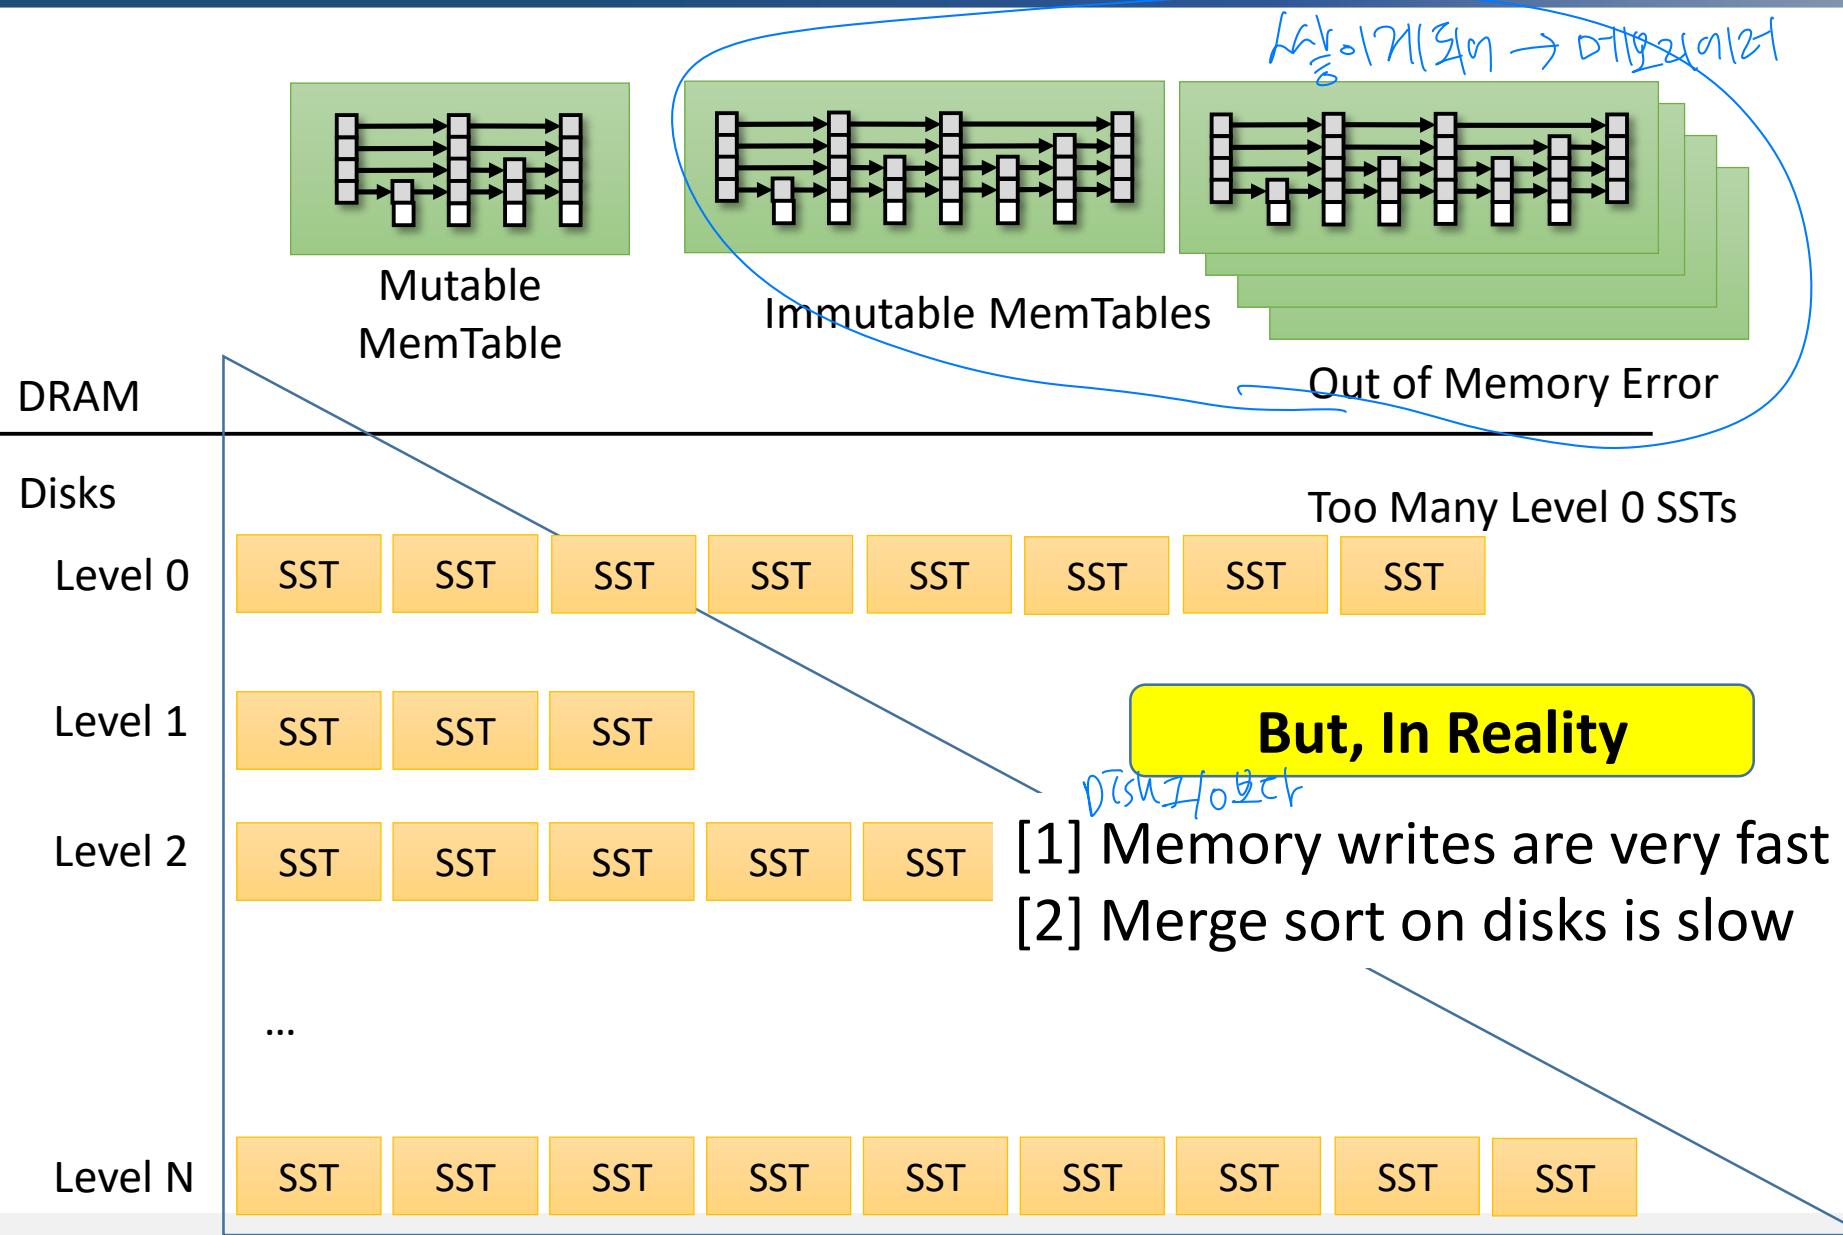


Log-Structured Merge-Tree

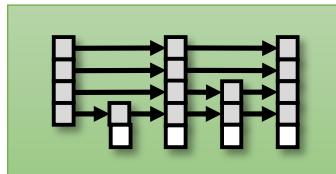


This is what LSM-Tree should be like

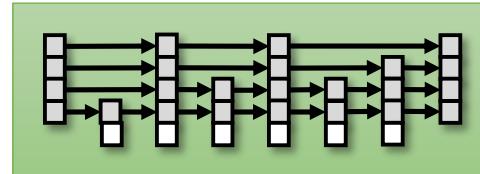
Log-Structured Merge-Tree



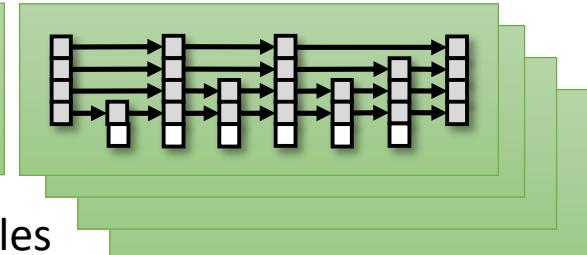
Log-Structured Merge-Tree



Mutable
MemTable



Immutable MemTables



Out of Memory Error

DRAM

Disks

Level 0



Too Many Level 0 SSTs

Level 1



Level 2

If there are too many MemTables or Level 0 SSTables,
Block Incoming Writes!!
This problem is known as Write Stall Problem

...





RDB ACID to NoSQL BASE

RDBMS ↓ 원칙

Atomicity

Consistency

Isolation

Durability



Mongo DB ↓ 원칙

Basically

Available (CP)

Soft-state

(State of system may change over time)

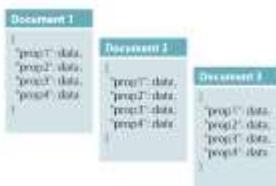
**Eventually
consistent**

(Asynchronous propagation)

NoSQL Databases



Graph database



Document-oriented
JSON

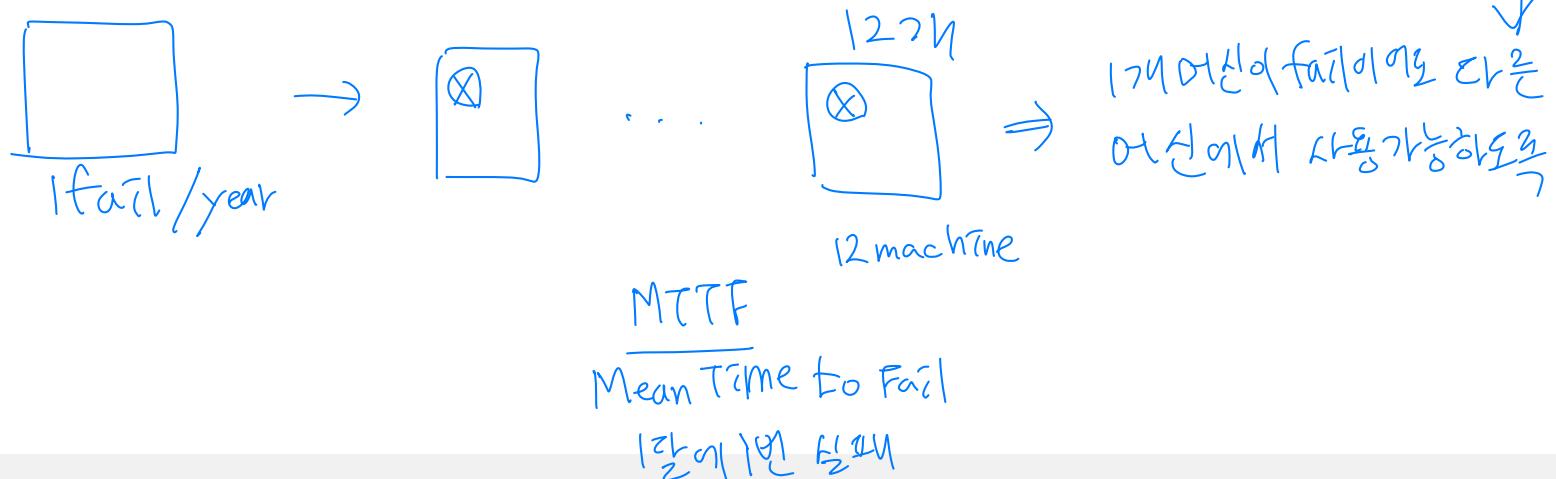


Column family



Key Value Storage Systems

- Key-value storage systems store **large numbers** (billions or even more) of **small** (KB-MB) sized records
- Records are **partitioned** across multiple machines
- Queries are **routed** by the system to appropriate machine
- **Records** are also **replicated** across **multiple machines**, to ensure availability even if a **machine fails**
 - Key-value stores ensure that updates are applied to all replicas, to ensure that their values are **consistent**



What is MongoDB?

- Developed by 10gen
 - Founded in 2007
- A document-oriented, NoSQL database
 - Hash-based, *schema-less database*
 - No Data Definition Language (*DDL*)
 - In practice, this means you can store hashes with any keys and values that you choose
 - Keys are stored as **strings**
 - **Document Identifiers** (`_id`) will be created for each document, field name reserved by system
 - Application tracks the schema and mapping
 - Uses **BSON** format
 - Based on JSON – B stands for Binary

Disk space
디스크 공간
↓

JSON : text → binary 변환
(BSON)

Functionality of MongoDB

- Dynamic schema
 - No DDL
- Document-based database
- Secondary indexes (값을 탐색할 때 유리함)
- Query language via an API
consistency 를 보장하지 X
- Atomic writes and fully-consistent reads
 - If system configured that way
- Master-slave replication with automated failover (replica sets)
*설정해놓고 끝나는 게
문제*
- Built-in horizontal scaling via automated range-based partitioning of data (sharding)
- No joins nor transactions

Why use MongoDB?

- Simple queries
- Functionality provided applicable to most web applications
- Easy and fast integration of data
 - No ERD diagram
- Not well suited for heavy and complex transactions systems

↳ schema less DB

consistency를 보장하지 않기

Data Model

- Stores data in form of **BSON (Binary JSON) documents**
- Group of related *documents* with a shared common index is a *collection*

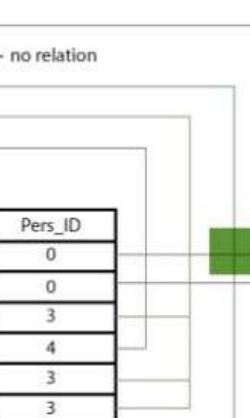
Relational

Person:

Pers_ID	Surname	First_Name	City
0	Miller	Paul	London
1	Ortega	Alvaro	Valencia
2	Huber	Urs	Zurich
3	Blanc	Gaston	Paris
4	Bertolini	Fabrizio	Rom

Car:

Car_ID	Model	Year	Value	Pers_ID
101	Bentley	1973	100000	0
102	Rolls Royce	1965	330000	0
103	Peugeot	1993	500	3
104	Ferrari	2005	150000	4
105	Renault	1998	2000	3
106	Renault	2001	7000	3
107	Smart	1999	2000	2



MongoDB Document

{

first_name: 'Paul',
surname: 'Miller'
city: 'London',
location: [45.123,47.232],
cars: [

{ model: 'Bentley',
year: 1973,
value: 100000, ... },
{ model: 'Rolls Royce',
year: 1965,
value: 330000, ... }

데이터
제작장치
(Join이 필요없다)

여러정보를 모으면
Join해야한다

→ Join 연산이 필요

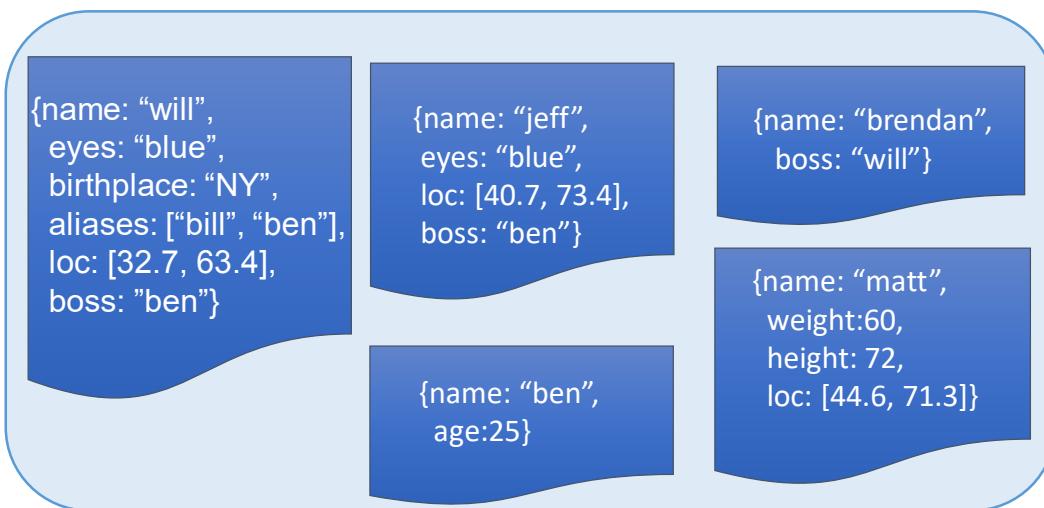
RDBMS에서 허용되지
(value만 가능하므로)

nested 구조로
구조화된 것을 고려한 설계

→ RDBMS 고려(차이점)

RDB Concepts to NO SQL

- MongoDB does not need any pre-defined data schema
- Every document could have different data
 - Schema Free



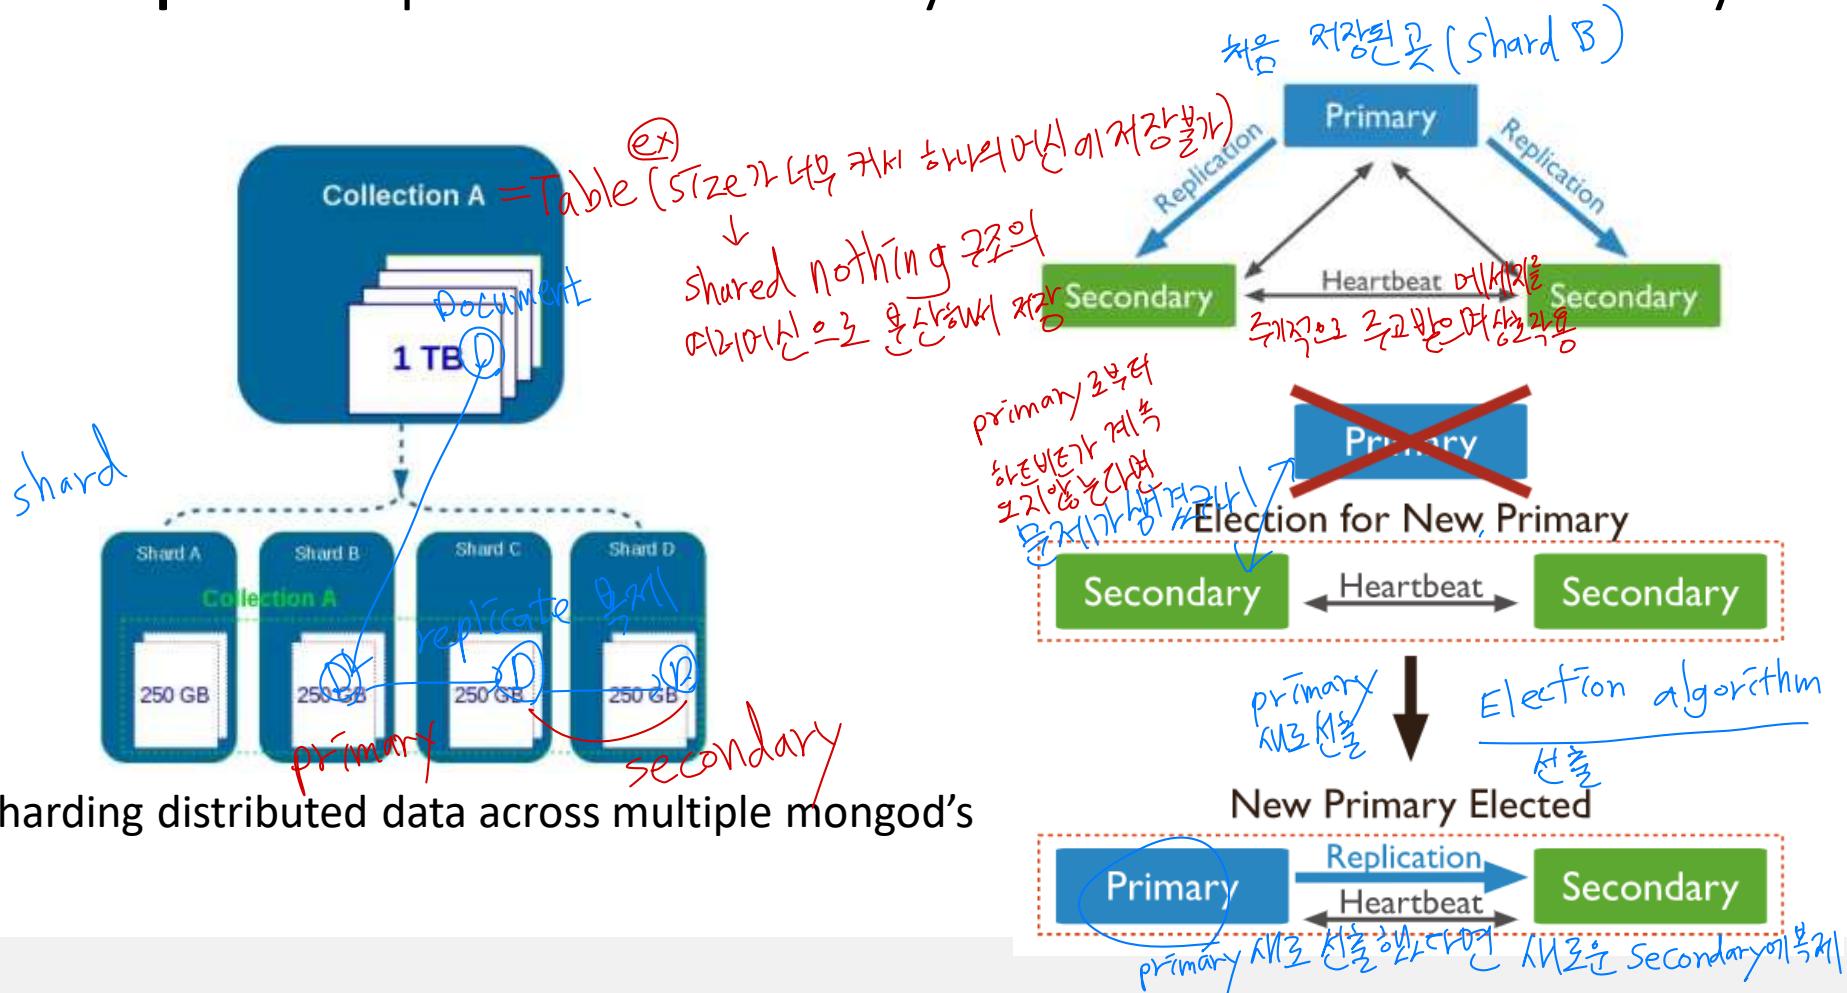
- RDBMS vs MongoDB

RDBMS	MongoDB
Database	Database
Table	Collection or Column Family
Row	Document (JSON, BSON)
Column	Field
Index	Index
Join	Embedded Document
Partition	Shard

- A MongoDB instance may have zero or more 'databases'
- A database may have zero or more 'collections'.
- A collection may have zero or more 'documents'.
- A document may have one or more 'fields'.
- MongoDB 'Indexes' function much like their RDBMS counterparts.

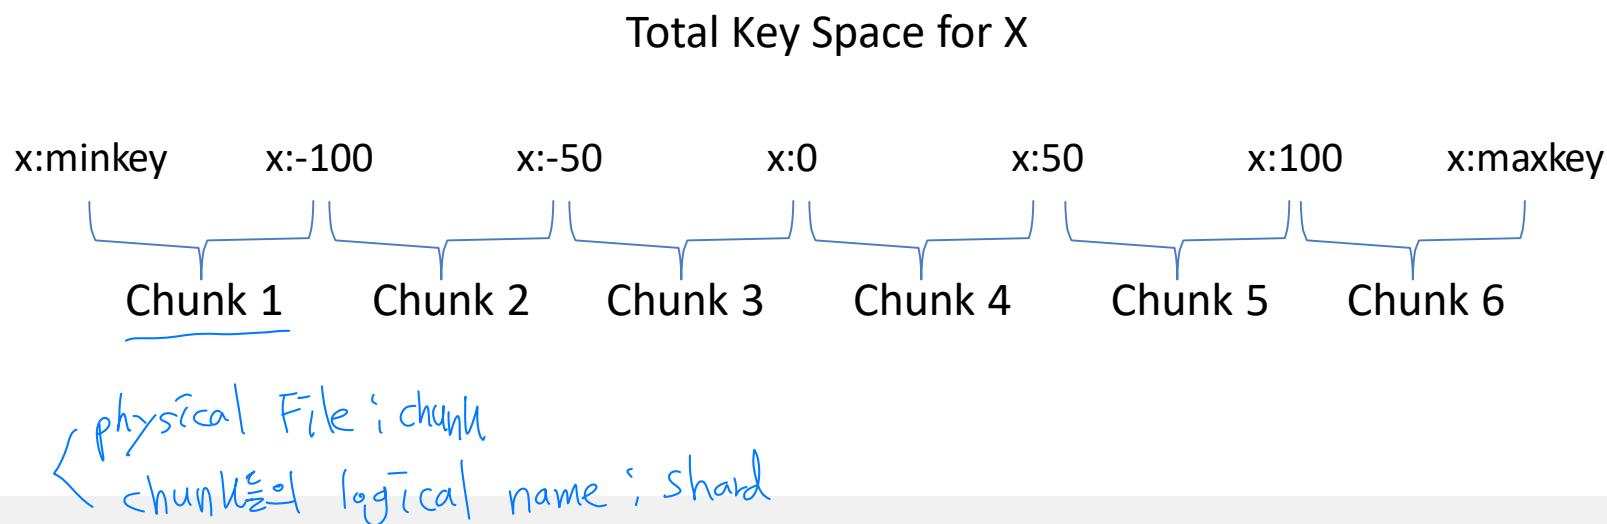
Sharding and Replication

- **Shard** (`mongod`) is a Mongo instance to handle a subset of original data.
 - **Replication** provides redundancy and increases data availability



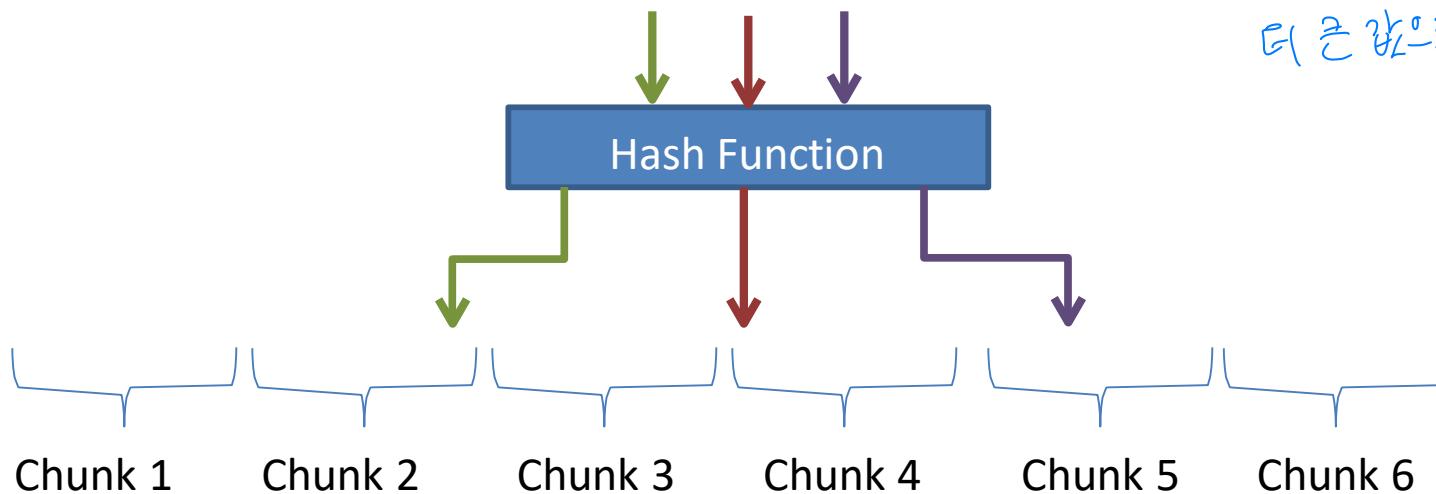
Range Sharding

- Shard Key: Single or compound field in schema used for data partitioning
- Partitions are called *chunks*. Two strategies:
 - [1] Range based: Shard Key Values are partitioned into ranges



Hash Sharding

- [2] Hash based: Hash of shard key values are partitioned into ranges



Hash % ?
? 같은 예선수요
더 큰 값으로!

- Range Queries are efficient for the first strategy
- Hash Scheme leads to better data balancing

↳ load balancing이 좋다
골고루 분배

100 < λ < 200
range query 를 나누는
원인 모든 shard를 탐색해야 함
해시는 지역성이 있어서

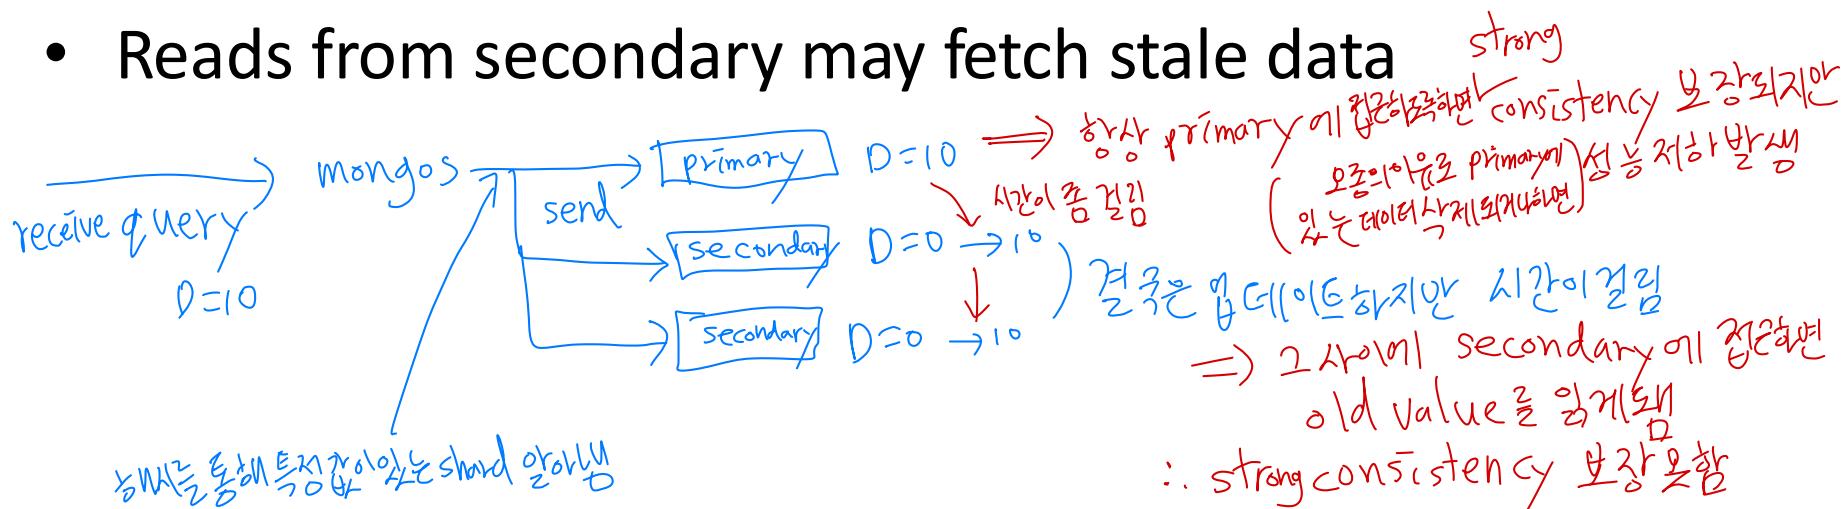
Balancing

- Splitting: Background process which splits when a chunks grows beyond a threshold
- Balancing: Migrates chunks among shards if there is an uneven distribution

특정 shard의 데이터가 올라다연 몇개의 key를 주변 shard로 옮겨서 데이터 분포의 균형을 맞춘다.

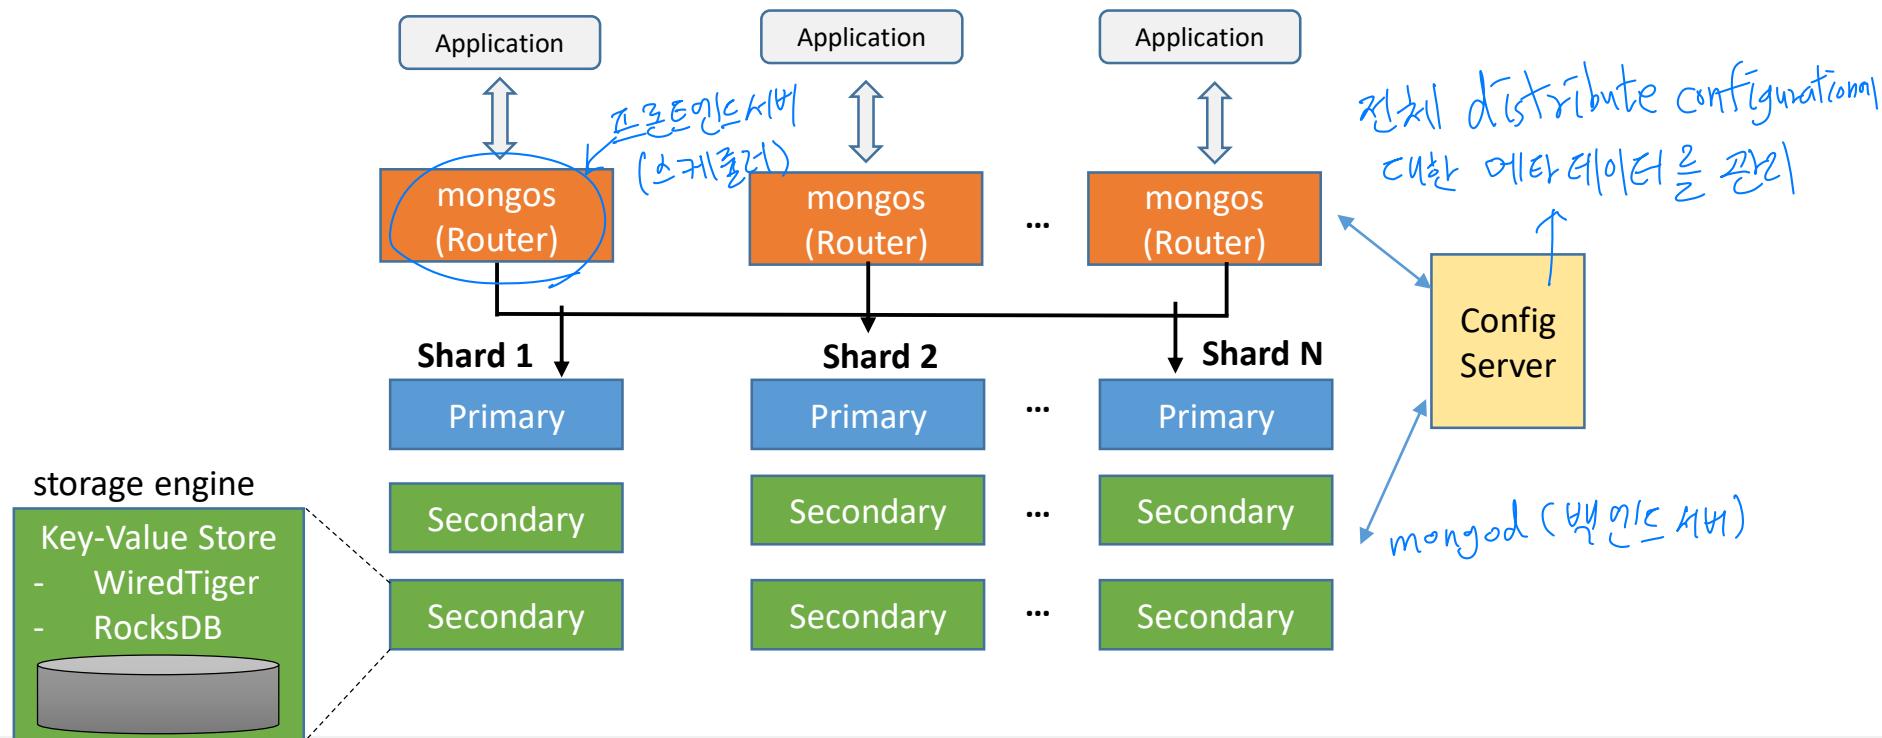
Read Preference with Replication

- Determine where to route read operation
- Default is primary. Possible options are secondary, primary-preferred, etc.
- Helps reduce latency, improve throughput
- Reads from secondary may fetch stale data

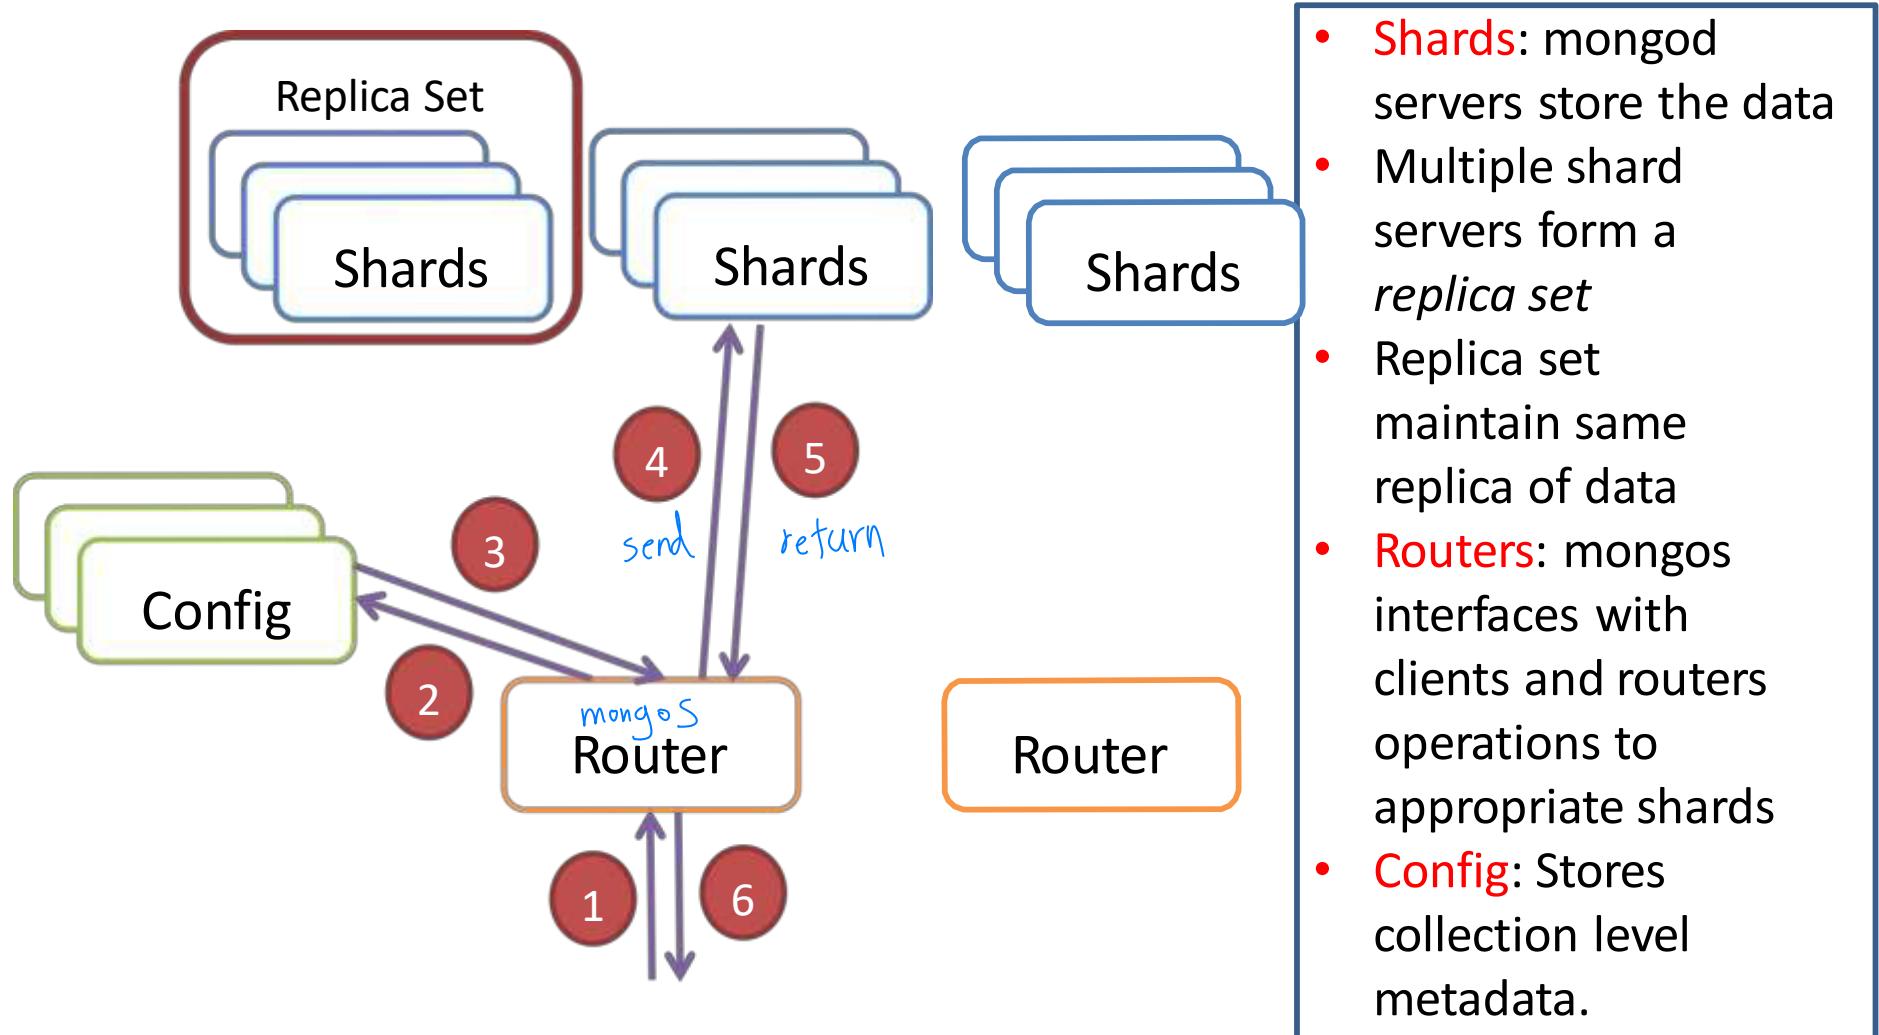


MongoDB Architecture

- **Mongos** is a query router to shards
 - Analogous to a database router.
 - Decides how many and which mongods should receive the query
 - Mongos collates the results, and sends it back to the client.
- **Config Server** is a Mongo instance which stores metadata



Typical MongoDB Query Processing



Consistency

- Strongly Consistent: Read Preference is ~~Master~~ primary *strong consistency* 
- Eventually Consistent: Read Preference is ~~Slave~~ Secondary
- CAP Theorem: Under partition, MongoDB becomes write unavailable thereby ensuring consistency

option

Query

Mongo DB → SQL 지원 X 하지만 비슷한건 있다.

- Query all employee names with salary greater than 18000 sorted in ascending order

db.users.find({**salary:{\$gt:18000}**}, {**name:1**}).sort(**{salary:1}**)

Collection =table
Condition
Projection
Modifier
inclusion
ascending
정렬 조건

{salary:25000, ...}
{salary:10000, ...}
{salary:20000, ...}
{salary:2000, ...}
{salary:30000, ...}
{salary:21000, ...}
{salary:5000, ...}
{salary:50000, ...}

→ {salary:25000, ...}
→ {salary:20000, ...}
→ {salary:30000, ...}
→ {salary:21000, ...}
→ {salary:50000, ...}

{salary:20000, ...}
{salary:21000, ...}
{salary:25000, ...}
{salary:30000, ...}
{salary:50000, ...}

JSON 파일

기본 고급 알고 DB 처리 적용하는 문제 안내용!!

Query Operators

Name	Description
\$eq	Matches value that are equal to a specified value
\$gt, \$gte	Matches values that are greater than (or equal to) a specified value
\$lt, \$lte	Matches values less than or (equal to) a specified value
\$ne	Matches values that are not equal to a specified value
\$in	Matches any of the values specified in an array
\$nin	Matches none of the values specified in an array
\$or	Joins query clauses with a logical OR returns all
\$and	Join query clauses with a logical AND
\$not	Inverts the effect of a query expression
\$nor	Join query clauses with a logical NOR
\$exists	Matches documents that have a specified field

Insert

- Insert a row entry for new employee Sally

```
db.users.insert({  
    name: "sally",  
    salary: 15000,  
    designation: "MTS",  
    teams: [ "cluster-management" ]  
})
```

JSON
key
obj

[:] array

Update

- All employees with salary greater than 18000 get a designation of Manager

```
db.users.update(  
    Update Criteria      {salary: {$gt:18000}},  
    Update Action        {$set: {designation: "Manager"}},  
    Update Option        {multi: true}  
)
```

- Multi option allows multiple document update

Delete

- Remove all employees who earn less than 10000

Remove Criteria

```
db.users.remove(  
    {salary:{$lt:10000}},  
)
```

- Can accept a flag to limit the number of document removal