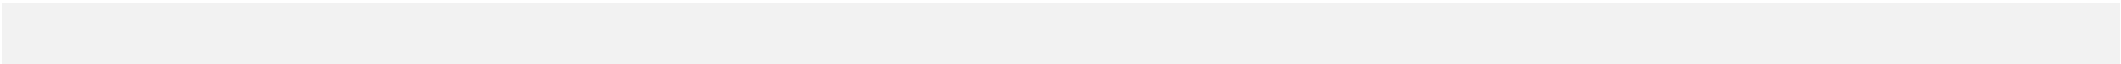


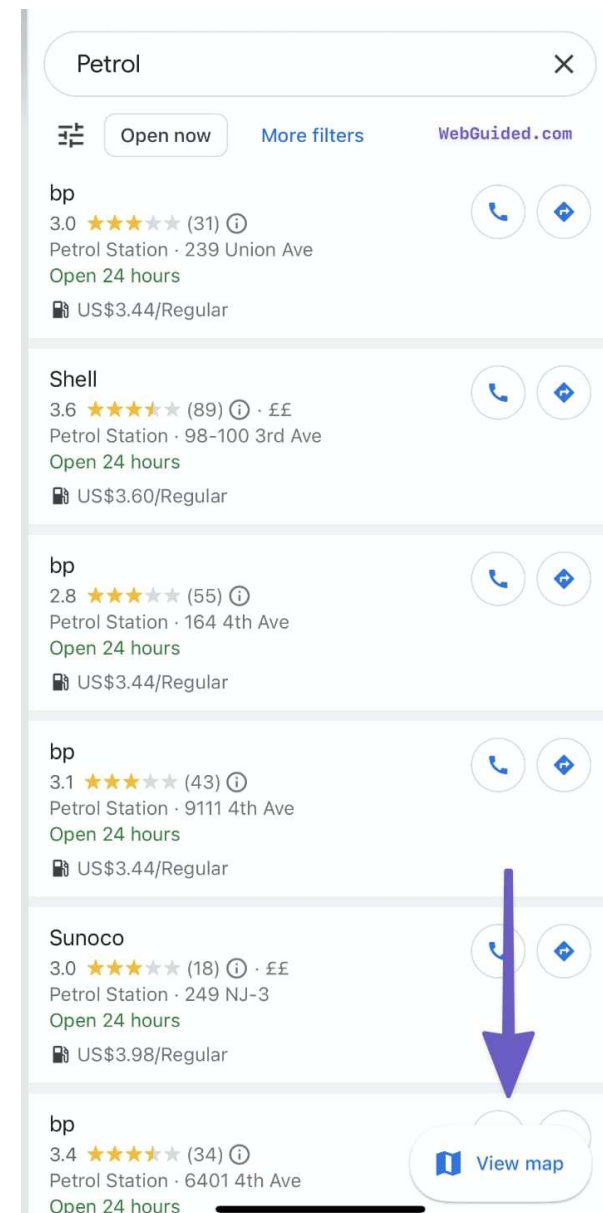
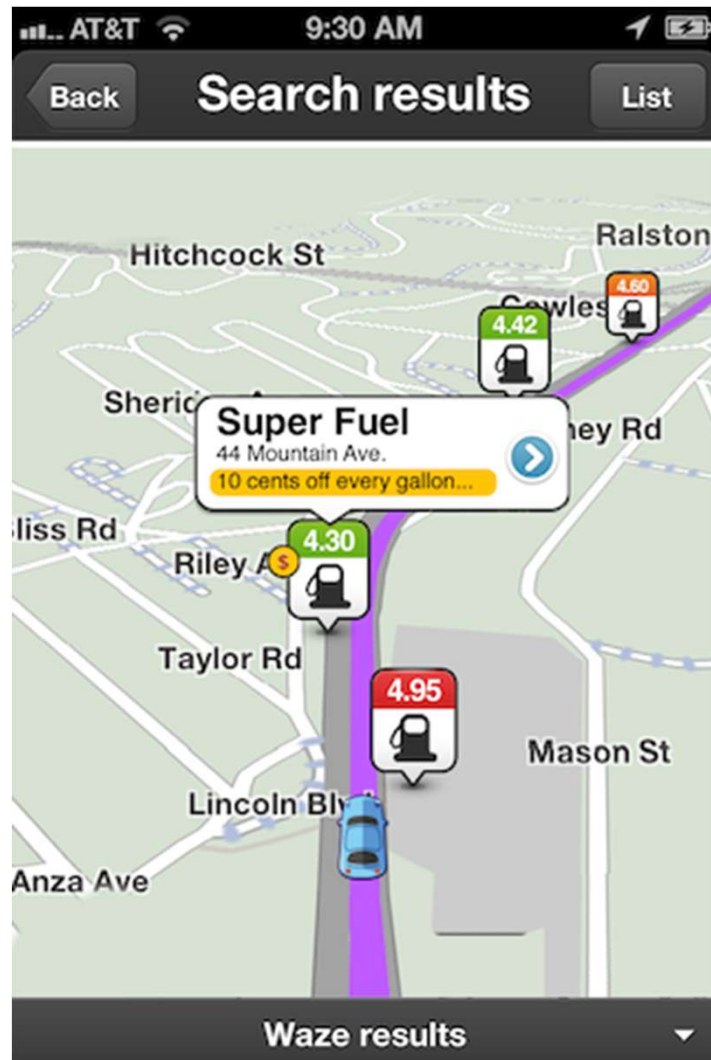
# Database Systems

## Lecture23 – Multi-Dimensional Index & Vector Database

Beomseok Nam (남범석)  
bnam@skku.edu




# Finding the closest gas station near me



# Spatial Data

multi dimension Data


- Data types such as points, lines, and polygons
- **Nearest neighbor queries**, given a point or an object, find the nearest object that satisfies given conditions.
- **Range queries** deal with spatial regions. e.g., ask for objects that lie partially or fully inside a specified region.




# K-D trees: Space-partitioning Method

BST의 확장 버전


- Each level of a K-D tree partitions the space into two.
  - In each node, choose one dimension for partitioning, cycling through the dimensions.



(a)




(b)




# Quadrees (space partitioning method)

4개 B 나누는 것

- The root node represents the entire target space.
- Each non-leaf node divides its region into four equal sized quadrants




Range query → 이걸 하나씩 열거해나가는 것이 아니라 backtracking이 필요!



# Range search


- Near neighbor (range search):
    - put the root on the stack *backtracking을 위해*
    - repeat
      - pop the next node  $T$  from the stack
      - for each child  $C$  of  $T$ :
        - if  $C$  is a leaf, examine point(s) in  $C$
        - if  $C$  intersects with the ball of radius  $r$  around  $q$ , add  $C$  to the stack
- 범위가 원하는 값과 overlap 된다면

# Nearest neighbor



# R-trees: Data-partitioning Method



- N-dimensional extension of B<sup>+</sup>-trees
- The **bounding box** of a node is a minimum sized rectangle that contains all the rectangles/polygons associated with the node
  - *Bounding boxes of children of a node are allowed to overlap*





# R-trees: Data-partitioning Method

- N-dimensional extension of B<sup>+</sup>-trees
- The **bounding box** of a node is a minimum sized rectangle that contains all the rectangles/polygons associated with the node
  - *Bounding boxes of children of a node are allowed to overlap*
  - *Range Query → Visit all overlapping child nodes → Backtracking*




# Nearest Neighbor Query (NN-Query)

- Start range search with  $r = \infty$ 
  - Or, guess a range  $r$  that contains at least one object say  $O$ 
    - if the current guess does not include any object, increase range size until an object found.
- put the root on the stack
- Repeat
  - pop the next node  $T$  from the stack
  - for each child  $C$  of  $T$ :
    - if  $C$  is a leaf, examine object(s) in  $C$
    - Whenever an object with smaller distance is found, update  $r$ ; Only investigate nodes with respect to current  $r$
    - if  $C$  intersects with the ball of radius  $r$  around  $q$ , add  $C$  to the stack


# R-trees: Data-partitioning Method

- Example: NN-Query Processing with R-tree



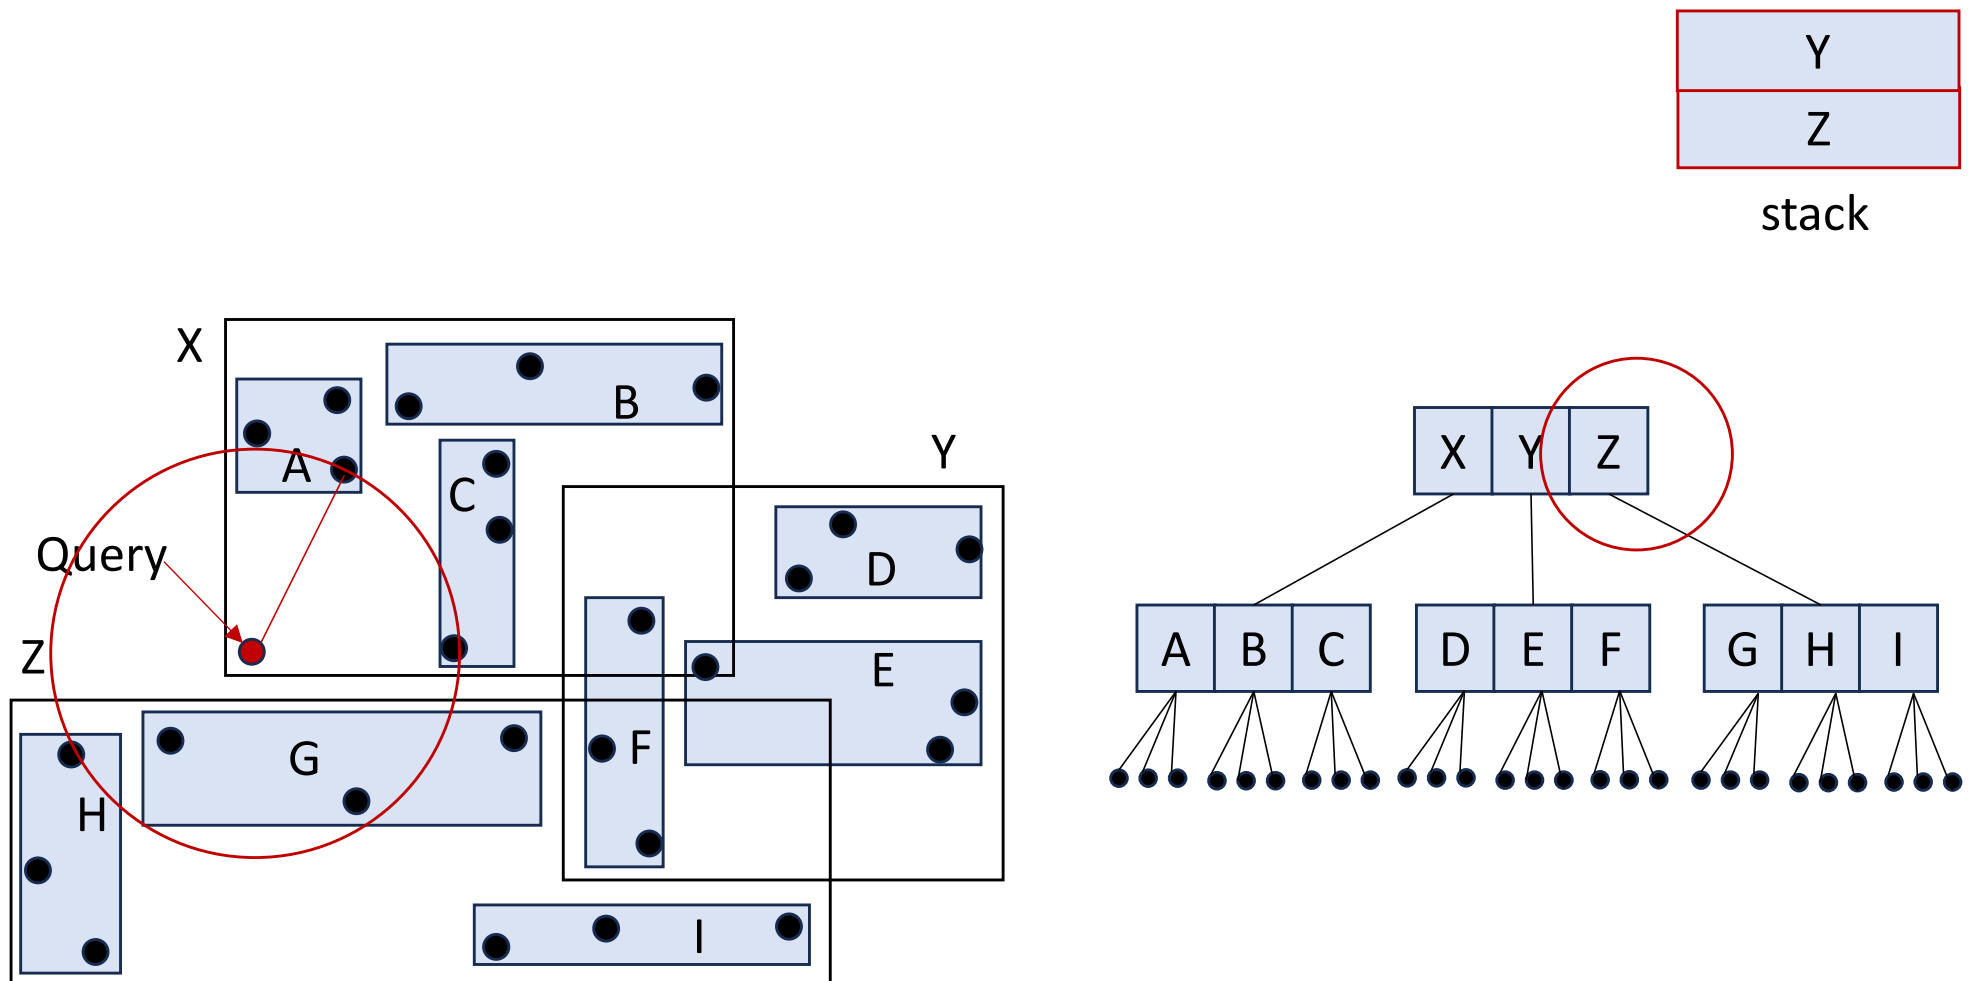
# R-trees: Data-partitioning Method

- Example: NN-Query Processing with R-tree



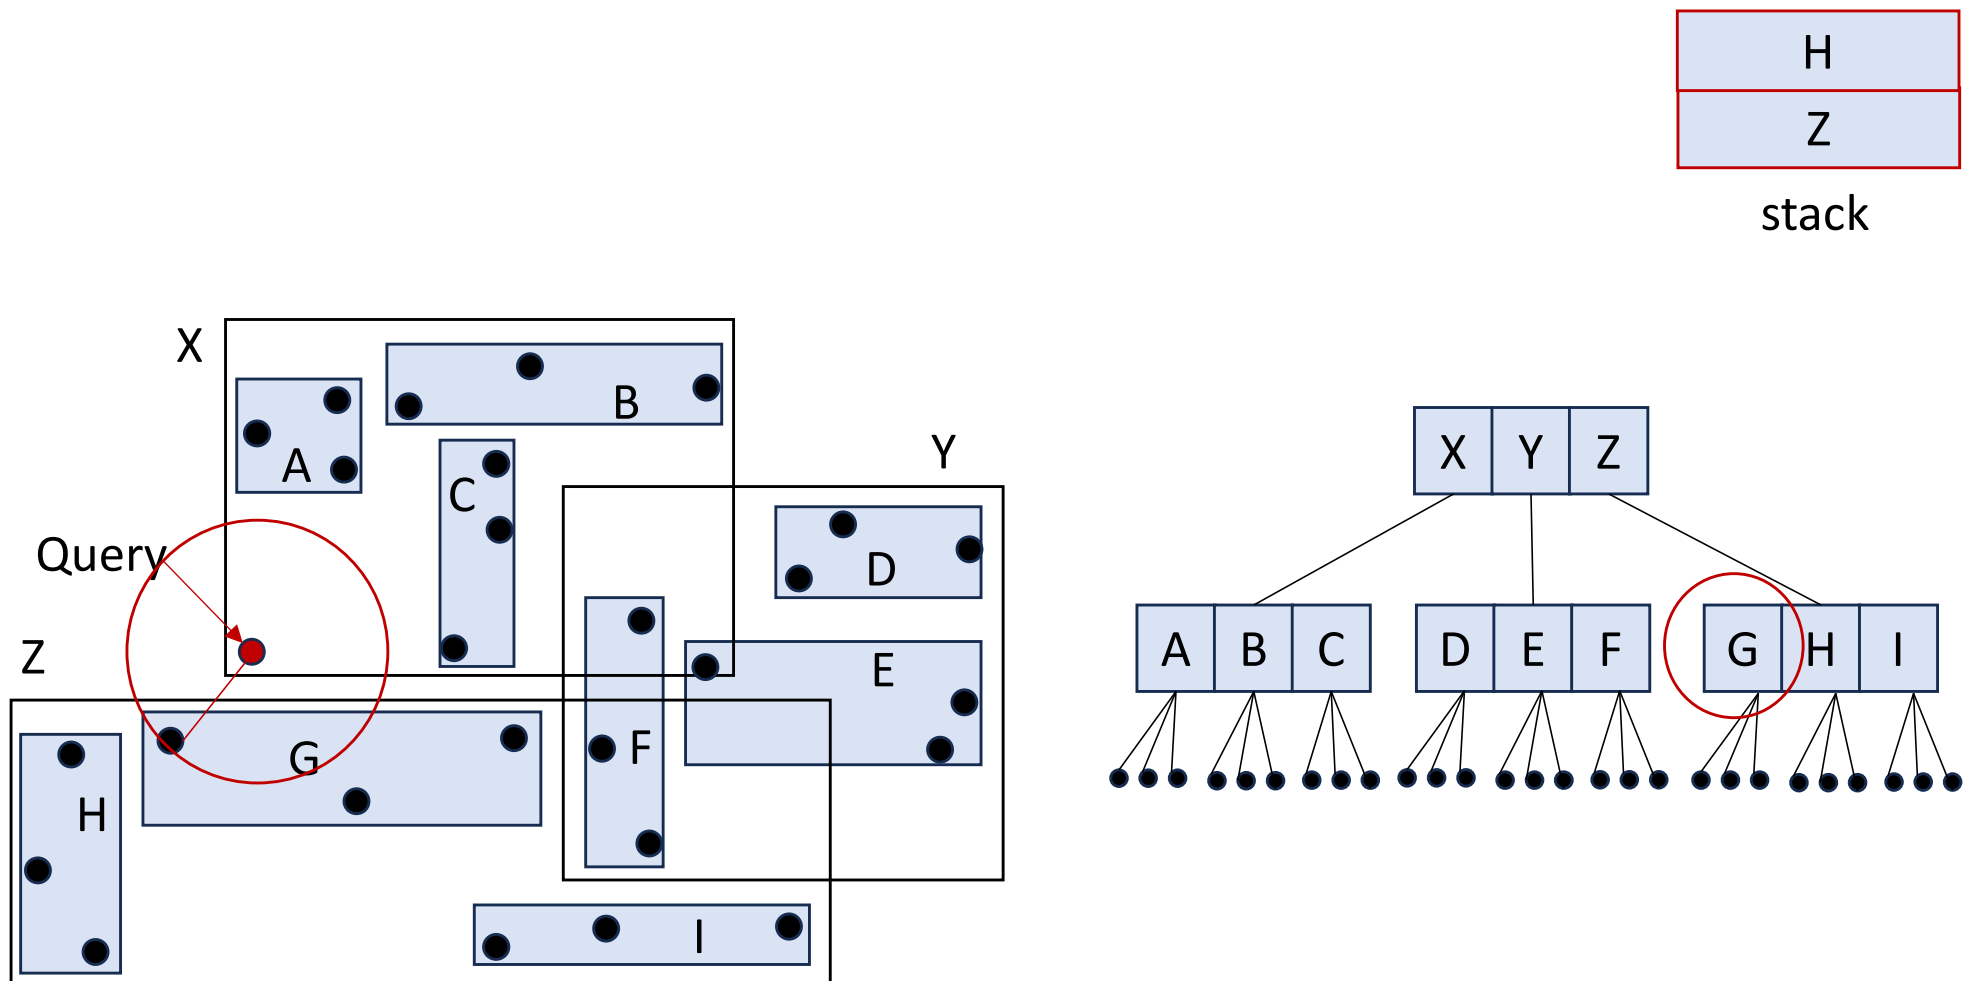
# R-trees: Data-partitioning Method

- Example: NN-Query Processing with R-tree



# R-trees: Data-partitioning Method

- Example: NN-Query Processing with R-tree



# Insert

RTree-Insert(T, entry)

leaf  $\leftarrow$  ChooseLeaf(T.root, entry) // locate place to insert

Insert entry into leaf

IF leaf overflows THEN

    newNode  $\leftarrow$  SplitNode(leaf)

    AdjustTree(leaf, newNode) // propagate changes upward

ELSE


    AdjustTree(leaf, NIL) // propagate MBB changes upward

IF root was split THEN


    create new root with children = previous root and newNode

# Split

- How to partition the  $M+1$  MBBs into two nodes?
  - 1. The total area of the two nodes is minimized
    - Large dead space hurts search performance
  - 2. The overlapping of the two nodes is minimized
    - Overlap causes backtracking and hurts performance
- Sometimes the two goals are conflicting



No overlap, large MBB



Overlap, small MBB



# Split

- Optimal solution: check every possible partition, complexity  $O(2^{M+1})$
- A quadratic algorithm:
  - Pick two “seed” entries  $e1$  and  $e2$  far from each other, that is to maximize  $\text{area}(\text{MBB}(e1,e2)) - \text{area}(e1) - \text{area}(e2)$ 
    - Here  $\text{MBB}(e1,e2)$  is the *minimum bounding box* containing both  $e1$  and  $e2$
    - complexity =  $O((M+1)^2)$
  - Insert the remaining  $(M-1)$  entries into the two groups
    - Continued on the next slide

# Split

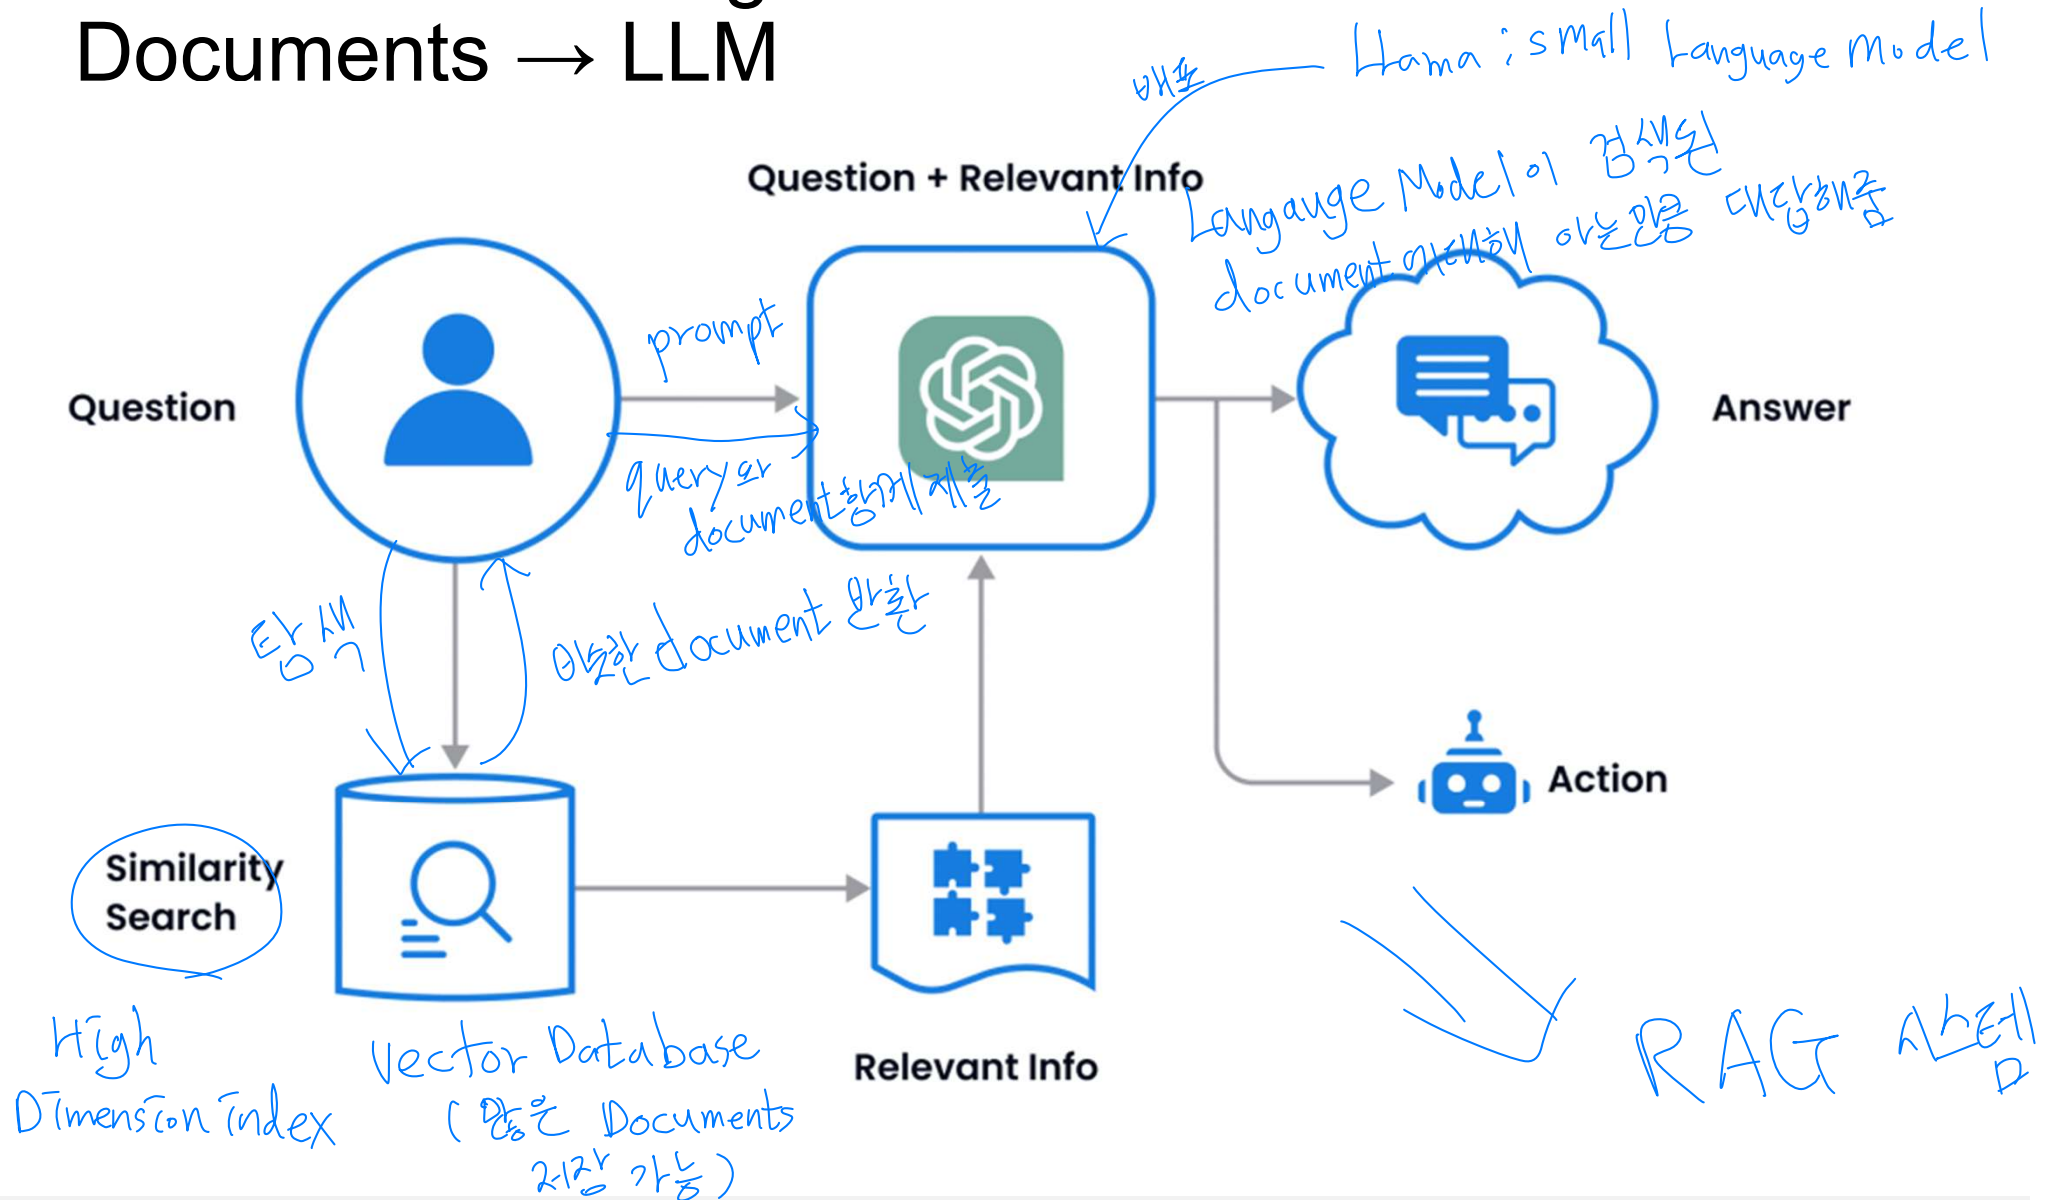
- A greedy method
  - At each step, pick an unassigned entry and assign it to one of the two groups based on:
    - Minimum area enlargement caused by adding the entry
    - If tied:
      - Select the group with smaller area
    - If still tied:
      - Select the group with fewer elements
  - Loop Until...
    - All entries are assigned, or
    - One group reaches  $(M - m + 1)$  entries
      - All remaining entries go to the other group
  - If the parent is also full, split the parent as well.

# LLM and Vector DB (Vector Store)

- **LLMs are context-limited** → Need external knowledge
- Vector DBs help retrieve **relevant documents** based on **semantic similarity**
- A Vector Store is a specialized database that stores and retrieves data using **vector embeddings** — numerical representations of text, images, or other unstructured data.
  - Enables **semantic search** (meaning-based, not exact keyword match)
  - Supports **context retrieval** in Retrieval-Augmented Generation (**RAG**)


# LLM and Vector DB (Vector Store)

- Text → Embedding → Vector Store → Similar Documents → LLM



# High-Dimensional Nearest Neighbor Search

- Example application: Reverse image search
  - Represent image by a vector
  - Pixel values arranged in a vector
  - More advanced features (SIFT, SURF, ORB)
  - Similar vectors  $\leftrightarrow$  similar images



# High-Dimensional Vectors

- 100-1000 dimensions
- **Curse of dimensionality**
  - Many methods scale poorly as the dimension increases
  - Considering one coordinate at a time is no longer enough

# Vector Indexes

Index Type	Description	Pros	Cons	Example Libraries
<b>Brute Force (Flat)</b>	Compares query against all vectors	100% accuracy	Very slow for large data	FAISS (IndexFlatL2)
<b>IVF (Inverted File Index)</b>	Clusters vectors, searches in relevant subsets	Fast, scalable	Slight accuracy drop	FAISS (IndexIVFFlat)
<b>HNSW (Hierarchical Navigable Small World)</b>	Graph-based approximate search	Very fast, high accuracy	Expensive to build	hnswlib, FAISS, Qdrant
<b>PQ (Product Quantization)</b>	Compresses vectors to save memory	Memory-efficient, scalable	Loss of precision	FAISS (IndexIVFPQ)
<b>Annoy</b>	Uses random projection trees	Lightweight, fast	Lower accuracy	Spotify Annoy
<b>Ball Tree / KD-Tree</b>	Traditional tree structures	Good for low dimensions	Poor performance in high-dim	Scikit-learn
<b>ScaNN (Google)</b>	Learned indexing for high recall	Fast and accurate	More complex to tune	ScaNN library

# ANN (Approximate Nearest Neighbors)

- Exact Nearest Neighbor (ENN):
  - Becomes computationally expensive in high-dimensional spaces (as dimensions increase, *distance calculations become less meaningful* and more costly).
  - Even brute-force methods with  $O(n)$  complexity outperforms in high-dimensional spaces.
- Approximate Nearest Neighbor (ANN):
  - Provides a trade-off between speed and accuracy.
  - Provides near-accurate results in a fraction of the time, making it viable for ML applications (e.g., 90-95% accurate results in milliseconds).



# Performance of ANN : Recall

- **Recall** is a metric to measure the **ability of correctly identifying all relevant instances** (i.e., all true positives).

- $$\text{Recall} = \frac{\text{True Positives (TP)}}{\text{True Positives (TP)} + \text{False Negatives (FN)}}$$

- **True Positives (TP)**: Correctly predicted positive cases.

- **False Negatives (FN)**: Actual positive cases that were incorrectly predicted as negative.


- Intuition

- **Recall** answers the question: *"Out of all the actual nearest neighbors, how many does it correctly find?"*
- It focuses on **minimizing missed positives**.

# ANN (Approximate Nearest Neighbors)


## ■ Proximity graph

- A **proximity graph** is a graph where each node represents a data point, and edges connect nodes based on their **proximity** (closeness) according to a specific distance metric (e.g., Euclidean distance, cosine similarity).
- Each node is connected to its **k-nearest neighbors**.




# Proximity Graph (2NN)

- edges connect vertices that are close to each other based on distance




# Proximity Graph (3NN)

- Full proximity graphs become too large in high-dimensions
  - Too many edges may introduce noise



# Sparse Neighborhood Graph (SNG)


- Proximity graph, where only a **subset of edges** are retained to reduce memory or computational cost



In undirected SNG, some nodes may have fewer than  $K$  edges


# Sparse Neighborhood Graph (SNG)

- Nearest nodes are not always selected as neighbors




# HNSW – Hierarchical Navigable-Small World

- **Hierarchical Navigable Small World (HNSW)** is an algorithm used for Approximate Nearest Neighbor (ANN) search.
- HNSW constructs hierarchical graphs where each node is connected to a set of nearby neighbors, creating a "small world" with short paths between any two points



# HNSW – Hierarchical Navigable-Small World

- **Multiple Layers:** Nodes are organized in a hierarchy of layers of proximity graphs (similar to SkipLists).
  - **Lower layers** have denser connections and capture local neighborhoods.
  - **Higher layers** are sparsely connected and provide global navigation across the graph.
- Greedy search in each layer
- Elements inserted one by one by searching in so far constructed index





# HNSW – Hierarchical Navigable-Small World

## ■ Index Construction:

1. Insert data points into multiple layers of the graph.
2. High-level layers capture global relationships; lower-level layers store local neighborhood information.
3. Each node connects to a subset of the most similar points in its layer.

## ■ Search Process:

1. Start at the **topmost layer** with a randomly selected node.
2. Traverse the graph, moving to closer neighbors until reaching the bottom layer.
3. In the **lowest layer**, perform a local search to find the nearest neighbors.