

SWE3003 Introduction to Database Systems - Midterm Spring 2024

Student ID	Name

For Instructor/TA only,

Q1	Q2	Q3	Q4	Q5	Q6	Q7	Total

Academic Honor Pledge

I affirm that I will **not** at any time be involved with **cheating** or **plagiarism** while enrolled as a student of **Introduction to Database Systems** class at Sungkyunkwan University.

I understand that violation of this code will result in penalties as severe as indefinite suspension from the university.

Your signature: _____

1. [30 pts] For each of the following statements, indicate whether it is TRUE or FALSE. You will get 3 points for each correct answer, -3 points for each incorrect answer, and 0 points for each blank answer or both marked answers.

	T	F
(a) A attribute is a property that describes the relationship between data and the entities that use it.	F <input type="checkbox"/>	<input type="checkbox"/>
(b) A candidate key is also a superkey, but there may be superkeys that are not candidate keys.	T <input type="checkbox"/>	<input type="checkbox"/>
(c) A weak entity set and its strong entity set must be stored as a single table.	F <input type="checkbox"/>	<input type="checkbox"/>
(d) One-to-one relationship set cannot be stored as a single table.	F <input type="checkbox"/>	<input type="checkbox"/>
(e) Data dictionary manages statistics about table accesses, such as the number of times a table is accessed, the frequency of particular queries, and other usage patterns.	T <input type="checkbox"/>	<input type="checkbox"/>
(f) In the WHERE clause of a SQL query, the condition NULL = NULL is evaluated to be true.....	F <input type="checkbox"/>	<input type="checkbox"/>
(g) Records cannot be updated in the underlying table through a view.	F <input type="checkbox"/>	<input type="checkbox"/>
(h) Using 'Statement' objects helps prevent SQL injection attacks.	F <input type="checkbox"/>	<input type="checkbox"/>
(i) You can only create one primary index (clustered index) on a table.	T R <input type="checkbox"/>	<input type="checkbox"/>
(j) Secondary indexes (unclustered indexes) are not effective when executing range queries	F <input type="checkbox"/>	<input type="checkbox"/>

2. [20 pts] Consider the following schema for a medical appointment scheduling program.

- Doctors(doctorid, dr_name)
- Patients(patientid, name)
- Appointments(doctorid, patientid, date)

Write a relational algebra for the following queries.

- (a) Find the names of patients who have appointments with 'Dr. Kim' on '2024-04-22':

answer:

$$\Pi_{\text{name}} (\text{Patients} \bowtie (\sigma_{\text{date}='2024-04-22'} (\text{Appointments}) \bowtie \sigma_{\text{dr_name}='Dr.Kim'} (\text{Doctors})))$$

- (b) Find the names of doctors who have appointments with other doctors on '2024-04-22'. Assume everyone's name is unique.

answer:

$$\Pi_{\text{dr_name}} ((\sigma_{\text{date}='2024-04-22'} (\text{Appointments}) \bowtie \text{Patients}) \bowtie_{\text{name} = \text{dr_name}} \text{Doctors})$$

- (c) Find the names of doctors who do not have appointments on '2024-04-22'.

answer:

$$\Pi_{\text{dr_name}} (\text{Doctors}) - \Pi_{\text{dr_name}} (\text{Doctors} \bowtie \sigma_{\text{date}='2024-04-22'} (\text{Appointments}))$$

- (d) Find the names of patients who have appointments with all doctors.

answer:

$$\Pi_{\text{name}} ((\text{Patients} \bowtie \text{Appointments}) \div \Pi_{\text{doctorid}} (\text{Doctors}))$$

Or

$$\begin{aligned} & \Pi_{\text{name}} (\text{Patients} \bowtie \text{Appointments}) - \\ & \Pi_{\text{name}} (\Pi_{\text{name}} (\text{Patients} \bowtie \text{Appointments}) \times \Pi_{\text{doctorid}} (\text{Doctors}) - \\ & \quad (\Pi_{\text{name,doctorid}} (\text{Patients} \bowtie \text{Appointments}))) \end{aligned}$$

3. [20 pts] Write each of the following queries in SQL for the given relations.

- Doctors(doctorid, dr_name)
- Patients(patientid, name)
- Appointments(doctorid, patientid, date)

(a) Find the dates on which Doctor ID 1 and Patient ID 1 have appointments but Doctor ID 2 and Patient ID 2 do not have appointments.

```
SELECT date
FROM Appointments
WHERE doctorid='1' AND patientid='1'
      AND date NOT IN (SELECT date
                        FROM Appointments
                        WHERE doctorid='2' AND patientid='2')
```

.

(b) Find the names of patients who have appointments with doctors on days when the doctors have 20 or more appointments.

```
With BusyDoctors(doctorid, date) AS
  (SELECT doctorid, date
   FROM Appointments
   GROUP BY doctorid, date
   HAVING count(*) >= 20
  )
SELECT name
FROM Patients NATURAL JOIN Appointments
      NATURAL JOIN BusyDoctors
```

.

(Cont'd) Write each of the following queries in SQL for the given relations.

- Doctors(doctorid, dr_name)
- Patients(patientid, name)
- Appointments(doctorid, patientid, date)


(c) Find the names of patients who have not had any appointments with doctors before '2024-04-22'.

```
SELECT name
FROM Patients
WHERE name NOT IN (
    SELECT name
      FROM Patients NATURAL JOIN Appointments
     WHERE date < '2024-04-22'
);
.
```


(d) Find the names of patients who have appointments with all doctors.

```
SELECT name
FROM Patients AS P
WHERE NOT EXISTS (
    (SELECT doctorid
      FROM Doctors)
  EXCEPT
  (SELECT A.doctorid
    FROM Appointments AS A
   WHERE A.patientid = P.patientid
  )
);
.
```

4. [15 pts] The figure below illustrates the relationships between the following entity sets and constraints about cardinality and total/partial participation.
- professor** entity set has professor_id and **name** attributes.
- project** entity set has proj_no, and **budget** attributes.
- student** entity set has student_id, **name**, and **department_name** attributes.



- (a) Draw an ER diagram for the three entity sets and two relationship sets.



- (b) How many database tables will be needed at minimum? Justify your answer

'supervise' relationship set can be merged into the 'project' entity set since it is one-to-many relationship set. 'work' relationship set can not be merged since it is many-to-many relationship set. So, the answer is 4.

- (c) Write a SQL DDL statement to create the 'project' table. All attributes are of VARCHAR(10) type.

```
CREATE TABLE project (
    professor_id VARCHAR(10) ,
    proj_no VARCHAR(10),
    budget VARCHAR(10)
    PRIMARY KEY (professor_id, proj_no)
    FOREIGN KEY (professor_id) REFERENCES professor(professor_id)
);
```

partial credits give for the following answer

```
CREATE TABLE project (
    proj_no VARCHAR(10) PRIMARY KEY,
    budget VARCHAR(10)
);
```

5. [15 pts] Consider the following relation and the functional dependency set:

- $R(A, B, C, D, E, F, G, H)$
- $FD = \{BC \rightarrow GH, AD \rightarrow E, A \rightarrow H, E \rightarrow BCF, G \rightarrow H\}$

(a) Find a candidate key of R.

answer: AD

AD

$\rightarrow ADE$

$\rightarrow ADEH$

$\rightarrow ABCDEF$

$\rightarrow ABCDEFGH$

(b) Decompose the relation R into BCNF.

answer: ADE, BCEF, GH, BCG

(c) Is the decomposition lossless? Justify your answer.


answer:

This is a lossless decomposition. It seems we lost $BC \rightarrow H$ and $A \rightarrow H$.


But, $BC \rightarrow G$ and $G \rightarrow H$.

So, we can reconstruct the original tuples when we join GH and BCG. Similarly, $A \rightarrow H$ can be reconstructed by joining decomposed tables.

6. [10 pts] Consider the following B+tree with degree 3, i.e, the maximum number of child nodes is 3.



- (a) Draw a tree structure after inserting 35 into the tree.



- (b) Suppose a single B+ tree node occupies 64 bytes. Each key is 8 bytes, and each pointer is 8 bytes as well. In a B+ tree containing 8 keys, what is the minimum and maximum number of nodes a query might access to find a specific key? Show how you calculated the answer.

Minimum: 2

Maximum: 3

Leaf nodes can have maximum 3 keys.

Internal nodes can have maximum 3 keys and 4 child pointers.

Due to 50% utilization, each leaf must have 2 keys at minimum and each internal node must have one key and 2 child pointers.

So, in the best case, there will be only two levels.

$$\begin{array}{c} 8 \text{ key } 8 \text{ key } 8 \text{ (key } 8) \text{ (8)} \\ 16 \text{ } 16 \text{ } 16 \text{ (8) } 8 \quad 16 \text{ } 16 \text{ } 16 \text{ (8) } 8 \quad 16 \text{ } 16 \text{ } 16 \text{ (8) } 8 \end{array}$$


In the worst case, there will be three levels.

$$\begin{array}{c} 8 \text{ key } 8 \text{ (key } 8 \text{ key } 8) \text{ (8)} \\ 8 \text{ key } 8 \text{ (key } 8 \text{ key } 8) \text{ (8)} \quad 8 \text{ key } 8 \text{ (key } 8 \text{ key } 8) \text{ (8)} \\ 16 \text{ } 16 \text{ (16) (8) } 8 \quad 16 \text{ } 16 \text{ (16) (8) } 8 \quad 16 \text{ } 16 \text{ (16) (8) } 8 \quad 16 \text{ } 16 \text{ (16) (8) } 8 \end{array}$$

7. [10 pts] Suppose each bucket can hold up to 2 records, and the least significant bits are used as the hash key. Suppose you insert the following 10 keys.

0001, 0010, 0011, 0100, 0101, 0110, 0111, 1000, 1001, 1010

- (1) Draw an extendable hash table after inserting the 10 keys.



- (2) Draw a linear hash table after inserting the 10 keys.

